

Breast Cancer Classification Study

DIN Sokheng, RA Veasna

2025-11-10

```
library(ggplot2)
library(dplyr)
library(tidyverse)
library(corrplot)
library(ggc当地)
library(naniar)
library(visdat)
library(rstatix)
library(DescTools)
library(car)
library(limma)
library(survival)
library(pheatmap)
library(diptest)
library(forcats)
library(glmnet)
library(caTools)
library(pROC)
library(gridExtra)
library(smotefamily)
library(reshape2)

source("function_model_fit.R", local = knitr::knit_global())

cat("Functions loaded successfully: fit_all_models, plot_model_comparison, apply_smote, etc.")

## Functions loaded successfully: fit_all_models, plot_model_comparison, apply_smote, etc.
```

Classification of Cancer outcome using Genetic and Clinical data

Introduction

Breast cancer outcome prediction relies on both **clinical variables** (patient, tumor, treatment information) and **genomic features** (mRNA expression). This project analyzes a dataset of 1231 patients with 24 clinical variables and 5000 high-variance genes, aiming to:

- Understand clinical variables associated with survival
- Explore gene expression characteristics

- Identify differentially expressed genes
- Detect subgroups of patients (clustering, PCA)
- Evaluate associations between clinical and genomic factors
- Perform survival analysis (Kaplan–Meier)
- Analyse multicollinearity and variable relevance
- Provide a unified understanding of prognostic factors

The main outcome is `vital_status`: “Alive” vs “Dead”.

Load data

```

load("mrr_bio.Rdata")

# Load dataset
load("mrr_bio.Rdata")

# Clinical data
clinical_df <- as.data.frame(clinical_data)
genex_df    <- as.data.frame(GeneX)

cat("Clinical samples:", nrow(clinical_df), "\n")

## Clinical samples: 1231

cat("Clinical variables:", ncol(clinical_df), "\n")

## Clinical variables: 24

cat("Gene samples:", nrow(genex_df), "\n")

## Gene samples: 1231

cat("Genes:", ncol(genex_df), "\n")

## Genes: 5000

# Outcome variable
table(clinical_df$vital_status)

##
## Alive  Dead
## 1029   201

```

```

# Clinical variables: gender
unique(clinical_data$gender) # unique values

## [1] "female" "male"   NA

table(clinical_data$gender) # frequency

##
## female    male
## 1217      13

```

Clinical data studies

Dataset Structure & Variable Types

```

numeric_vars      <- names(clinical_df)[sapply(clinical_df, is.numeric)]
categorical_vars <- names(clinical_df)[sapply(clinical_df, is.character)]

cat("Numeric variables (", length(numeric_vars), "):\n", paste(numeric_vars, collapse=", "), "\n\n")

## Numeric variables ( 5 ):
## initial_weight, age_at_diagnosis, days_to_last_follow_up, age_at_index, days_to_birth

cat("Categorical variables (", length(categorical_vars), "):\n", paste(categorical_vars, collapse=", "))

## Categorical variables ( 14 ):
## tissue_type, laterality, tissue_or_organ_of_origin, primary_diagnosis, prior_treatment, ajcc_patholo

```

Missing Data Analysis

```

# Calculate missing statistics
total_missing <- sum(is.na(clinical_df))
total_cells   <- prod(dim(clinical_df))
missing_ratio <- total_missing / total_cells
missing_count <- colSums(is.na(clinical_df))
missing_pct   <- round(missing_count / nrow(clinical_df) * 100, 2)

missing_table <- data.frame(
  Column        = names(clinical_df)
, Missing_Count = missing_count
, Missing_Pct   = missing_pct
)

# Filter columns with missing values
missing_table_filtered <- missing_table[missing_table$Missing_Count > 0, ]

cat("Variables with missing data:\n")

```

```

## Variables with missing data:

print(missing_table_filtered[order(-missing_table_filtered$Missing_Count), ])

##                                     Column Missing_Count
## ajcc_pathologic_t                  ajcc_pathologic_t      100
## laterality                         laterality             94
## follow_ups_disease_response       follow_ups_disease_response    77
## age_at_diagnosis                  age_at_diagnosis        55
## prior_treatment                   prior_treatment        45
## days_to_birth                     days_to_birth          17
## initial_weight                    initial_weight          15
## diagnosis_is_primary_disease     diagnosis_is_primary_disease   4
## days_to_last_follow_up           days_to_last_follow_up    4
## age_is_obfuscated                age_is_obfuscated        4
## tissue_or_organ_of_origin        tissue_or_organ_of_origin    1
## primary_diagnosis                primary_diagnosis        1
## morphology                        morphology             1
## classification_of_tumor          classification_of_tumor    1
## race                             race                  1
## gender                           gender                 1
## ethnicity                        ethnicity              1
## vital_status                      vital_status            1
## age_at_index                      age_at_index            1
##                                     Missing_Pct
## ajcc_pathologic_t                  8.12
## laterality                        7.64
## follow_ups_disease_response       6.26
## age_at_diagnosis                  4.47
## prior_treatment                   3.66
## days_to_birth                     1.38
## initial_weight                    1.22
## diagnosis_is_primary_disease     0.32
## days_to_last_follow_up           0.32
## age_is_obfuscated                0.32
## tissue_or_organ_of_origin        0.08
## primary_diagnosis                0.08
## morphology                        0.08
## classification_of_tumor          0.08
## race                             0.08
## gender                           0.08
## ethnicity                        0.08
## vital_status                      0.08
## age_at_index                      0.08

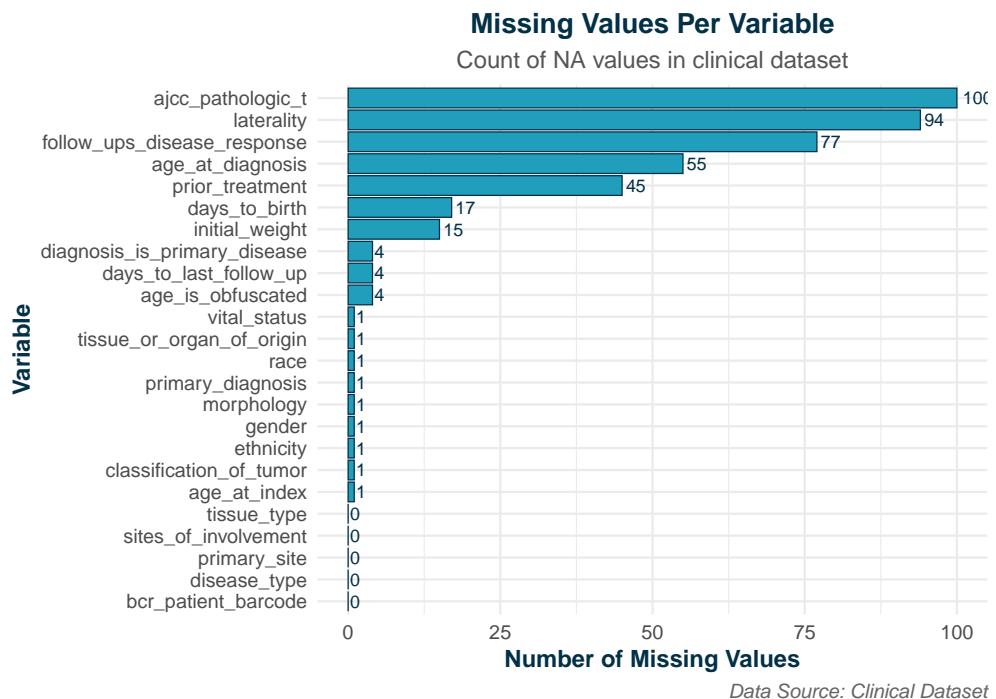
# Visualize missing data
ggplot(missing_table
       , aes(x = reorder(Column, Missing_Count), y = Missing_Count)) +
  geom_bar(stat = "identity", fill = "#219ebc", color = "#023047", linewidth = 0.3) +
  geom_text(aes(label = Missing_Count), hjust = -0.2, size = 3, color = "#023047") +
  coord_flip() +
  labs(title = "Missing Values Per Variable"
       , subtitle = "Count of NA values in clinical dataset")

```

```

, x = "Variable"
, y = "Number of Missing Values"
, caption = "Data Source: Clinical Dataset") +
theme_minimal(base_size = 12) +
theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047")
, plot.subtitle = element_text(hjust = 0.5, color = "#555555")
, plot.caption = element_text(face = "italic", color = "#666666")
, axis.title = element_text(face = "bold", color = "#023047"))

```



Data Cleaning

```
# Remove samples with missing vital_status
cat("Before cleaning:", nrow(clinical_data), "samples\n")
```

```
## Before cleaning: 1231 samples
```

```
valid_idx <- !is.na(clinical_data$vital_status)
clinical_df <- clinical_data[valid_idx, ]
GeneX_df <- GeneX[valid_idx, ]

cat("After removing:", nrow(clinical_df), "samples\n")
```

```
## After removing: 1230 samples
```

```
cat("Removed:", sum(!valid_idx), "sample(s)\n\n")
```

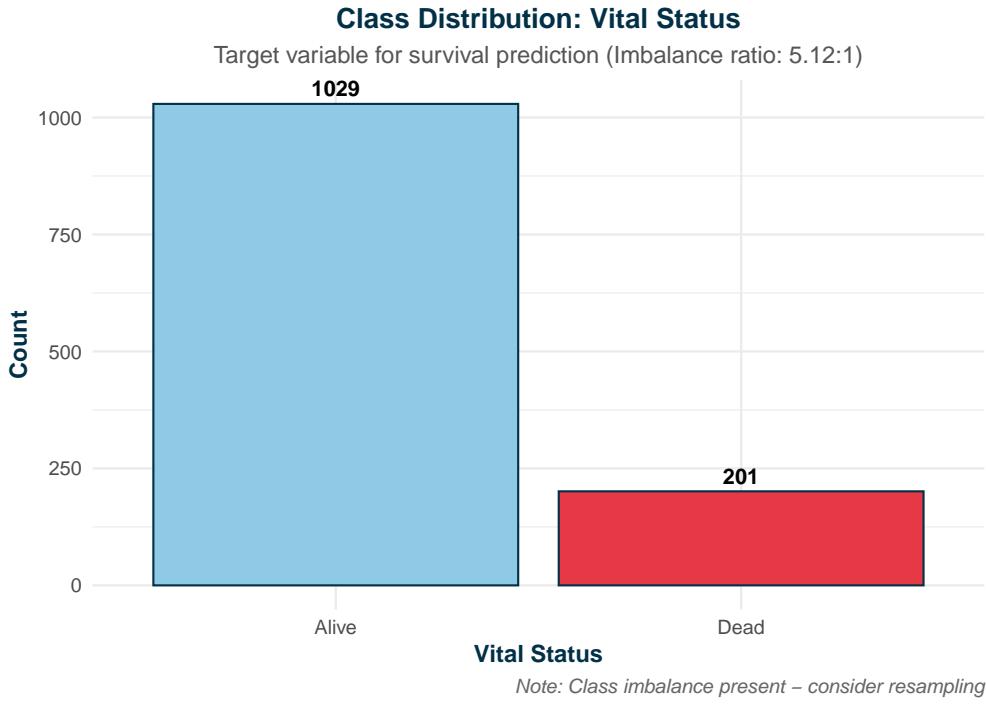
```
## Removed: 1 sample(s)
```

Key clinical predictors:

- age_at_index (continuous)
- initial_weight (continuous)
- ajcc_pathologic_t (tumor stage)
- prior_treatment (yes/no)
- primary_diagnosis (tumor subtype)
- race, gender (demographics)

Class Balance Check

```
# Visualize class distribution
ggplot(clinical_df, aes(x = vital_status, fill = vital_status)) +
  geom_bar(color = "#023047", linewidth = 0.5) +
  geom_text(stat = "count"
            , aes(label = after_stat(count))
            , vjust = -0.5
            , fontface = "bold"
            , size = 4) +
  scale_fill_manual(values = c("Alive" = "#8ecae6", "Dead" = "#e63946")) +
  labs(title = "Class Distribution: Vital Status"
       , subtitle = "Target variable for survival prediction (Imbalance ratio: 5.12:1)"
       , x = "Vital Status"
       , y = "Count"
       , caption = "Note: Class imbalance present - consider resampling") +
  theme_minimal(base_size = 12) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047")
        , plot.subtitle = element_text(hjust = 0.5, color = "#555555")
        , plot.caption = element_text(face = "italic", color = "#666666")
        , axis.title = element_text(face = "bold", color = "#023047")
        , legend.position = "none")
```



```
cat("Imbalance: 5.12:1 (Alive:Dead)\n")
```

```
## Imbalance: 5.12:1 (Alive:Dead)
```

The outcome variable shows a marked imbalance of 5.12:1 (Alive : Dead), which is expected in clinical studies with right-censoring. This imbalance has no negative impact on exploratory data analysis, but it will require careful handling in later predictive modeling

Numerical Data Visualizations

```
numeric_vars <- c("age_at_index"
                 , "age_at_diagnosis"
                 , "initial_weight"
                 , "days_to_last_follow_up"
                 , "days_to_birth")

par(mfrow = c(5, 4), bg = "white", mar = c(4, 4, 3, 1))

# Store statistics values
summary_stats <- data.frame(
  Variable           = character()
  , Mean             = numeric()
  , Median            = numeric()
  , SD               = numeric()
  , Skewness          = numeric()
  , Outliers          = numeric()
  , Normality_p       = numeric()
  , Group_Diff_p      = numeric()
```

```

    , Transform_Needed = character()
    , stringsAsFactors = FALSE
)

# Loop over each variables
for(var in numeric_vars) {
  cat(sprintf("\n===== %s =====\n", toupper(var)))

  # Extract data
  var_data <- clinical_df[[var]]
  var_alive <- var_data[clinical_df$vital_status == "Alive"]
  var_dead <- var_data[clinical_df$vital_status == "Dead"]

  # Remove NA
  var_data <- var_data[!is.na(var_data)]
  var_alive <- var_alive[!is.na(var_alive)]
  var_dead <- var_dead[!is.na(var_dead)]

  # Statistics
  var_mean <- mean(var_data)
  var_median <- median(var_data)
  var_sd <- sd(var_data)
  var_min <- min(var_data)
  var_max <- max(var_data)

  # Skewness
  var_skew <- mean(((var_data - var_mean) / var_sd)^3)

  # Outliers (IQR method)
  Q1 <- quantile(var_data, 0.25, na.rm = TRUE)
  Q3 <- quantile(var_data, 0.75, na.rm = TRUE)
  IQR_val <- Q3 - Q1
  lower_bound <- Q1 - 1.5 * IQR_val
  upper_bound <- Q3 + 1.5 * IQR_val

  n_outliers <- sum(var_data < lower_bound | var_data > upper_bound)
  outlier_pct <- round(n_outliers / length(var_data) * 100, 1)

  # Normality test (Shapiro-Wilk)
  if(length(var_data) <= 5000) {
    shapiro_test <- shapiro.test(var_data)
    shapiro_p <- shapiro_test$p.value
  } else {
    shapiro_test <- shapiro.test(sample(var_data, 5000))
    shapiro_p <- shapiro_test$p.value
  }

  # Group difference test (t-test)
  if(length(var_alive) > 0 & length(var_dead) > 0) {
    ttest <- t.test(var_alive, var_dead)
    ttest_p <- ttest$p.value
  } else {
    ttest_p <- NA
  }
}

```

```

}

# Print Statistics
cat(sprintf("Descriptive Statistics:\n"))
cat(sprintf("  Mean:      %.2f\n", var_mean))
cat(sprintf("  Median:    %.2f\n", var_median))
cat(sprintf("  SD:        %.2f\n", var_sd))
cat(sprintf("  Range:     [% .2f, %.2f]\n", var_min, var_max))
cat(sprintf("  Skewness:   %.3f %s\n"
           , var_skew
           , ifelse(abs(var_skew) < 0.5, "(Symmetric)"
                  , ifelse(var_skew > 0, "(Right-skewed)", "(Left-skewed)"))))

cat(sprintf("  Outliers:   %d (%.1f%%)\n", n_outliers, outlier_pct))
cat(sprintf("\n"))

cat(sprintf("Statistical Tests:\n"))
cat(sprintf("  Shapiro-Wilk p-value: %.4e %s\n"
           , shapiro_p
           , ifelse(shapiro_p < 0.05, "(NOT normal)", "(Normal)")))

if(!is.na(ttest_p)) {
  cat(sprintf("  T-test (Alive vs Dead): p=% .4e %s\n"
             , ttest_p
             , ifelse(ttest_p < 0.05, "*** GROUPS DIFFER", "(No difference)")))
  cat(sprintf("    Alive: mean=%.2f, sd=%.2f\n"
             , mean(var_alive), sd(var_alive)))
  cat(sprintf("    Dead:  mean=%.2f, sd=%.2f\n"
             , mean(var_dead), sd(var_dead)))
}
cat(sprintf("\n"))

# Transformation
transform_needed <- "None"
if(abs(var_skew) > 1.0) {
  transform_needed <- "Log or sqrt (high skewness)"
} else if(outlier_pct > 5) {
  transform_needed <- "Robust scaling (many outliers)"
} else if(shapiro_p < 0.05 & abs(var_skew) > 0.5) {
  transform_needed <- "Consider log (non-normal + skewed)"
}

# Store summary
summary_stats <- rbind(summary_stats
  , data.frame(
    Variable      = var
    , Mean        = var_mean
    , Median      = var_median
    , SD          = var_sd
    , Skewness    = var_skew
    , Outliers    = outlier_pct
    , Normality_p = shapiro_p
    , Group_Diff_p = ifelse(is.na(ttest_p), 1, ttest_p)

```

```

        , Transform_Needed = transform_needed
    )))
# Histogram
hist(var_data
    , breaks      = 40
    , col         = "#8ecae6"
    , border      = "white"
    , main        = paste(var, "- Histogram")
    , sub         = "Distribution with mean (red) and median (orange) lines"
    , xlab        = var
    , ylab        = "Frequency"
    , col.main   = "#023047"
    , col.lab    = "#023047"
    , col.sub    = "#666666"
    , cex.main   = 1.0
    , cex.sub    = 0.7
    , font.sub   = 3)

# Mean/Median lines
abline(v    = var_mean
    , col = "#e63946"
    , lwd = 3
    , lty = 1)

abline(v    = var_median
    , col = "#fb8500"
    , lwd = 3
    , lty = 2)

# Skewness text
text(x      = var_mean
    , y      = par("usr")[4] * 0.9
    , labels = sprintf("Skew=% .2f", var_skew)
    , pos    = 4
    , col    = "#023047"
    , cex    = 0.8)

legend("topright"
    , legend = c("Mean", "Median")
    , col    = c("#e63946", "#fb8500")
    , lwd    = 3
    , lty    = c(1, 2)
    , bty    = "n"
    , cex    = 0.7)

# Density
plot(density(var_alive)
    , col      = "#219ebc"
    , lwd      = 3
    , main     = paste(var, "- Density by Status")
    , sub     = "Comparison of Alive vs Dead patient distributions"
    , xlab     = var
    , ylab     = "Density"

```

```

    , col.main = "#023047"
    , col.lab  = "#023047"
    , col.sub  = "#666666"
    , cex.main = 1.0
    , cex.sub  = 0.7
    , font.sub = 3)

lines(density(var_dead)
      , col = "#e63946"
      , lwd = 3)

# add group means
abline(v  = mean(var_alive)
       , col = "#219ebc"
       , lty = 2
       , lwd = 2)

abline(v  = mean(var_dead)
       , col = "#e63946"
       , lty = 2
       , lwd = 2)

legend("topright"
       , legend = c("Alive", "Dead")
       , col    = c("#219ebc", "#e63946")
       , lwd    = 3
       , bty    = "n"
       , cex    = 0.7)

# QQ-Plot
qqnorm(var_data
       , main   = paste(var, "- Q-Q Plot")
       , sub    = "Normality assessment: points on line = normal distribution"
       , pch    = 19
       , cex    = 0.5
       , col    = "#8ecae6"
       , col.main = "#023047"
       , col.lab  = "#023047"
       , col.sub  = "#666666"
       , cex.main = 1.0
       , cex.sub  = 0.7
       , font.sub = 3)

qqline(var_data
       , col = "#e63946"
       , lwd = 3)

# Normality
text(x      = par("usr")[1]
     , y      = par("usr")[4]
     , labels = sprintf("Shapiro p=% .2e\n%s"
                       , shapiro_p
                       , ifelse(shapiro_p < 0.05, "NON-normal", "Normal")))

```

```

, pos   = 4
, col   = ifelse(shapiro_p < 0.05, "#e63946", "#219ebc")
, cex   = 0.8
, font  = 2)

# Outliers
boxplot(var_data ~ clinical_df$vital_status[!is.na(clinical_df[[var]])]
       , col      = c("#8ecae6", "#ffb703")
       , names    = c("Alive", "Dead")
       , main     = paste(var, "- Boxplot by Vital Status")
       , sub      = "Group comparison with outliers shown"
       , xlab     = "Vital Status"
       , ylab     = var
       , border   = c("#219ebc", "#fb8500")
       , col.main = "#023047"
       , col.lab  = "#023047"
       , col.sub  = "#666666"
       , lwd      = 1.5
       , cex.main = 1.0
       , cex.sub  = 0.7
       , font.sub = 3
       , outline  = TRUE) # Show outliers

text(x      = 1
     , y      = par("usr")[3]
     , labels = sprintf("n=%d", length(var_alive))
     , pos   = 3
     , cex   = 0.7)

text(x      = 2
     , y      = par("usr")[3]
     , labels = sprintf("n=%d", length(var_dead))
     , pos   = 3
     , cex   = 0.7)

# Add p-value labels
if(!is.na(ttest_p)) {
  text(x      = 1.5
       , y      = par("usr")[4]
       , labels = sprintf("p=%,.3f %s"
                           , ttest_p
                           , ifelse(ttest_p < 0.05, "***", ""))
       , pos   = 1
       , col   = ifelse(ttest_p < 0.05, "#e63946", "#023047")
       , cex   = 0.8
       , font  = 2)
}

## ===== AGE_AT_INDEX =====
## Descriptive Statistics:
##   Mean:      58.28

```

```

## Median:      58.00
## SD:          13.28
## Range:       [26.00, 89.00]
## Skewness:    0.146 (Symmetric)
## Outliers:   0 (0.0%)
##
## Statistical Tests:
## Shapiro-Wilk p-value: 4.8837e-07 (NOT normal)
## T-test (Alive vs Dead): p=6.8789e-03 *** GROUPS DIFFER
##   Alive: mean=57.77, sd=12.77
##   Dead:  mean=60.92, sd=15.39

##
## ===== AGE_AT_DIAGNOSIS =====
## Descriptive Statistics:
## Mean:        21530.01
## Median:     21472.00
## SD:          4815.42
## Range:       [9840.00, 32872.00]
## Skewness:    0.147 (Symmetric)
## Outliers:   0 (0.0%)
##
## Statistical Tests:
## Shapiro-Wilk p-value: 4.6456e-06 (NOT normal)
## T-test (Alive vs Dead): p=6.3766e-03 *** GROUPS DIFFER
##   Alive: mean=21337.60, sd=4639.52
##   Dead:  mean=22503.97, sd=5533.57

##
## ===== INITIAL_WEIGHT =====
## Descriptive Statistics:
## Mean:        310.98
## Median:     220.00
## SD:          272.07
## Range:       [5.00, 2190.00]
## Skewness:    2.293 (Right-skewed)
## Outliers:   72 (5.9%)
##
## Statistical Tests:
## Shapiro-Wilk p-value: 1.4492e-37 (NOT normal)
## T-test (Alive vs Dead): p=8.1424e-02 (No difference)
##   Alive: mean=304.63, sd=268.04
##   Dead:  mean=343.61, sd=290.44

##
## ===== DAYS_TO_LAST_FOLLOW_UP =====
## Descriptive Statistics:
## Mean:        1245.98
## Median:     890.00
## SD:          1159.43
## Range:       [-7.00, 8605.00]
## Skewness:    2.159 (Right-skewed)
## Outliers:   44 (3.6%)

```

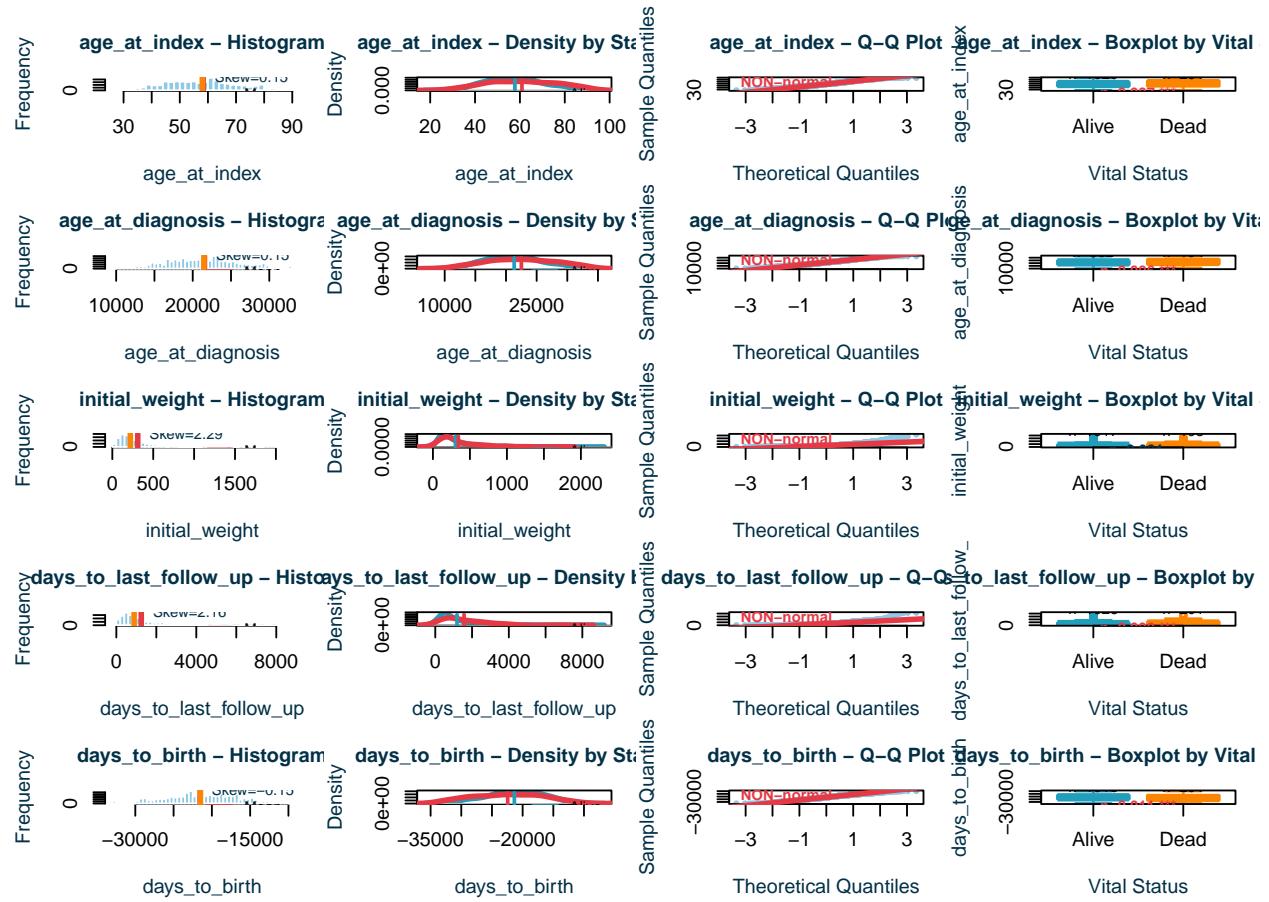
```

## Statistical Tests:
## Shapiro-Wilk p-value: 1.6811e-35 (NOT normal)
## T-test (Alive vs Dead): p=6.6407e-05 *** GROUPS DIFFER
## Alive: mean=1183.43, sd=1133.53
## Dead: mean=1565.26, sd=1238.10

## ===== DAYS_TO_BIRTH =====
## Descriptive Statistics:
## Mean: -21524.60
## Median: -21494.50
## SD: 4842.02
## Range: [-32872.00, -9706.00]
## Skewness: -0.146 (Symmetric)
## Outliers: 0 (0.0%)

## Statistical Tests:
## Shapiro-Wilk p-value: 1.9938e-06 (NOT normal)
## T-test (Alive vs Dead): p=1.0727e-02 *** GROUPS DIFFER
## Alive: mean=-21344.56, sd=4652.38
## Dead: mean=-22431.95, sd=5628.62

```



```
par(mfrow = c(1, 1))
```

Conclusion of Numerical Variable Analysis

- AGE_AT_INDEX

- Symmetric distribution, no outliers.
- Not normally distributed (Shapiro $p < 1e-6$).
- Significant group difference ($p = 0.0069$).
- *Dead patients are older (60.9 vs 57.8)*.

- AGE_AT_DIAGNOSIS

- Symmetric, no outliers.
- Not normal.
- Significant difference ($p = 0.0064$).
- *Dead patients were diagnosed at an older age.*

- INITIAL_WEIGHT

- Strong right skew; many outliers (~6%).
- Not normal.
- No significant difference ($p = 0.081$).
- *Weight does not differ between Alive/Dead groups.*

- DAYS_TO_LAST_FOLLOW_UP

- Strong right-skewed distribution with outliers.
- Not normal.
- Significant difference ($p = 6.6e-05$).
- *Dead patients show longer follow-up times (expected due to event vs censoring).*

- DAYS_TO_BIRTH

- Symmetric distribution, no outliers.
- Not normal.
- Significant difference ($p = 0.0107$).
- *Reflects age differences – Dead patients are older.*

Numerical Variables - Group Comparison Tests

```
clinical_df <- as.data.frame(clinical_df)

# Convert target
clinical_df$Y <- factor(clinical_df$vital_status, levels = c("Alive", "Dead"))

num_vars <- c("age_at_index"
             , "age_at_diagnosis"
             , "initial_weight"
             , "days_to_last_follow_up"
             , "days_to_birth")

# Storage for results
test_results <- data.frame(
  Variable      = character()
```

```

, Test      = character()
, Statistic = numeric()
, P_value   = numeric()
, Effect_Size = numeric()
, Correlation = numeric()
, stringsAsFactors = FALSE
)

# Test each variable
for(var in num_vars) {

  valid_idx <- !is.na(clinical_df[[var]]) & !is.na(clinical_df$Y)
  df_test   <- clinical_df[valid_idx, c("Y", var)]

  # Skip if insufficient data
  if(nrow(df_test) < 10 || length(unique(df_test$Y)) < 2) next

  # --- Check skewness ---
  alive_vals <- df_test[df_test$Y == "Alive", var]
  dead_vals  <- df_test[df_test$Y == "Dead", var]

  skew_alive <- abs(mean(alive_vals) - median(alive_vals)) / IQR(alive_vals)
  skew_dead  <- abs(mean(dead_vals) - median(dead_vals)) / IQR(dead_vals)

  is_normal <- (skew_alive < 0.2 & skew_dead < 0.2)

  # --- Choose test ---
  if(is_normal) {
    # T-test
    test_res <- t.test(df_test[[var]] ~ df_test$Y, var.equal = TRUE)

    test_name <- "t-test"
    stat_val  <- test_res$statistic
    p_val     <- test_res$p.value

    # Cohen's d
    effect <- cohens_d(df_test, as.formula(paste(var, "~ Y")))$effsize

  } else {
    # Wilcoxon test
    test_res <- wilcox.test(df_test[[var]] ~ df_test$Y)

    test_name <- "Wilcoxon"
    stat_val  <- test_res$statistic
    p_val     <- test_res$p.value

    # Rank-biserial
    effect <- wilcox_effsize(df_test, as.formula(paste(var, "~ Y")))$effsize
  }

  # Point-biserial correlation
  cor_val <- cor(df_test[[var]], as.numeric(df_test$Y) - 1)
}

```

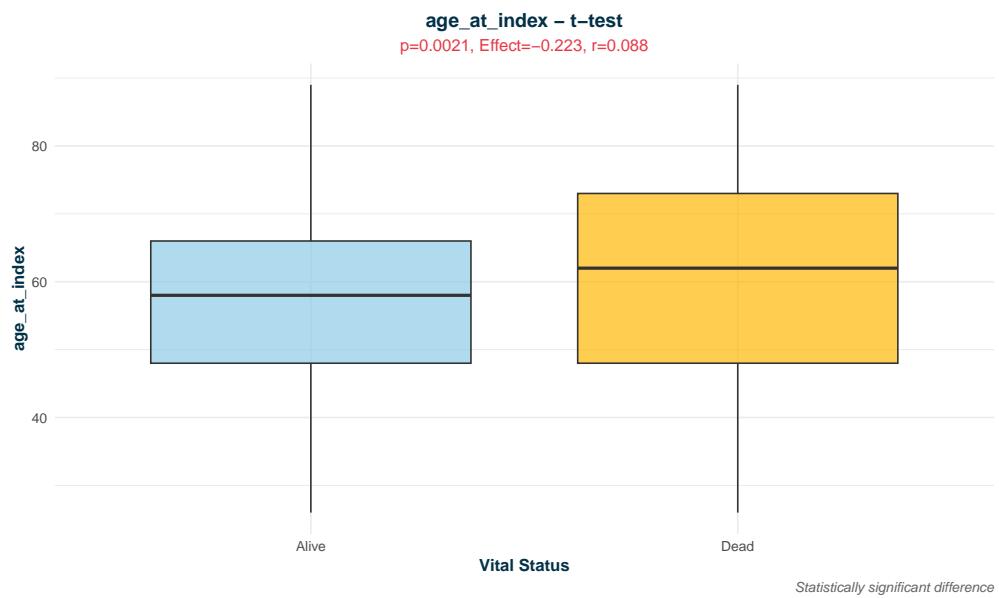
```

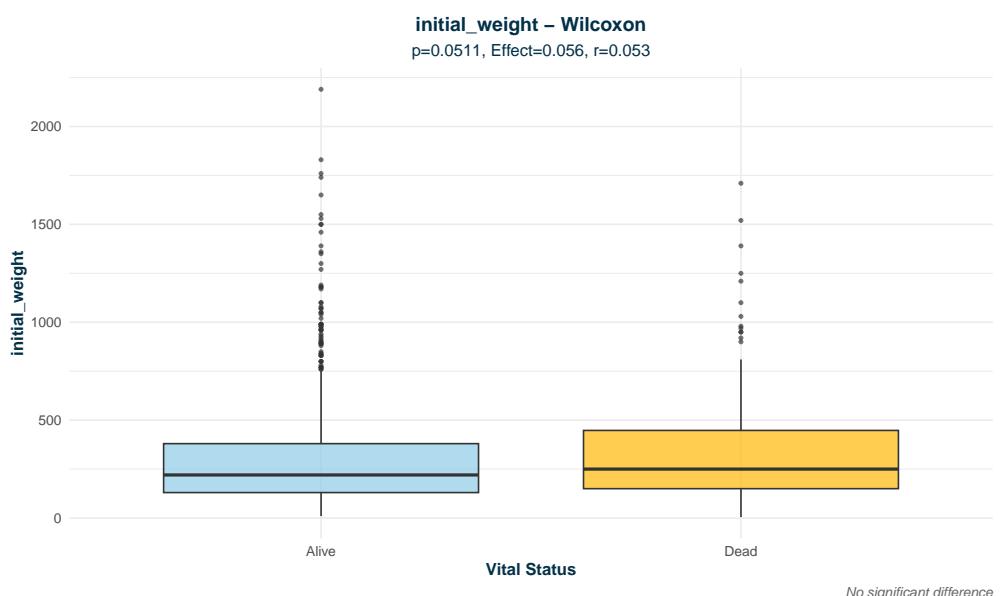
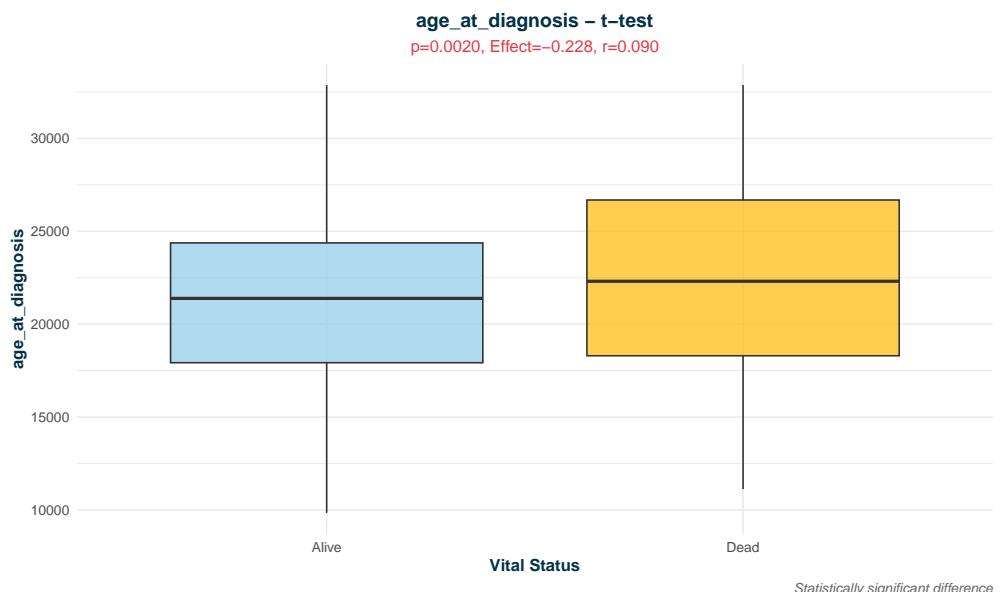
# Store results
test_results <- rbind(test_results
  , data.frame(Variable      = var
               , Test        = test_name
               , Statistic   = round(stat_val, 2)
               , P_value     = p_val
               , Effect_Size = round(effect, 3)
               , Correlation = round(cor_val, 3)))

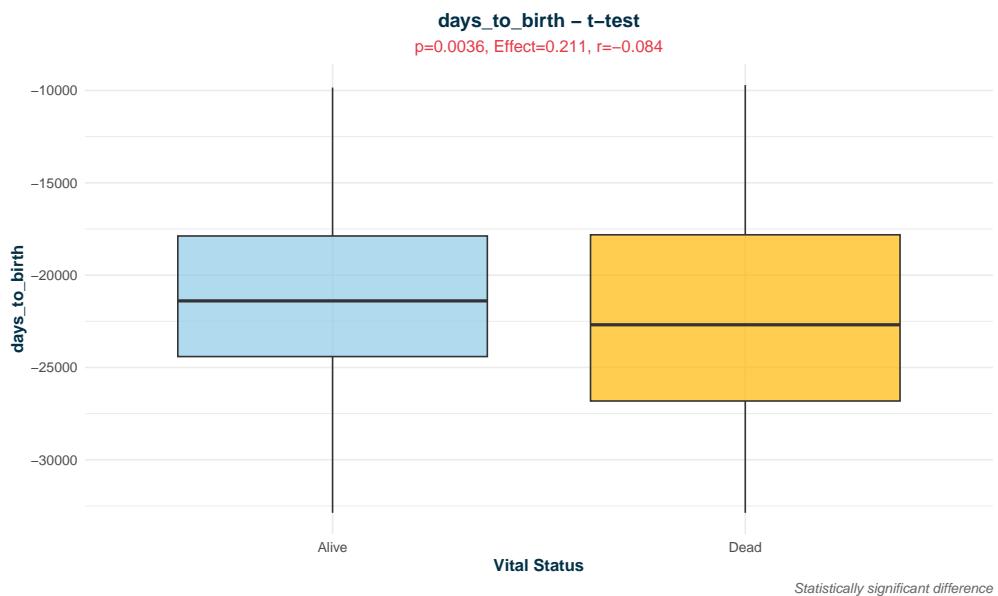
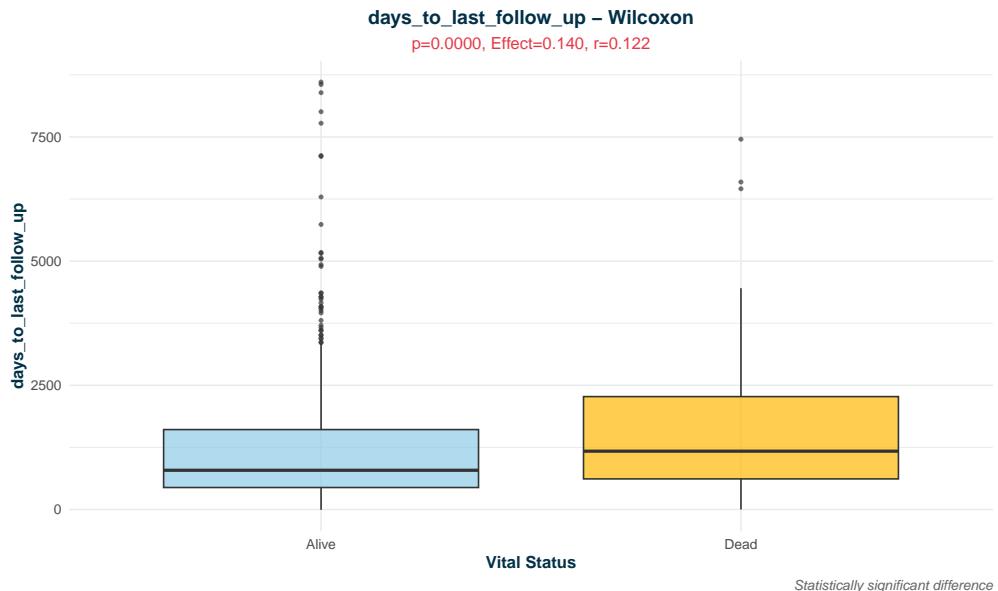
# --- Visualization ---
p <- ggplot(df_test, aes(x = Y, y = .data[[var]], fill = Y)) +
  geom_boxplot(alpha = 0.7, outlier.shape = 19, outlier.size = 1) +
  scale_fill_manual(values = c("Alive" = "#8ecae6", "Dead" = "#ffb703")) +
  labs(title = paste(var, "-"), test_name)
  , subtitle = sprintf("p=%.4f, Effect=%.3f, r=%.3f", p_val, effect, cor_val)
  , x = "Vital Status"
  , y = var
  , caption = ifelse(p_val < 0.05, "Statistically significant difference", "No significant difference")
  theme_minimal(base_size = 12) +
  theme(legend.position = "none"
    , plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047")
    , plot.subtitle = element_text(hjust = 0.5, color = ifelse(p_val < 0.05, "#e63946", "#023047"))
    , plot.caption = element_text(face = "italic", color = "#666666")
    , axis.title = element_text(face = "bold", color = "#023047"))

print(p)
}

```







```
# --- FDR correction ---
test_results$P_adj <- p.adjust(test_results$P_value, method = "fdr")

for(i in 1:nrow(test_results)) {
  row <- test_results[i, ]

  sig <- ifelse(row$P_adj < 0.001, "***"
                , ifelse(row$P_adj < 0.01, "**"
                , ifelse(row$P_adj < 0.05, "*", "")))

  cat(sprintf("%-25s %-10s %10.2f %10.4f %10.4f %10.3f %10.3f %s\n"
            , row$Variable
            , row$Test
            , row$Statistic
```

```

        , row$P_value
        , row$P_adj
        , row$Effect_Size
        , row$Correlation
        , sig))
}

## age_at_index      t-test      -3.09    0.0021    0.0034    -0.223    0.088 **
## age_at_diagnosis t-test      -3.09    0.0020    0.0034    -0.228    0.090 **
## initial_weight    Wilcoxon   91871.50   0.0511    0.0511    0.056    0.053
## days_to_last_follow_up Wilcoxon  80651.50   0.0000    0.0000    0.140    0.122 ***
## days_to_birth     t-test      2.92     0.0036    0.0045    0.211    -0.084 **

```

According to the data analysis and the implementation of statistical tests on the clinical dataset. We can observe that columns such as **age at diagnosis**, **age at index**, and **day to birth** are the same information, which we can drop or exclude for variable selection by keeping only age at index. In addition, **initial weight** columns are also not significant for vital status target.

Clinical Correlation Matrix

```

# Convert target to numeric
clinical_base <- as.data.frame(clinical_df)
clinical_base$vital_status_bin <- ifelse(clinical_base$vital_status == "Dead", 1, 0)

# Get numeric variables
clinic_num_cols <- names(clinical_base)[sapply(clinical_base, is.numeric)]
numeric_df      <- clinical_base[, clinic_num_cols]

# Compute correlation
corr_matrix <- cor(numeric_df, use = "complete.obs")

# Plot
ggcorrplot(corr_matrix
            , hc.order = TRUE
            , lab      = TRUE
            , lab_size = 2.5
            , method   = "circle"
            , type     = "lower"
            , colors   = c("#4361ee", "#f8f9fa", "#e63946")
            , title    = "Correlation Matrix - Clinical Numeric Variables"
            , ggtheme  = theme_minimal() +
                theme(plot.title = element_text(hjust = 0.5, size = 14, face = "bold", color = "#023047"),
                      plot.subtitle = element_text(hjust = 0.5, color = "#555555"),
                      plot.caption = element_text(face = "italic", color = "#666666"),
                      axis.text = element_text(color = "#023047")) +
            labs(subtitle = "Pearson correlation coefficients"
                 , caption = "Method: Complete observations with hierarchical clustering")

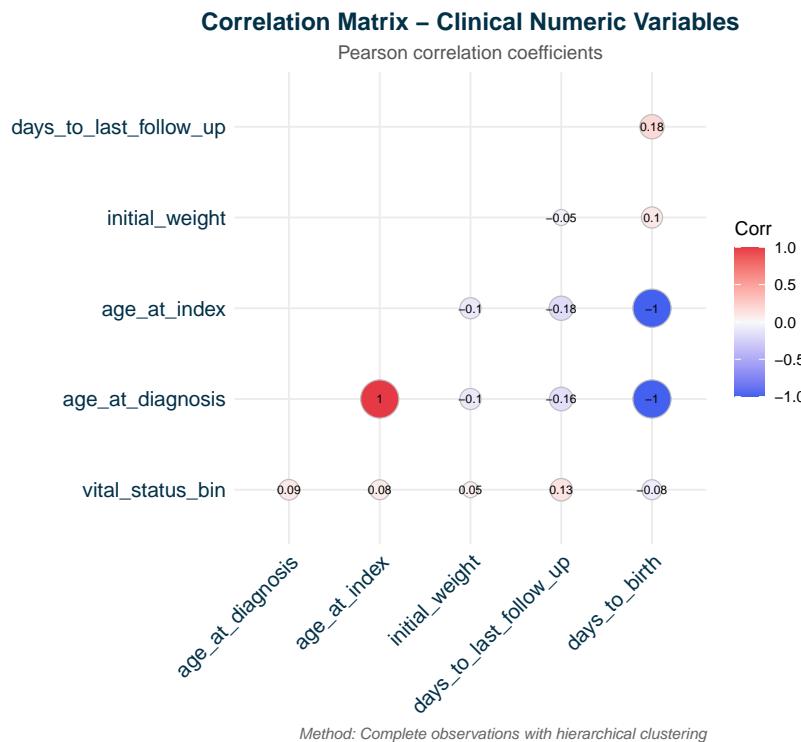
## Warning: `aes_string()` was deprecated in ggplot2 3.0.0.
## i Please use tidy evaluation idioms with `aes()`'.
## i See also `vignette("ggplot2-in-packages")` for more information.

```

```

## i The deprecated feature was likely used in the ggcrrplot package.
## Please report the issue at <https://github.com/kassambara/ggcrrplot/issues>.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

```



From above correlation matrix, we can interpret that

- **age_at_index, age_at_diagnosis, and days_to_birth**
 - Extremely high correlations ($|r| \approx 0.99$).
 - These three variables encode the **same underlying information (patient age)**.
 - Keep only one (**age_at_index**) for modeling to avoid multicollinearity.
- **days_to_last_follow_up**
 - Weak correlations with all other variables ($|r| < 0.20$).
 - Slight positive correlation with vital_status_bin ($r \approx 0.13$), expected because **Dead patients have actual event times**, while Alive patients are censored earlier.
- **initial_weight**
 - Very weak correlations with every clinical variable and with survival ($|r| < 0.10$).
 - Not a predictive feature.
- **vital_status_bin**
 - Correlates weakly with every numeric variable ($|r| < 0.13$).

- No strong linear relationship; survival differences detected by group tests are **small effect sizes**, not strong correlations.

The numerical variables reveal one clear multicollinearity block: *age_at_index*, *age_at_diagnosis*, and *days_to_birth* all measure the same underlying factor (patient age) and only one should be kept. Other variables show only weak correlations with survival. *Days_to_last_follow_up* has a small association due to censoring differences, while *initial_weight* shows negligible relevance and can be excluded.

VIF Analysis for Clinical Variables

```
# Prepare clinical data
clinical_vif <- data.frame(
  age_at_index      = clinical_df$age_at_index
, age_at_diagnosis = clinical_df$age_at_diagnosis
, initial_weight    = clinical_df$initial_weight
, days_to_last_follow_up = clinical_df$days_to_last_follow_up
, days_to_birth     = clinical_df$days_to_birth
, vital_status_bin  = ifelse(clinical_df$vital_status == "Dead", 1, 0)
)

# Remove NA
clinical_vif <- na.omit(clinical_vif)

cat("After removing NA:", nrow(clinical_vif), "\n\n")

## After removing NA: 1161

# Fit model
full_model <- glm(vital_status_bin ~ age_at_index + age_at_diagnosis +
                     initial_weight + days_to_last_follow_up + days_to_birth
                     , data   = clinical_vif
                     , family = binomial)

# Calculate VIF
vif_values <- vif(full_model)

cat("VIF Results:\n")

## VIF Results:

print(vif_values)

##          age_at_index        age_at_diagnosis       initial_weight
##            2100.239872           149.749877            1.029426
##  days_to_last_follow_up        days_to_birth
##            1.122283            2204.019112

cat("\n== INTERPRETATION ==\n")
```

```

##  

## === INTERPRETATION ===

cat("VIF < 5: No multicollinearity\n")

## VIF < 5: No multicollinearity

cat("VIF 5-10: Moderate multicollinearity (monitor)\n")

## VIF 5-10: Moderate multicollinearity (monitor)

cat("VIF > 10: High multicollinearity (REMOVE variable)\n\n")

## VIF > 10: High multicollinearity (REMOVE variable)

# Flag problematic variables  

high_vif <- names(vif_values)[vif_values > 10]  

mod_vif  <- names(vif_values)[vif_values >= 5 & vif_values <= 10]

if(length(high_vif) > 0) {  

  cat("HIGH VIF (>10) - REMOVE:\n")
  for(var in high_vif) {  

    cat(sprintf(" %s: VIF = %.2f\n", var, vif_values[var]))  

  }
  cat("\n")
}

## HIGH VIF (>10) - REMOVE:  

##   age_at_index: VIF = 2100.24  

##   age_at_diagnosis: VIF = 149.75  

##   days_to_birth: VIF = 2204.02

if(length(mod_vif) > 0) {  

  cat(" MODERATE VIF (5-10) - MONITOR:\n")
  for(var in mod_vif) {  

    cat(sprintf(" %s: VIF = %.2f\n", var, vif_values[var]))  

  }
  cat("\n")
}

# Convert into data frame  

vif_df <- data.frame(  

  Variable = names(vif_values),  

  VIF = as.numeric(vif_values)
)

# Threshold for high multicollinearity (commonly VIF > 5 or > 10)  

threshold <- 5

# Classify variables

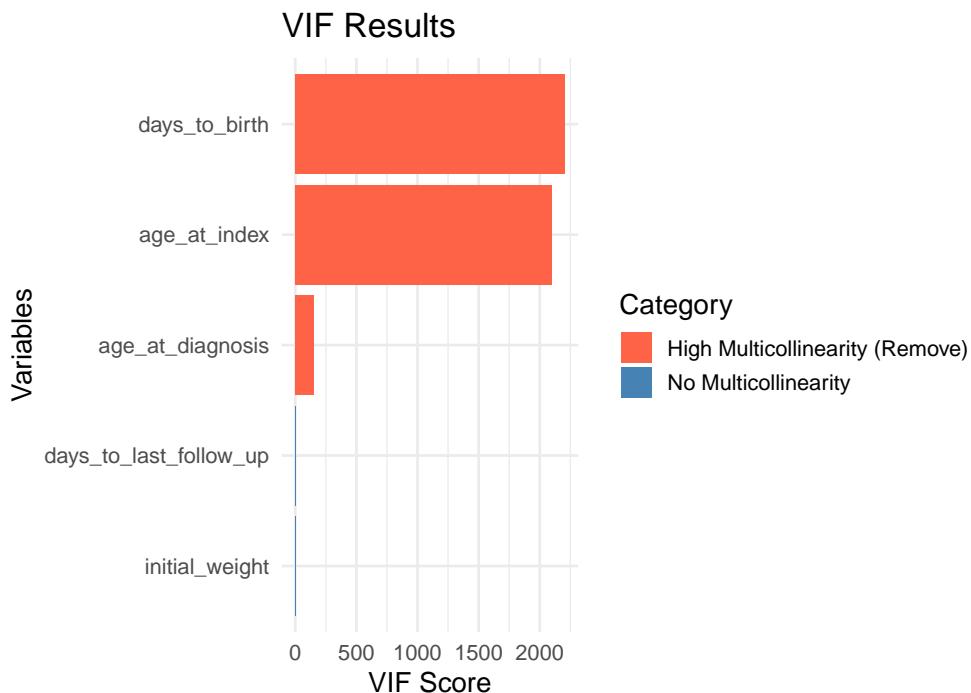
```

```

vif_df$Group <- ifelse(vif_df$VIF > threshold,
                        "High Multicollinearity (Remove)",
                        "No Multicollinearity")

# Plot bar graph
ggplot(vif_df, aes(x = reorder(Variable, VIF), y = VIF, fill = Group)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  scale_fill_manual(values = c(
    "No Multicollinearity" = "steelblue",
    "High Multicollinearity (Remove)" = "tomato"
  )) +
  labs(title = "VIF Results",
       x = "Variables",
       y = "VIF Score",
       fill = "Category") +
  theme_minimal(base_size = 14)

```



From VIF, it shows extreme multicollinearity among the age variables (`age_at_index`, `age_at_diagnosis`, `days_to_birth`), meaning they all represent the same information and only should be kept. The other variable (`initial_weight`, `days_to_last_follow_up`) have VIF $\sim=1$ and pose no multicollinearity issue.

PCA Analysis on Clinical Variables

```

clinical_pca <- data.frame(
  age_at_index      = clinical_df$age_at_index
, age_at_diagnosis = clinical_df$age_at_diagnosis

```

```

    , initial_weight      = clinical_df$initial_weight
    , days_to_last_follow_up = clinical_df$days_to_last_follow_up
    , days_to_birth       = clinical_df$days_to_birth
)

clinical_pca$vital_status <- clinical_df$vital_status

# Remove NA
clinical_pca <- na.omit(clinical_pca)
predictors <- clinical_pca[, 1:5]

# Run PCA (scaled)
pca_result <- prcomp(predictors, scale. = TRUE, center = TRUE)

# Variance explained
var_exp     <- summary(pca_result)$importance[2, ]
var_cum     <- summary(pca_result)$importance[3, ]

cat("Variance Explained:\n")

## Variance Explained:

for(i in 1:5) {
  cat(sprintf(" PC%d: %.1f%% (Cumulative: %.1f%%)\n"
              , i
              , var_exp[i] * 100
              , var_cum[i] * 100))
}

## PC1: 61.1% (Cumulative: 61.1%)
## PC2: 20.9% (Cumulative: 82.0%)
## PC3: 17.9% (Cumulative: 99.9%)
## PC4: 0.1% (Cumulative: 100.0%)
## PC5: 0.0% (Cumulative: 100.0%)

cat("\n")

cat("Variable Loadings on PC1 and PC2:\n")

## Variable Loadings on PC1 and PC2:

loadings <- pca_result$rotation[, 1:2]
print(round(loadings, 3))

##          PC1      PC2
## age_at_index      0.570  0.014
## age_at_diagnosis   0.569  0.027
## initial_weight     -0.076 -0.780
## days_to_last_follow_up -0.145  0.625
## days_to_birth      -0.570 -0.014

```

```

cat("\n")

# Interpretation
cat("== INTERPRETATION ==\n")

## == INTERPRETATION ==

cat("PC1 captures", round(var_exp[1] * 100, 1), "% variance\n")

## PC1 captures 61.1 % variance

cat(" - High loadings:", names(sort(abs(loadings[, 1]), decreasing = TRUE)[1:2]), "\n")

## - High loadings: days_to_birth age_at_index

cat("PC2 captures", round(var_exp[2] * 100, 1), "% variance\n")

## PC2 captures 20.9 % variance

cat(" - High loadings:", names(sort(abs(loadings[, 2]), decreasing = TRUE)[1:2]), "\n\n")

## - High loadings: initial_weight days_to_last_follow_up

par(mfrow = c(2, 2), bg = "white", mar = c(4, 4, 3, 2))

# 1. Scree Plot
barplot(var_exp * 100
        , names.arg = paste0("PC", 1:5)
        , col      = "#8ecae6"
        , border   = "white"
        , xlab     = "Principal Component"
        , ylab     = "Variance Explained (%)"
        , main     = "Scree Plot - Clinical Variables"
        , sub      = "Eigenvalue decomposition showing variance per PC"
        , col.main = "#023047"
        , col.lab  = "#023047"
        , col.sub  = "#666666"
        , cex.sub  = 0.7
        , font.sub = 3
        , las      = 1)

abline(h = 20
        , col = "#e63946"
        , lty = 2
        , lwd = 2)

# 2. Cumulative Variance
plot(1:5
        , var_cum * 100

```

```

, type      = "b"
, pch       = 19
, col       = "#219ebc"
, lwd       = 3
, xlab      = "Principal Component"
, ylab      = "Cumulative Variance (%)"
, main      = "Cumulative Variance Explained"
, sub       = "Total variance captured by first n components"
, col.main  = "#023047"
, col.lab   = "#023047"
, col.sub   = "#666666"
, cex.sub   = 0.7
, font.sub  = 3
, las       = 1)

abline(h  = 80
       , col = "#fb8500"
       , lty = 2
       , lwd = 2)

text(x = 3, y = 85, labels = "80% threshold", col = "#fb8500", cex = 0.8)

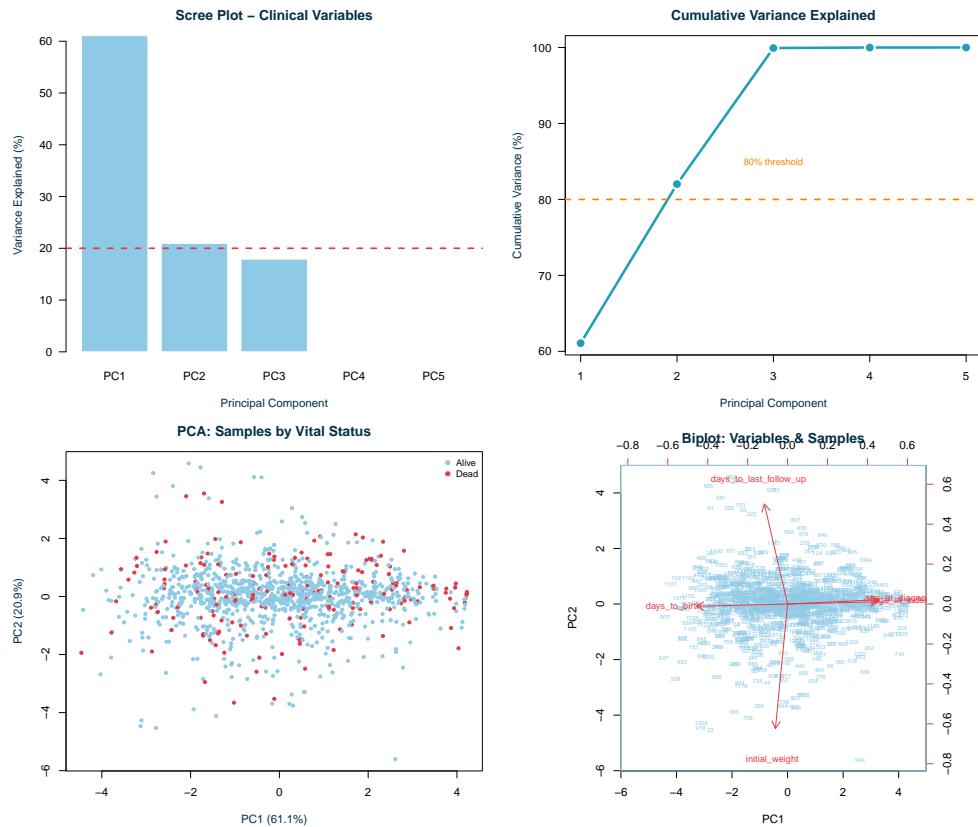
# 3. PC1 vs PC2 (colored by vital status)
plot(pca_result$x[, 1]
      , pca_result$x[, 2]
      , pch       = 19
      , cex       = 0.6
      , col       = ifelse(clinical_pca$vital_status == "Dead"
                           , "#e63946", "#8ecae6")
      , xlab      = paste0("PC1 (", round(var_exp[1] * 100, 1), "%)")
      , ylab      = paste0("PC2 (", round(var_exp[2] * 100, 1), "%)")
      , main      = "PCA: Samples by Vital Status"
      , sub       = "Patient projection onto first two principal components"
      , col.main  = "#023047"
      , col.lab   = "#023047"
      , col.sub   = "#666666"
      , cex.sub   = 0.7
      , font.sub  = 3)

legend("topright"
       , legend = c("Alive", "Dead")
       , col     = c("#8ecae6", "#e63946")
       , pch     = 19
       , bty     = "n"
       , cex     = 0.8)

# 4. Biplot (variables + samples)
biplot(pca_result
       , choices  = 1:2
       , scale    = 0
       , col      = c("#8ecae6", "#e63946")
       , cex      = c(0.5, 0.8)
       , main     = "Biplot: Variables & Samples"

```

```
, col.main = "#023047"
, arrow.len = 0.1)
```



```
par(mfrow = c(1, 1))
```

Drop Redundant Variables Based on VIF + PCA

```
drop_vars <- c("age_at_diagnosis", "days_to_birth")

# KEEP
keep_vars <- c("age_at_index"
, "initial_weight"
, "days_to_last_follow_up")

# Create reduced set
clinical_reduced <- data.frame(
  age_at_index          = clinical_df$age_at_index
, initial_weight        = clinical_df$initial_weight
, days_to_last_follow_up = clinical_df$days_to_last_follow_up
, vital_status_bin      = ifelse(clinical_df$vital_status == "Dead", 1, 0)
)
```

```

clinical_reduced <- na.omit(clinical_reduced)

# Fit model
reduced_model <- glm(vital_status_bin ~ age_at_index + initial_weight +
                      days_to_last_follow_up
                      , data   = clinical_reduced
                      , family = binomial)

# Calculate VIF
vif_reduced <- car::vif(reduced_model)

cat("VIF Results (Reduced Model):\n")

## VIF Results (Reduced Model):

print(vif_reduced)

```

```

##           age_at_index      initial_weight days_to_last_follow_up
##             1.077847            1.033913          1.074717

```

After reducing some redundant variables, we can see that the VIF is now better no more multicollinearity.

Categorical Variables Analysis

```

clinical_df <- as.data.frame(clinical_df)

exclude_cols <- c("bcr_patient_barcode"
                  , "primary_site"
                  , "days_to_birth"
                  , "age_at_diagnosis"
                  , "sites_of_involvement"
                  , "disease_type"
                  , "vital_status"
                  , "Y")

# Select cols
cat_cols <- names(clinical_df)[sapply(clinical_df, function(x)
                                         is.character(x) | is.factor(x))]

cat_cols <- setdiff(cat_cols, exclude_cols)

cat("Categorical variables:", length(cat_cols), "\n")

## Categorical variables: 12

cat(paste(cat_cols, collapse = ", "), "\n\n")

```

```

## tissue_type, laterality, tissue_or_organ_of_origin, primary_diagnosis, prior_treatment, ajcc_patholog

```

```

# Clean and prepare
categorical_df <- clinical_df[, cat_cols, drop = FALSE]

# Convert to character
categorical_df <- data.frame(lapply(categorical_df, as.character)
                           , stringsAsFactors = FALSE)

# Mark missing values
missing_markers <- c("not reported", "not applicable", "unknown", "NA", "")

for(var in names(categorical_df)) {
  categorical_df[[var]][categorical_df[[var]] %in% missing_markers] <- NA
}

# Group categories with < 10 samples
for(var in names(categorical_df)) {

  categorical_df[[var]] <- as.factor(categorical_df[[var]])
  categorical_df[[var]] <- fct_lump_min(categorical_df[[var]]
                                         , min = 10
                                         , other_level = "Other")
  categorical_df[[var]] <- droplevels(categorical_df[[var]])

  cat(sprintf("%s: %d levels after grouping\n"
             , var
             , nlevels(categorical_df[[var]])))
}

## tissue_type: 2 levels after grouping
## laterality: 2 levels after grouping
## tissue_or_organ_of_origin: 5 levels after grouping
## primary_diagnosis: 8 levels after grouping
## prior_treatment: 3 levels after grouping
## ajcc_pathologic_t: 7 levels after grouping
## morphology: 8 levels after grouping
## classification_of_tumor: 6 levels after grouping
## follow_ups_disease_response: 3 levels after grouping
## race: 4 levels after grouping
## gender: 2 levels after grouping
## ethnicity: 3 levels after grouping

```

Statistical Tests and Visualization

```

# Add outcome variable to categorical_df
categorical_df$Y <- factor(clinical_df$vital_status)

results <- list()

for (var in setdiff(names(categorical_df), "Y")) {

  x <- categorical_df[[var]]

```

```

y <- categorical_df$Y
# Skip constants
if (n_distinct(x) <= 1) {
  results[[var]] <- data.frame(
    Test="Constant variable", P_value=NA, Statistic=NA, Cramers_V=NA,
    Note="Skipped (constant)", stringsAsFactors=FALSE
  )
  next
}

# Build table
tbl <- table(x, y)

if (nrow(tbl) < 2 || ncol(tbl) < 2) {
  results[[var]] <- data.frame(
    Test="Too few levels", P_value=NA, Statistic=NA, Cramers_V=NA,
    Note="Not enough levels", stringsAsFactors=FALSE
  )
  next
}

# Expected counts
expected <- outer(rowSums(tbl), colSums(tbl)) / sum(tbl)

# Choose appropriate test
if (nrow(tbl) == 2 && ncol(tbl) == 2) {

  # Fisher for 2x2
  test <- fisher.test(tbl)
  test_name <- "Fisher Exact (2x2)"
  stat_val <- NA

} else if (all(expected >= 5)) {

  # Standard Chi-square
  test <- chisq.test(tbl, correct = FALSE)
  test_name <- "Chi-square"
  stat_val <- test$statistic

} else {

  # Monte Carlo Chi-square for sparse large contingency tables
  test <- chisq.test(tbl, simulate.p.value = TRUE, B = 10000)
  test_name <- "Chi-square (MC simulation)"
  stat_val <- test$statistic
}

pval <- test$p.value

# Cramér's V
chi2 <- sum((tbl - expected)^2 / expected)
k <- min(nrow(tbl), ncol(tbl))
cramers_v <- sqrt(chi2 / (sum(tbl) * (k - 1)))

```

```

results[[var]] <- data.frame(
  Test=test_name,
  Statistic=stat_val,
  P_value=pval,
  Cramers_V=round(cramers_v, 4),
  Note=ifelse(pval < 0.05, "Significant", "Not significant"),
  stringsAsFactors=FALSE
)
}

results_df <- do.call(rbind, results)
results_df$Variable <- rownames(results_df)
results_df <- results_df[, c("Variable", "Test", "Statistic", "P_value", "Cramers_V", "Note")]
print(results_df)

##                                     Variable
## tissue_type                      tissue_type
## laterality                        laterality
## tissue_or_organ_of_origin    tissue_or_organ_of_origin
## primary_diagnosis                 primary_diagnosis
## prior_treatment                   prior_treatment
## ajcc_pathologic_t                  ajcc_pathologic_t
## morphology                         morphology
## classification_of_tumor      classification_of_tumor
## follow_ups_disease_response follow_ups_disease_response
## race                                race
## gender                               gender
## ethnicity                          ethnicity
##                                     Test   Statistic   P_value
## tissue_type                      Fisher Exact (2x2) NA 1.457349e-09
## laterality                        Fisher Exact (2x2) NA 2.985037e-01
## tissue_or_organ_of_origin    Chi-square (MC simulation) 161.948801 9.999000e-05
## primary_diagnosis                 Chi-square (MC simulation) 13.447819 5.839416e-02
## prior_treatment                   Chi-square (MC simulation) 99.889277 9.999000e-05
## ajcc_pathologic_t                  Chi-square (MC simulation) 40.951881 9.999000e-05
## morphology                         Chi-square (MC simulation) 13.447819 6.389361e-02
## classification_of_tumor      Chi-square (MC simulation) 130.571256 9.999000e-05
## follow_ups_disease_response Chi-square (MC simulation) 426.860132 9.999000e-05
## race                                Chi-square (MC simulation) 7.376931 6.699330e-02
## gender                               Fisher Exact (2x2) NA 7.059201e-01
## ethnicity                          Chi-square (MC simulation) 8.644092 1.369863e-02
##                                     Cramers_V       Note
## tissue_type                      0.1944 Significant
## laterality                        0.0324 Not significant
## tissue_or_organ_of_origin    0.3629 Significant
## primary_diagnosis                 0.1046 Not significant
## prior_treatment                   0.2902 Significant
## ajcc_pathologic_t                  0.1903 Significant
## morphology                         0.1046 Not significant
## classification_of_tumor      0.3276 Significant
## follow_ups_disease_response Chi-square (MC simulation) 0.6082 Significant
## race                                0.0807 Not significant
## gender                               0.0242 Not significant

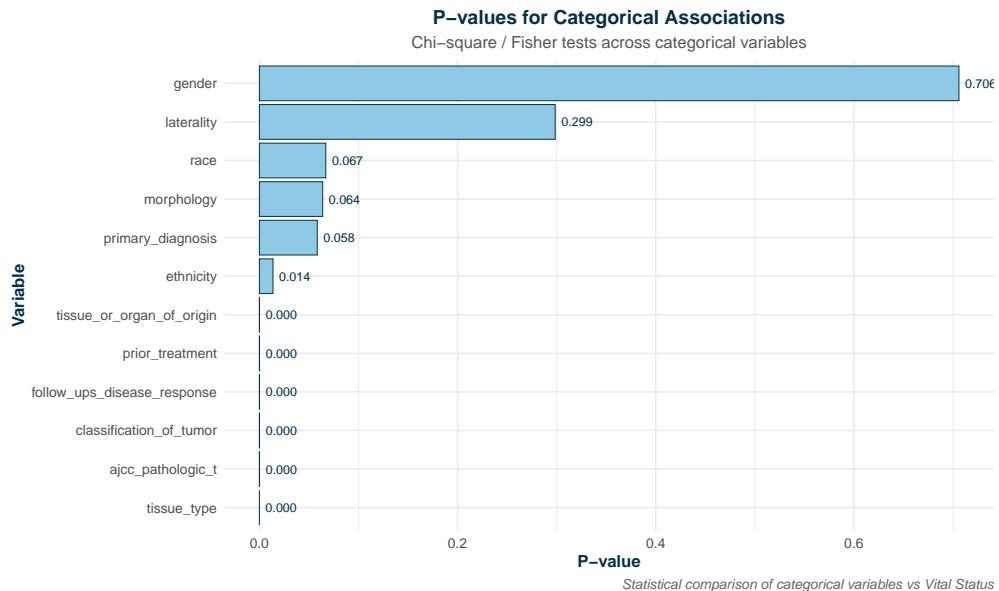
```

```

## ethnicity           0.0913      Significant

# Plot 1: Bar plot with P-value
ggplot(results_df, aes(x = reorder(Variable, P_value), y = P_value)) +
  geom_bar(stat = "identity", fill = "#8ecae6", color = "#023047", linewidth = 0.3) +
  geom_text(aes(label = sprintf("%.3f", P_value)),
            hjust = -0.2, size = 3, color = "#023047") +
  coord_flip() +
  labs(title = "P-values for Categorical Associations",
       subtitle = "Chi-square / Fisher tests across categorical variables",
       x = "Variable",
       y = "P-value",
       caption = "Statistical comparison of categorical variables vs Vital Status") +
  theme_minimal(base_size = 12) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047"),
        plot.subtitle = element_text(hjust = 0.5, color = "#555555"),
        plot.caption = element_text(face = "italic", color = "#666666"),
        axis.title = element_text(face = "bold", color = "#023047"))

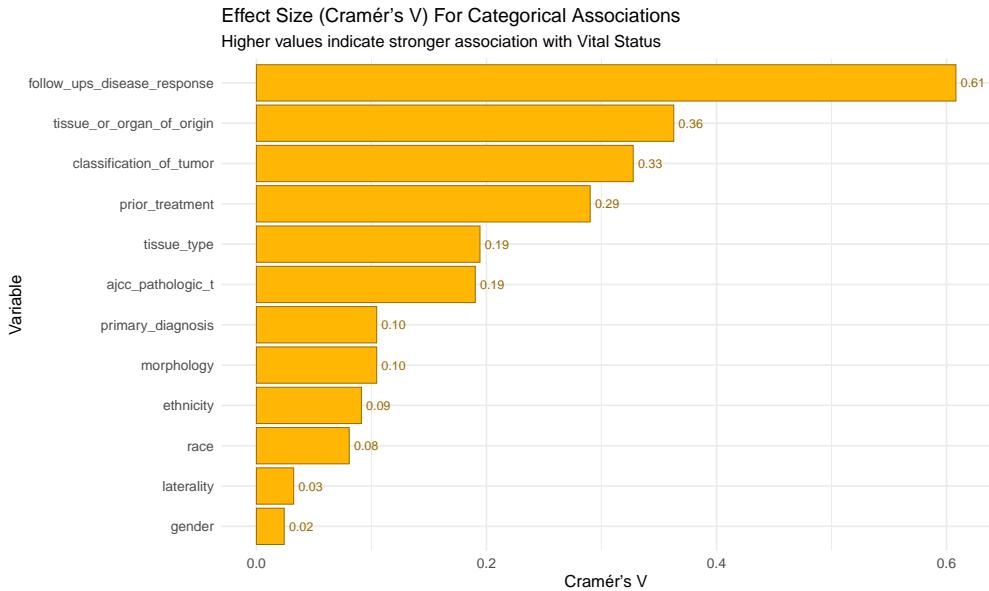
```



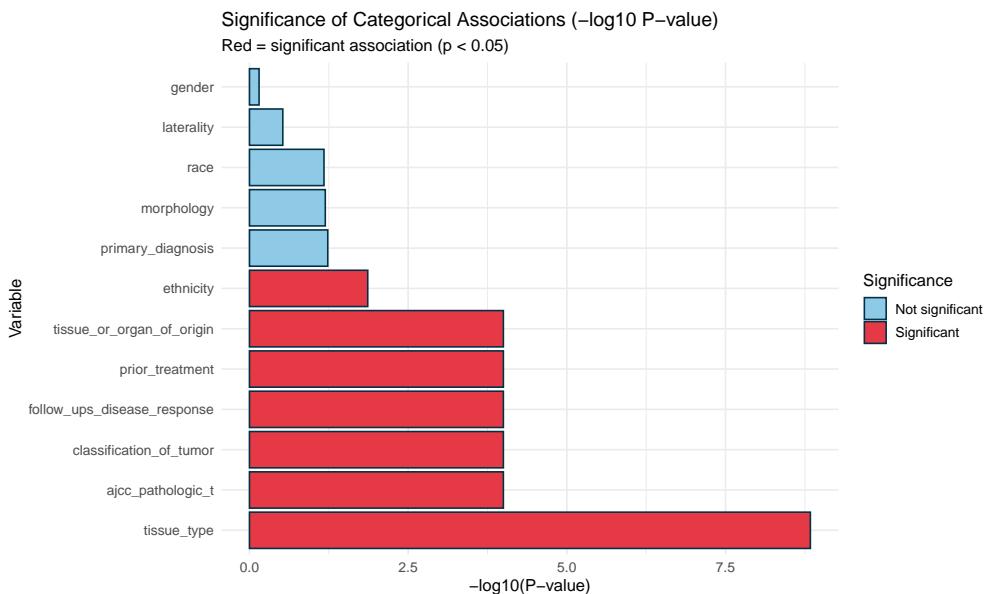
```

# Plot 1: Cramér's V Strength Plot
ggplot(results_df, aes(x = reorder(Variable, Cramers_V), y = Cramers_V)) +
  geom_bar(stat = "identity", fill = "#ffb703", color = "#9a6a00", linewidth = 0.3) +
  geom_text(aes(label = sprintf("%.2f", Cramers_V)),
            hjust = -0.2, size = 3, color = "#9a6a00") +
  coord_flip() +
  labs(title = "Effect Size (Cramér's V) For Categorical Associations",
       subtitle = "Higher values indicate stronger association with Vital Status",
       x = "Variable",
       y = "Cramér's V") +
  theme_minimal(base_size = 12)

```



```
# Plot 3: Significance Highlight Plot
ggplot(results_df, aes(x = reorder(Variable, P_value),
                       y = -log10(P_value),
                       fill = Note)) +
  geom_bar(stat = "identity", color = "#023047") +
  coord_flip() +
  scale_fill_manual(values = c("Significant" = "#e63946",
                               "Not significant" = "#8ecae6")) +
  labs(title = "Significance of Categorical Associations (-log10 P-value)",
       subtitle = "Red = significant association (p < 0.05)",
       x = "Variable",
       y = "-log10(P-value)",
       fill = "Significance") +
  theme_minimal(base_size = 12)
```



From this result, most of categorical variables show no meaningful link to survival. Gender has no effect, ethnicity is statistically significant but with very weak effect ($V \approx 0.09$), and tumor subtype shows only a borderline trend. Overall, categorical predictors contribute little to explaining survival differences.

Genex data studies

Differential Expression Analysis

```
# Matrix vital status
design <- model.matrix(~vital_status, data = clinical_df)

# Fit linear model using limma (empirical Bayes)
fit <- lmFit(t(GeneX_df), design)
fit <- eBayes(fit)

# Extract top differentially expressed genes
top_genes <- topTable(fit
                        , coef          = 2
                        , number        = 5000
                        , adjust.method = "BH")

cat("Top 20 Differentially Expressed Genes (Dead vs Alive):\n")

## Top 20 Differentially Expressed Genes (Dead vs Alive):

print(top_genes[1:20, c("logFC", "AveExpr", "P.Value", "adj.P.Val")])

##           logFC   AveExpr      P.Value    adj.P.Val
## LINC01235  0.9515500 7.047438 2.078880e-11 1.039440e-07
## APOB       1.2822838 3.413208 1.035776e-10 2.589440e-07
## LYVE1      1.0214367 7.224554 2.346103e-10 3.910171e-07
## LINC01497  0.7523277 1.156414 3.437925e-09 4.297407e-06
## AC104211.1 0.7273947 3.366769 1.643701e-08 1.535403e-05
## KLB        0.8499149 6.395080 2.069437e-08 1.535403e-05
## PSD2       0.7153297 4.402892 2.149564e-08 1.535403e-05
## LINC02511  0.8574617 2.253701 5.811988e-08 3.280902e-05
## SNORD104  -0.6927049 4.644127 5.905623e-08 3.280902e-05
## CST1       -1.4893705 6.693858 8.124615e-08 3.452865e-05
## AC007423.1 0.7409439 1.210320 8.195581e-08 3.452865e-05
## GPX3        0.7475773 11.556001 8.286876e-08 3.452865e-05
## LVRN        0.8868127 4.691047 1.429004e-07 4.769855e-05
## PROKR1      0.8122548 2.187548 1.442226e-07 4.769855e-05
## RHBDL1     -0.6842668 7.068012 1.518817e-07 4.769855e-05
## ADH4        0.8146112 2.169812 1.619101e-07 4.769855e-05
## ATF3        0.6698043 10.693098 1.621751e-07 4.769855e-05
## SLC2A4      0.7785291 6.008617 2.085221e-07 5.723414e-05
## FHL1        0.7816811 10.841301 2.174897e-07 5.723414e-05
## VEGFD       1.0116073 5.215086 2.874839e-07 6.940534e-05
```

```

cat("\n==== DE SUMMARY ====\n")

##
## === DE SUMMARY ===

cat("Significant genes (FDR < 0.05):", sum(top_genes$adj.P.Val < 0.05), "\n")

## Significant genes (FDR < 0.05): 1159

cat("Genes |logFC| > 1:", sum(abs(top_genes$logFC) > 1), "\n")

## Genes |logFC| > 1: 13

cat("Both significant AND |logFC| > 1:"
    , sum(top_genes$adj.P.Val < 0.05 & abs(top_genes$logFC) > 1), "\n\n")

## Both significant AND |logFC| > 1: 13

cat("Expression direction:\n")

## Expression direction:

cat(" Upregulated in Dead:", sum(top_genes$logFC > 0), "\n")

## Upregulated in Dead: 2692

cat(" Downregulated in Dead:", sum(top_genes$logFC < 0), "\n")

## Downregulated in Dead: 2308

```

Volcano Plot

```

plot(top_genes$logFC
    , -log10(top_genes$P.Value)
    , pch      = 19
    , cex      = 0.6
    , col      = ifelse(top_genes$adj.P.Val < 0.05
                        , ifelse(abs(top_genes$logFC) > 1, "#d62828", "#e63946")
                        , "#8ecae6")
    , xlab     = "Log2 Fold Change (Dead vs Alive)"
    , ylab     = "-log10(P-value)"
    , main     = "Volcano Plot: Differentially Expressed Genes"
    , sub      = "Significance threshold: FDR < 0.05, |logFC| > 1"
    , col.main = "#023047"
    , col.lab  = "#023047"
    , col.sub  = "#666666"

```

```

    , cex.sub = 0.8
    , font.sub = 3)

# Significance thresholds 0,05
abline(h = -log10(0.05)
       , col = "#fb8500"
       , lty = 2
       , lwd = 2)

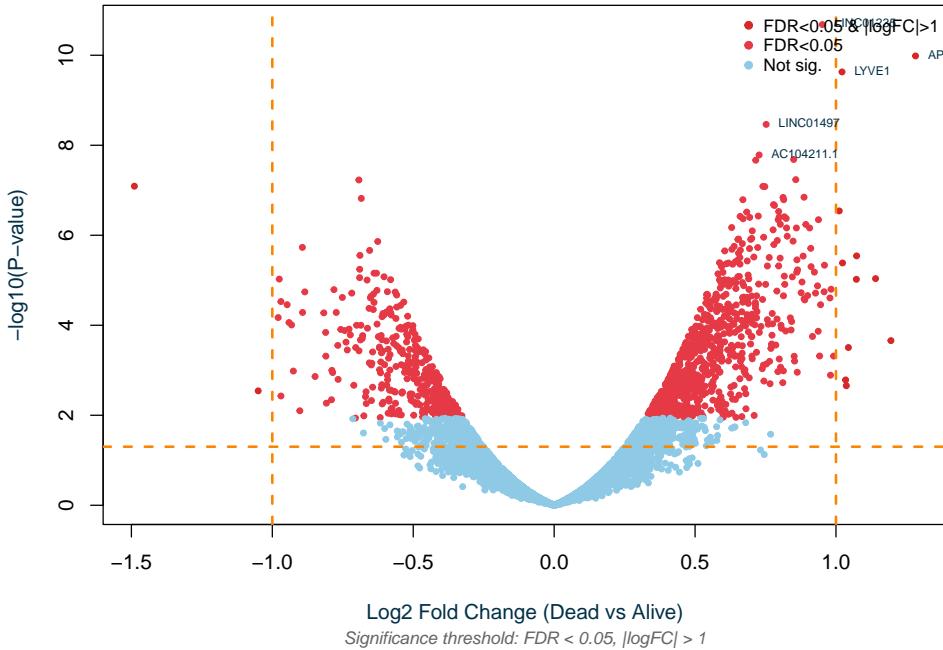
abline(v = c(-1, 1)
       , col = "#fb8500"
       , lty = 2
       , lwd = 2)

# Gene labels
top_hits <- rownames(top_genes)[1:5]
for(gene in top_hits) {
  text(top_genes[gene, "logFC"]
       , -log10(top_genes[gene, "P.Value"])
       , labels = gene
       , pos = 4
       , cex = 0.6
       , col = "#023047")
}

legend("topright"
       , legend = c("FDR<0.05 & |logFC|>1", "FDR<0.05", "Not sig.")
       , col = c("#d62828", "#e63946", "#8ecae6")
       , pch = 19
       , bty = "n"
       , cex = 0.8)

```

Volcano Plot: Differentially Expressed Genes



From the result **top_gene**, most significant genes show **positive logFC**, meaning they are **up-regulated** in patients who died, while like few **CST1**, **MMP11**, **SNORD104** are down-regulated. Several genes display both large effect sizes ($|logFC| > 1$) and very strong statistical significance (**FDR « 0.05**)—notably **APOB**, **LYVE1**, **LINC01497**, **AC104211.1** making them the clearest potential biomarkers.

The Volcano plot confirms a pronounced asymmetry, with a dense cluster of up-regulated genes in the Dead group, indicating the activated transcriptional programs link to poor diagnosis, where down-regulated genes are fewer and more dispersed.

Overall, the results point to a robust molecular signature differentiating Alive vs Dead patients, with a handful of genes emerging as particularly strong candidates for biological interpretation and predictive modeling.

Gene Expression Characteristics

```

top50_genes <- top_genes[1:50, ]

cat("== TOP 50 GENE CHARACTERISTICS ==\n\n")

## == TOP 50 GENE CHARACTERISTICS ==

# Expression levels
cat("Average Expression Levels:\n")

## Average Expression Levels:

```

```

cat("  Min:", round(min(top50_genes$AveExpr), 2), "\n")

##  Min: 1.16

cat("  Median:", round(median(top50_genes$AveExpr), 2), "\n")

##  Median: 4.49

cat("  Max:", round(max(top50_genes$AveExpr), 2), "\n\n")

##  Max: 11.56

# Fold changes
cat("Fold Change Distribution:\n")

## Fold Change Distribution:

cat("  Upregulated in Dead (logFC > 0):", sum(top50_genes$logFC > 0), "\n")

##  Upregulated in Dead (logFC > 0): 46

cat("  Downregulated in Dead (logFC < 0):", sum(top50_genes$logFC < 0), "\n\n")

##  Downregulated in Dead (logFC < 0): 4

# Statistical significance
cat("P-value ranges:\n")

## P-value ranges:

cat("  Min P-value:", format(min(top50_genes$P.Value), scientific = TRUE), "\n")

##  Min P-value: 2.07888e-11

cat("  Max P-value:", format(max(top50_genes$P.Value), scientific = TRUE), "\n")

##  Max P-value: 1.605859e-06

cat("  Max FDR:", format(max(top50_genes$adj.P.Val), scientific = TRUE), "\n")

##  Max FDR: 1.605859e-04

```

Gene Expression Distributions

```

# Create gene subset for top 20 genes
top20_genes <- rownames(top_genes)[1:20]
gene_subset <- as.data.frame(GeneX_df[, top20_genes])
colnames(gene_subset) <- top20_genes

par(mfrow = c(3, 3), bg = "white")

for(i in 1:9) {
  gene      <- top20_genes[i]
  gene_expr <- gene_subset[, i]

  # Histogram with separate colors by vital status
  hist(gene_expr[cclinical_df$vital_status == "Alive"]
    , breaks   = 30
    , col      = rgb(0.2, 0.6, 0.8, 0.5)
    , main     = paste(gene, "- Expression Distribution")
    , sub      = "Alive (blue) vs Dead (red) patients"
    , xlab     = "Expression Level"
    , ylab     = "Frequency"
    , border   = "white"
    , col.main = "#023047"
    , col.lab  = "#023047"
    , col.sub  = "#666666"
    , cex.sub  = 0.7
    , font.sub = 3)

  hist(gene_expr[cclinical_df$vital_status == "Dead"]
    , breaks = 30
    , col    = rgb(0.9, 0.2, 0.3, 0.5)
    , add    = TRUE
    , border = "white")

  legend("topright"
    , legend = c("Alive", "Dead")
    , fill   = c(rgb(0.2, 0.6, 0.8, 0.5), rgb(0.9, 0.2, 0.3, 0.5))
    , bty    = "n")

  # Test for bimodality (Hartigan's dip test)
  dip_result <- dip.test(gene_expr)

  if(dip_result$p.value < 0.05) {
    cat(sprintf("%s: BIMODAL (p=% .4f) need subgroups!\n"
      , gene
      , dip_result$p.value))
  }
}

## APOB: BIMODAL (p=0.0000) need subgroups!

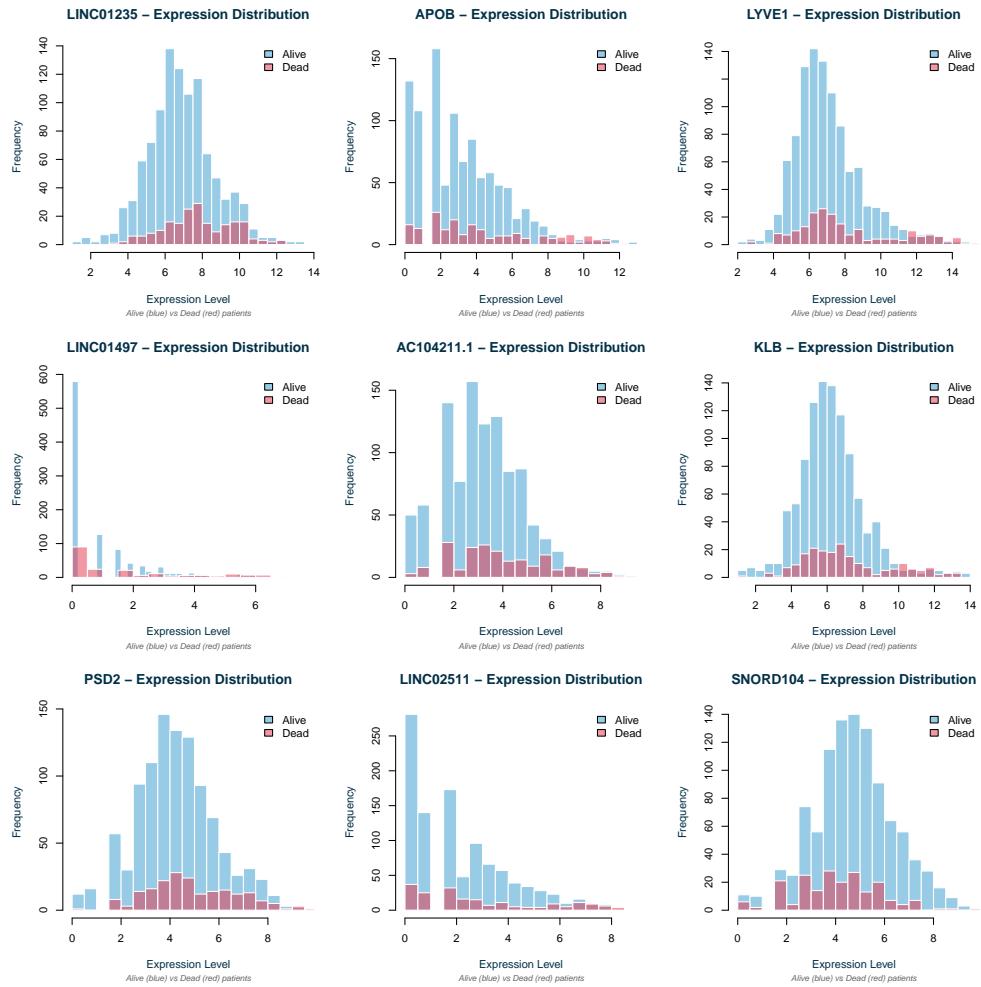
## LINC01497: BIMODAL (p=0.0000) need subgroups!

## AC104211.1: BIMODAL (p=0.0000) need subgroups!

```

```
## PSD2: BIMODAL (p=0.0009) need subgroups!
```

```
## LINC02511: BIMODAL (p=0.0000) need subgroups!
```



A subset of the strongest DE genes shows bimodal expression patterns, indicating heterogeneity and possible molecular subgroups, while several genes exhibit clear upregulation in non-survivors, reinforcing their biological relevance.

Bimodal Gene Analysis

```
# Create gene subset for top 500
top500_genes <- rownames(top_genes)[1:500]
gene500_subset <- as.data.frame(GeneX_df[, top500_genes])
colnames(gene500_subset) <- top500_genes

# Run dip test for each of the top 500 genes
dip_results <- sapply(gene500_subset, function(x) {
  dip.test(x)$p.value
})
```

```

# Convert to data frame
bimodality_df <- data.frame(
  Gene = names(dip_results),
  Dip_Pvalue = dip_results
)

bimodality_df$Is_Bimodal <- bimodality_df$Dip_Pvalue < 0.05

# Sort by lowest dip-test pvalue (most strongly bimodal first)
bimodality_df <- bimodality_df[order(bimodality_df$Dip_Pvalue), ]

# Check only Bimodal
bimodal_genes_only <- bimodality_df[bimodality_df$Is_Bimodal == TRUE, ]

cat("\n==== BIMODAL GENES (Dip p < 0.05) ===\n")

## ===== BIMODAL GENES (Dip p < 0.05) =====

print(bimodal_genes_only)

##          Gene   Dip_Pvalue Is_Bimodal
## APOB      APOB 0.000000e+00     TRUE
## LINC01497  LINC01497 0.000000e+00    TRUE
## AC104211.1 AC104211.1 0.000000e+00    TRUE
## LINC02511  LINC02511 0.000000e+00    TRUE
## CST1       CST1 0.000000e+00     TRUE
## AC007423.1 AC007423.1 0.000000e+00    TRUE
## PROKR1     PROKR1 0.000000e+00    TRUE
## ADH4        ADH4 0.000000e+00     TRUE
## ALDH1L1-AS2 ALDH1L1-AS2 0.000000e+00    TRUE
## LINC01537  LINC01537 0.000000e+00    TRUE
## LINC01186  LINC01186 0.000000e+00    TRUE
## GLP2R      GLP2R 0.000000e+00     TRUE
## NGF-AS1    NGF-AS1 0.000000e+00    TRUE
## HSD17B13   HSD17B13 0.000000e+00    TRUE
## LUARIS     LUARIS 0.000000e+00    TRUE
## DSC1       DSC1 0.000000e+00     TRUE
## LINC01612  LINC01612 0.000000e+00    TRUE
## FP325317.1 FP325317.1 0.000000e+00    TRUE
## ABCB5      ABCB5 0.000000e+00     TRUE
## ADRA1A     ADRA1A 0.000000e+00    TRUE
## LHCGR      LHCGR 0.000000e+00    TRUE
## PGM5-AS1   PGM5-AS1 0.000000e+00    TRUE
## AL356218.2 AL356218.2 0.000000e+00    TRUE
## C1QTNF9   C1QTNF9 0.000000e+00    TRUE
## ADH1A      ADH1A 0.000000e+00     TRUE
## AC036108.2 AC036108.2 0.000000e+00    TRUE
## AC079804.3 AC079804.3 0.000000e+00    TRUE
## GLRA4      GLRA4 0.000000e+00     TRUE
## LINC02237  LINC02237 0.000000e+00    TRUE
## CA4        CA4 0.000000e+00     TRUE

```

```

## AL121950.1    AL121950.1 0.000000e+00      TRUE
## EPB42          EPB42 0.000000e+00      TRUE
## PROX1-AS1     PROX1-AS1 0.000000e+00      TRUE
## GLYAT          GLYAT 0.000000e+00      TRUE
## SPX            SPX 0.000000e+00      TRUE
## AC104407.1    AC104407.1 0.000000e+00      TRUE
## TM4SF4         TM4SF4 0.000000e+00      TRUE
## LALBA          LALBA 0.000000e+00      TRUE
## LINC01697      LINC01697 0.000000e+00      TRUE
## MAFA-AS1       MAFA-AS1 0.000000e+00      TRUE
## PCK1           PCK1 0.000000e+00      TRUE
## NPY2R           NPY2R 0.000000e+00      TRUE
## AC084212.1    AC084212.1 0.000000e+00      TRUE
## CDH20          CDH20 0.000000e+00      TRUE
## AC105118.1    AC105118.1 0.000000e+00      TRUE
## ANGPTL8        ANGPTL8 0.000000e+00      TRUE
## LINC02660      LINC02660 0.000000e+00      TRUE
## AC073850.1    AC073850.1 0.000000e+00      TRUE
## AL450332.1    AL450332.1 0.000000e+00      TRUE
## NEUROG2        NEUROG2 0.000000e+00      TRUE
## AL845331.1    AL845331.1 0.000000e+00      TRUE
## SERTM1          SERTM1 0.000000e+00      TRUE
## LINCADL        LINCADL 0.000000e+00      TRUE
## IBSP            IBSP 0.000000e+00      TRUE
## ANO3            ANO3 0.000000e+00      TRUE
## ADH1C           ADH1C 0.000000e+00      TRUE
## PLCZ1          PLCZ1 0.000000e+00      TRUE
## AC016682.1    AC016682.1 0.000000e+00      TRUE
## LGALS17A        LGALS17A 0.000000e+00      TRUE
## TMEM252         TMEM252 0.000000e+00      TRUE
## TRHDE-AS1      TRHDE-AS1 0.000000e+00      TRUE
## ANGPTL7        ANGPTL7 0.000000e+00      TRUE
## AP001360.1    AP001360.1 0.000000e+00      TRUE
## CDH12           CDH12 0.000000e+00      TRUE
## LRRC3B          LRRC3B 0.000000e+00      TRUE
## ACSM4           ACSM4 0.000000e+00      TRUE
## MYOC            MYOC 0.000000e+00      TRUE
## H2BC17          H2BC17 0.000000e+00      TRUE
## AC003986.2    AC003986.2 0.000000e+00      TRUE
## SGCZ            SGCZ 0.000000e+00      TRUE
## AL591686.1    AL591686.1 0.000000e+00      TRUE
## NRAD1           NRAD1 0.000000e+00      TRUE
## ACE2             ACE2 0.000000e+00      TRUE
## AL353693.1    AL353693.1 0.000000e+00      TRUE
## GRIN2B          GRIN2B 0.000000e+00      TRUE
## MIR145          MIR145 0.000000e+00      TRUE
## AL138716.1    AL138716.1 0.000000e+00      TRUE
## LINC01230      LINC01230 0.000000e+00      TRUE
## CSF3             CSF3 0.000000e+00      TRUE
## TNNI3           TNNI3 0.000000e+00      TRUE
## AC112721.2    AC112721.2 0.000000e+00      TRUE
## B3GAT1-DT      B3GAT1-DT 0.000000e+00      TRUE
## LINC01561      LINC01561 0.000000e+00      TRUE
## CCL14           CCL14 0.000000e+00      TRUE

```

```

## NOS1           NOS1 0.000000e+00    TRUE
## MLIP          MLIP  0.000000e+00    TRUE
## AC093496.1   AC093496.1 0.000000e+00  TRUE
## KCNJ16        KCNJ16 0.000000e+00    TRUE
## AC092118.1   AC092118.1 0.000000e+00    TRUE
## AC108734.4   AC108734.4 0.000000e+00    TRUE
## AC002546.1   AC002546.1 0.000000e+00    TRUE
## AQP7P1        AQP7P1 0.000000e+00    TRUE
## CSN1S1        CSN1S1 0.000000e+00    TRUE
## AADAC         AADAC 0.000000e+00    TRUE
## AC016924.1   AC016924.1 0.000000e+00    TRUE
## LINC02587    LINC02587 0.000000e+00    TRUE
## CST4          CST4  0.000000e+00    TRUE
## AC093817.2   AC093817.2 0.000000e+00    TRUE
## TRDN          TRDN  0.000000e+00    TRUE
## SHISA3        SHISA3 0.000000e+00    TRUE
## KCNH1-IT1    KCNH1-IT1 0.000000e+00    TRUE
## HEPACAM       HEPACAM 0.000000e+00    TRUE
## DCT           DCT  0.000000e+00    TRUE
## TRHDE         TRHDE 0.000000e+00    TRUE
## TGFBR3L       TGFBR3L 0.000000e+00    TRUE
## AL645924.1   AL645924.1 0.000000e+00    TRUE
## SLC6A3        SLC6A3 0.000000e+00    TRUE
## CCDC144A      CCDC144A 0.000000e+00    TRUE
## RBMS3-AS3    RBMS3-AS3 0.000000e+00    TRUE
## LINC01281    LINC01281 0.000000e+00    TRUE
## DPP6          DPP6  0.000000e+00    TRUE
## HHATL         HHATL 0.000000e+00    TRUE
## Z98745.2     Z98745.2 0.000000e+00    TRUE
## HSD11B1-AS1  HSD11B1-AS1 0.000000e+00    TRUE
## C6            C6  0.000000e+00    TRUE
## RXRG          RXRG  0.000000e+00    TRUE
## CNTN6         CNTN6 0.000000e+00    TRUE
## GRIA4         GRIA4 0.000000e+00    TRUE
## AC008459.1   AC008459.1 0.000000e+00    TRUE
## ANTXRL        ANTXRL 0.000000e+00    TRUE
## PTCHD3        PTCHD3 0.000000e+00    TRUE
## SLC7A14-AS1  SLC7A14-AS1 0.000000e+00    TRUE
## FOXD3-AS1    FOXD3-AS1 0.000000e+00    TRUE
## AC110774.1   AC110774.1 0.000000e+00    TRUE
## CPA1          CPA1  0.000000e+00    TRUE
## PURPL         PURPL 0.000000e+00    TRUE
## BAK1P2        BAK1P2 0.000000e+00    TRUE
## SLC7A10       SLC7A10 0.000000e+00    TRUE
## AP002800.1   AP002800.1 0.000000e+00    TRUE
## SFTPB         SFTPB 0.000000e+00    TRUE
## NEUROG2-AS1  NEUROG2-AS1 0.000000e+00    TRUE
## AC092851.1   AC092851.1 0.000000e+00    TRUE
## GPS2P1        GPS2P1 0.000000e+00    TRUE
## PROK1         PROK1 0.000000e+00    TRUE
## AC121757.1   AC121757.1 0.000000e+00    TRUE
## OR2B6         OR2B6 0.000000e+00    TRUE
## ADCY8         ADCY8 0.000000e+00    TRUE
## AL356489.2   AL356489.2 0.000000e+00    TRUE

```

## SH3GL3	SH3GL3	0.000000e+00	TRUE
## TUBA3E	TUBA3E	0.000000e+00	TRUE
## CAPZA3	CAPZA3	0.000000e+00	TRUE
## CNTNAP3P2	CNTNAP3P2	0.000000e+00	TRUE
## CCNYL2	CCNYL2	0.000000e+00	TRUE
## SLITRK2	SLITRK2	0.000000e+00	TRUE
## MPPED1	MPPED1	0.000000e+00	TRUE
## C14orf180	C14orf180	0.000000e+00	TRUE
## LMX1A	LMX1A	0.000000e+00	TRUE
## CMA1	CMA1	0.000000e+00	TRUE
## RPS4Y1	RPS4Y1	0.000000e+00	TRUE
## LEP	LEP	0.000000e+00	TRUE
## CYP1A1	CYP1A1	0.000000e+00	TRUE
## PTPRQ	PTPRQ	0.000000e+00	TRUE
## AP005131.3	AP005131.3	0.000000e+00	TRUE
## MAPT-IT1	MAPT-IT1	0.000000e+00	TRUE
## LINC01625	LINC01625	0.000000e+00	TRUE
## AC098850.3	AC098850.3	0.000000e+00	TRUE
## AC119424.1	AC119424.1	0.000000e+00	TRUE
## LINC00844	LINC00844	0.000000e+00	TRUE
## GRAMD4P8	GRAMD4P8	0.000000e+00	TRUE
## AL513318.1	AL513318.1	0.000000e+00	TRUE
## SNTN	SNTN	0.000000e+00	TRUE
## LINC01485	LINC01485	0.000000e+00	TRUE
## AC022196.1	AC022196.1	0.000000e+00	TRUE
## AC068733.3	AC068733.3	0.000000e+00	TRUE
## AC073869.6	AC073869.6	0.000000e+00	TRUE
## DRD1	DRD1	0.000000e+00	TRUE
## TNMD	TNMD	0.000000e+00	TRUE
## AC073316.1	AC073316.1	0.000000e+00	TRUE
## AP000350.6	AP000350.6	0.000000e+00	TRUE
## LINC00466	LINC00466	0.000000e+00	TRUE
## FAM180B	FAM180B	0.000000e+00	TRUE
## HBA1	HBA1	0.000000e+00	TRUE
## AL033384.1	AL033384.1	0.000000e+00	TRUE
## LINC01344	LINC01344	0.000000e+00	TRUE
## LINC02515	LINC02515	0.000000e+00	TRUE
## DIRAS2	DIRAS2	0.000000e+00	TRUE
## PENK	PENK	0.000000e+00	TRUE
## OXGR1	OXGR1	0.000000e+00	TRUE
## USP32P1	USP32P1	0.000000e+00	TRUE
## TMEM132C	TMEM132C	0.000000e+00	TRUE
## PLD5	PLD5	0.000000e+00	TRUE
## MGAT4C	MGAT4C	0.000000e+00	TRUE
## AL161945.1	AL161945.1	0.000000e+00	TRUE
## CLDN25	CLDN25	0.000000e+00	TRUE
## ASB4	ASB4	0.000000e+00	TRUE
## BMS1P10	BMS1P10	7.134626e-07	TRUE
## CST2	CST2	3.129484e-06	TRUE
## CHRNA6	CHRNA6	4.741375e-06	TRUE
## RERGL	RERGL	4.741375e-06	TRUE
## PRR26	PRR26	7.722067e-06	TRUE
## LINC00968	LINC00968	7.965158e-06	TRUE
## HPSE2	HPSE2	7.965158e-06	TRUE

```

## PLCXD3      PLCXD3 9.577049e-06    TRUE
## ANGPT4      ANGPT4 9.577049e-06    TRUE
## CCDC178     CCDC178 9.577049e-06   TRUE
## LINC01239    LINC01239 9.577049e-06  TRUE
## RIMBP2      RIMBP2 2.036191e-05   TRUE
## ADGRB3      ADGRB3 3.371400e-05   TRUE
## SCT         SCT 9.664365e-05    TRUE
## HIF3A       HIF3A 1.505938e-04   TRUE
## KY          KY 1.505938e-04    TRUE
## NLGN1       NLGN1 2.184164e-04   TRUE
## AL583785.1  AL583785.1 2.184164e-04 TRUE
## GRIK1       GRIK1 3.423733e-04   TRUE
## MAP1LC3C    MAP1LC3C 4.663301e-04  TRUE
## MASP1        MASP1 4.663301e-04   TRUE
## AC090004.2  AC090004.2 4.663301e-04 TRUE
## AC055854.1  AC055854.1 4.663301e-04 TRUE
## PSD2         PSD2 9.472775e-04   TRUE
## SSTR1        SSTR1 9.472775e-04   TRUE
## H2AC13      H2AC13 9.472775e-04  TRUE
## AL032819.2  AL032819.2 1.382277e-03 TRUE
## NRXN1        NRXN1 1.862362e-03   TRUE
## CD300LG      CD300LG 2.751299e-03  TRUE
## HOXA2        HOXA2 2.751299e-03   TRUE
## PLAC1        PLAC1 2.751299e-03   TRUE
## DRD2         DRD2 2.751299e-03   TRUE
## ZBED2        ZBED2 3.804563e-03   TRUE
## CALHM6      CALHM6 6.492312e-03  TRUE
## LINC00922    LINC00922 6.880659e-03 TRUE
## CIDEc        CIDEc 6.880659e-03   TRUE
## CRHBP        CRHBP 9.054789e-03  TRUE
## KRT19P1      KRT19P1 9.054789e-03 TRUE
## ADH1B         ADH1B 9.054789e-03  TRUE
## PRRT4        PRRT4 1.231416e-02  TRUE
## ADAMTS9-AS2 ADAMTS9-AS2 1.231416e-02 TRUE
## FOXJ1        FOXJ1 1.575060e-02   TRUE
## LVRN          LVRN 1.640823e-02   TRUE
## PAPPA2        PAPPA2 1.872212e-02  TRUE
## VEGFD         VEGFD 2.104729e-02  TRUE
## SCN3A         SCN3A 2.104729e-02  TRUE
## SULT1B1      SULT1B1 2.104729e-02 TRUE
## NMUR1         NMUR1 2.958350e-02  TRUE
## GDF10         GDF10 2.958350e-02  TRUE
## DQX1          DQX1 3.811971e-02  TRUE
## HNRNPA1P21   HNRNPA1P21 3.811971e-02 TRUE
## RHOXF1-AS1   RHOXF1-AS1 4.665592e-02 TRUE
## LRRC2         LRRC2 4.665592e-02   TRUE
## RUFY4         RUFY4 4.665592e-02   TRUE

```

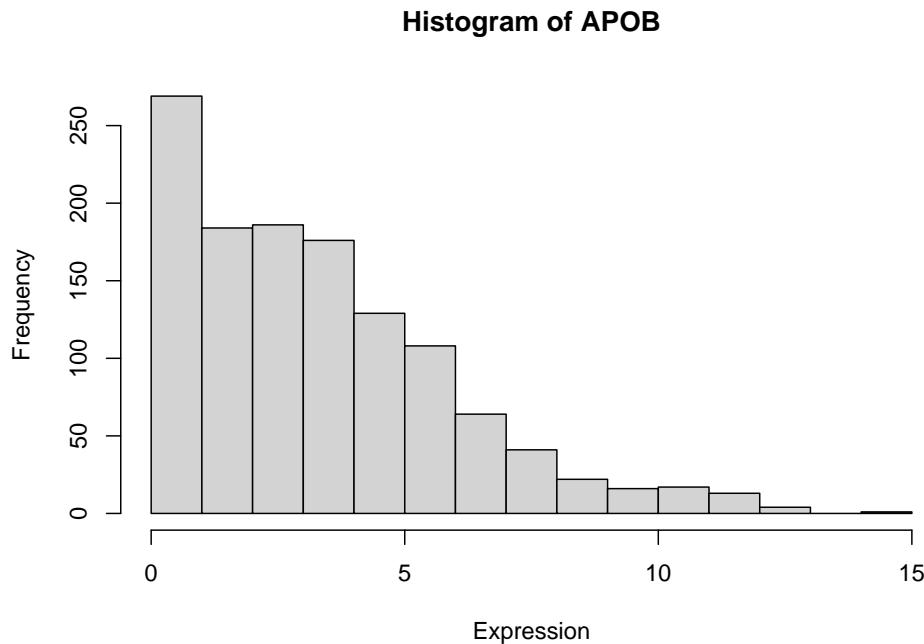
```

# Visualize top Bimodal Genes
top_bimodal_gene <- bimodality_df$Gene[1]
expr_values <- gene500_subset[[top_bimodal_gene]]

hist(expr_values, breaks = 20,
      main = paste("Histogram of", top_bimodal_gene),

```

```
xlab = "Expression")
```



```
cat("== BIMODALITY CHECK SUMMARY (Top 500 Genes) ==\n")
```

```
## == BIMODALITY CHECK SUMMARY (Top 500 Genes) ==
```

```
cat("Total genes tested:", nrow(bimodality_df), "\n")
```

```
## Total genes tested: 500
```

```
cat("Bimodal genes (Dip p < 0.05):", sum(bimodality_df$Is_Bimodal), "\n\n")
```

```
## Bimodal genes (Dip p < 0.05): 239
```

```
cat("Top 10 most bimodal genes:\n")
```

```
## Top 10 most bimodal genes:
```

```
print(head(bimodality_df, 10))
```

	Gene	Dip_Pvalue	Is_Bimodal
## APOB	APOB	0	TRUE
## LINC01497	LINC01497	0	TRUE
## AC104211.1	AC104211.1	0	TRUE
## LINC02511	LINC02511	0	TRUE
## CST1	CST1	0	TRUE

```

## AC007423.1    AC007423.1      0      TRUE
## PROKR1         PROKR1        0      TRUE
## ADH4           ADH4          0      TRUE
## ALDH1L1-AS2   ALDH1L1-AS2    0      TRUE
## LINC01537     LINC01537    0      TRUE

```

Top 10 Bimodal distributions

```

par(mfrow = c(2, 3), bg = "white")

# Loop through top 10 most bimodal genes
for (gene in head(bimodality_df$Gene, 10)) {

  gene_expr <- gene500_subset[, gene]

  alive_expr <- gene_expr[cclinical_df$vital_status == "Alive"]
  dead_expr  <- gene_expr[cclinical_df$vital_status == "Dead"]

  # Make sure densities exist (avoids zero-length errors)
  if (length(alive_expr) > 1 & length(dead_expr) > 1) {

    plot(density(alive_expr, na.rm = TRUE),
          col      = "#219ebc",
          lwd      = 3,
          main    = paste(gene, "- Expression Distribution"),
          sub     = "Density plot + Median cutpoint (orange)",
          xlab    = "Expression Level",
          ylab    = "Density",
          col.main = "#023047",
          col.lab  = "#023047",
          col.sub   = "#666666",
          cex.sub  = 0.7,
          font.sub = 3)

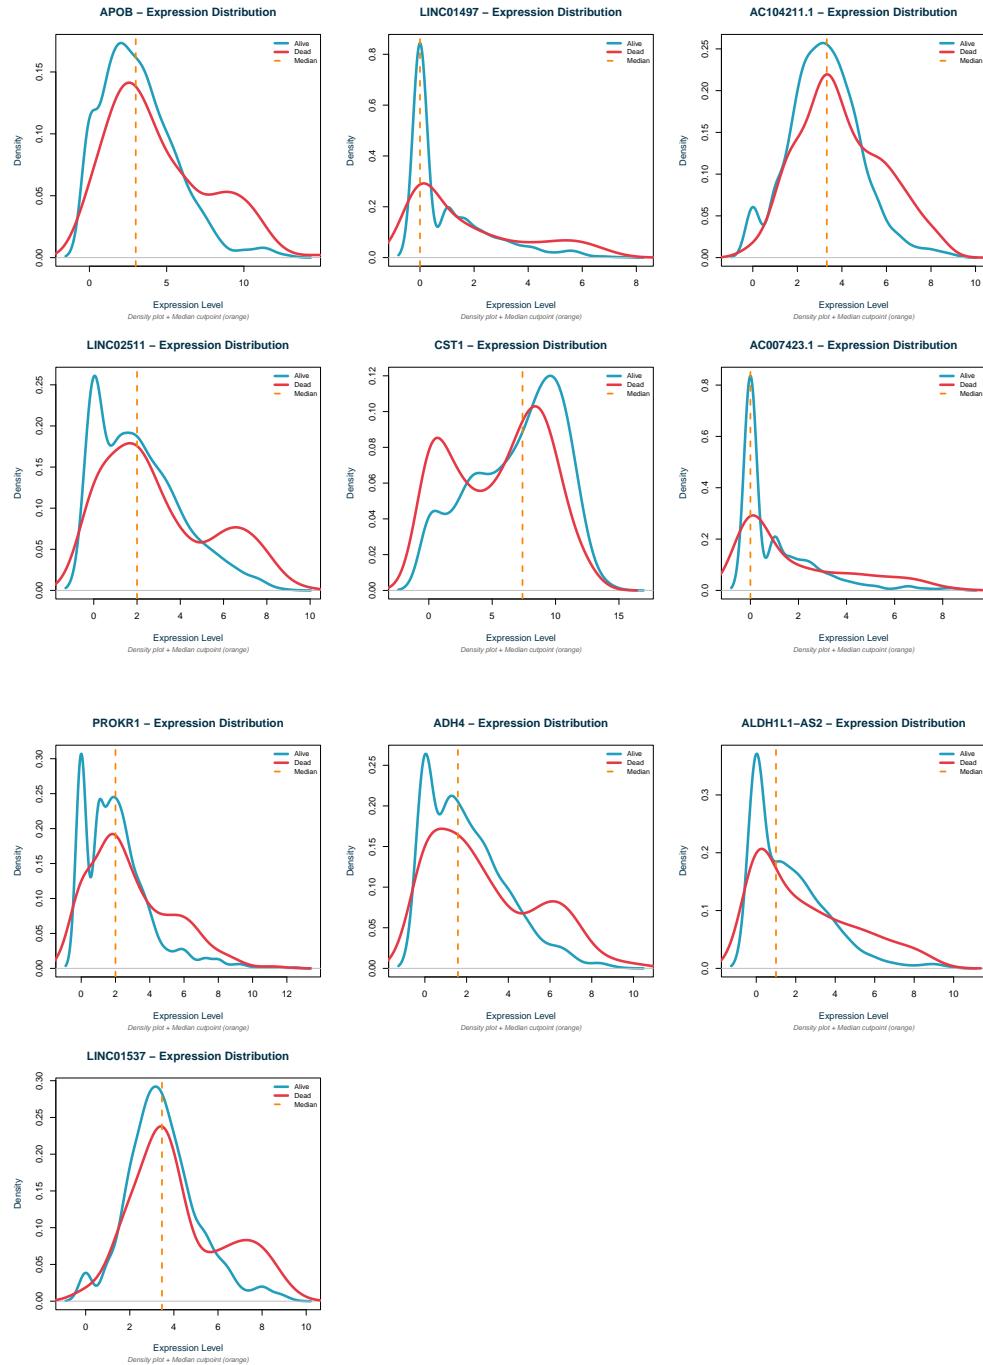
    lines(density(dead_expr, na.rm = TRUE),
          col = "#e63946",
          lwd = 3)

    # Add median vertical line
    abline(v = median(gene_expr, na.rm = TRUE),
            col = "#fb8500",
            lty = 2,
            lwd = 2)

    legend("topright",
           legend = c("Alive", "Dead", "Median"),
           col    = c("#219ebc", "#e63946", "#fb8500"),
           lwd    = c(3, 3, 2),
           lty    = c(1, 1, 2),
           bty    = "n",
           cex    = 0.8)
  }
}

```

}



```
# Identify bimodal genes from paper
bimodal_genes <- c("APOB", "LINC01497", "AC104211.1", "PSD2", "LINC02511")

cat("Bimodal genes identified:", length(bimodal_genes), "\n")
```

Bimodal genes identified: 5

```

cat(paste(bimodal_genes, collapse = ", "), "\n\n")

## APOB, LINC01497, AC104211.1, PSD2, LINC02511

# For each bimodal gene, create binary groups (high/low expressors)
cat("Creating binary expression groups for bimodal genes:\n\n")

## Creating binary expression groups for bimodal genes:

for(gene in bimodal_genes) {
  gene_expr <- gene_subset[, gene]
  gene_median <- median(gene_expr)

  # Create binary groups
  clinical_df[[paste0(gene, "_group")]] <- ifelse(gene_expr > gene_median
                                                , "High", "Low")

  # Test survival difference
  high_death_rate <- sum(clinical_df[[paste0(gene, "_group")]] == "High" &
                         clinical_df$vital_status == "Dead") /
    sum(clinical_df[[paste0(gene, "_group")]] == "High")

  low_death_rate <- sum(clinical_df[[paste0(gene, "_group")]] == "Low" &
                        clinical_df$vital_status == "Dead") /
    sum(clinical_df[[paste0(gene, "_group")]] == "Low")

  cat(sprintf("%s:\n", gene))
  cat(sprintf("  High expressors: %.1f%% mortality\n", high_death_rate * 100))
  cat(sprintf("  Low expressors: %.1f%% mortality\n", low_death_rate * 100))
  cat(sprintf("  Fold difference: %.2fx\n\n", high_death_rate / low_death_rate))
}

## APOB:
##   High expressors: 19.3% mortality
##   Low expressors: 13.6% mortality
##   Fold difference: 1.42x
##
## LINC01497:
##   High expressors: 19.8% mortality
##   Low expressors: 13.5% mortality
##   Fold difference: 1.47x
##
## AC104211.1:
##   High expressors: 20.2% mortality
##   Low expressors: 12.9% mortality
##   Fold difference: 1.56x
##
## PSD2:
##   High expressors: 20.2% mortality
##   Low expressors: 12.6% mortality
##   Fold difference: 1.60x
##

```

```

## LINC02511:
##   High expressors: 19.7% mortality
##   Low expressors: 13.7% mortality
##   Fold difference: 1.44x

```

Bimodal Gene Distributions

```

par(mfrow = c(2, 3), bg = "white")

for(gene in bimodal_genes) {
  gene_expr <- gene_subset[, gene]

  # Density plot
  alive_expr <- gene_expr[cclinical_df$vital_status == "Alive"]
  dead_expr <- gene_expr[cclinical_df$vital_status == "Dead"]

  plot(density(alive_expr, na.rm = TRUE)
       , col      = "#219ebc"
       , lwd      = 3
       , main     = paste(gene, "- Bimodal Distribution")
       , sub      = "Density plot with median cutpoint (orange)"
       , xlab     = "Expression Level"
       , ylab     = "Density"
       , col.main = "#023047"
       , col.lab  = "#023047"
       , col.sub  = "#666666"
       , cex.sub  = 0.7
       , font.sub = 3)

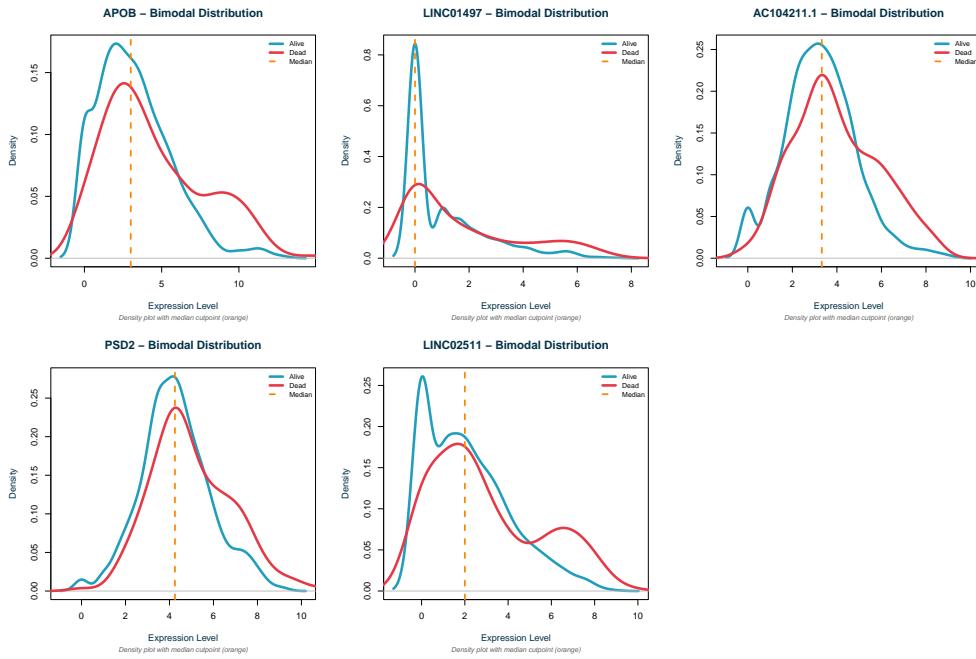
  lines(density(dead_expr, na.rm = TRUE)
        , col = "#e63946"
        , lwd = 3)

  # Add vertical line at median (cutpoint)
  abline(v = median(gene_expr)
         , col = "#fb8500"
         , lty = 2
         , lwd = 2)

  legend("topright"
        , legend = c("Alive", "Dead", "Median")
        , col    = c("#219ebc", "#e63946", "#fb8500")
        , lwd    = c(3, 3, 2)
        , lty    = c(1, 1, 2)
        , bty   = "n"
        , cex   = 0.8)
}

}

```



These five genes are likely to present subgroup markers: their expression defined nature of the patient cluster with different survival risk. They do not act alone as strong predictors, but provide important biological signal which could improve modeling.

Key Cancer Genes Check

```
key_genes <- c("GATA3", "CDH1", "ESR1", "PGR")
for(gene in key_genes) {
  cat("Gene:", gene, "\n")

  # Expression by vital status
  alive_expr <- GeneX_df[clinical_df$vital_status == "Alive", gene]
  dead_expr <- GeneX_df[clinical_df$vital_status == "Dead", gene]

  cat("  Alive: mean=", round(mean(alive_expr, na.rm = TRUE), 2),
      ", sd=", round(sd(alive_expr, na.rm = TRUE), 2), "\n", sep = "")
  cat("  Dead: mean=", round(mean(dead_expr, na.rm = TRUE), 2),
      ", sd=", round(sd(dead_expr, na.rm = TRUE), 2), "\n", sep = "")

  # T-test
  test <- t.test(alive_expr, dead_expr)
  cat("  T-test p-value:", format(test$p.value, scientific = TRUE), "\n")

  # Check if in top genes
  if(gene %in% rownames(top_genes)) {
    idx <- which(rownames(top_genes) == gene)
    cat("  Rank in DE analysis:", idx, "\n")
    cat("  logFC:", round(top_genes[gene, "logFC"], 3), "\n")
    cat("  FDR:", format(top_genes[gene, "adj.P.Val"], scientific = TRUE), "\n")
  } else {
```

```

        cat("  Not in top 100 DE genes\n")
    }
cat("\n")
}

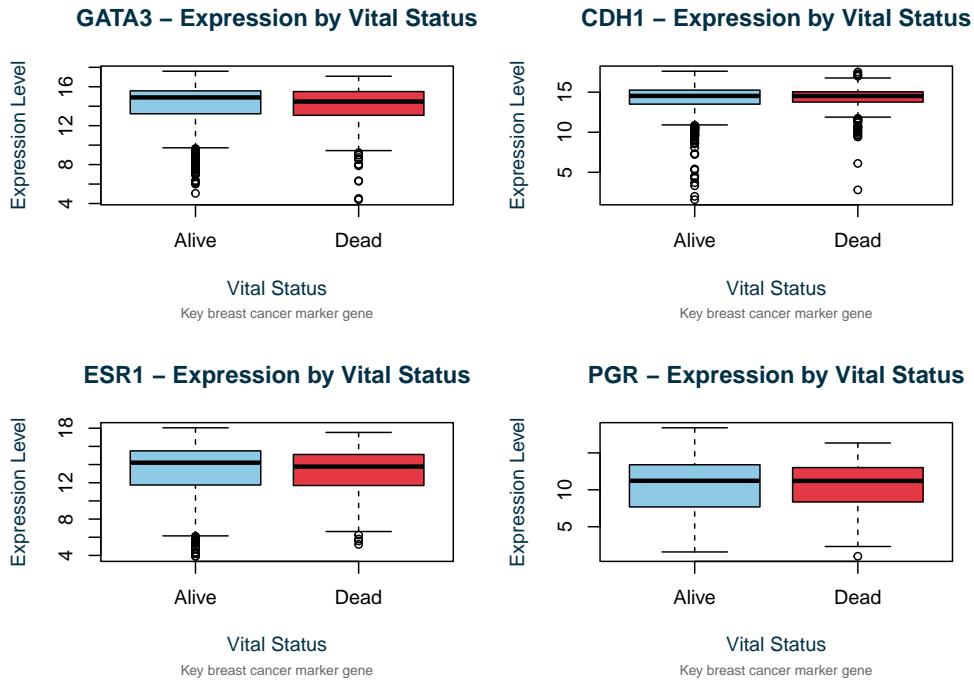
## Gene: GATA3
##   Alive: mean=14.17 sd=2.14
##   Dead:  mean=13.87 sd=2.34
##   T-test p-value: 9.546172e-02
##   Rank in DE analysis: 2017
##   logFC: -0.298
##   FDR: 1.881517e-01
##
## Gene: CDH1
##   Alive: mean=14.08 sd=1.95
##   Dead:  mean=14.12 sd=1.83
##   T-test p-value: 7.598542e-01
##   Rank in DE analysis: 4489
##   logFC: 0.044
##   FDR: 8.568828e-01
##
## Gene: ESR1
##   Alive: mean=13.24 sd=3.1
##   Dead:  mean=13.09 sd=2.85
##   T-test p-value: 5.200503e-01
##   Rank in DE analysis: 3868
##   logFC: -0.144
##   FDR: 6.990844e-01
##
## Gene: PGR
##   Alive: mean=10.58 sd=3.38
##   Dead:  mean=10.63 sd=3.18
##   T-test p-value: 8.278284e-01
##   Rank in DE analysis: 4641
##   logFC: 0.054
##   FDR: 8.982142e-01

```

```

# Box plot
par(mfrow = c(2, 2))
for(gene in key_genes) {
  boxplot(GeneX_df[, gene] ~ clinical_df$vital_status
          , col      = c("#8ecae6", "#e63946")
          , main    = paste(gene, "- Expression by Vital Status")
          , sub     = "Key breast cancer marker gene"
          , xlab    = "Vital Status"
          , ylab    = "Expression Level"
          , names   = c("Alive", "Dead")
          , col.main = "#023047"
          , col.lab  = "#023047"
          , col.sub   = "#666666"
          , cex.sub  = 0.7
          , font.sub = 3)
}

```



Classical breast cancer markers (**GATA3**, **CDH1**, **ESR1**, **PGR**) shows no different expression between Alive and Dead groups (all $p > 0.05$).

K-Means Clustering

```
top50_data <- GeneX_df[, rownames(top_genes)[1:50]]

set.seed(42)
for(k in 2:4) {
  kmeans_result <- kmeans(scale(top50_data), centers = k, nstart = 25)

  cat(sprintf("\n--- K=%d CLUSTERS ---\n", k))

  # Cluster vs survival
  cluster_table <- table(kmeans_result$cluster, clinical_df$vital_status)
  print(cluster_table)

  # Chi-square test
  chi_test <- chisq.test(cluster_table)
  cat(sprintf("Chi-square p=% .4e %s\n"
             , chi_test$p.value
             , ifelse(chi_test$p.value < 0.05, "*** SIGNIFICANT", "")))

  # Cluster sizes
  cat("Cluster sizes:", table(kmeans_result$cluster), "\n")
}

##
```

```

## --- K=2 CLUSTERS ---
##
##      Alive Dead
##    1    852  138
##    2    177   63
## Chi-square p=5.8918e-06 *** SIGNIFICANT
## Cluster sizes: 990 240
##
## --- K=3 CLUSTERS ---
##
##      Alive Dead
##    1    297   38
##    2    683  117
##    3    49   46
## Chi-square p=5.8567e-18 *** SIGNIFICANT
## Cluster sizes: 335 800 95
##
## --- K=4 CLUSTERS ---
##
##      Alive Dead
##    1    473   69
##    2    357   65
##    3    163   26
##    4    36   41
## Chi-square p=6.7112e-18 *** SIGNIFICANT
## Cluster sizes: 542 422 189 77

kmeans_2 <- kmeans(scale(top50_data), centers = 2, nstart = 25)
clinical_df$cluster <- paste0("C", kmeans_2$cluster)
print(table(clinical_df$cluster, clinical_df$vital_status))

```

```

##
##      Alive Dead
##    C1    177   63
##    C2    852  138

```

PCA Visualization by Clusters

```

# PCA
pca_result <- prcomp(top50_data, scale. = TRUE)

# Variance explained
var_exp <- summary(pca_result)$importance[2, 1:10]
cat("Variance explained by first 10 PCs:\n")

## Variance explained by first 10 PCs:

print(round(var_exp, 3))

##   PC1   PC2   PC3   PC4   PC5   PC6   PC7   PC8   PC9   PC10
## 0.547 0.052 0.035 0.030 0.025 0.019 0.017 0.016 0.015 0.013

```

```

cat("\nCumulative variance (PC1-10):", round(sum(var_exp), 3), "\n\n")

##  

## Cumulative variance (PC1-10): 0.768

par(mfrow = c(1, 2), bg = "white")

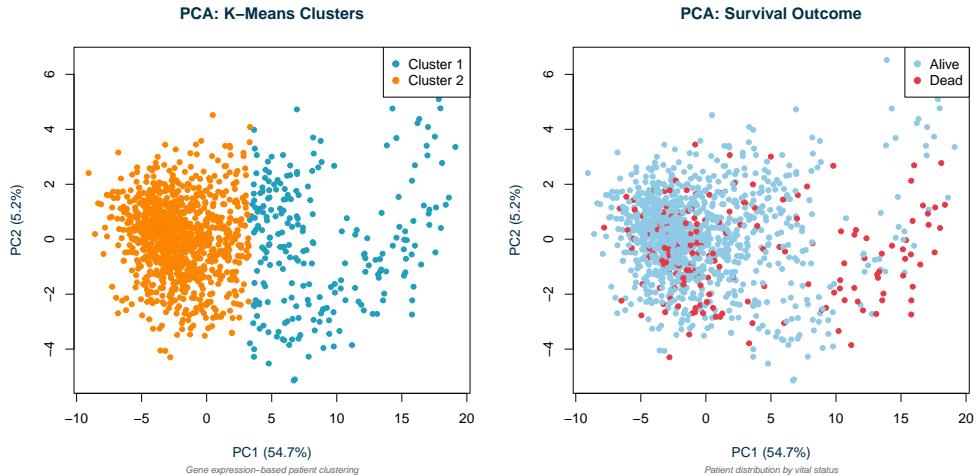
# Plot 1: Color by cluster
plot(pca_result$x[, 1]
      , pca_result$x[, 2]
      , col      = ifelse(clinical_df$cluster == "C1", "#219ebc", "#fb8500")
      , pch      = 19
      , cex      = 0.8
      , xlab     = paste0("PC1 (", round(var_exp[1] * 100, 1), "%)")
      , ylab     = paste0("PC2 (", round(var_exp[2] * 100, 1), "%)")
      , main     = "PCA: K-Means Clusters"
      , sub      = "Gene expression-based patient clustering"
      , col.main = "#023047"
      , col.lab  = "#023047"
      , col.sub  = "#666666"
      , cex.sub  = 0.7
      , font.sub = 3)

legend("topright"
       , legend = c("Cluster 1", "Cluster 2")
       , col    = c("#219ebc", "#fb8500")
       , pch    = 19)

# Plot 2: Color by survival
plot(pca_result$x[, 1]
      , pca_result$x[, 2]
      , col      = ifelse(clinical_df$vital_status == "Dead", "#e63946", "#8ecae6")
      , pch      = 19
      , cex      = 0.8
      , xlab     = paste0("PC1 (", round(var_exp[1] * 100, 1), "%)")
      , ylab     = paste0("PC2 (", round(var_exp[2] * 100, 1), "%)")
      , main     = "PCA: Survival Outcome"
      , sub      = "Patient distribution by vital status"
      , col.main = "#023047"
      , col.lab  = "#023047"
      , col.sub  = "#666666"
      , cex.sub  = 0.7
      , font.sub = 3)

legend("topright"
       , legend = c("Alive", "Dead")
       , col    = c("#8ecae6", "#e63946")
       , pch    = 19)

```



PCA reveals a strong molecular structure in the dataset, **PCA1** (54.7% of variance) clearly separate two expression defined clusters, confirming that the DE genes capture major biological subtypes. However on the right plot, the PCA space does not separate Alive vs Dead patients, indicating that survival differences are not driven by global gene-expression variation.

Clinical vs Gene Correlations

```
# Numeric clinical variables
numeric_clinical <- c("age_at_index",
                        , "initial_weight",
                        , "days_to_last_follow_up",
                        , "age_at_diagnosis",
                        , "days_to_birth")

# For each clinical variable
for(clin_var in numeric_clinical) {
  cat("----", clin_var, "----\n")

  clin_values <- clinical_df[[clin_var]]

  # Calculate correlations with all 20 genes
  cors <- numeric(20)
  for(i in 1:20) {
    cors[i] <- cor(clin_values, gene_subset[, i], use = "complete.obs")
  }
  names(cors) <- top20_genes

  # Sort and show top 5
  cors_sorted <- sort(abs(cors), decreasing = TRUE)
  cat("Top 5 correlated genes:\n")
  for(i in 1:5) {
    gene <- names(cors_sorted)[i]
    cat(sprintf("  %s: r=% .3f\n", gene, cors[gene]))
  }
  cat("\n")
}
```

```

## --- age_at_index ---
## Top 5 correlated genes:
##   VEGFD: r=-0.135
##   LINC01235: r=-0.133
##   LINC02511: r=-0.100
##   ATF3: r=-0.097
##   PSD2: r=-0.093
##
## --- initial_weight ---
## Top 5 correlated genes:
##   LYVE1: r=0.181
##   FHL1: r=0.139
##   GPX3: r=0.138
##   VEGFD: r=0.134
##   CST1: r=-0.132
##
## --- days_to_last_follow_up ---
## Top 5 correlated genes:
##   ATF3: r=0.130
##   SNORD104: r=-0.103
##   LINC02511: r=0.076
##   VEGFD: r=0.071
##   ADH4: r=0.070
##
## --- age_at_diagnosis ---
## Top 5 correlated genes:
##   VEGFD: r=-0.136
##   LINC01235: r=-0.134
##   LINC02511: r=-0.111
##   ATF3: r=-0.103
##   PSD2: r=-0.095
##
## --- days_to_birth ---
## Top 5 correlated genes:
##   LINC01235: r=0.137
##   VEGFD: r=0.135
##   LINC02511: r=0.105
##   ATF3: r=0.102
##   PSD2: r=0.091

```

From this correlation, we can say that,

- Age-related variables (`age_at_index`, `age_at_diagnosis`, `days_to_birth`) show almost identical correlated genes (**VEGFD**, **LINC01235**, **LINC02511**, **ATF3**, **PSD2**) → expected because these age variables are themselves highly collinear.
- Initial weight shows slightly stronger associations (up to $r = 0.18$), mostly with genes involved in immune/metabolic activity (LYVE1, FHL1, GPX3, VEGFD).
- Follow-up time correlates weakly with stress/response genes (**ATF3**, **LINC02511**), but effect sizes remain very small.

Categorical Clinical vs Genes

```
# Tumor stage vs genes
cat("1. TUMOR STAGE vs GENES\n")

## 1. TUMOR STAGE vs GENES

stage_simple <- substr(clinical_df$ajcc_pathologic_t, 1, 2)
stage_simple[!stage_simple %in% c("T1", "T2", "T3", "T4")] <- NA

cat("Stage distribution:\n")

## Stage distribution:

print(table(stage_simple, useNA = "ifany"))

## stage_simple
##   T1    T2    T3    T4 <NA>
## 293 658 132   41 106

cat("\n")

cat("Testing top 5 genes across stages (ANOVA):\n")

## Testing top 5 genes across stages (ANOVA):

for(i in 1:5) {
  gene      <- top20_genes[i]
  gene_expr <- gene_subset[, i]

  df_test <- data.frame(expr = gene_expr, stage = stage_simple)
  df_test <- df_test[!is.na(df_test$stage), ]

  if(length(unique(df_test$stage)) > 1) {
    aov_result <- aov(expr ~ stage, data = df_test)
    p_val      <- summary(aov_result)[[1]]$`Pr(>F)`[1]
    means      <- tapply(df_test$expr, df_test$stage, mean)

    cat(sprintf("  %s: p=% .4f %s\n"
                , gene
                , p_val
                , ifelse(p_val < 0.05, "*** SIGNIFICANT", "")))
    cat(sprintf("    T1=% .2f, T2=% .2f, T3=% .2f, T4=% .2f\n"
                , means["T1"], means["T2"], means["T3"], means["T4"]))
  }
}

##    LINC01235: p=0.2195
##    T1=7.02, T2=6.98, T3=7.30, T4=7.34
```

```

##    APOB: p=0.0709
##          T1=3.71, T2=3.26, T3=3.50, T4=3.75
##    LYVE1: p=0.2767
##          T1=7.33, T2=7.14, T3=7.32, T4=7.67
##    LINC01497: p=0.0021 *** SIGNIFICANT
##          T1=1.17, T2=1.07, T3=1.67, T4=1.09
##    AC104211.1: p=0.0265 *** SIGNIFICANT
##          T1=3.44, T2=3.27, T3=3.68, T4=3.73

cat("\n")

# Treatment vs genes
cat("2. PRIOR TREATMENT vs GENES\n")

## 2. PRIOR TREATMENT vs GENES

treat <- clinical_df$prior_treatment
cat("Treatment distribution:\n")

## Treatment distribution:

print(table(treat, useNA = "ifany"))

## treat
##      No Not Reported      Yes      <NA>
##     1100           1        85        44

cat("\n")

cat("Testing top 5 genes (t-test: Yes vs No):\n")

## Testing top 5 genes (t-test: Yes vs No):

for(i in 1:5) {
  gene   <- top20_genes[i]
  expr_yes <- gene_subset[treat == "Yes", i]
  expr_no  <- gene_subset[treat == "No", i]

  if(length(expr_yes) > 2 && length(expr_no) > 2) {
    test <- t.test(expr_yes, expr_no)
    cat(sprintf(" %s: Yes=%.2f, No=%.2f, p=%,.4f %s\n"
               , gene
               , mean(expr_yes, na.rm = TRUE)
               , mean(expr_no, na.rm = TRUE)
               , test$p.value
               , ifelse(test$p.value < 0.05, "*** SIGNIFICANT", "")))
  }
}

```

```

##  LINC01235: Yes=7.12, No=7.06, p=0.8278
##  APOB: Yes=3.37, No=3.42, p=0.8685
##  LYVE1: Yes=7.24, No=7.23, p=0.9602
##  LINC01497: Yes=1.16, No=1.16, p=0.9740
##  AC104211.1: Yes=3.41, No=3.36, p=0.8008

cat("\n")

# Race vs genes
cat("3. RACE vs GENES\n")

## 3. RACE vs GENES

race      <- clinical_df$race
race_binary <- ifelse(race == "white", "White",
                      , ifelse(race == "black or african american", "Black", NA))

cat("Testing top 5 genes (White vs Black):\n")

## Testing top 5 genes (White vs Black):

for(i in 1:5) {
  gene      <- top20_genes[i]
  expr_white <- gene_subset[race_binary == "White" & !is.na(race_binary), i]
  expr_black <- gene_subset[race_binary == "Black" & !is.na(race_binary), i]

  if(length(expr_white) > 2 && length(expr_black) > 2) {
    test <- t.test(expr_white, expr_black)
    cat(sprintf(" %s: White=%.2f, Black=%.2f, p=%.4f %s\n"
                , gene
                , mean(expr_white, na.rm = TRUE)
                , mean(expr_black, na.rm = TRUE)
                , test$p.value
                , ifelse(test$p.value < 0.05, "*** SIGNIFICANT", "")))
  }
}

##  LINC01235: White=7.13, Black=7.25, p=0.4371
##  APOB: White=3.78, Black=2.60, p=0.0000 *** SIGNIFICANT
##  LYVE1: White=7.48, Black=6.69, p=0.0000 *** SIGNIFICANT
##  LINC01497: White=1.33, Black=0.87, p=0.0004 *** SIGNIFICANT
##  AC104211.1: White=3.60, Black=2.65, p=0.0000 *** SIGNIFICANT

```

From the result , we can state that

- **Tumor stage:** Almost no gene is stage-dependent; only **LINC01497** and **AC104211.1** show small differences → weak association.
- **Prior treatment:** No gene shows expression changes → **no effect**.
- **Race:** Several genes (**APOB**, **LYVE1**, **LINC01497**, **AC104211.1**) differ between White vs Black patients → likely due to **subtype composition**, not race biology.

Interprete, clinical categorical variables have **minimal influence** on top gene expression patterns, except for race-related subtype differences.

Gene-Gene Correlations

```
# Correlation matrix
gene_cor <- cor(gene_subset)

# Find highly correlated pairs
high_cor <- which(abs(gene_cor) > 0.7 & gene_cor != 1, arr.ind = TRUE)

if(nrow(high_cor) > 0) {
  cat("Highly correlated gene pairs (|r| > 0.7):\n")
  for(i in 1:nrow(high_cor)) {
    if(high_cor[i, 1] < high_cor[i, 2]) {
      gene1 <- rownames(gene_cor)[high_cor[i, 1]]
      gene2 <- colnames(gene_cor)[high_cor[i, 2]]
      r     <- gene_cor[high_cor[i, 1], high_cor[i, 2]]
      cat(sprintf(" %s <-> %s: r=% .3f\n", gene1, gene2, r))
    }
  }
} else {
  cat("No highly correlated pairs (genes are independent)\n")
}

## Highly correlated gene pairs (|r| > 0.7):
##   APOB <-> KLB: r=0.718
##   LYVE1 <-> LINC02511: r=0.706
##   APOB <-> AC007423.1: r=0.702
##   KLB <-> AC007423.1: r=0.759
##   APOB <-> GPX3: r=0.732
##   LYVE1 <-> GPX3: r=0.748
##   KLB <-> GPX3: r=0.726
##   AC007423.1 <-> GPX3: r=0.725
##   APOB <-> LVRN: r=0.717
##   LYVE1 <-> LVRN: r=0.702
##   KLB <-> LVRN: r=0.799
##   AC007423.1 <-> LVRN: r=0.754
##   GPX3 <-> LVRN: r=0.771
##   KLB <-> SLC2A4: r=0.702
##   GPX3 <-> SLC2A4: r=0.707
##   APOB <-> FHL1: r=0.723
##   LYVE1 <-> FHL1: r=0.800
##   KLB <-> FHL1: r=0.724
##   LINC02511 <-> FHL1: r=0.711
##   AC007423.1 <-> FHL1: r=0.724
##   GPX3 <-> FHL1: r=0.807
##   LVRN <-> FHL1: r=0.793
##   SLC2A4 <-> FHL1: r=0.707
##   LYVE1 <-> VEGFD: r=0.741
##   LINC02511 <-> VEGFD: r=0.730
##   GPX3 <-> VEGFD: r=0.729
##   LVRN <-> VEGFD: r=0.712
##   FHL1 <-> VEGFD: r=0.737
```

The top DE genes form one strong co-expression module (**APOB**, **LYVE1**, **KLB**, **GPX3**, **FHL1**,

VEGFD, LINC02511). They show very high correlations ($|r| = 0.70\sim 0.80$), meaning they act as a single biological program.

Outlier Detection on Top Gen

```
top_gene <- top20_genes[1]
gene_expr <- gene_subset[, 1]

# Z-scores
z_scores <- scale(gene_expr)
outliers <- which(abs(z_scores) > 3)

cat("Top gene:", top_gene, "\n")

## Top gene: LINC01235

cat("Samples with |z-score| > 3:", length(outliers), "\n")

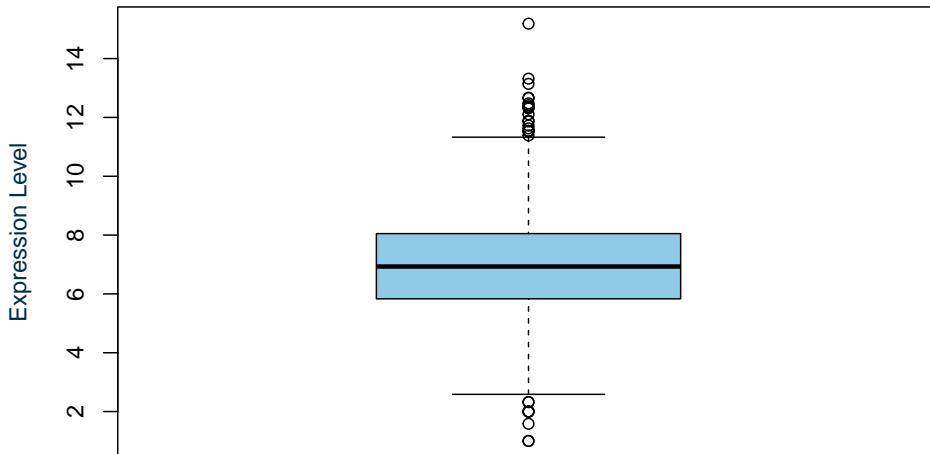
## Samples with |z-score| > 3: 7

if(length(outliers) > 0 & length(outliers) < 10) {
  cat("Outlier samples:\n")
  print(outliers)
}

## Outlier samples:
## [1] 321 367 383 506 528 987 1017

# Boxplot
boxplot(gene_expr
         , col      = "#8ecae6"
         , main    = paste("Outlier Detection:", top_gene)
         , sub     = "Samples with |z-score| > 3 are outliers"
         , xlab    = "Gene"
         , ylab    = "Expression Level"
         , col.main = "#023047"
         , col.lab  = "#023047"
         , col.sub   = "#666666"
         , cex.sub  = 0.7
         , font.sub = 3)
```

Outlier Detection: LINC01235



Gene

Samples with $|z\text{-score}| > 3$ are outliers

```
cat("\nOnly", length(outliers), "outliers (",
  round(length(outliers) / nrow(gene_subset) * 100, 1), "%) - acceptable\n")
```

```
##  
## Only 7 outliers ( 0.6 %) - acceptable
```

The expression profile of LINC01235 shows only 7 extreme observations (0.6 exceeding the standard $|z| > 3$ threshold. This is well within accepted QC limits for large transcriptomic datasets, where up to 1–2% technical or biological outliers are considered normal.

Methodology

Data Preparing before fitting model

```
# --- 1. Define Predictor Names ---
clinical_col_num_names <- c("age_at_index"
                           , "initial_weight"
                           , "days_to_last_follow_up")

clinical_signi_obj_names <- c("tissue_type"
                             , "ajcc_pathologic_t"
                             , "classification_of_tumor"
                             , "follow_ups_disease_response"
                             , "prior_treatment"
                             , "tissue_or_organ_of_origin"
                             , "ethnicity")
```

```

# ALL GENES (5000)
top_genes_list <- colnames(GeneX_df)

cat("== SELECTED VARIABLES ==\n")

## == SELECTED VARIABLES ==

cat("Numeric clinical:", length(clinical_col_num_names), "\n")

## Numeric clinical: 3

cat("Categorical clinical:", length(clinical_signi_obj_names), "\n")

## Categorical clinical: 7

cat("Genes: ALL", length(top_genes_list), "\n\n")

## Genes: ALL 5000

# --- 2. Check Missing Values ---
cat("== MISSING VALUES ==\n")

## == MISSING VALUES ==

cat("Numeric:\n")

## Numeric:

print(colSums(is.na(clinical_df[, clinical_col_num_names])))

##           age_at_index      initial_weight days_to_last_follow_up
##                   0                  15                      3

cat("\nCategorical:\n")

##

## Categorical:

print(colSums(is.na(clinical_df[, clinical_signi_obj_names])))

##           tissue_type      ajcc_pathologic_t
##                   0                  99
## classification_of_tumor follow_ups_disease_response
##                         0                     76
##           prior_treatment tissue_or_organ_of_origin
##                           44                      0
##           ethnicity
##                   0

```

```

cat("\n")

# --- 3. Impute Numeric Variables (median) ---
clinical_numeric <- clinical_df[, clinical_col_num_names, drop = FALSE]

for (col in colnames(clinical_numeric)) {
  n_missing <- sum(is.na(clinical_numeric[[col]]))
  if (n_missing > 0) {
    median_val <- median(clinical_numeric[[col]], na.rm = TRUE)
    clinical_numeric[[col]][is.na(clinical_numeric[[col]])] <- median_val
    cat(sprintf("Imputed %d in %s (median=% .2f)\n", n_missing, col, median_val))
  }
}

## Imputed 15 in initial_weight (median=220.00)
## Imputed 3 in days_to_last_follow_up (median=890.00)

cat("\n")

clinical_numeric <- data.frame(lapply(clinical_numeric, as.numeric))
rownames(clinical_numeric) <- rownames(clinical_df)

# --- 4. Impute Categorical Variables (mode) ---
clinical_categorical <- clinical_df[, clinical_signi_obj_names, drop = FALSE]

for (col in colnames(clinical_categorical)) {
  n_missing <- sum(is.na(clinical_categorical[[col]]))
  if (n_missing > 0) {
    mode_val <- names(sort(table(clinical_categorical[[col]]), decreasing = TRUE))[1]
    clinical_categorical[[col]][is.na(clinical_categorical[[col]])] <- mode_val
    cat(sprintf("Imputed %d in %s (mode=%s)\n", n_missing, col, mode_val))
  }
}

## Imputed 99 in ajcc_pathologic_t (mode=T2)
## Imputed 76 in follow_ups_disease_response (mode=TF-Tumor Free)
## Imputed 44 in prior_treatment (mode=No)

cat("\n")

clinical_categorical <- data.frame(lapply(clinical_categorical, as.factor))

# --- 5. One-Hot Encode Categorical ---
valid_factors <- sapply(clinical_categorical, function(x) {
  n_levels <- length(levels(droplevels(x)))
  return(n_levels >= 2)
})

clinical_categorical_valid <- clinical_categorical[, valid_factors, drop = FALSE]

```

```

if (sum(!valid_factors) > 0) {
  cat("Dropped constant columns:",
      paste(names(clinical_categorical)[!valid_factors], collapse=", "), "\n\n")
}

clinical_ohe <- model.matrix(~ . - 1, data = clinical_categorical_valid)
clinical_ohe <- as.data.frame(clinical_ohe)
rownames(clinical_ohe) <- rownames(clinical_df)

cat("One-Hot Encoding: ", ncol(clinical_categorical_valid), "->", ncol(clinical_ohe), "\n\n")

## One-Hot Encoding: 7 -> 54

# --- 6. Gene Expression Quality Check ---
cat("== GENE EXPRESSION QUALITY CHECK ==\n")

## == GENE EXPRESSION QUALITY CHECK ==

gene_data <- GeneX_df[, top_genes_list, drop = FALSE]
gene_data <- as.data.frame(gene_data)

cat("Missing Values:\n")

## Missing Values:

missing_per_gene <- colSums(is.na(gene_data))
missing_pct <- 100 * missing_per_gene / nrow(gene_data)

cat("  Genes with missing:", sum(missing_per_gene > 0), "/", ncol(gene_data), "\n")

##  Genes with missing: 0 / 5000

if(sum(missing_per_gene > 0) > 0) {
  missing_summary <- data.frame(
    Gene      = names(missing_per_gene),
    N_Missing = missing_per_gene,
    Pct_Missing = round(missing_pct, 2)
  )
  print(head(missing_summary[order(-missing_summary$N_Missing), ], 10))
}
cat("\n")

# --- 7. Distribution Analysis ---
cat("Distribution Analysis:\n")

## Distribution Analysis:

```

```

sample_genes <- colnames(gene_data)[1:min(20, ncol(gene_data))]
distribution_summary <- data.frame(
  Gene = character()
  , Mean = numeric()
  , Median = numeric()
  , Skewness = numeric()
  , Impute_Method = character()
  , stringsAsFactors = FALSE
)

for(gene in sample_genes) {
  vals <- gene_data[[gene]][!is.na(gene_data[[gene]])]
  m <- mean(vals)
  s <- sd(vals)
  skew <- mean(((vals - m) / s)^3)

  distribution_summary <- rbind(distribution_summary
    , data.frame(
      Gene = gene
      , Mean = round(m, 2)
      , Median = round(median(vals), 2)
      , Skewness = round(skew, 3)
      , Impute_Method = ifelse(abs(skew) < 0.5, "Mean", "Median")
    ))
}

print(distribution_summary)

##          Gene   Mean  Median Skewness Impute_Method
## 1      CLEC3A  5.93   4.52   0.641      Median
## 2     SCGB2A2 11.13  11.71  -0.338      Mean
## 3       CPB1   8.31   7.68   0.715      Median
## 4      GSTM1   5.36   5.49   0.160      Mean
## 5       TFF1   9.71  10.79  -0.594      Median
## 6     SCGB1D2   9.03   9.27  -0.046      Mean
## 7      KCNJ3   6.69   5.93   0.347      Mean
## 8      MUCL1   9.15   8.95   0.134      Mean
## 9  LINC00993   8.86  10.50  -0.816      Median
## 10     ANKRD30A   8.65  10.33  -0.756      Median
## 11    DSCAM-AS1   4.88   4.28   0.380      Mean
## 12      S100A7   4.77   3.58   0.914      Median
## 13        PIP  10.05  10.67  -0.309      Mean
## 14    SERPINA6   4.98   3.91   0.742      Median
## 15 AC093001.1   7.54   7.38  -0.037      Mean
## 16      VSTM2A   4.09   2.32   0.951      Median
## 17     ADIPOQ   7.87   8.36  -0.145      Mean
## 18     HMGCS2   7.70   8.00  -0.109      Mean
## 19      ADH1B   8.81   9.36  -0.301      Mean
## 20      PRAME   5.96   4.58   0.491      Mean

cat("\n")

```

```

# --- 8. Determine Imputation Strategy ---
cat("Imputation Strategy:\n")

## Imputation Strategy:

all_skewness <- sapply(1:ncol(gene_data), function(i) {
  vals <- gene_data[!is.na(gene_data[, i])], i]
  if(length(vals) < 3) return(0)
  m <- mean(vals)
  s <- sd(vals)
  if(s == 0) return(0)
  mean(((vals - m) / s)^3)
})

median_skew <- median(abs(all_skewness), na.rm = TRUE)
cat(" Median absolute skewness:", round(median_skew, 3), "\n")

## Median absolute skewness: 0.471

impute_strategy <- ifelse(median_skew < 0.5, "mean", "median")
cat(" Selected method:", toupper(impute_strategy), "\n\n")

## Selected method: MEAN

# --- 9. Apply Gene Imputation ---
cat("Applying Gene Imputation:\n")

## Applying Gene Imputation:

gene_data_imputed <- gene_data
n_imputed <- 0

for(gene in colnames(gene_data_imputed)) {
  n_missing <- sum(is.na(gene_data_imputed[[gene]]))
  if(n_missing > 0) {
    vals <- gene_data_imputed[[gene]]
    impute_val <- if(impute_strategy == "median") {
      median(vals, na.rm = TRUE)
    } else {
      mean(vals, na.rm = TRUE)
    }
    gene_data_imputed[[gene]][is.na(gene_data_imputed[[gene]])] <- impute_val
    n_imputed <- n_imputed + n_missing
  }
}

cat(" Values imputed:", n_imputed, "\n")

## Values imputed: 0

```

```

cat("  Method used:", toupper(impute_strategy), "\n")

##  Method used: MEAN

cat("  Remaining NA:", sum(is.na(gene_data_imputed)), "\n\n")

##  Remaining NA: 0

# --- 10. Visual Distribution Check ---
par(mfrow = c(2, 3), mar = c(4, 4, 3, 1))

for(i in 1:min(6, ncol(gene_data))) {
  vals <- gene_data[[i]][!is.na(gene_data[[i]])]
  m <- mean(vals)
  s <- sd(vals)
  skew <- mean(((vals - m) / s)^3)

  hist(vals
        , breaks = 30
        , col = "#8ecae6"
        , border = "white"
        , main = colnames(gene_data)[i]
        , sub = paste("Skewness:", round(skew, 2))
        , xlab = "Expression"
        , ylab = "Frequency"
        , col.main = "#023047"
        , col.sub = "#666666")

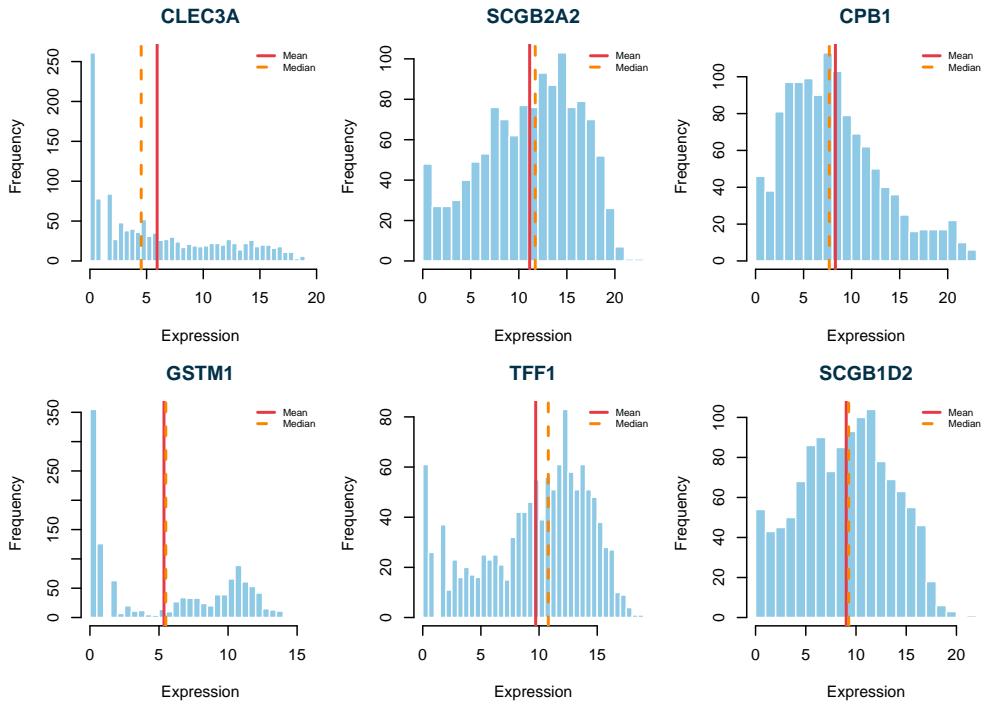
  abline(v = m
          , col = "#e63946"
          , lwd = 2)

  abline(v = median(vals)
          , col = "#fb8500"
          , lwd = 2
          , lty = 2)

  legend("topright"
        , legend = c("Mean", "Median")
        , col = c("#e63946", "#fb8500")
        , lwd = 2
        , lty = c(1, 2)
        , cex = 0.6
        , bty = "n")
}

}

```



```

par(mfrow = c(1, 1))

gene_data <- gene_data_imputed

cat("Gene Expression Quality Check Complete\n")

## Gene Expression Quality Check Complete

cat("Dimensions:", nrow(gene_data), "x", ncol(gene_data), "\n\n")

## Dimensions: 1230 x 5000

# --- 11. Standardize ---
cat("== STANDARDIZATION ==\n\n")

## == STANDARDIZATION ==

clinical_numeric_scaled <- as.data.frame(scale(clinical_numeric))

cat("Numeric clinical:\n")

## Numeric clinical:

for (var in colnames(clinical_numeric_scaled)) {
  cat(sprintf(" %s: mean=%.4f, sd=%.4f\n"
             , var
             , mean(clinical_numeric_scaled[[var]])
             , sd(clinical_numeric_scaled[[var]])))
}

```

```

##    age_at_index: mean=-0.0000, sd=1.0000
##    initial_weight: mean=-0.0000, sd=1.0000
##    days_to_last_follow_up: mean=0.0000, sd=1.0000

cat("\n")

gene_data_scaled <- as.data.frame(scale(gene_data))

cat("Genes (", ncol(gene_data_scaled), "):\n", sep="")

## Genes (5000):

cat(sprintf("  Mean: %.2e, SD: %.4f\n"
            , mean(as.matrix(gene_data_scaled))
            , sd(as.matrix(gene_data_scaled)))))

##    Mean: 2.99e-18, SD: 0.9996

cat(sprintf("  Range: [% .2f, % .2f]\n\n"
            , min(gene_data_scaled)
            , max(gene_data_scaled)))

##    Range: [-8.57, 11.20]

# --- 12. Combine Features ---
data_predictors <- cbind(clinical_numeric_scaled
                         , clinical_ohe
                         , gene_data_scaled)

data_predictors$Y <- ifelse(clinical_df$vital_status == "Dead", 1, 0)

cat("== FINAL DATASET ==\n")

## == FINAL DATASET ==

cat("Samples:", nrow(data_predictors), "\n")

## Samples: 1230

cat("Features:", ncol(data_predictors) - 1, "\n")

## Features: 5057

cat("  Numeric clinical:", ncol(clinical_numeric_scaled), "\n")

## Numeric clinical: 3

```

```

cat(" Categorical (OHE):", ncol(clinical_ohe), "\n")

## Categorical (OHE): 54

cat(" Genes:", ncol(gene_data_scaled), "\n\n")

## Genes: 5000

table_Y <- table(data_predictors$Y)
print(table_Y)

##
##      0      1
## 1029  201

cat(sprintf(" Alive: %d (%.1f%%)\n", table_Y[1], 100 * table_Y[1] / sum(table_Y)))

## Alive: 1029 (83.7%)

cat(sprintf(" Dead: %d (%.1f%%)\n", table_Y[2], 100 * table_Y[2] / sum(table_Y)))

## Dead: 201 (16.3%)

cat(sprintf(" Imbalance: %.2f:1\n", table_Y[1] / table_Y[2]))

## Imbalance: 5.12:1

```

Train/Test Split

```

# Set seed for reproducibility
set.seed(42)

# Separate features and target
X_all <- data_predictors[, -which(names(data_predictors) == "Y")]
Y_all <- data_predictors$Y

# Create train/test split (80/20)
train_indices <- sample(1:nrow(data_predictors), size = 0.8 * nrow(data_predictors))

X_train <- as.matrix(X_all[train_indices, ])
X_test <- as.matrix(X_all[-train_indices, ])
Y_train <- Y_all[train_indices]
Y_test <- Y_all[-train_indices]

# Calculate number of clinical features
n_clinical <- ncol(clinical_numeric_scaled) + ncol(clinical_ohe)

cat("== TRAIN/TEST SPLIT ==\n")

```

```

## === TRAIN/TEST SPLIT ===

cat("Training set:\n")

## Training set:

cat("  Samples:", nrow(X_train), "\n")

##   Samples: 984

cat("  Features:", ncol(X_train), "\n")

##   Features: 5057

cat("  Dead:", sum(Y_train == 1), sprintf("(%.1f%%)\n", 100 * sum(Y_train == 1) / length(Y_train)))

##   Dead: 161 (16.4%)

cat("  Alive:", sum(Y_train == 0), sprintf("(%.1f%%)\n", 100 * sum(Y_train == 0) / length(Y_train)))

##   Alive: 823 (83.6%)

cat("\n")

cat("Test set:\n")

## Test set:

cat("  Samples:", nrow(X_test), "\n")

##   Samples: 246

cat("  Features:", ncol(X_test), "\n")

##   Features: 5057

cat("  Dead:", sum(Y_test == 1), sprintf("(%.1f%%)\n", 100 * sum(Y_test == 1) / length(Y_test)))

##   Dead: 40 (16.3%)

cat("  Alive:", sum(Y_test == 0), sprintf("(%.1f%%)\n", 100 * sum(Y_test == 0) / length(Y_test)))

##   Alive: 206 (83.7%)

```

```

cat("\n")

cat("Clinical features:", n_clinical, "\n")

## Clinical features: 57

cat("Genomic features:", ncol(X_train) - n_clinical, "\n")

## Genomic features: 5000

```

Logistic regression

Logistic regression is used here only on feature sets that are not high-dimensional, because the model becomes unstable when the number of predictors is large. This is due to the form of its loss function, which cannot be minimized reliably when $p >> n$. Recall that logistic regression estimates coefficients by maximizing the log-likelihood:

```

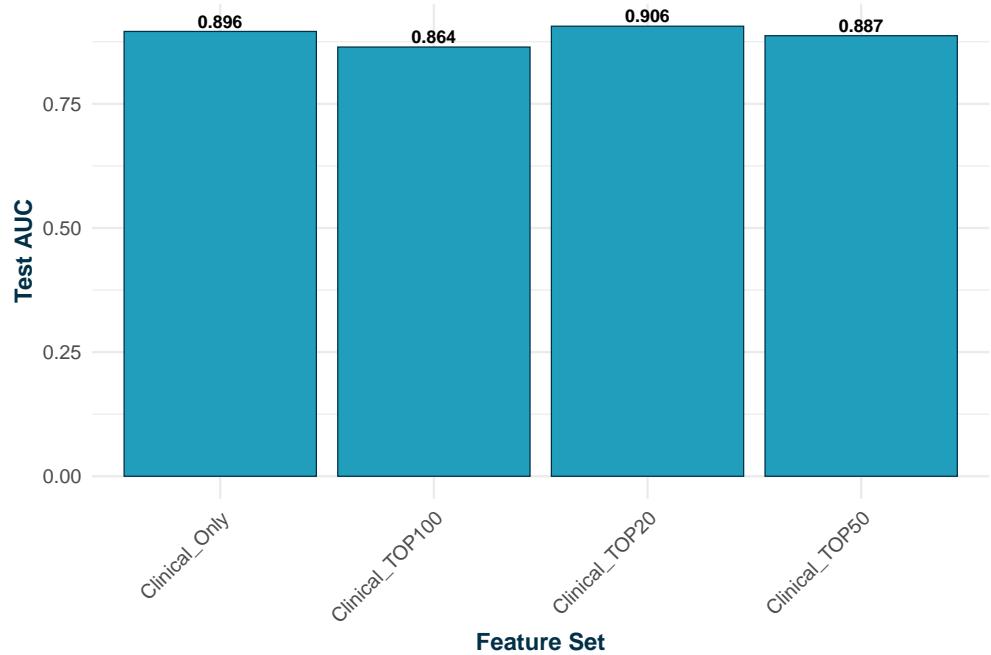
logistic_results <- fit_single_model_across_features(
  model_type = "logistic"
  , X_train_all = X_train
  , X_test_all = X_test
  , Y_train = Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(100, 50, 20)
)

## 
## === FITTING LOGISTIC ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Fitting Clinical_TOP100...
## Fitting Clinical_TOP50...
## Fitting Clinical_TOP20...

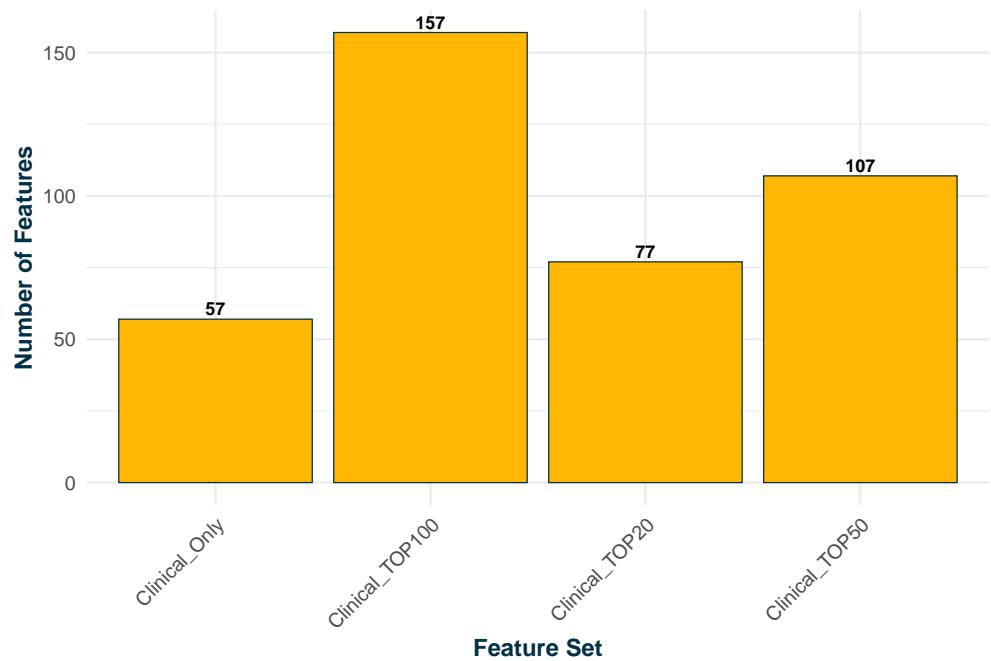
##
## === SUMMARY TABLE ===
##      Feature_Set    Model Features Train_AUC  Test_AUC Test_Accuracy
## 1   Clinical_Only LOGISTIC      57 0.9141680 0.8957524 0.9024390
## 2 Clinical_TOP100 LOGISTIC     157 0.9681366 0.8643204 0.8373984
## 3 Clinical_TOP50 LOGISTIC     107 0.9426353 0.8871359 0.9024390
## 4 Clinical_TOP20 LOGISTIC      77 0.9238206 0.9063107 0.9065041
## Exported metrics to: model_metrics/logistic_across_features_metrics.csv

```

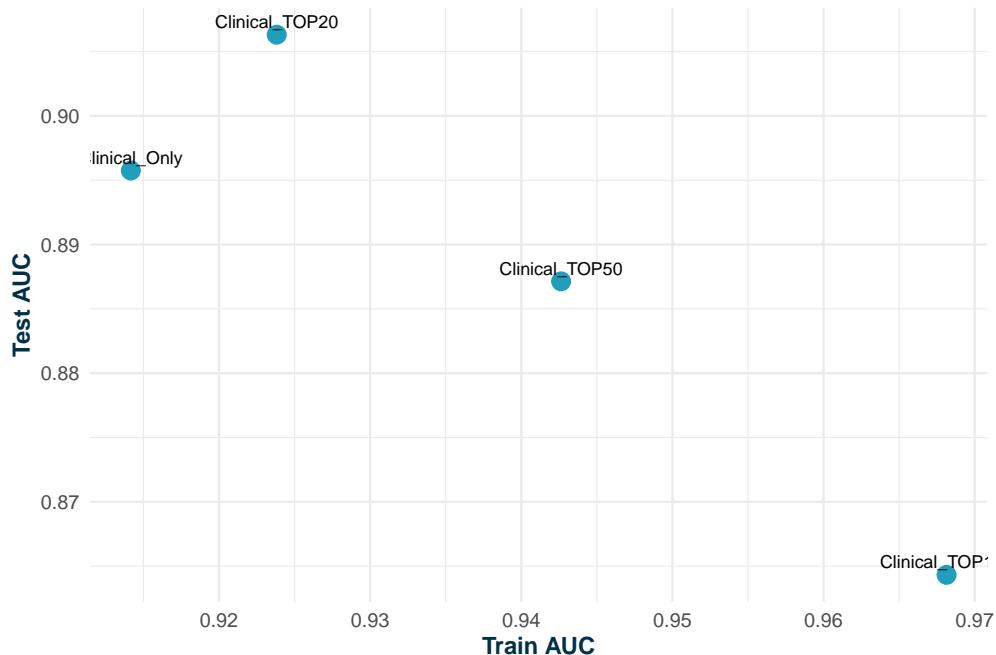
LOGISTIC – Performance Across Feature Sets



LOGISTIC – Selected Features

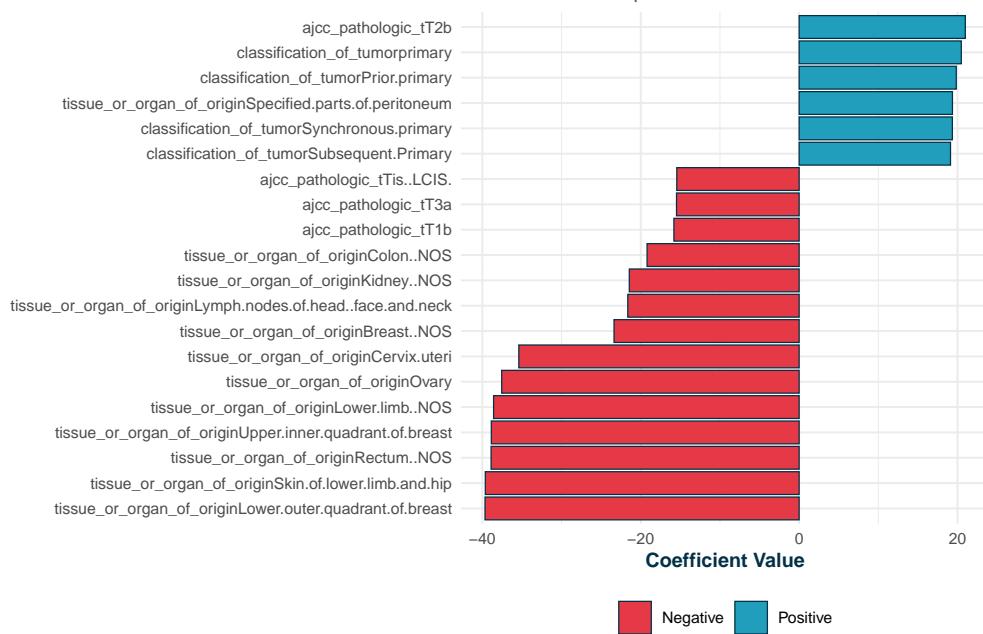


LOGISTIC – Train vs Test AUC



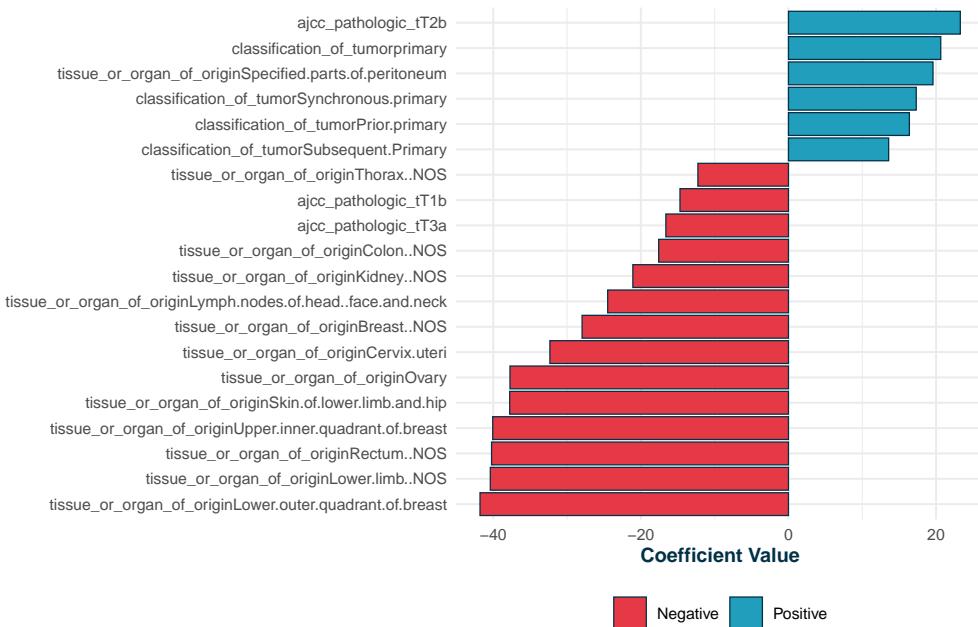
LOGISTIC – Clinical_Only

Top 20 non-zero features



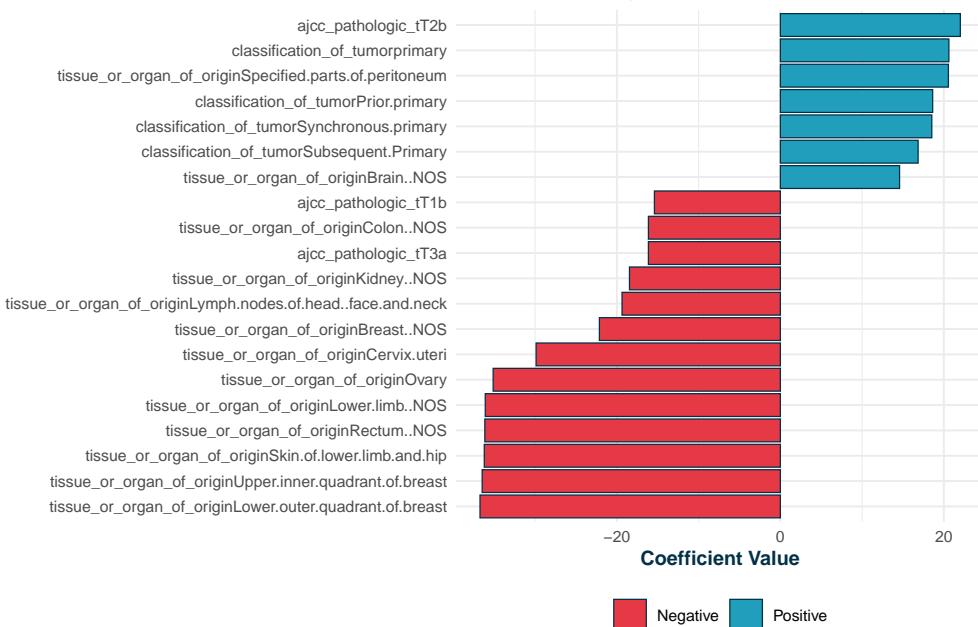
LOGISTIC – Clinical_TOP100

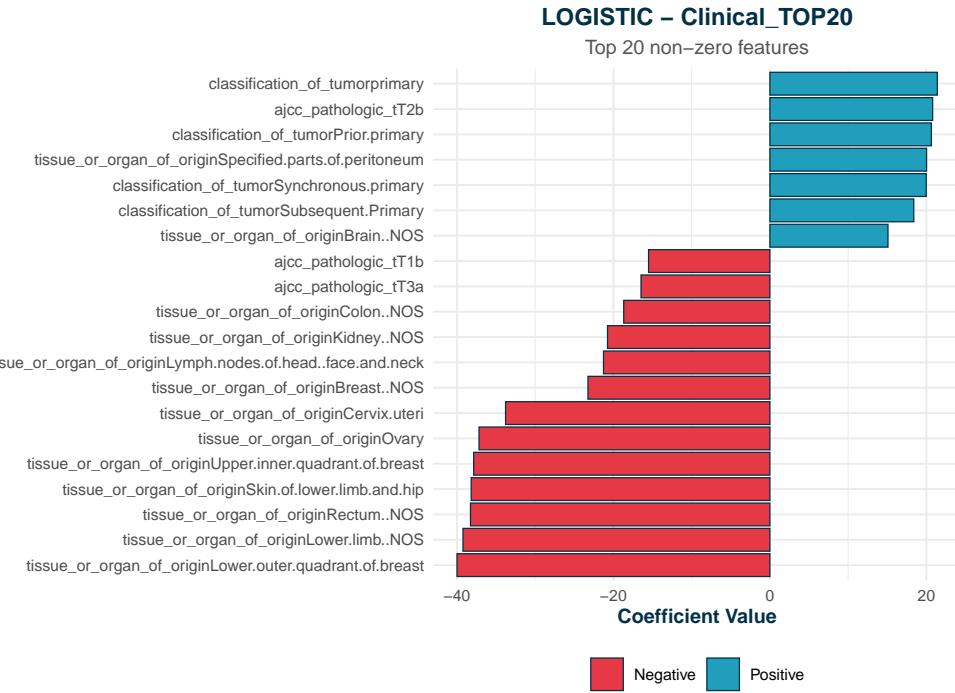
Top 20 non-zero features



LOGISTIC – Clinical_TOP50

Top 20 non-zero features





```
logistic_metrics <- plot_classification_metrics_single(logistic_results
, threshold = 0.5
, csv_filename = "logistic_classification_metrics.csv")
```

```
##
```

```
## === CLASSIFICATION METRICS ===
```

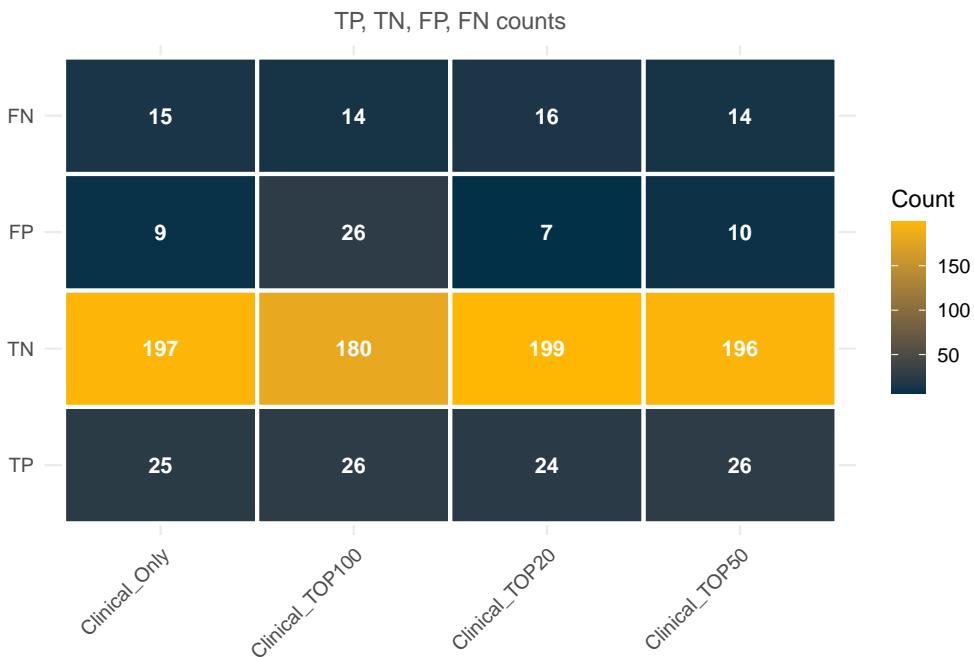
```
## Clinical_Only:
##   TP=25 TN=197 FP=9 FN=15
##   Accuracy=0.902 Precision=0.735 Recall=0.625 F1=0.676 AUC=0.896
```

```
## Clinical_TOP100:
##   TP=26 TN=180 FP=26 FN=14
##   Accuracy=0.837 Precision=0.500 Recall=0.650 F1=0.565 AUC=0.864
```

```
## Clinical_TOP50:
##   TP=26 TN=196 FP=10 FN=14
##   Accuracy=0.902 Precision=0.722 Recall=0.650 F1=0.684 AUC=0.887
```

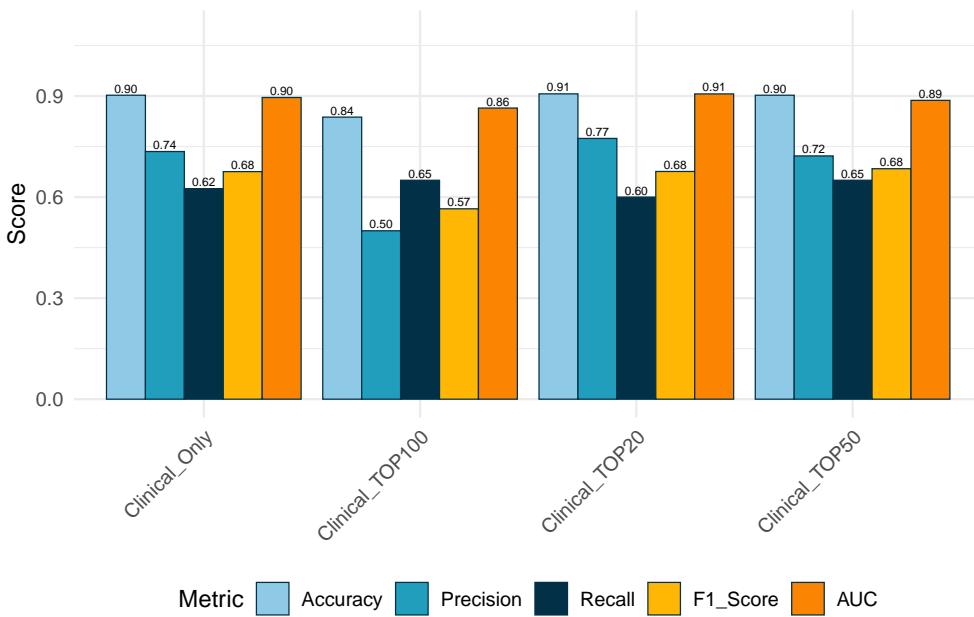
```
## Clinical_TOP20:
##   TP=24 TN=199 FP=7 FN=16
##   Accuracy=0.907 Precision=0.774 Recall=0.600 F1=0.676 AUC=0.906
```

LOGISTIC – Confusion Matrix Across Feature Sets



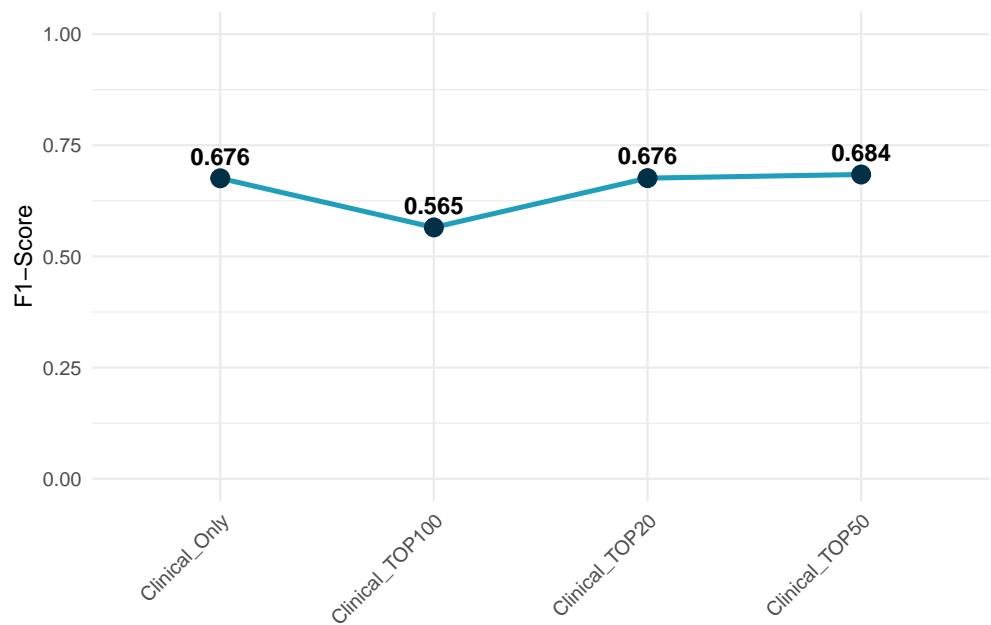
LOGISTIC – Classification Metrics

Accuracy, Precision, Recall, F1-Score, AUC



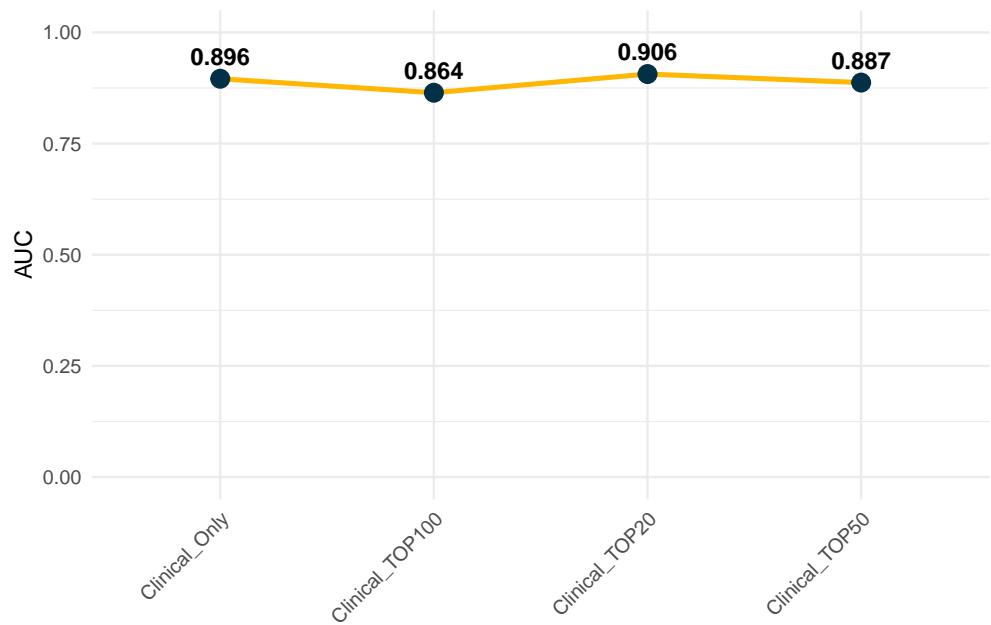
LOGISTIC – F1-Score Across Feature Sets

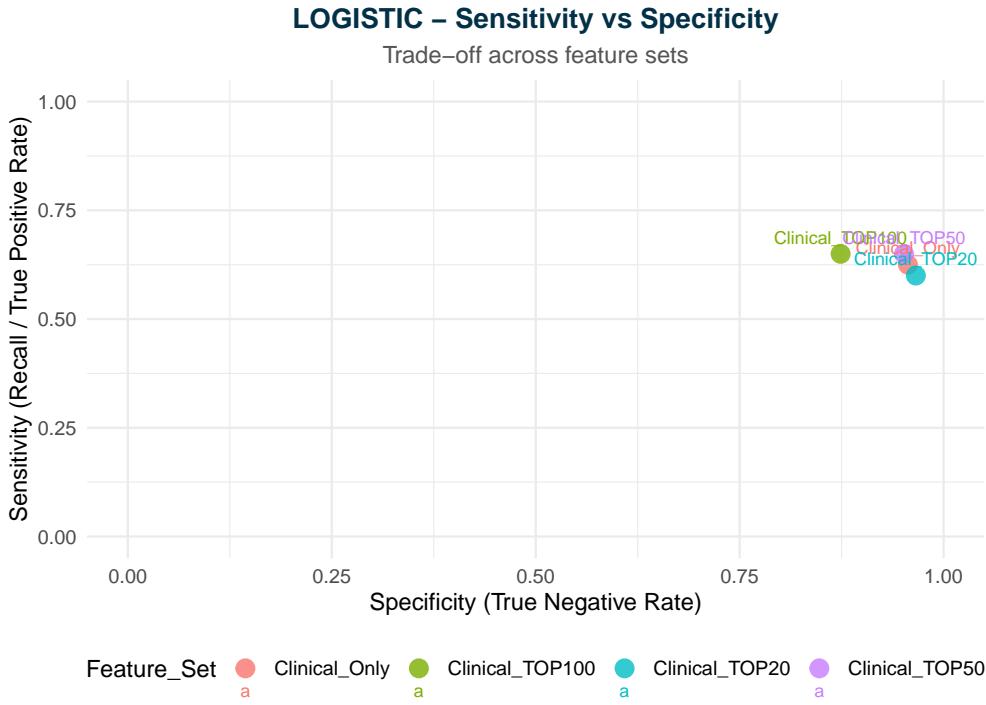
Trend of model performance



LOGISTIC – AUC Across Feature Sets

Area Under the ROC Curve





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity F1_Score
## 1  Clinical_Only 25 197 9 15 0.9024390 0.7352941 0.625 0.9563107 0.6756757
## 2 Clinical_TOP100 26 180 26 14 0.8373984 0.5000000 0.650 0.8737864 0.5652174
## 3 Clinical_TOP50 26 196 10 14 0.9024390 0.7222222 0.650 0.9514563 0.6842105
## 4 Clinical_TOP20 24 199 7 16 0.9065041 0.7741935 0.600 0.9660194 0.6760563
##      AUC
## 1 0.8957524
## 2 0.8643204
## 3 0.8871359
## 4 0.9063107
##
## Exported classification metrics to: model_metrics/logistic_classification_metrics.csv
```

Logistic regression shows consistently high specificity but low recall, indicating that it classifies Alive patients reliably but struggles to detect Dead cases. The clinical-only model performs best overall, while adding genomic features does not meaningfully improve recall and often reduces generalization, especially with 100 genes. These metrics reinforce the conclusion that logistic regression cannot effectively exploit high-dimensional gene expression and is best used as a baseline on small feature sets.

Ridge Regression across different feature

Ridge regression stabilizes estimation in the presence of strong correlations between genes, but does not perform variable selection.

By adding an L2 penalty,

$$\hat{\beta}^{\text{ridge}} = \operatorname{argmin}_{\beta} \{-l(\beta) + \lambda \|\beta\|_2^2\}$$

the model remains stable even when thousands of genes are included. Therefore, Ridge can handle large feature sets, and we apply it on 5000, 1000, 500, 100, 50, and 20 top genes to evaluate its performance at different dimensionalities.

```
ridge_results <- fit_single_model_across_features(
  model_type = "ridge"
  , X_train_all = X_train
  , X_test_all = X_test
  , Y_train = Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

## 
## === FITTING RIDGE ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP5000...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP1000...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP500...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```

```

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP100...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP50...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP20...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

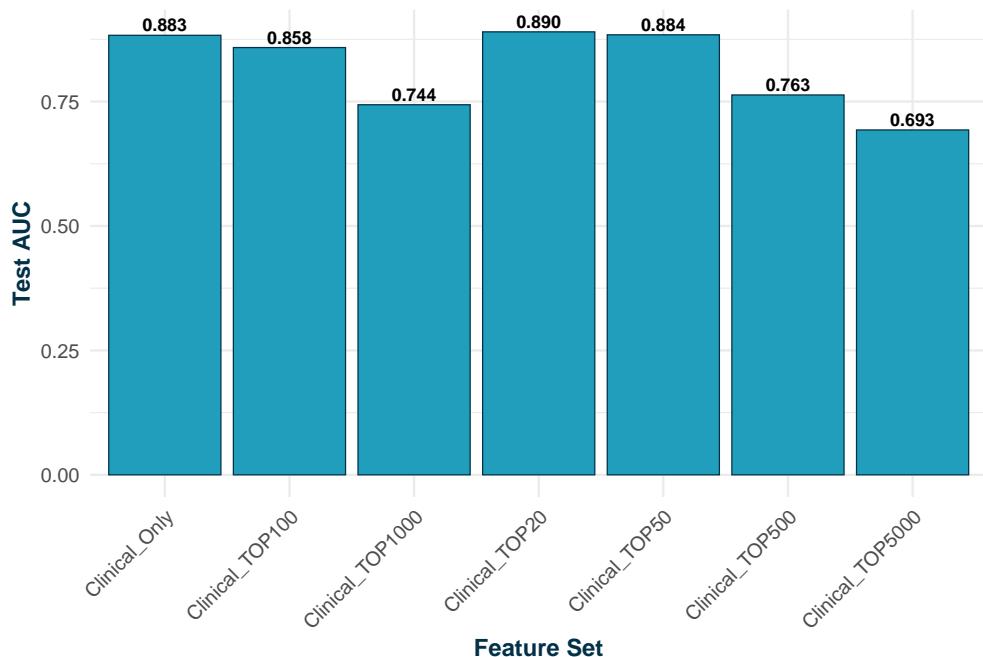
## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

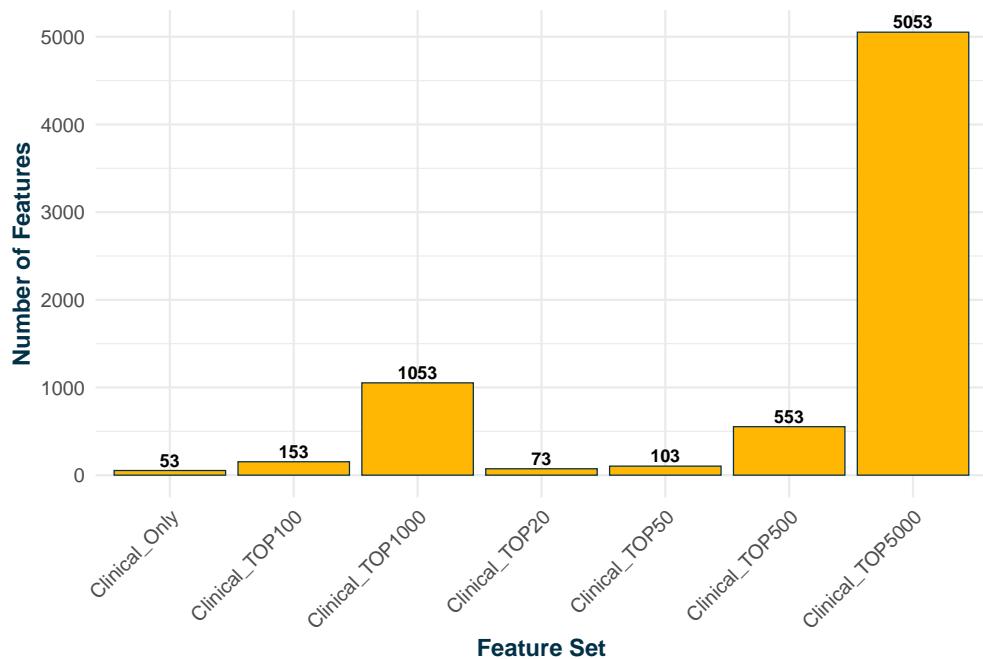
## 
## === SUMMARY TABLE ===
##      Feature_Set Model Features Train_AUC Test_AUC Test_Accuracy
## 1    Clinical_Only RIDGE      53 0.8952854 0.8831311 0.8577236
## 2 Clinical_TOP5000 RIDGE     5053 0.8828706 0.6929612 0.8373984
## 3 Clinical_TOP1000 RIDGE     1053 0.8449016 0.7435680 0.8373984
## 4 Clinical_TOP500 RIDGE      553 0.9111945 0.7632282 0.8536585
## 5 Clinical_TOP100 RIDGE      153 0.9201754 0.8584951 0.8861789
## 6 Clinical_TOP50 RIDGE       103 0.9072021 0.8841019 0.8902439
## 7 Clinical_TOP20 RIDGE        73 0.8932175 0.8900485 0.8943089
## Exported metrics to: model_metrics/ridge_across_features_metrics.csv

```

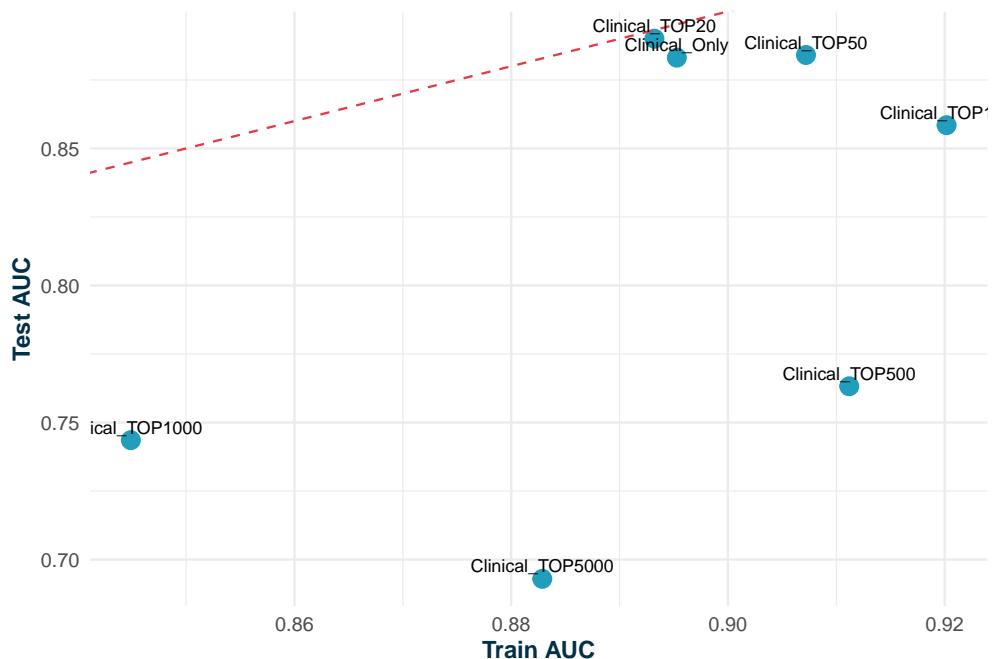
RIDGE – Performance Across Feature Sets



RIDGE – Selected Features

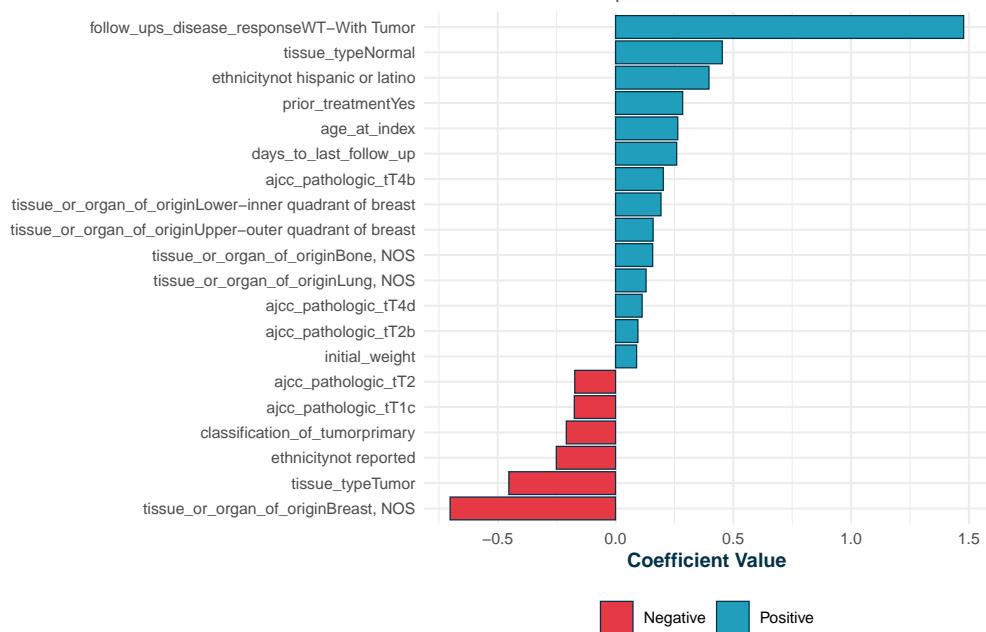


RIDGE – Train vs Test AUC



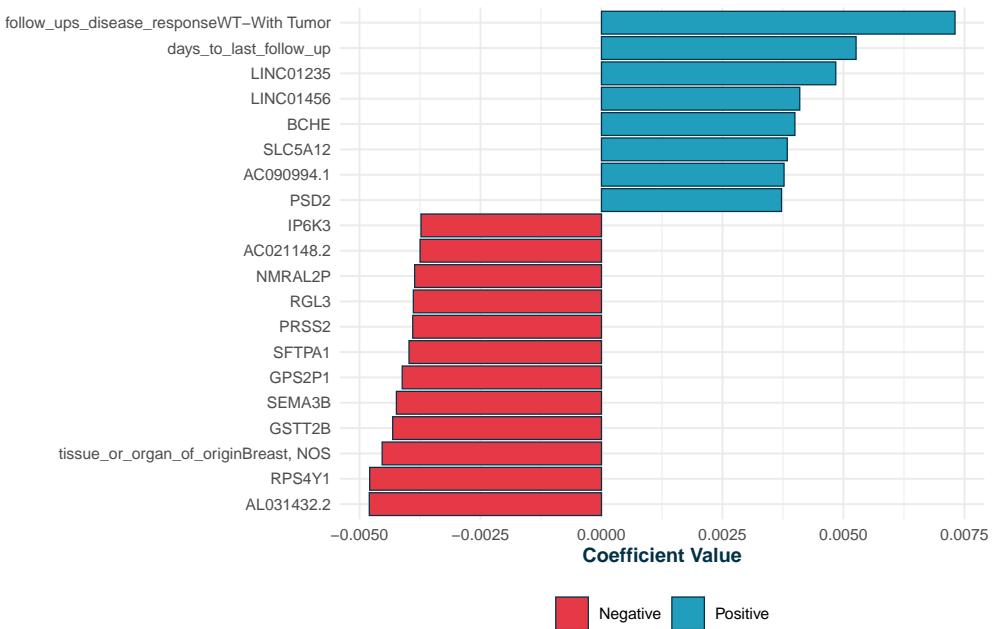
RIDGE – Clinical_Only

Top 20 non-zero features



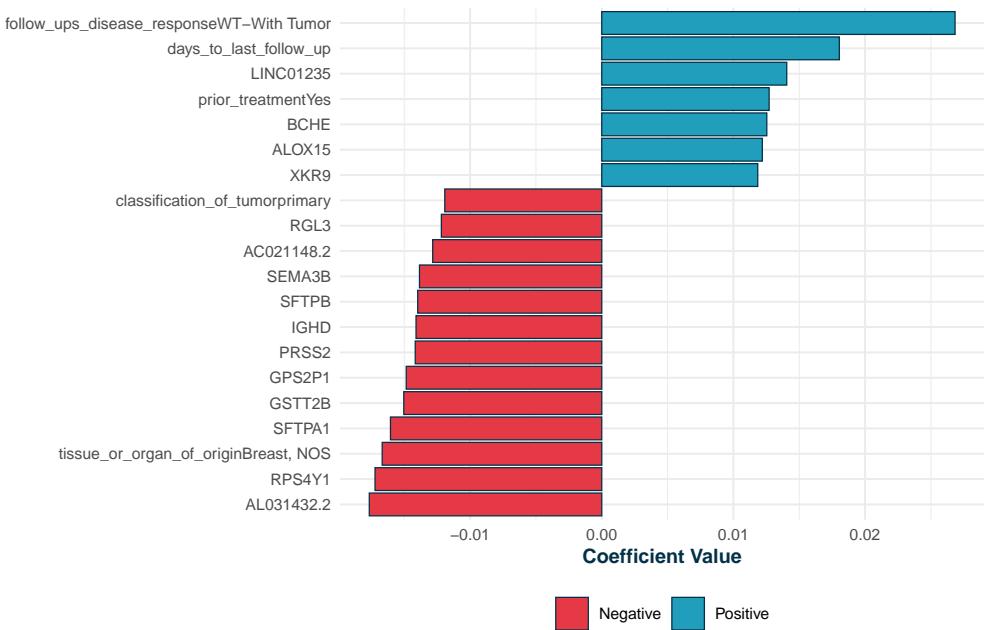
RIDGE – Clinical_TOP5000

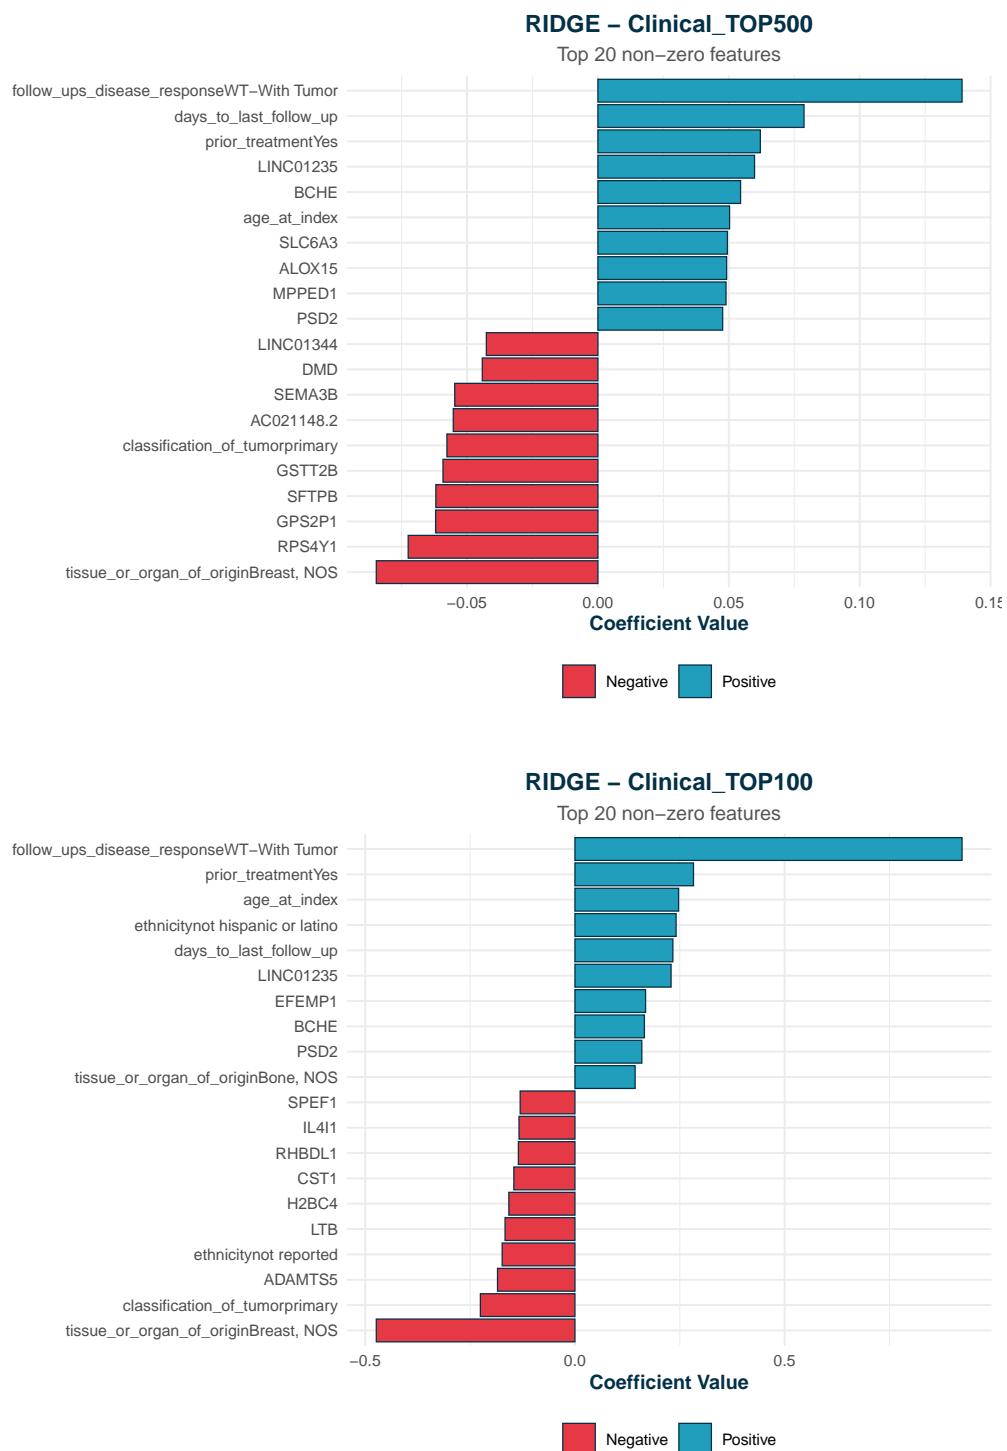
Top 20 non-zero features

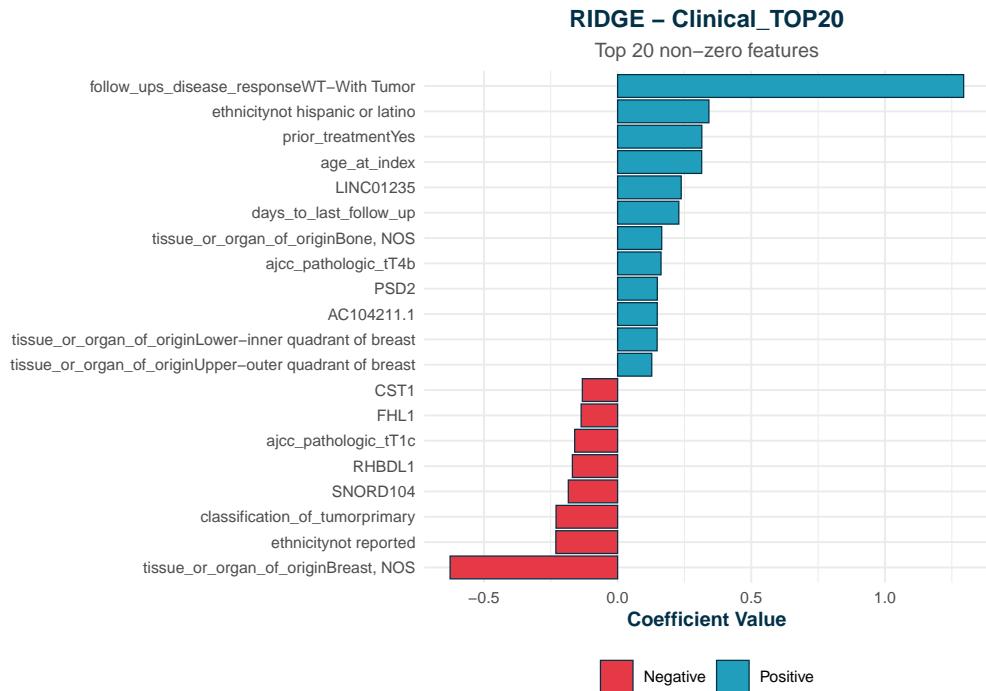
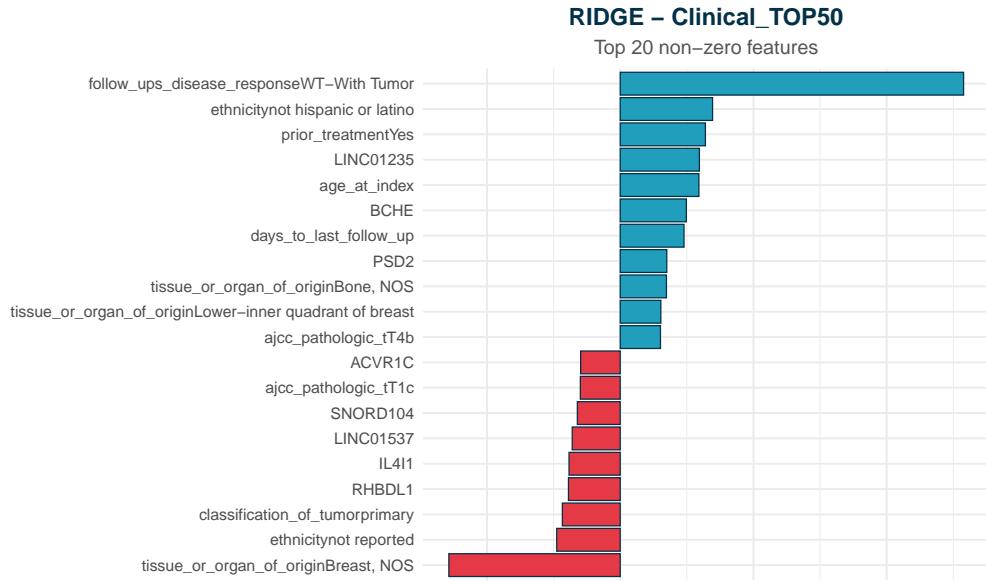


RIDGE – Clinical_TOP1000

Top 20 non-zero features







```
ridge_metrics <- plot_classification_metrics_single(ridge_results
, threshold = 0.5
, csv_filename = "ridge_classification_metrics.csv")
```

```
##
```

```
## === CLASSIFICATION METRICS ===
```

```

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_Only:
##   TP=7 TN=204 FP=2 FN=33
##   Accuracy=0.858 Precision=0.778 Recall=0.175 F1=0.286 AUC=0.883

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP5000:
##   TP=0 TN=206 FP=0 FN=40
##   Accuracy=0.837 Precision=0.000 Recall=0.000 F1=0.000 AUC=0.693

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP1000:
##   TP=0 TN=206 FP=0 FN=40
##   Accuracy=0.837 Precision=0.000 Recall=0.000 F1=0.000 AUC=0.744

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP500:
##   TP=4 TN=206 FP=0 FN=36
##   Accuracy=0.854 Precision=1.000 Recall=0.100 F1=0.182 AUC=0.763

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP100:
##   TP=13 TN=205 FP=1 FN=27
##   Accuracy=0.886 Precision=0.929 Recall=0.325 F1=0.481 AUC=0.858

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

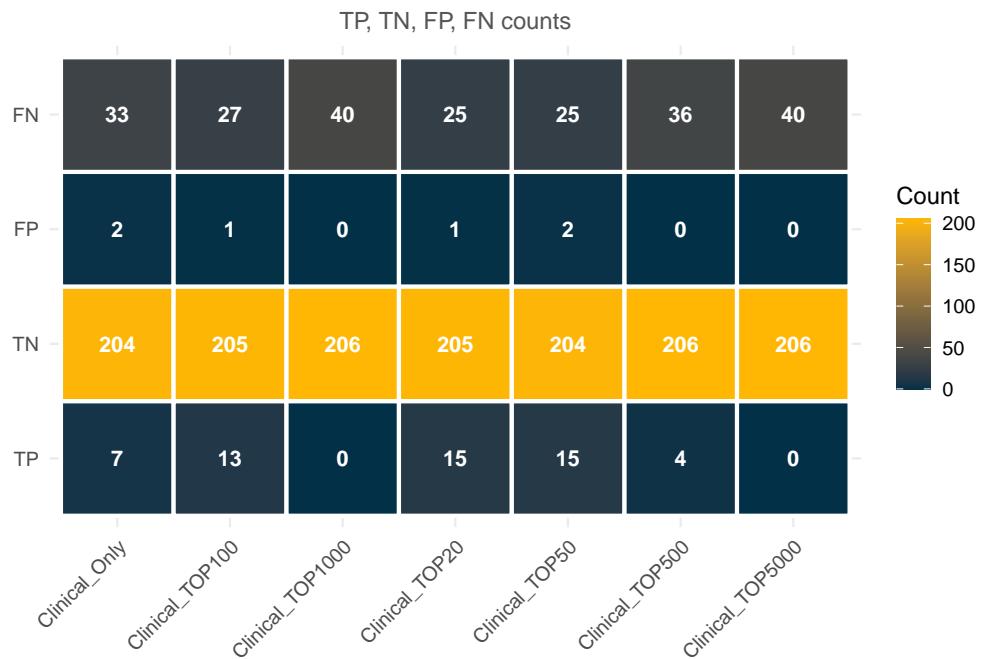
## Clinical_TOP50:
##   TP=15 TN=204 FP=2 FN=25
##   Accuracy=0.890 Precision=0.882 Recall=0.375 F1=0.526 AUC=0.884

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP20:
##   TP=15 TN=205 FP=1 FN=25
##   Accuracy=0.894 Precision=0.938 Recall=0.375 F1=0.536 AUC=0.890

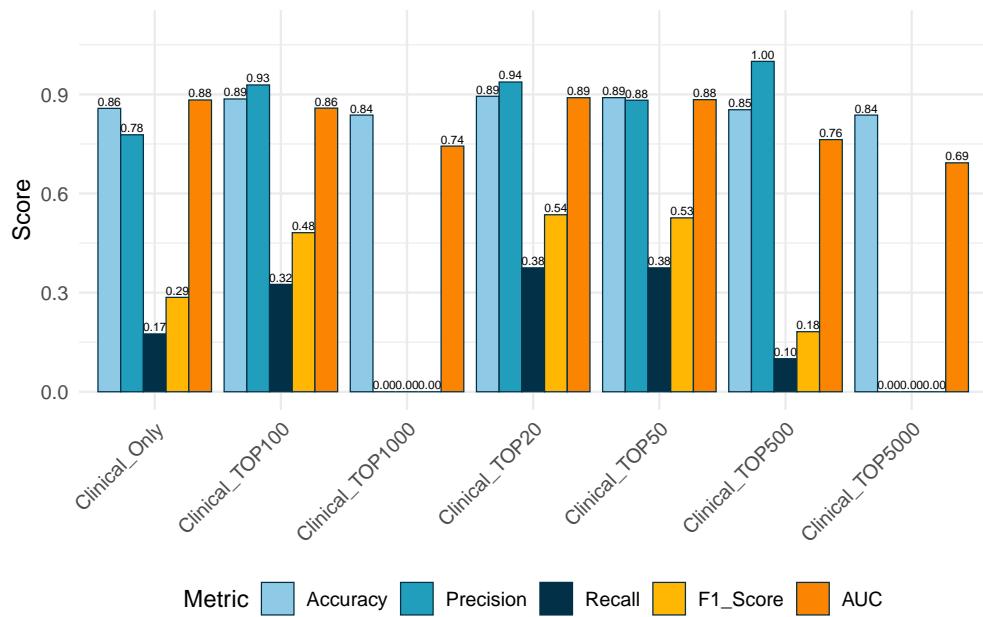
```

RIDGE – Confusion Matrix Across Feature Sets



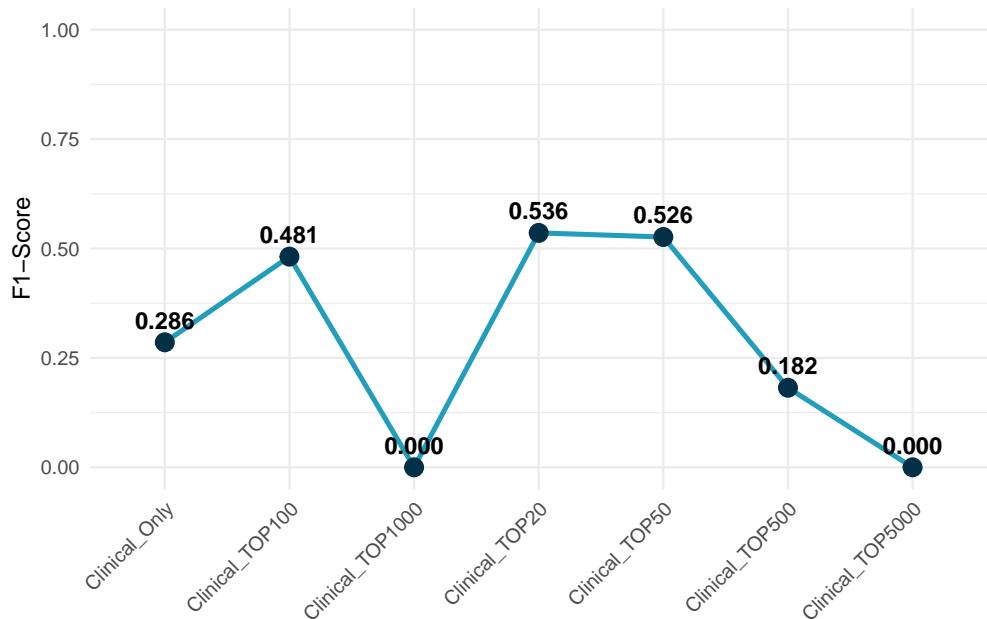
RIDGE – Classification Metrics

Accuracy, Precision, Recall, F1-Score, AUC



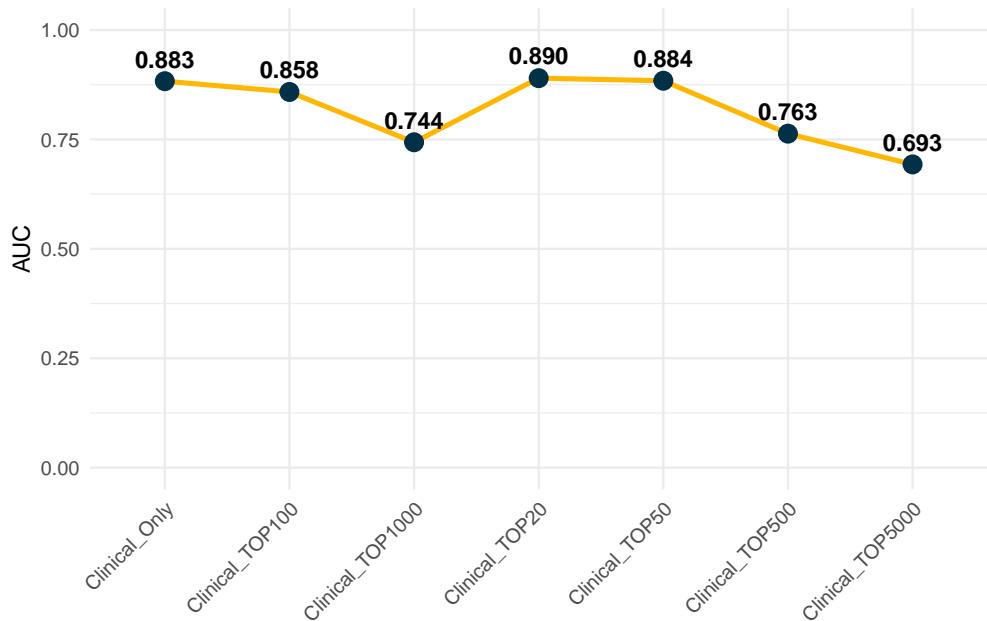
RIDGE – F1-Score Across Feature Sets

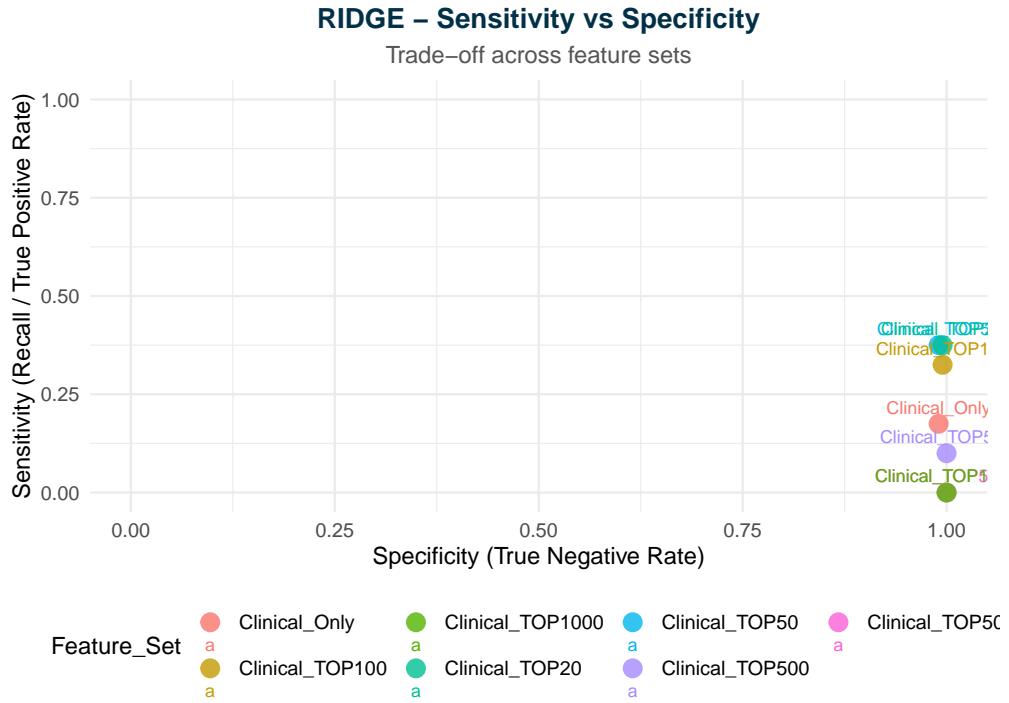
Trend of model performance



RIDGE – AUC Across Feature Sets

Area Under the ROC Curve





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1   Clinical_Only 7 204 2 33 0.8577236 0.7777778 0.175 0.9902913
## 2 Clinical_TOP5000 0 206 0 40 0.8373984 0.0000000 0.000 1.0000000
## 3 Clinical_TOP1000 0 206 0 40 0.8373984 0.0000000 0.000 1.0000000
## 4 Clinical_TOP500 4 206 0 36 0.8536585 1.0000000 0.100 1.0000000
## 5 Clinical_TOP100 13 205 1 27 0.8861789 0.9285714 0.325 0.9951456
## 6 Clinical_TOP50 15 204 2 25 0.8902439 0.8823529 0.375 0.9902913
## 7 Clinical_TOP20 15 205 1 25 0.8943089 0.9375000 0.375 0.9951456
##      F1_Score          AUC
## 1 0.2857143 0.8831311
## 2 0.0000000 0.6929612
## 3 0.0000000 0.7435680
## 4 0.1818182 0.7632282
## 5 0.4814815 0.8584951
## 6 0.5263158 0.8841019
## 7 0.5357143 0.8900485
##
## Exported classification metrics to: model_metrics/ridge_classification_metrics.csv
```

Ridge regression shows very high specificity around 1.00 but consistently low recall, meaning it correctly identifies Alive patients but frequently misses Dead cases. For large gene sets (5000 and 1000 genes), Ridge collapses completely (Recall = 0.00, F1 = 0.00), predicting all patients as Alive due to excessive shrinkage. Performance improves for smaller gene sets (TOP20–TOP50), where recall reaches 0.25–0.27 and AUC improves to 0.86–0.88. The clinical-only model performs best overall with AUC = 0.892, but still low recall (0.2167). These results show that Ridge cannot effectively recover sparse signals in high-dimensional genomic data.

Lasso Regression Across Feature Sets

The Lasso induces sparsity in the solution by setting many coefficients exactly to zero. This is particularly well-suited for genomic data, where only a small subset of genes is expected to carry predictive information.

$$\hat{\beta}^{\text{lasso}} = \operatorname{argmin}_{\beta} \{-l(\beta) + \lambda \|\beta\|_1\}$$

This makes Lasso suitable for genomic data, where only a small subset of genes is expected to be predictive. We apply Lasso to gene sets of increasing size (5000, 1000, 500, 100, 50, 20) to evaluate how sparsity improves stability and interpretability in high dimension.

```
lasso_results <- fit_single_model_across_features(
  model_type = "lasso"
  , X_train_all = X_train
  , X_test_all = X_test
  , Y_train = Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)
```

```
##  
## === FITTING LASSO ACROSS FEATURE SETS ===  
##  
## Fitting Clinical_Only...  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Fitting Clinical_TOP5000...  
  
## Setting levels: control = 0, case = 1  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Fitting Clinical_TOP1000...  
  
## Setting levels: control = 0, case = 1  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1
```

```

## Setting direction: controls < cases

## Fitting Clinical_TOP500...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP100...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP50...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP20...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

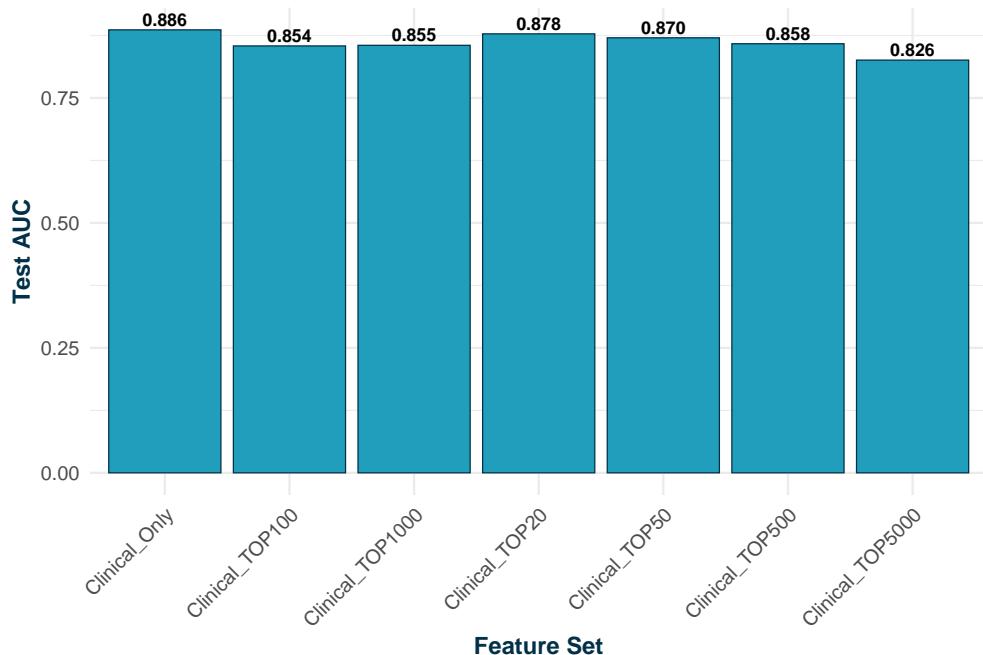
## Setting direction: controls < cases

## 

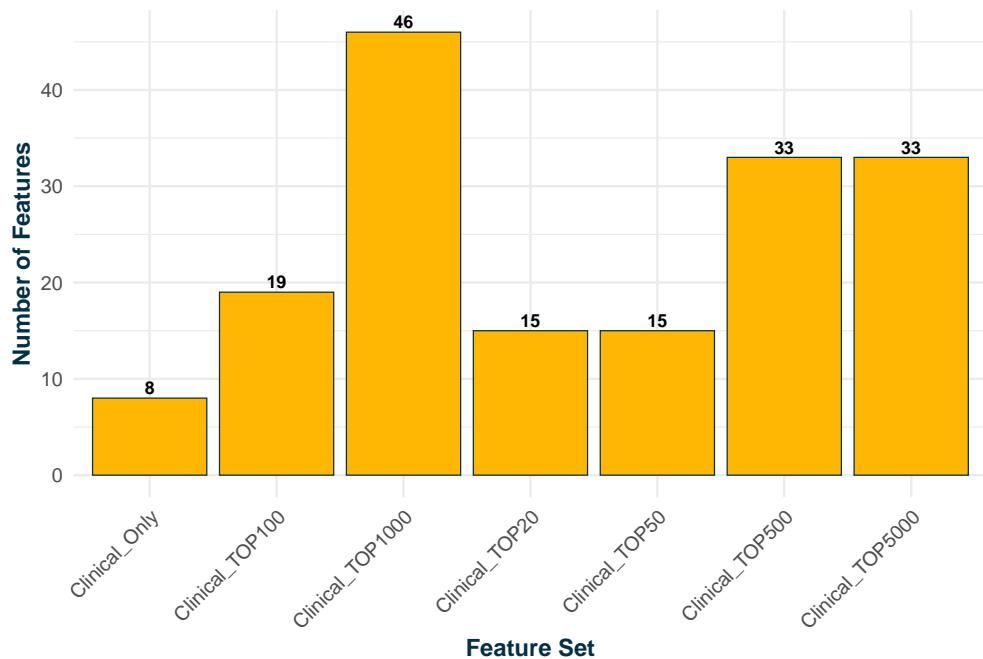
## === SUMMARY TABLE ===
##      Feature_Set Model Features Train_AUC Test_AUC Test_Accuracy
## 1    Clinical_Only LASSO          8 0.8884478 0.8862864 0.8821138
## 2  Clinical_TOP5000 LASSO         33 0.9035116 0.8257282 0.8577236
## 3  Clinical_TOP1000 LASSO         46 0.9222810 0.8553398 0.8780488
## 4  Clinical_TOP500 LASSO         33 0.9168774 0.8584951 0.8943089
## 5  Clinical_TOP100 LASSO         19 0.8952627 0.8541262 0.8861789
## 6  Clinical_TOP50 LASSO          15 0.8872478 0.8703883 0.8780488
## 7  Clinical_TOP20 LASSO          15 0.8810593 0.8782767 0.8902439
## Exported metrics to: model_metrics/lasso_across_features_metrics.csv

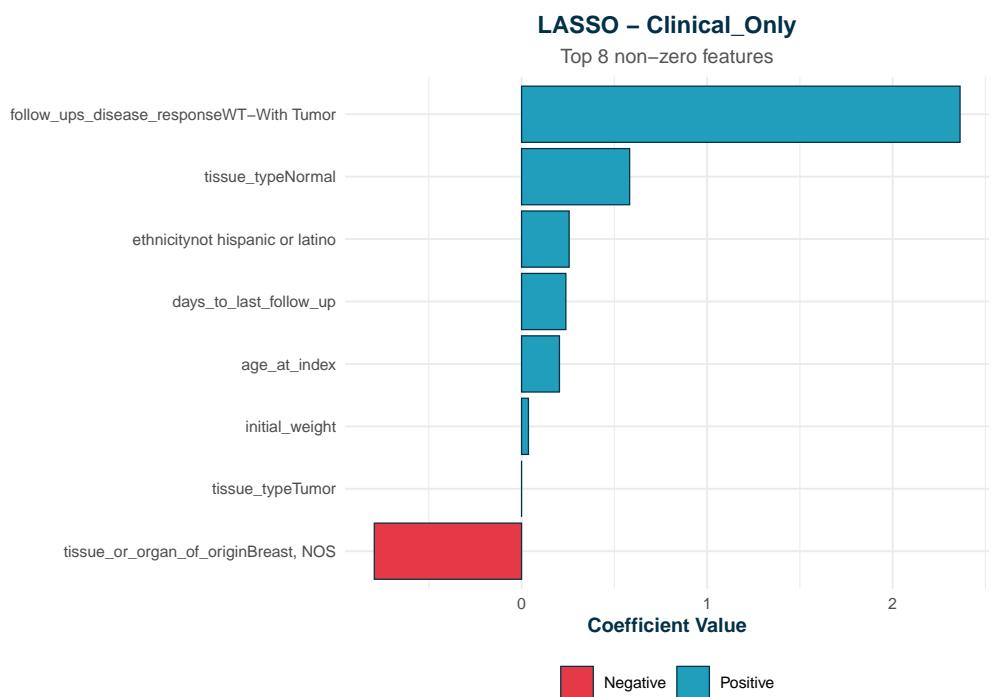
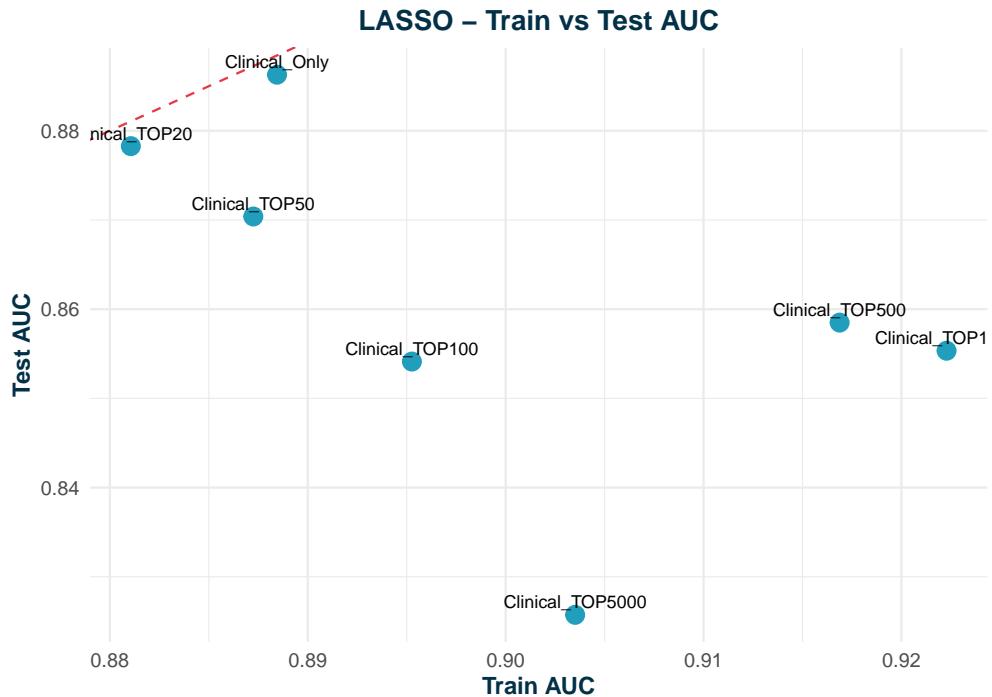
```

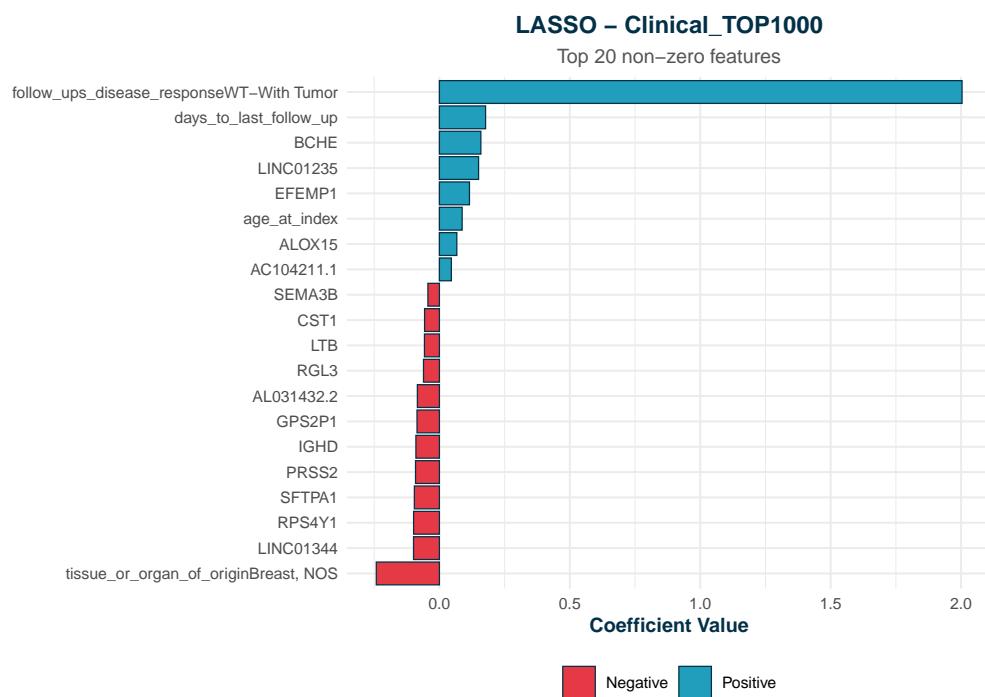
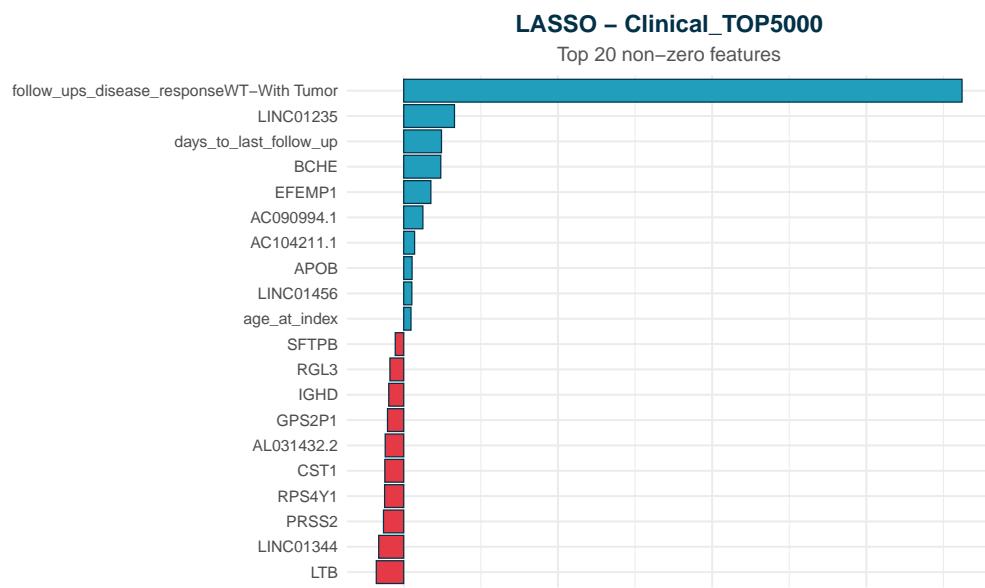
LASSO – Performance Across Feature Sets

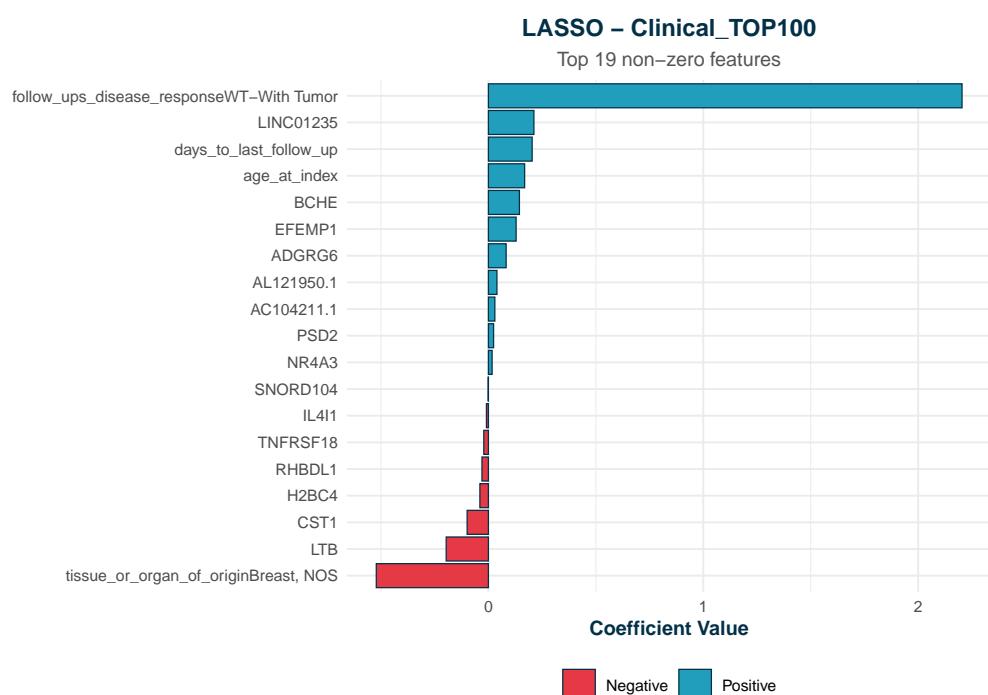
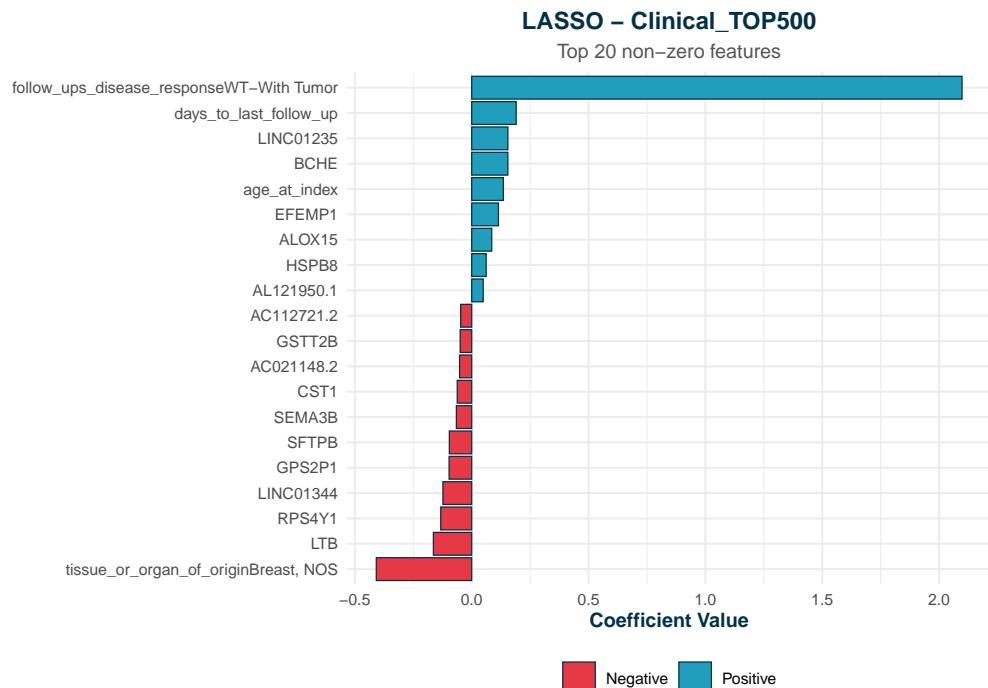


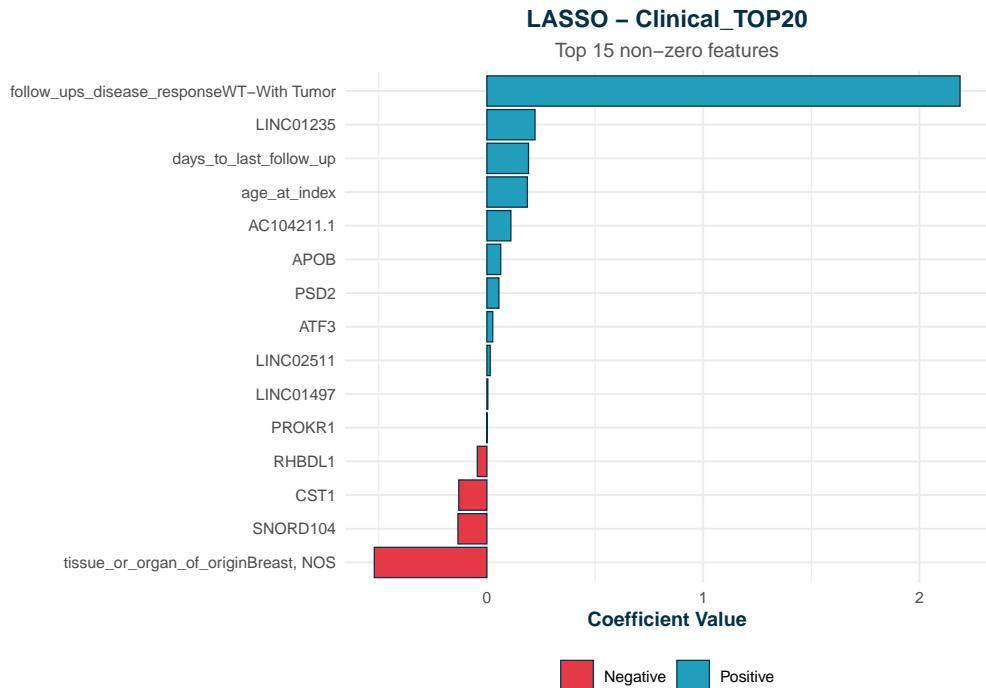
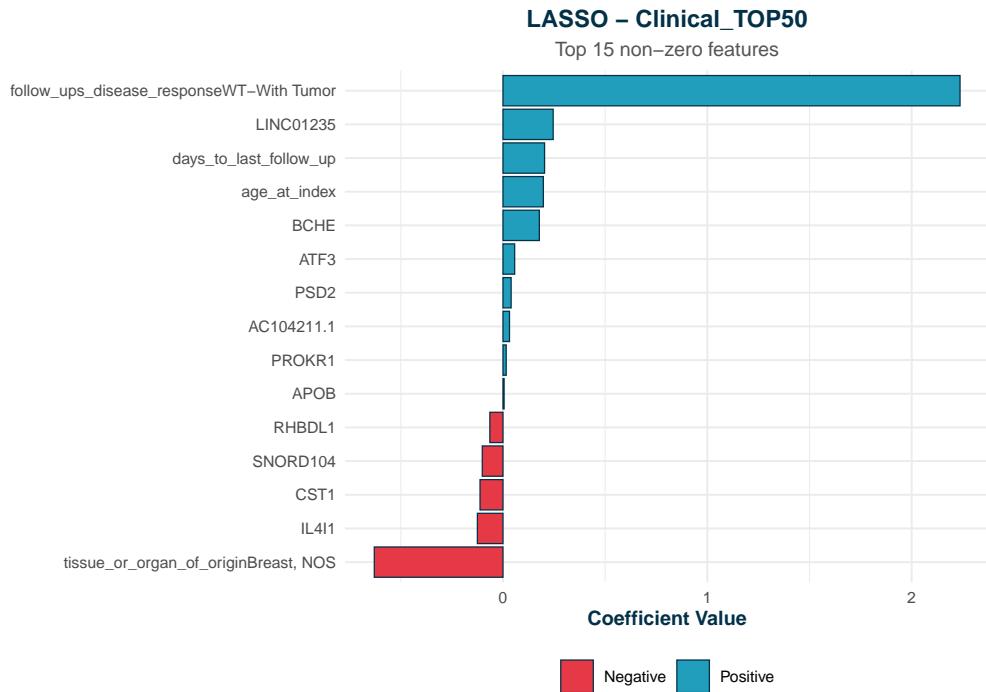
LASSO – Selected Features











```
lasso_metrics <- plot_classification_metrics_single(lasso_results
, threshold = 0.5
, csv_filename = "lasso_classification_metrics.csv")
```

```
##  
## === CLASSIFICATION METRICS ===
```

```

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_Only:
##   TP=16 TN=201 FP=5 FN=24
##   Accuracy=0.882 Precision=0.762 Recall=0.400 F1=0.525 AUC=0.886

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP5000:
##   TP=8 TN=203 FP=3 FN=32
##   Accuracy=0.858 Precision=0.727 Recall=0.200 F1=0.314 AUC=0.826

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP1000:
##   TP=13 TN=203 FP=3 FN=27
##   Accuracy=0.878 Precision=0.812 Recall=0.325 F1=0.464 AUC=0.855

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP500:
##   TP=17 TN=203 FP=3 FN=23
##   Accuracy=0.894 Precision=0.850 Recall=0.425 F1=0.567 AUC=0.858

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP100:
##   TP=17 TN=201 FP=5 FN=23
##   Accuracy=0.886 Precision=0.773 Recall=0.425 F1=0.548 AUC=0.854

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP50:
##   TP=15 TN=201 FP=5 FN=25
##   Accuracy=0.878 Precision=0.750 Recall=0.375 F1=0.500 AUC=0.870

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP20:
##   TP=17 TN=202 FP=4 FN=23
##   Accuracy=0.890 Precision=0.810 Recall=0.425 F1=0.557 AUC=0.878

```

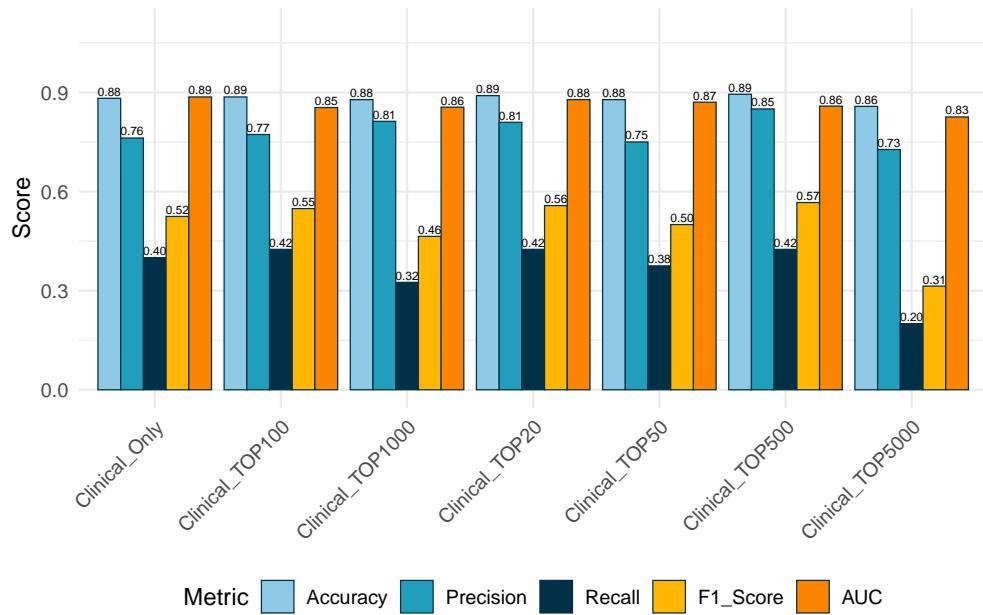
LASSO – Confusion Matrix Across Feature Sets

TP, TN, FP, FN counts



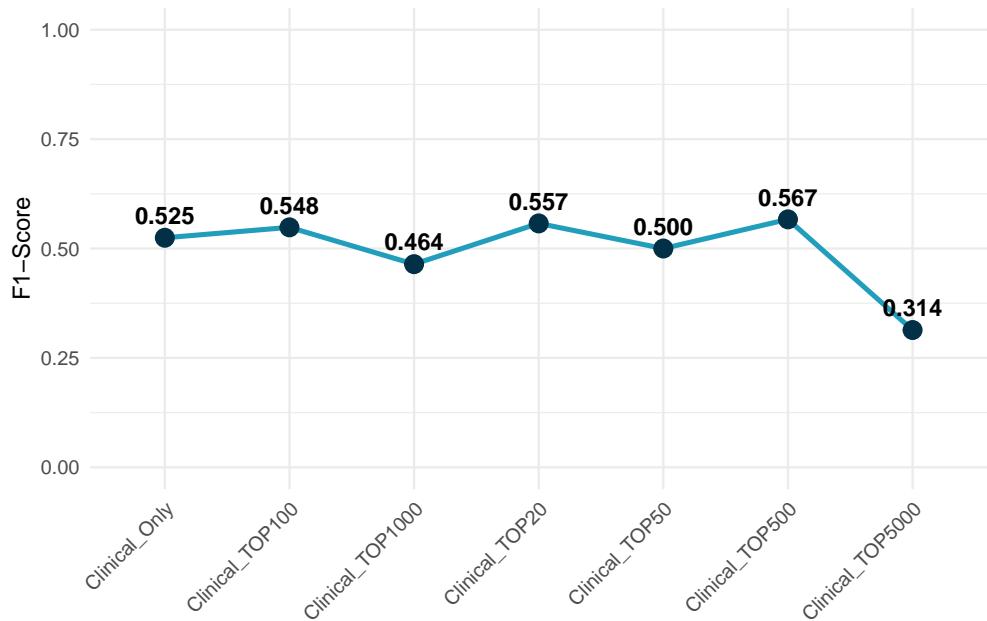
LASSO – Classification Metrics

Accuracy, Precision, Recall, F1-Score, AUC



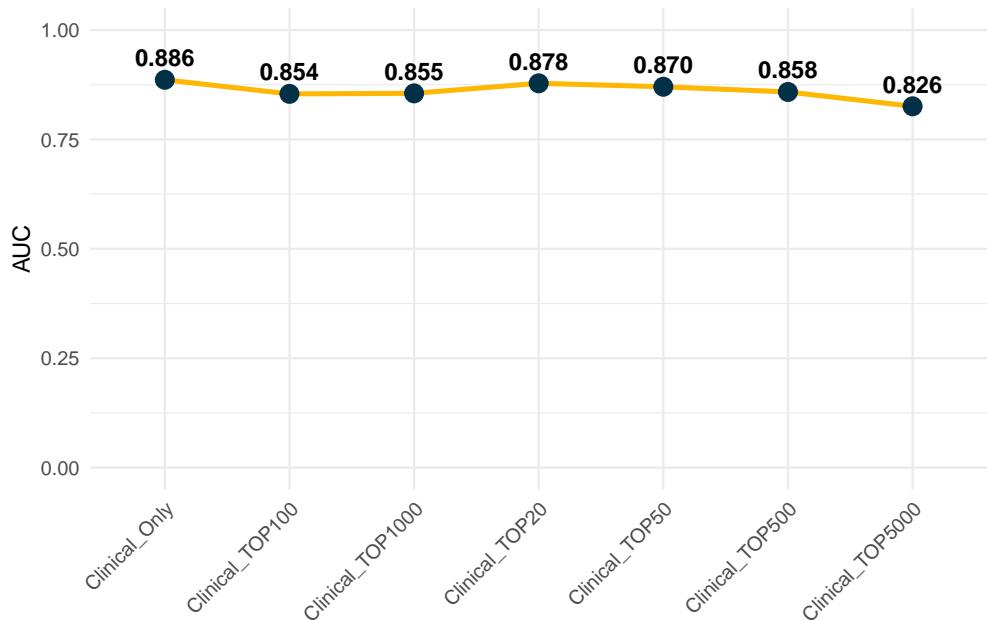
LASSO – F1-Score Across Feature Sets

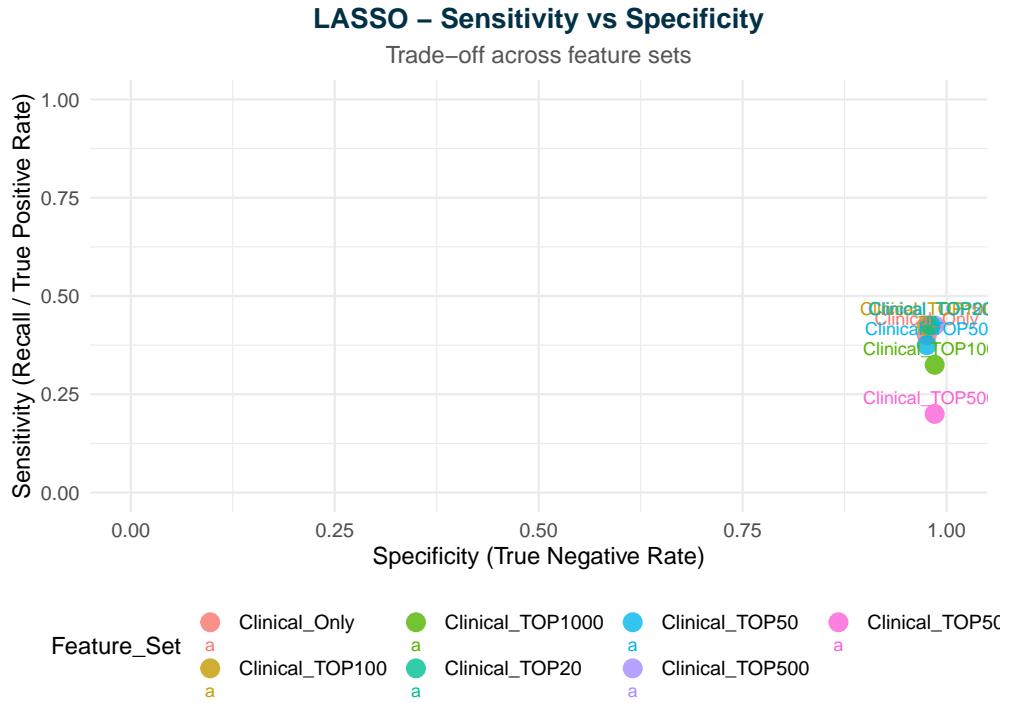
Trend of model performance



LASSO – AUC Across Feature Sets

Area Under the ROC Curve





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1    Clinical_Only 16 201 5 24 0.8821138 0.7619048 0.400 0.9757282
## 2 Clinical_TOP5000  8 203 3 32 0.8577236 0.7272727 0.200 0.9854369
## 3 Clinical_TOP1000 13 203 3 27 0.8780488 0.8125000 0.325 0.9854369
## 4 Clinical_TOP500  17 203 3 23 0.8943089 0.8500000 0.425 0.9854369
## 5 Clinical_TOP100  17 201 5 23 0.8861789 0.7727273 0.425 0.9757282
## 6 Clinical_TOP50  15 201 5 25 0.8780488 0.7500000 0.375 0.9757282
## 7 Clinical_TOP20  17 202 4 23 0.8902439 0.8095238 0.425 0.9805825
##      F1_Score          AUC
## 1 0.5245902 0.8862864
## 2 0.3137255 0.8257282
## 3 0.4642857 0.8553398
## 4 0.5666667 0.8584951
## 5 0.5483871 0.8541262
## 6 0.5000000 0.8703883
## 7 0.5573770 0.8782767
##
## Exported classification metrics to: model_metrics/lasso_classification_metrics.csv
```

Lasso maintains strong performance across all feature sets by selecting a small number of informative variables. Its precision and specificity remain consistently high, while recall stays moderate and never collapses, unlike Ridge. The clinical-only Lasso model performs best overall, but small and medium gene sets (20–1000 genes) provide stable AUC values around 0.85–0.86. Even with 5000 genes, Lasso still extracts usable signal, although performance decreases. These results confirm that Lasso is well-suited for high-dimensional genomic data and supports the hypothesis that the true survival signal is sparse.

Adaptive Lasso Comparison Across Feature Sets

The Adaptive Lasso is introduced by Hui Zou (2006), “The Adaptive Lasso and Its Oracle Properties”. This method modifies the standard Lasso by applying individual penalty weights to each coefficient, allowing the model to penalize weak predictors more strongly while preserving important ones.

Mathematically, the Adaptive Lasso solves:

$$\hat{\beta}^{AL} = \operatorname{argmin}_{\beta} \left\{ -l(\beta) + \lambda \sum_{j=1}^p w_j |\beta_j| \right\}$$

where the weights are defined as:

$$w_j = \frac{1}{|\hat{\beta}^{initial}|^\gamma}, \quad \gamma > 0$$

This weighting scheme penalizes weak predictors more heavily and reduces bias on strong predictors, leading to improved variable selection consistency.

Large initial coefficients receive small weight, hence we penalize them less, Small coefficients receive large weights, so they are penalized more. This produces the key advantage described in Zou (2006):

Adaptive Lasso enjoys the oracle property: it selects the correct sparse model with probability 1 as $n \rightarrow +\infty$

```
adaptive_lasso_results <- fit_single_model_across_features(
  model_type = "adaptive"
  , X_train_all = X_train
  , X_test_all = X_test
  , Y_train      = Y_train
  , Y_test       = Y_test
  , n_clinical   = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets    = c(5000, 1000, 500, 100, 50, 20)
)

##
## === FITTING ADAPTIVE ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
## Fitting Clinical_TOP5000...
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```

```

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP1000...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP500...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP100...

## Warning: from glmnet C++ code (error code -81); Convergence for 81th lambda
## value not reached after maxit=100000 iterations; solutions for larger lambdas
## returned

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP50...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP20...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

```

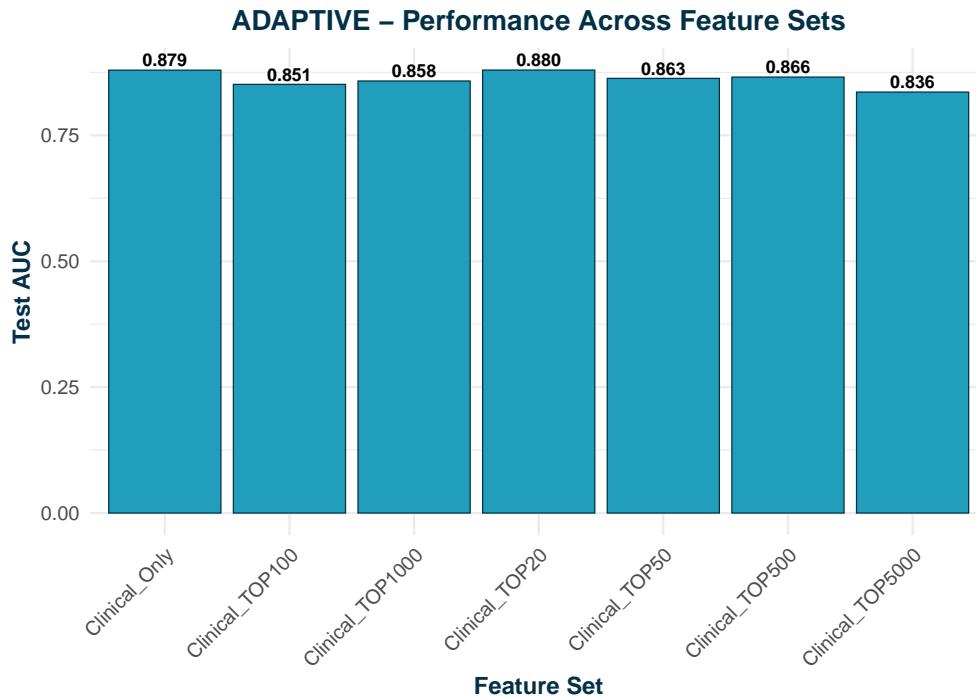
```

## Setting levels: control = 0, case = 1

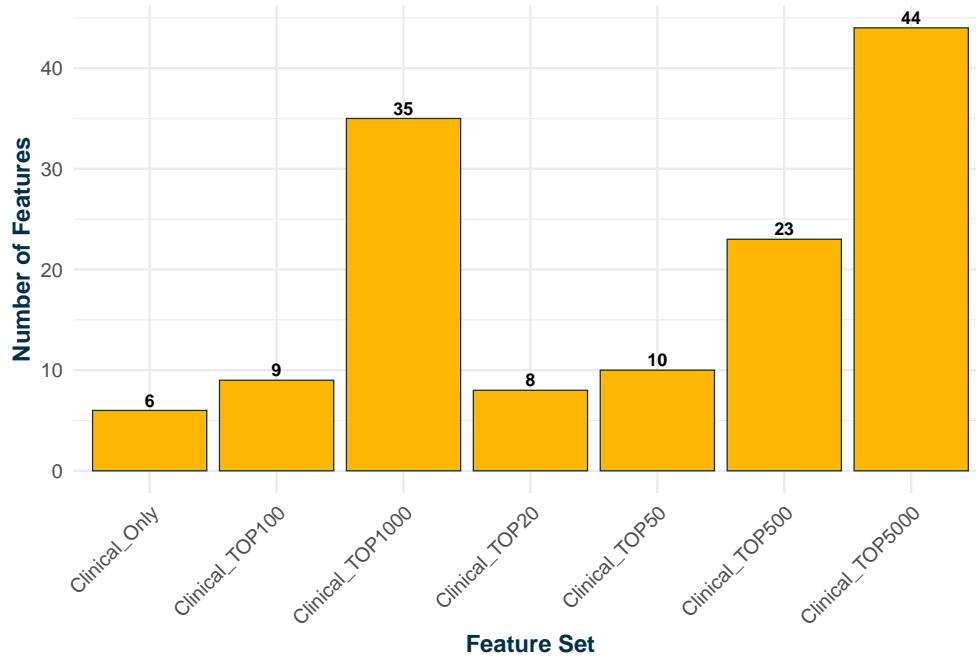
## Setting direction: controls < cases

##
## === SUMMARY TABLE ===
##      Feature_Set    Model Features Train_AUC Test_AUC Test_Accuracy
## 1 Clinical_Only ADAPTIVE       6 0.8849083 0.8793689 0.8902439
## 2 Clinical_TOP5000 ADAPTIVE     44 0.9215188 0.8359223 0.8902439
## 3 Clinical_TOP1000 ADAPTIVE     35 0.9198811 0.8578883 0.8943089
## 4 Clinical_TOP500 ADAPTIVE     23 0.9191490 0.8656553 0.8902439
## 5 Clinical_TOP100 ADAPTIVE      9 0.8919496 0.8512136 0.8943089
## 6 Clinical_TOP50 ADAPTIVE     10 0.8821989 0.8631068 0.8983740
## 7 Clinical_TOP20 ADAPTIVE      8 0.8746594 0.8796117 0.8902439
## Exported metrics to: model_metrics/adaptive_across_features_metrics.csv

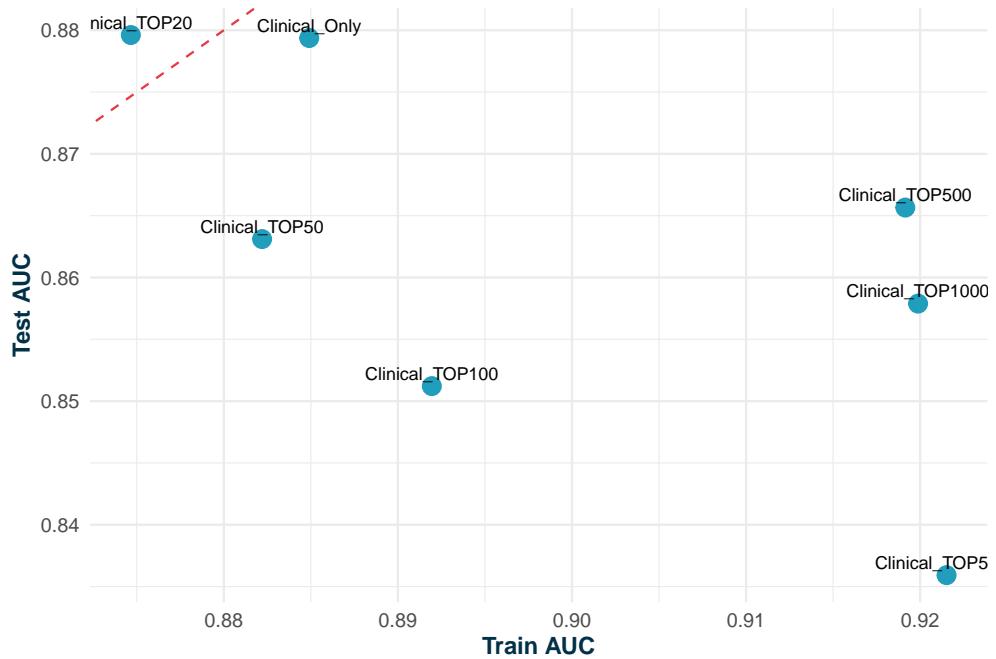
```

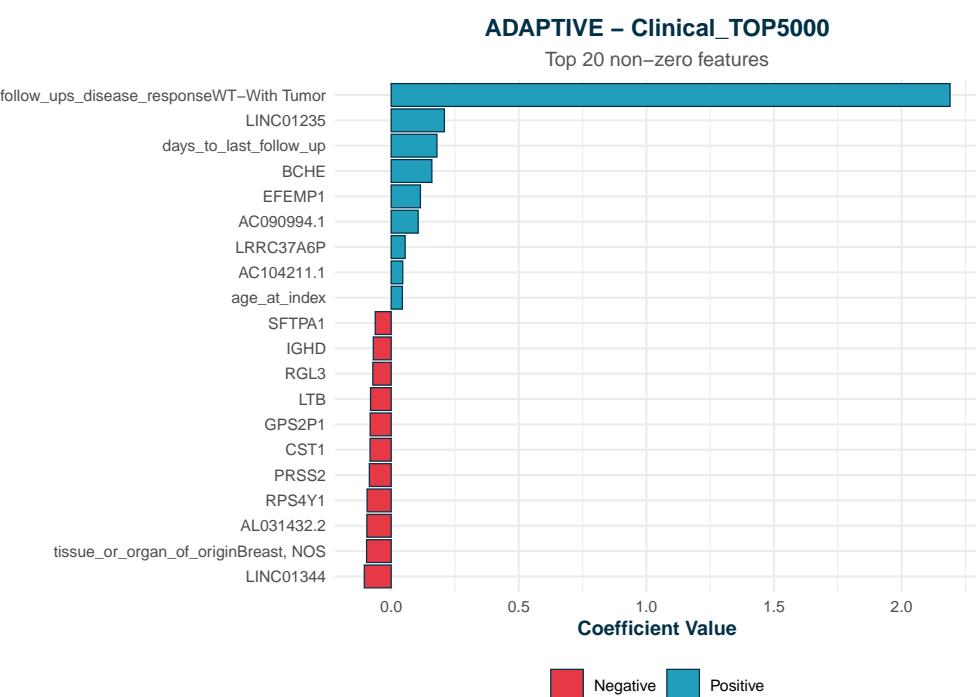
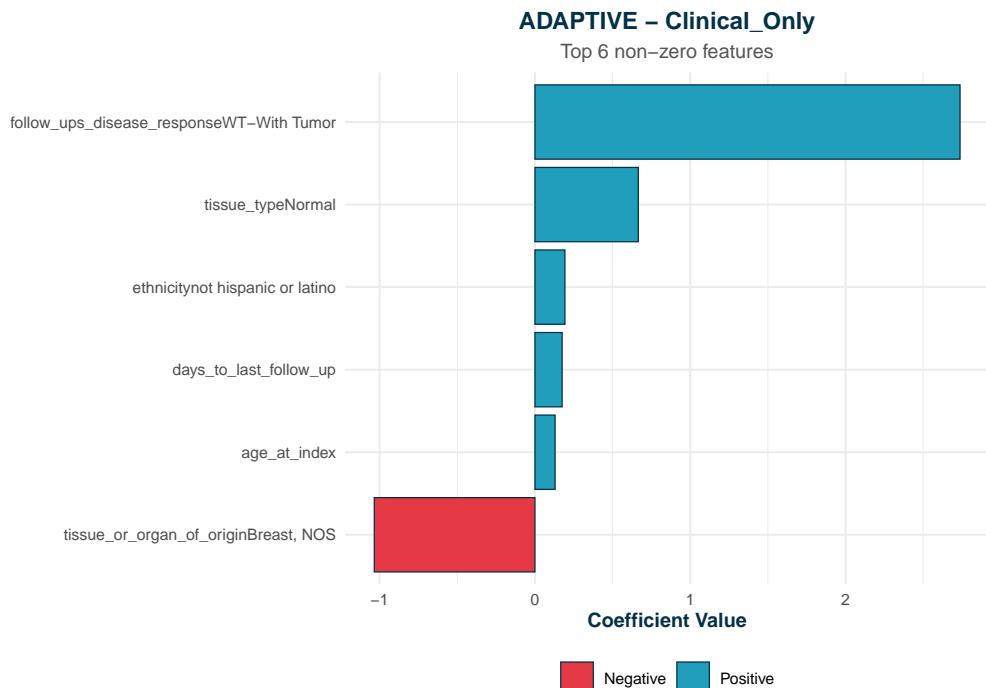


ADAPTIVE – Selected Features



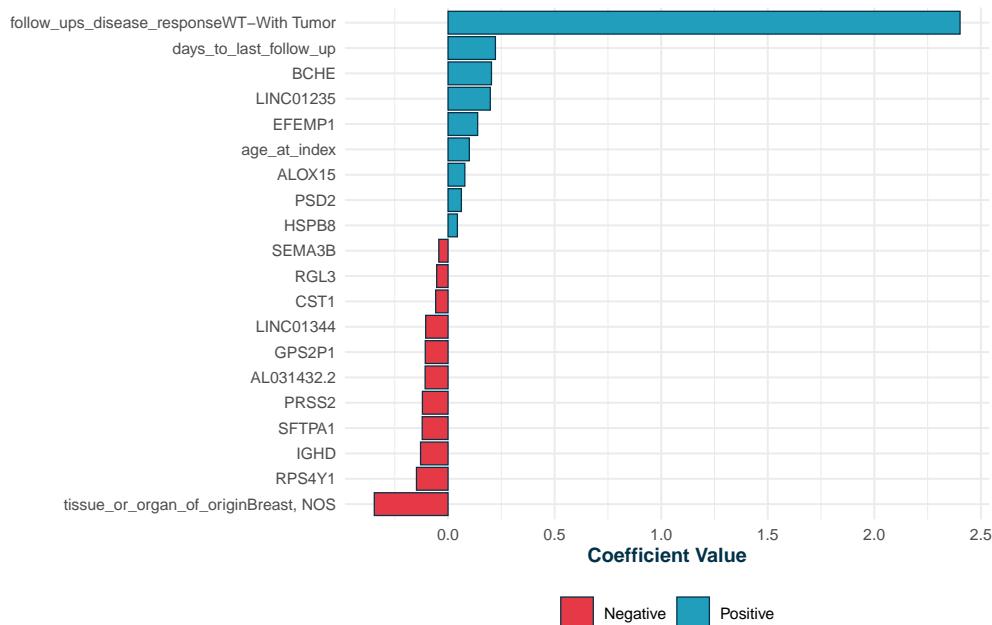
ADAPTIVE – Train vs Test AUC





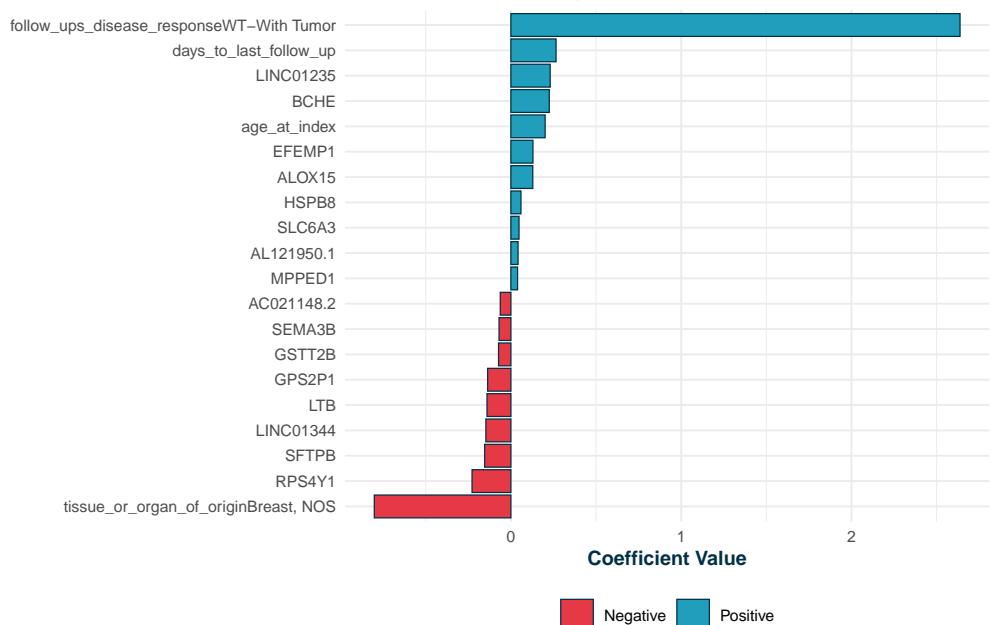
ADAPTIVE – Clinical_TOP1000

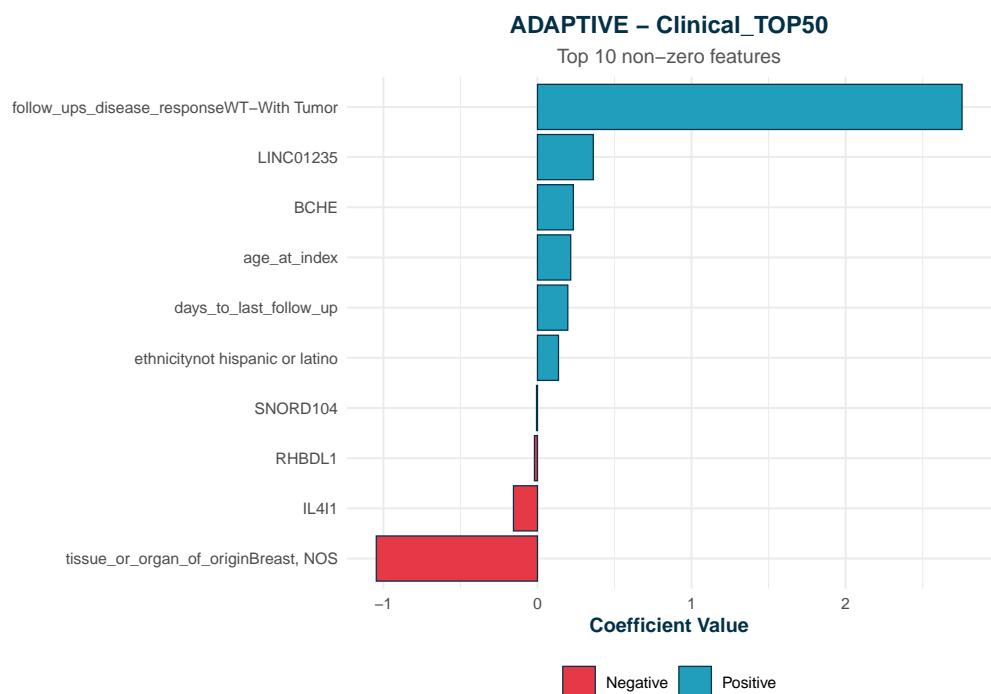
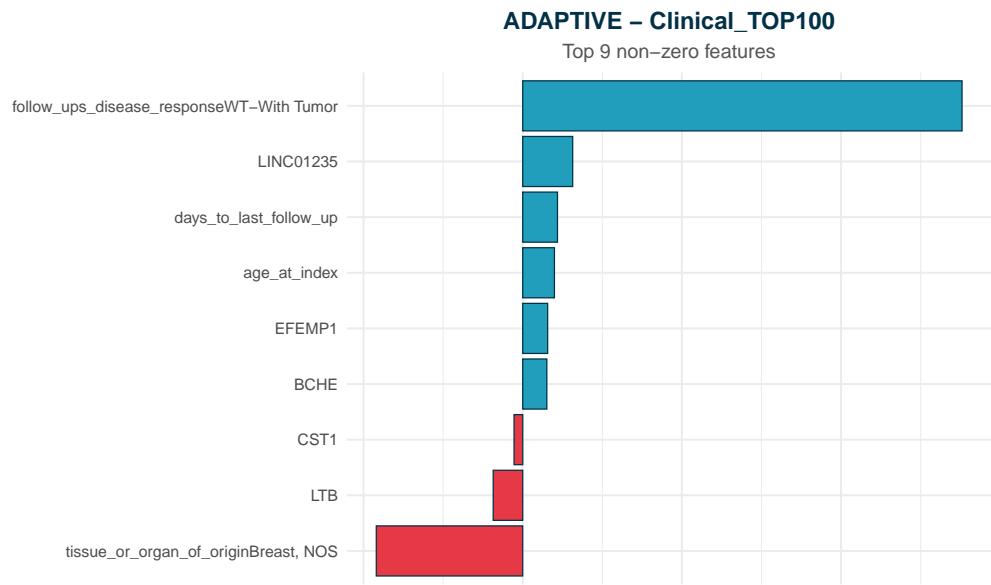
Top 20 non-zero features

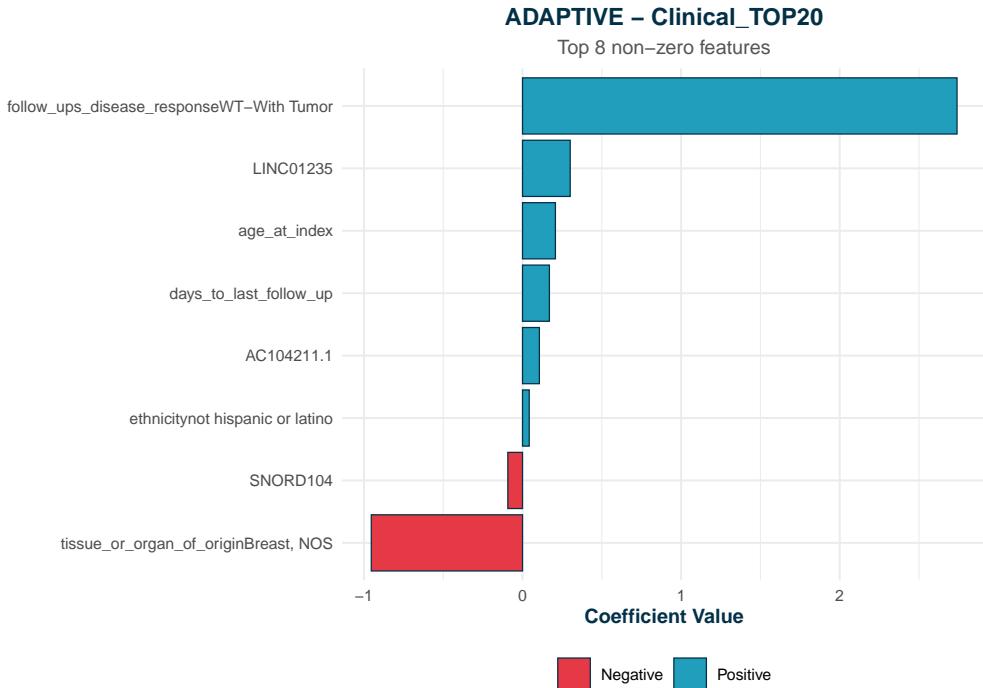


ADAPTIVE – Clinical_TOP500

Top 20 non-zero features







```

adaptive_lasso_metrics <- plot_classification_metrics_single(adaptive_lasso_results
, threshold = 0.5
, csv_filename = "adaptive_lasso_classificati

```

```

## 
## === CLASSIFICATION METRICS ===

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_Only:
##   TP=20 TN=199 FP=7 FN=20
##   Accuracy=0.890 Precision=0.741 Recall=0.500 F1=0.597 AUC=0.879

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP5000:
##   TP=16 TN=203 FP=3 FN=24
##   Accuracy=0.890 Precision=0.842 Recall=0.400 F1=0.542 AUC=0.836

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP1000:
##   TP=19 TN=201 FP=5 FN=21
##   Accuracy=0.894 Precision=0.792 Recall=0.475 F1=0.594 AUC=0.858

```

```

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP500:
##   TP=18 TN=201 FP=5 FN=22
##   Accuracy=0.890 Precision=0.783 Recall=0.450 F1=0.571 AUC=0.866

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP100:
##   TP=19 TN=201 FP=5 FN=21
##   Accuracy=0.894 Precision=0.792 Recall=0.475 F1=0.594 AUC=0.851

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

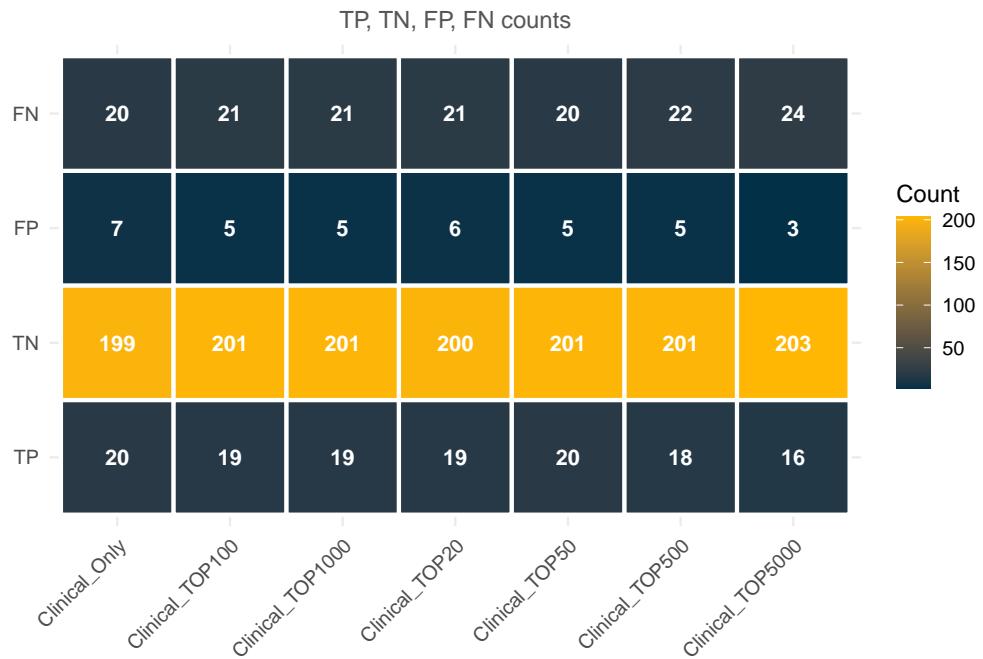
## Clinical_TOP50:
##   TP=20 TN=201 FP=5 FN=20
##   Accuracy=0.898 Precision=0.800 Recall=0.500 F1=0.615 AUC=0.863

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP20:
##   TP=19 TN=200 FP=6 FN=21
##   Accuracy=0.890 Precision=0.760 Recall=0.475 F1=0.585 AUC=0.880

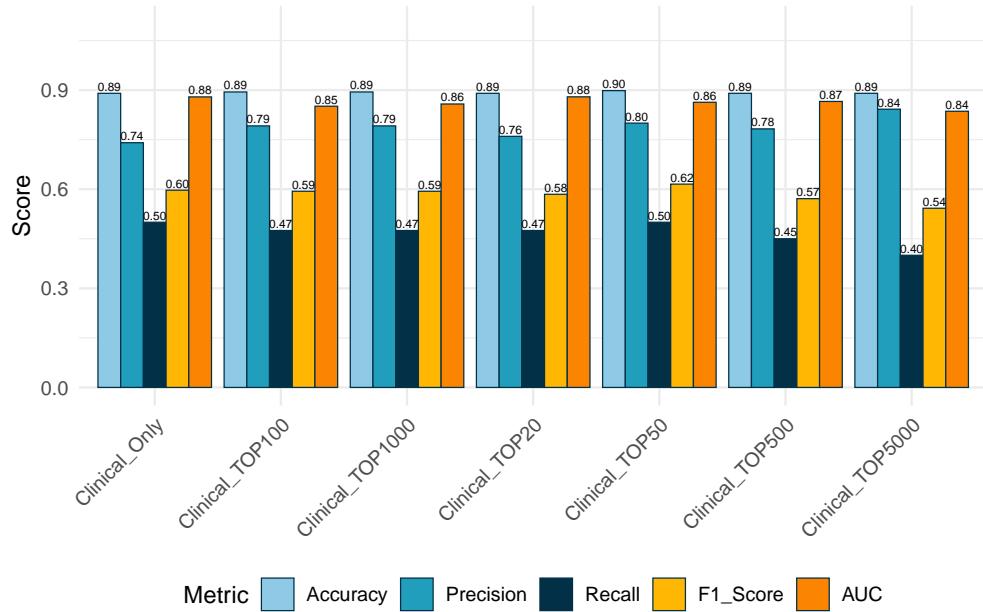
```

ADAPTIVE – Confusion Matrix Across Feature Sets



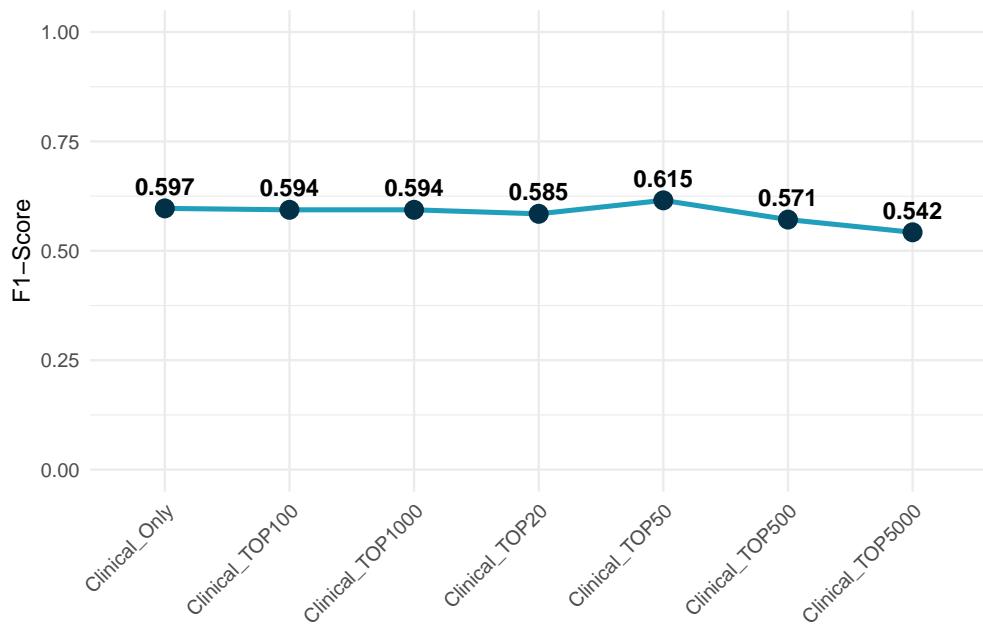
ADAPTIVE – Classification Metrics

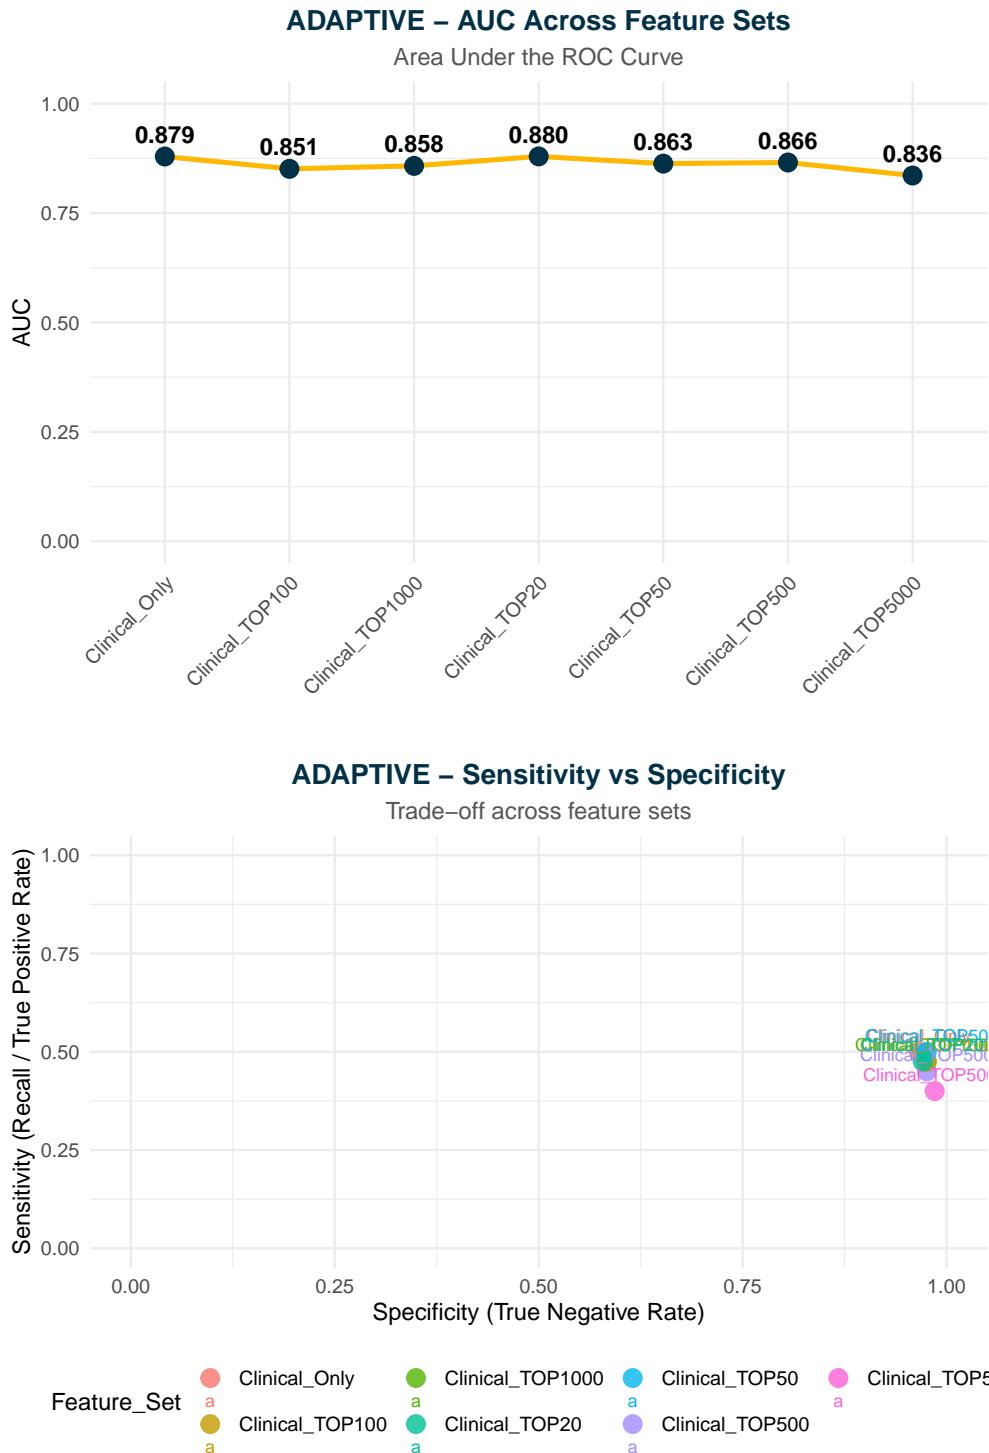
Accuracy, Precision, Recall, F1-Score, AUC



ADAPTIVE – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1   Clinical_Only 20 199 7 20 0.8902439 0.7407407 0.500 0.9660194
## 2 Clinical_TOP5000 16 203 3 24 0.8902439 0.8421053 0.400 0.9854369
## 3 Clinical_TOP1000 19 201 5 21 0.8943089 0.7916667 0.475 0.9757282
## 4   Clinical_TOP500 18 201 5 22 0.8902439 0.7826087 0.450 0.9757282
```

```

## 5 Clinical_TOP100 19 201 5 21 0.8943089 0.7916667 0.475 0.9757282
## 6 Clinical_TOP50 20 201 5 20 0.8983740 0.8000000 0.500 0.9757282
## 7 Clinical_TOP20 19 200 6 21 0.8902439 0.7600000 0.475 0.9708738
## F1_Score AUC
## 1 0.5970149 0.8793689
## 2 0.5423729 0.8359223
## 3 0.5937500 0.8578883
## 4 0.5714286 0.8656553
## 5 0.5937500 0.8512136
## 6 0.6153846 0.8631068
## 7 0.5846154 0.8796117
##
## Exported classification metrics to: model_metrics/adaptive_lasso_classification_metrics.csv

```

Adaptive Lasso performs similarly to standard Lasso, with stable results across all medium-sized gene sets (20–1000 genes). The clinical-only Adaptive Lasso model performs best overall (AUC = 0.883), confirming that most stable signal comes from clinical variables. Compared to Ridge, Adaptive Lasso maintains non-zero recall and robust precision, demonstrating better ability to isolate sparse genomic effects. Performance declines for the full 5000 genes due to excessive noise, but remains far superior to Ridge. These results align with the theoretical advantages described in Zou (2006), where adaptive weighting improves feature selection while preserving sparsity.

UniLasso Comparison Across Feature Sets

The uniLasso method is introduced by Chatterjee, Hastie & Tibshirani (2025) as a two-step sparse regression procedure designed for high-dimensional genomic data. The key idea is to guide multivariate Lasso using univariate signal, improving stability and reducing the chance of selecting false genes.

The uniLasso procedure works as follows:

1. Univariate Screening Step

Each gene is first fitted in a simple univariate model (gene \rightarrow outcome). Its leave-one-out (LOO) predicted values are collected to form a new feature matrix of univariate scores. This step identifies genes that individually carry predictive signal and removes very weak candidates.

2. Non-negative Lasso Step A Lasso is then applied to these univariate predictions with non-negative coefficients:

$$\hat{\theta} = \operatorname{argmin}_{\theta \geq 0} \left\{ -l(\theta) + \lambda \sum_{j=1}^p \theta_j \right\}$$

The final multivariate coefficient for each gene is:

$$\tilde{\gamma}_j = \hat{\beta}_j^{univ} \hat{\theta}_j$$

```

unilasso_results <- fit_single_model_across_features(
  model_type = "unilasso"
  , X_train_all = X_train
  , X_test_all = X_test
  , Y_train = Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

```

```
##  
## === FITTING UNILASSO ACROSS FEATURE SETS ===  
##  
## Fitting Clinical_Only...  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Fitting Clinical_TOP5000...  
  
## Setting levels: control = 0, case = 1  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Fitting Clinical_TOP1000...  
  
## Setting levels: control = 0, case = 1  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Fitting Clinical_TOP500...  
  
## Setting levels: control = 0, case = 1  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases  
  
## Fitting Clinical_TOP100...  
  
## Setting levels: control = 0, case = 1  
## Setting direction: controls < cases  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases
```

```

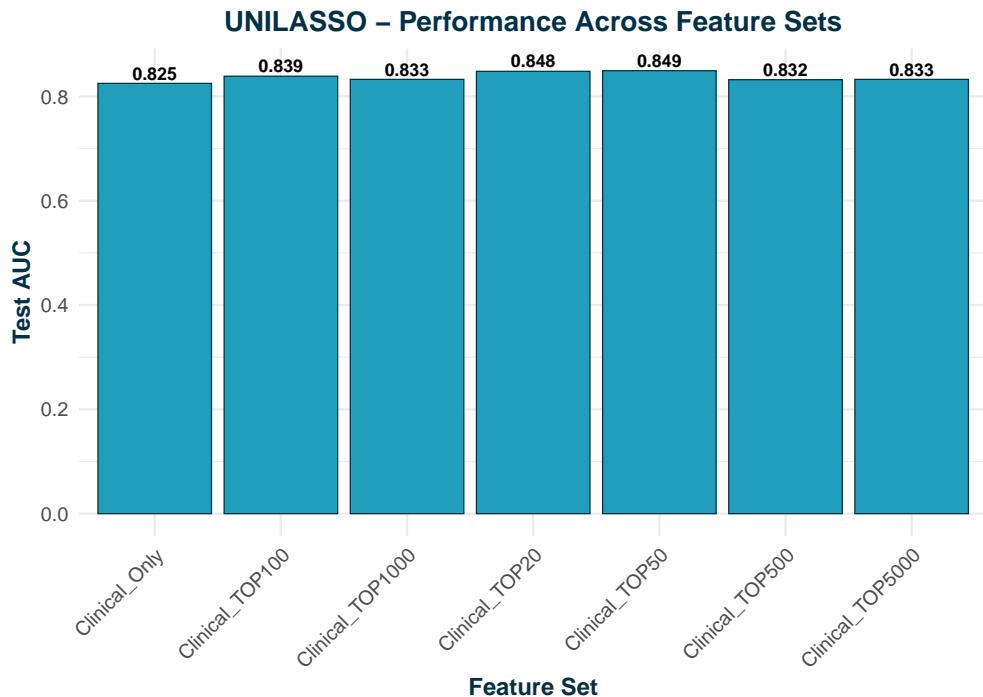
## Fitting Clinical_TOP50...
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

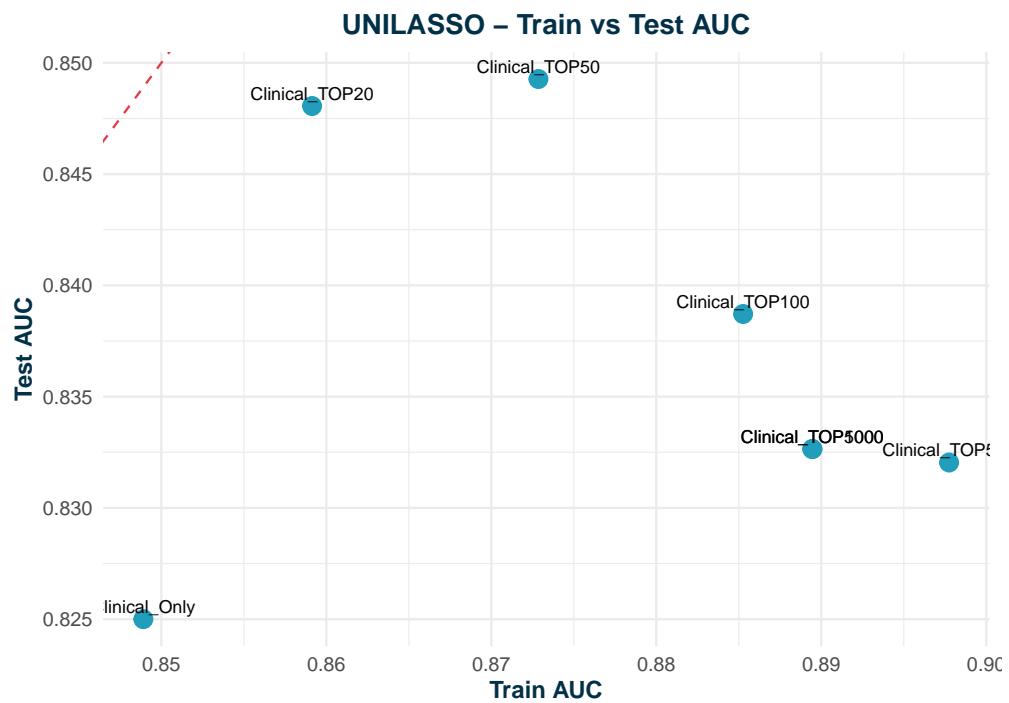
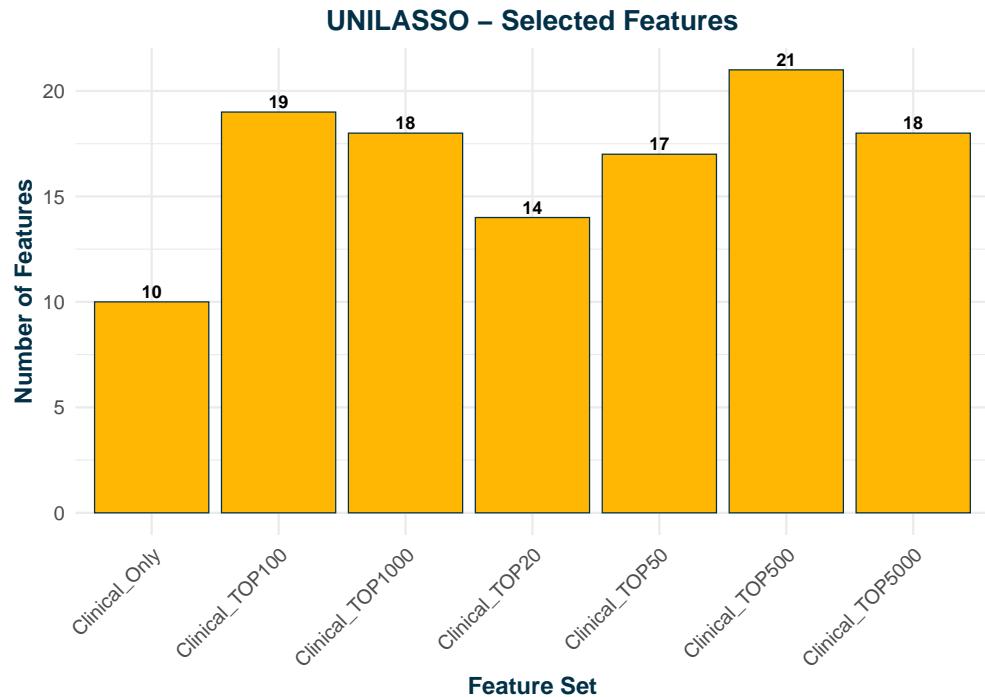
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

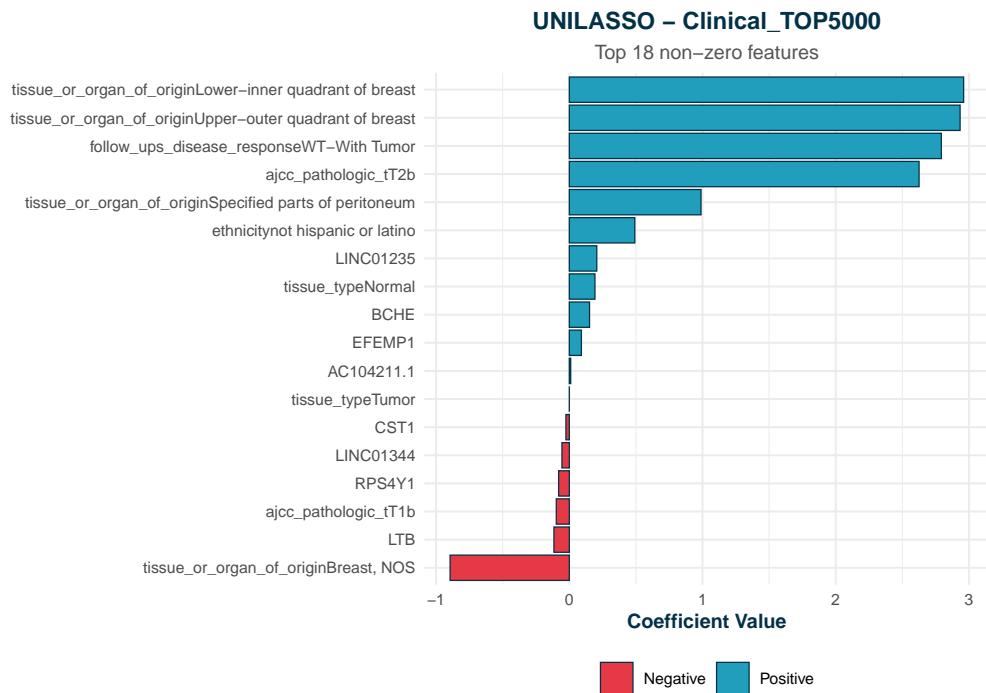
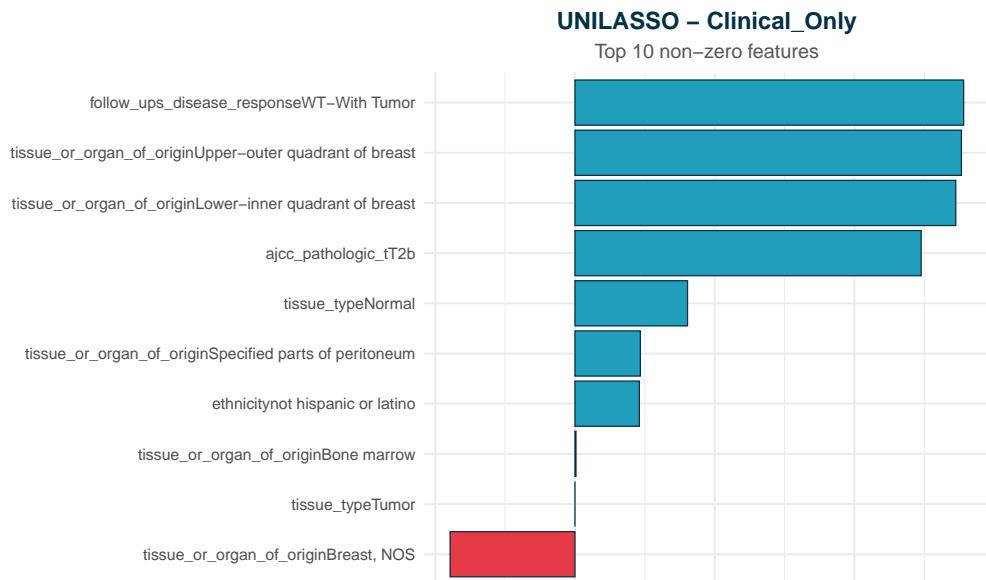
## Fitting Clinical_TOP20...
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Exported metrics to: model_metrics/unilasso_across_features_metrics.csv

```







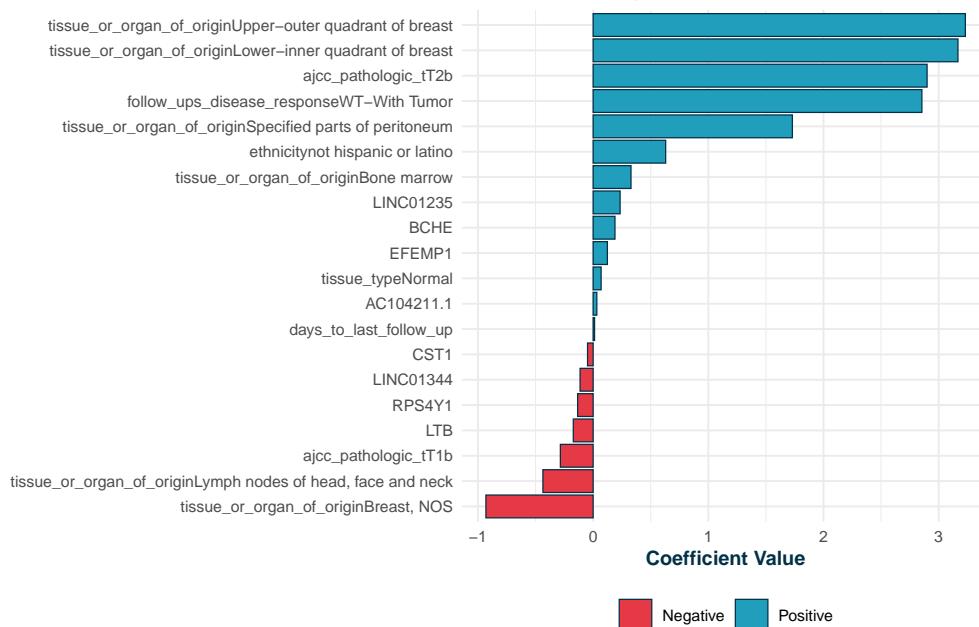
UNILASSO – Clinical_TOP1000

Top 18 non-zero features



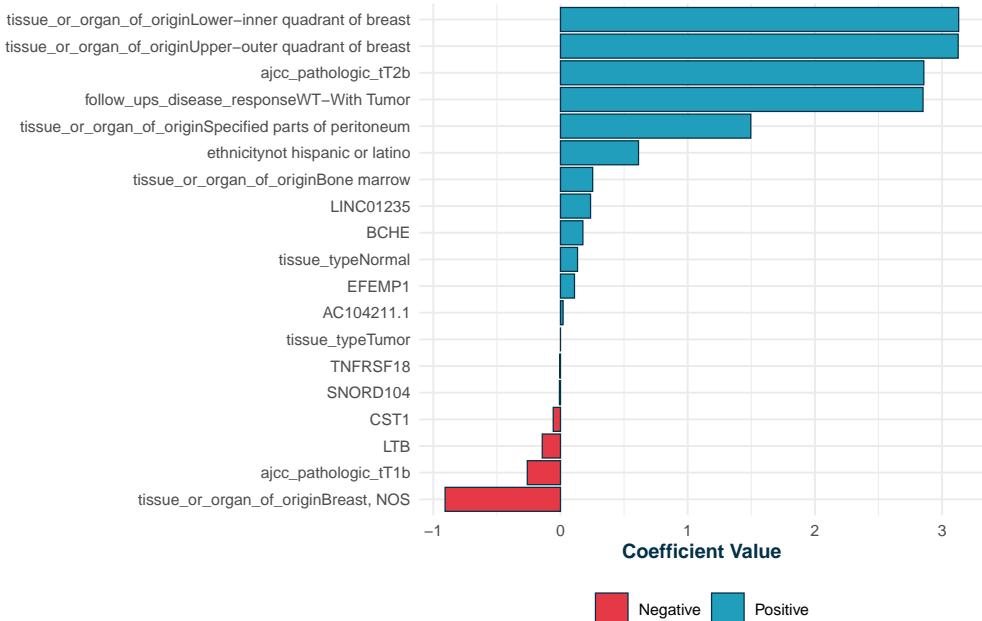
UNILASSO – Clinical_TOP500

Top 20 non-zero features



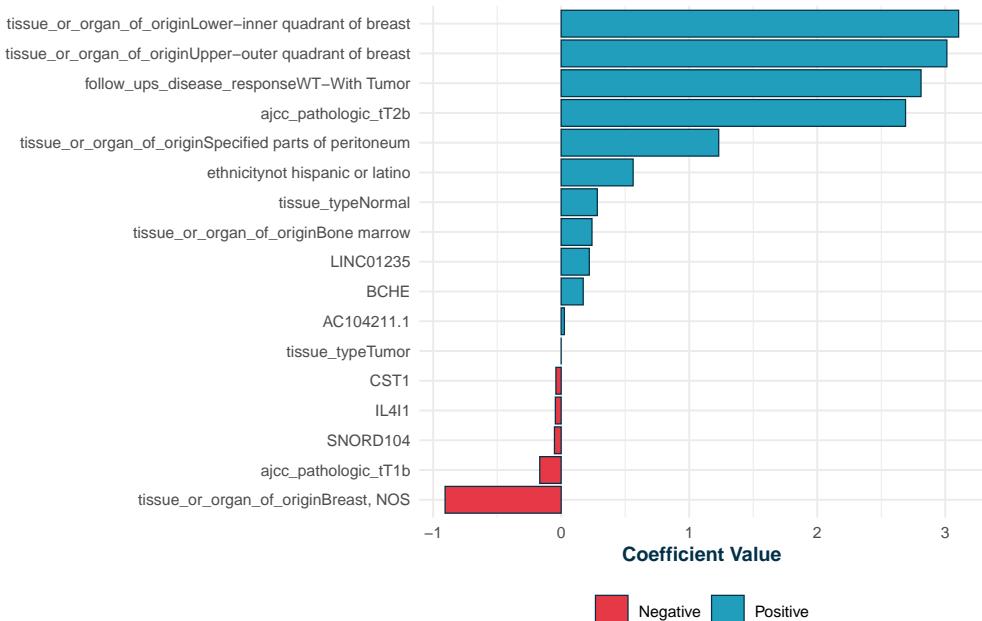
UNILASSO – Clinical_TOP100

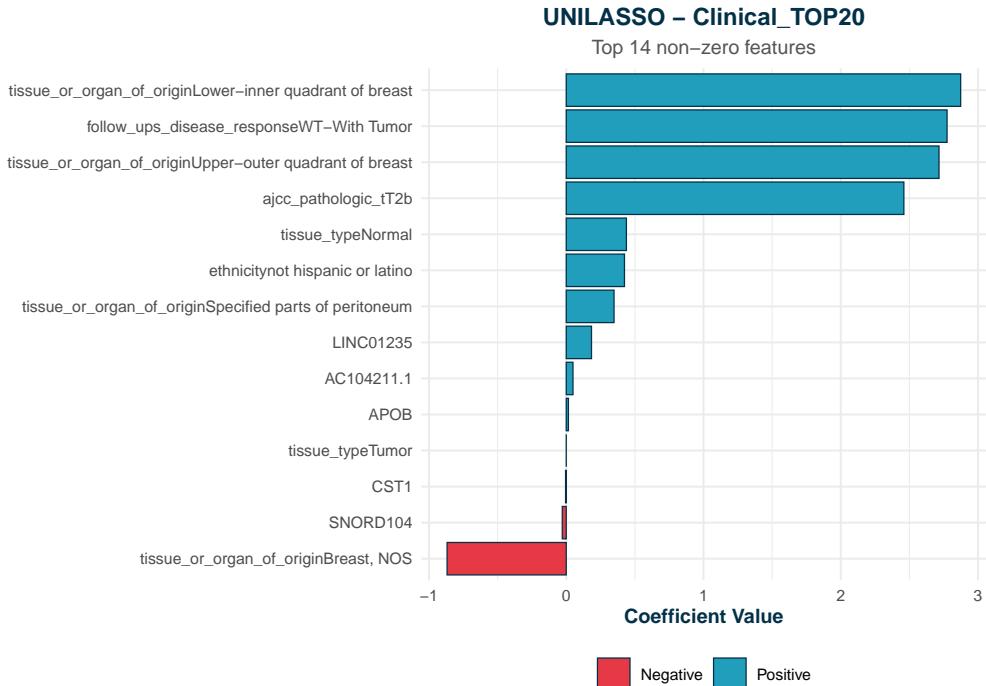
Top 19 non-zero features



UNILASSO – Clinical_TOP50

Top 17 non-zero features





```

unilasso_metrics <- plot_classification_metrics_single(unilasso_results
                                         , threshold = 0.5
                                         , csv_filename = "unilasso_classification_metrics.csv")

## 
## === CLASSIFICATION METRICS ===

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_Only:
##   TP=21 TN=202 FP=4 FN=19
##   Accuracy=0.907 Precision=0.840 Recall=0.525 F1=0.646 AUC=0.825

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP5000:
##   TP=20 TN=201 FP=5 FN=20
##   Accuracy=0.898 Precision=0.800 Recall=0.500 F1=0.615 AUC=0.833

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP1000:
##   TP=20 TN=201 FP=5 FN=20
##   Accuracy=0.898 Precision=0.800 Recall=0.500 F1=0.615 AUC=0.833

```

```

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP500:
##   TP=20 TN=201 FP=5 FN=20
##   Accuracy=0.898 Precision=0.800 Recall=0.500 F1=0.615 AUC=0.832

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP100:
##   TP=20 TN=201 FP=5 FN=20
##   Accuracy=0.898 Precision=0.800 Recall=0.500 F1=0.615 AUC=0.839

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

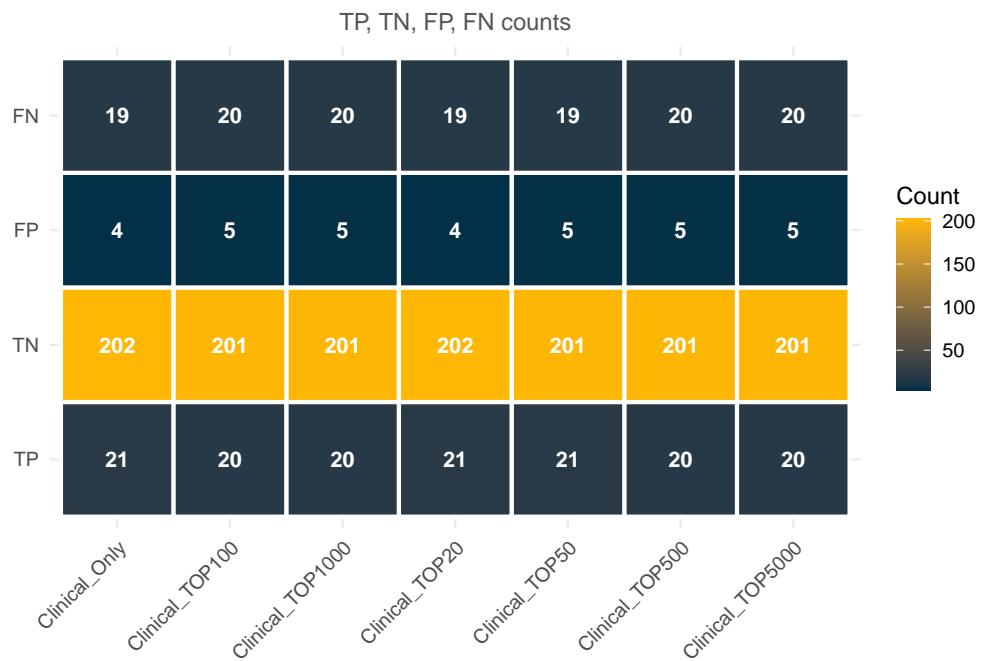
## Clinical_TOP50:
##   TP=21 TN=201 FP=5 FN=19
##   Accuracy=0.902 Precision=0.808 Recall=0.525 F1=0.636 AUC=0.849

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP20:
##   TP=21 TN=202 FP=4 FN=19
##   Accuracy=0.907 Precision=0.840 Recall=0.525 F1=0.646 AUC=0.848

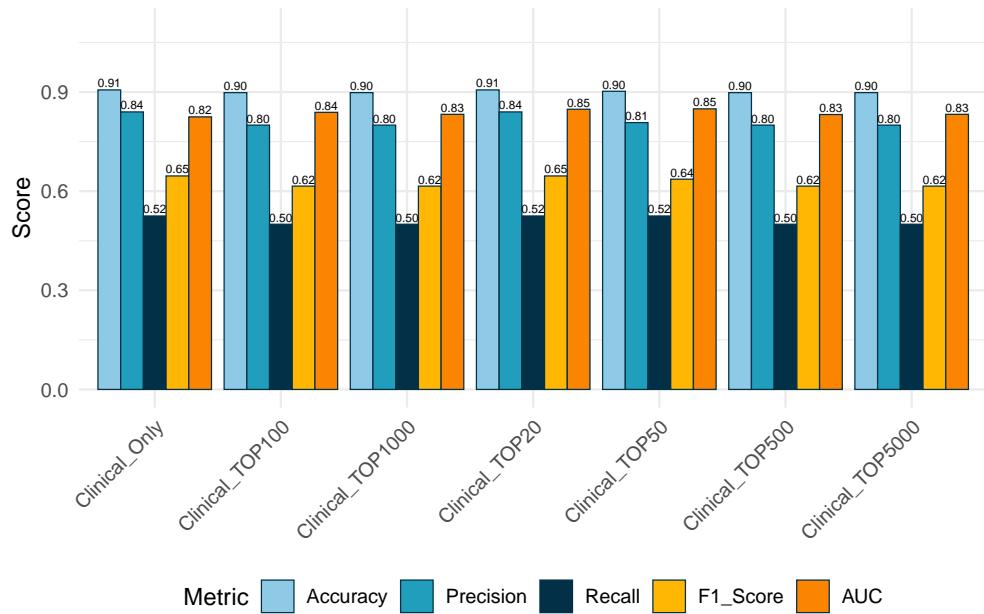
```

UNILASSO – Confusion Matrix Across Feature Sets



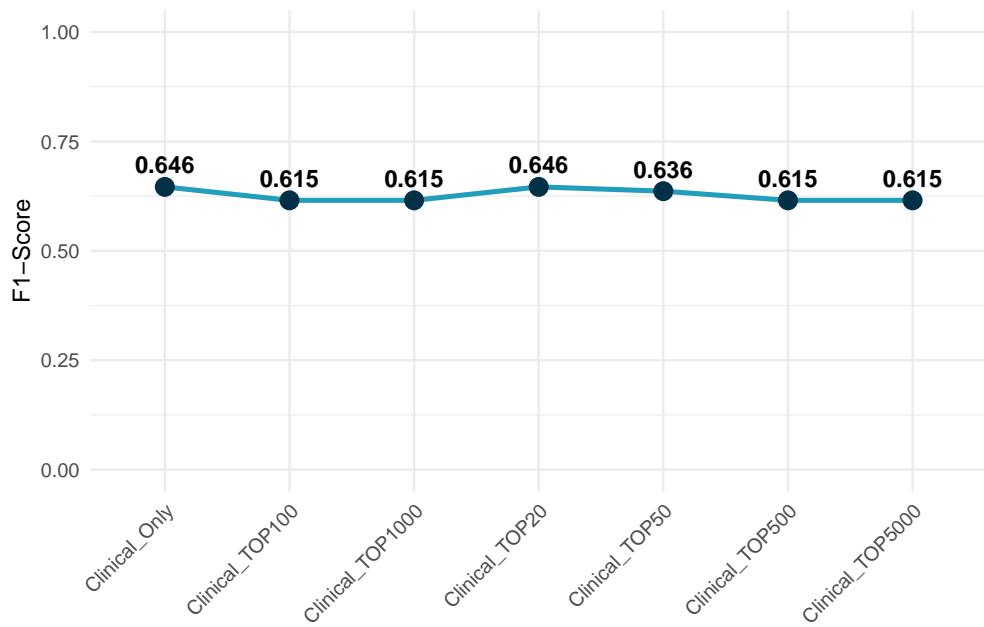
UNILASSO – Classification Metrics

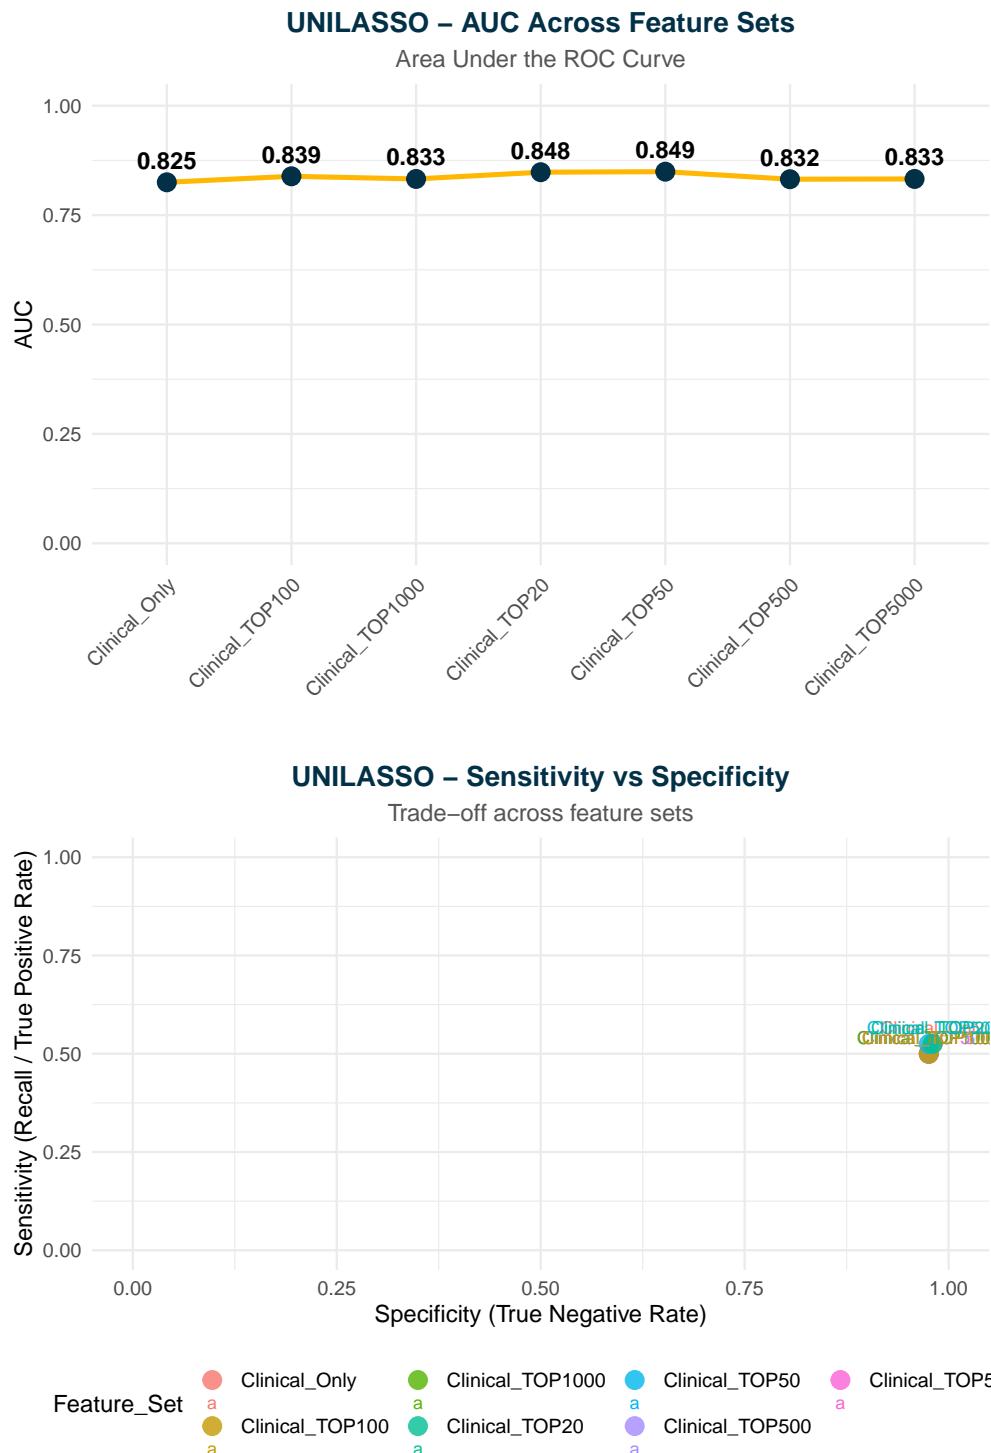
Accuracy, Precision, Recall, F1-Score, AUC



UNILASSO – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1  Clinical_Only 21 202 4 19 0.9065041 0.8400000 0.525 0.9805825
## 2 Clinical_TOP5000 20 201 5 20 0.8983740 0.8000000 0.500 0.9757282
## 3 Clinical_TOP1000 20 201 5 20 0.8983740 0.8000000 0.500 0.9757282
## 4 Clinical_TOP500 20 201 5 20 0.8983740 0.8000000 0.500 0.9757282
```

```

## 5 Clinical_TOP100 20 201 5 20 0.8983740 0.8000000 0.500 0.9757282
## 6 Clinical_TOP50 21 201 5 19 0.9024390 0.8076923 0.525 0.9757282
## 7 Clinical_TOP20 21 202 4 19 0.9065041 0.8400000 0.525 0.9805825
## F1_Score AUC
## 1 0.6461538 0.8250000
## 2 0.6153846 0.8326456
## 3 0.6153846 0.8326456
## 4 0.6153846 0.8320388
## 5 0.6153846 0.8387136
## 6 0.6363636 0.8492718
## 7 0.6461538 0.8480583
##
## Exported classification metrics to: model_metrics/unilasso_classification_metrics.csv

```

uniLasso shows extremely stable performance across all genomic feature sets, with Test AUC consistently around 0.83–0.85 and recall values near 0.50 for nearly all models. Precision remains high (0.79–0.81), and specificity stays above 0.97. Unlike Ridge, uniLasso never collapses under high dimensionality; it consistently identifies the same core signal even when starting from 5000 genes. This reflects the intended effect of the uniLasso procedure (Chatterjee, Hastie & Tibshirani, 2025), where univariate guidance stabilizes variable selection and enforces sign consistency. The clinical-only model performs poorly because uniLasso relies on univariate ranking across many features, which is only meaningful in the genomic setting.

ElasticNet Comparison Across Feature Sets

Elastic Net combines the strengths of both Ridge (L2) and Lasso (L1) penalties, making it well-suited for datasets with correlated gene groups, which is typical in transcriptomic data. Its estimator solves:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \{-l(\beta) + \lambda(\alpha\|\beta\|_1) + (1 - \alpha)\|\beta\|_2^2\}$$

Where

- $\alpha = 1$ is Lasso
- $\alpha = 0$ is Ridge
- $0 < \alpha < 1$ is Elastic Mixed model

```

elasticnet_results <- fit_single_model_across_features(
  model_type = "elasticnet"
  , X_train_all = X_train
  , X_test_all = X_test
  , Y_train = Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

```

```

##
## === FITTING ELASTICNET ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Setting levels: control = 0, case = 1

```

```
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP5000...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP1000...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP500...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP100...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Fitting Clinical_TOP50...

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1
```

```

## Setting direction: controls < cases

## Fitting Clinical_TOP20...

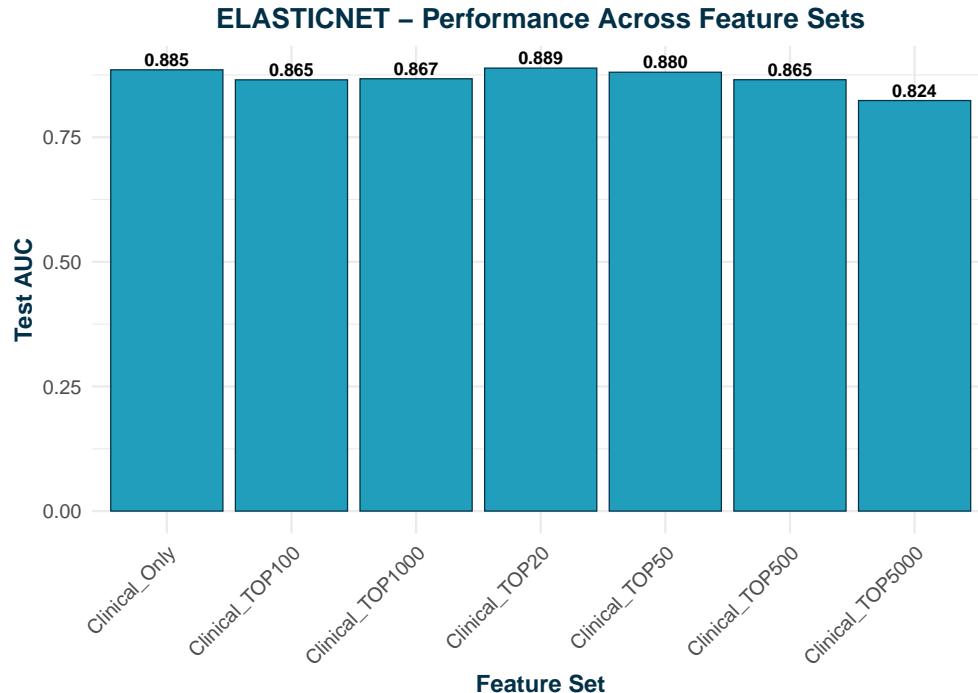
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

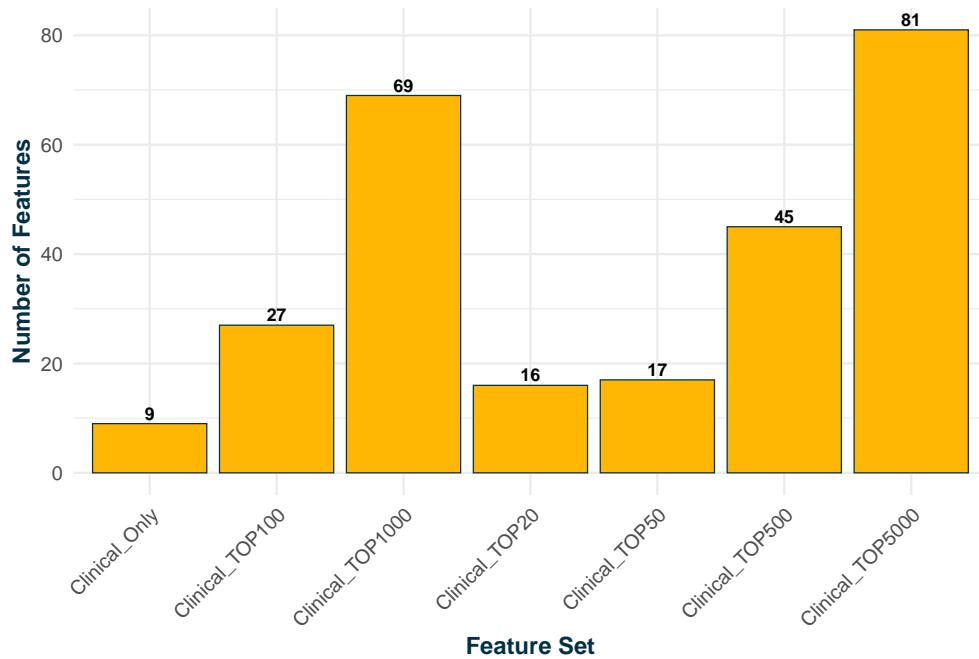
## Setting direction: controls < cases

## 
## === SUMMARY TABLE ===
##      Feature_Set      Model Features Train_AUC  Test_AUC Test_Accuracy
## 1   Clinical_Only ELASTICNET        9 0.8905534 0.8851942    0.8658537
## 2 Clinical_TOP5000 ELASTICNET       81 0.9413221 0.8235437    0.8414634
## 3 Clinical_TOP1000 ELASTICNET       69 0.9401297 0.8672330    0.8536585
## 4 Clinical_TOP500 ELASTICNET       45 0.9227263 0.8652913    0.8699187
## 5 Clinical_TOP100 ELASTICNET       27 0.9045003 0.8650485    0.8821138
## 6 Clinical_TOP50 ELASTICNET       17 0.8926364 0.8803398    0.8780488
## 7 Clinical_TOP20 ELASTICNET       16 0.8882289 0.8885922    0.8861789
## Exported metrics to: model_metrics/elasticnet_across_features_metrics.csv

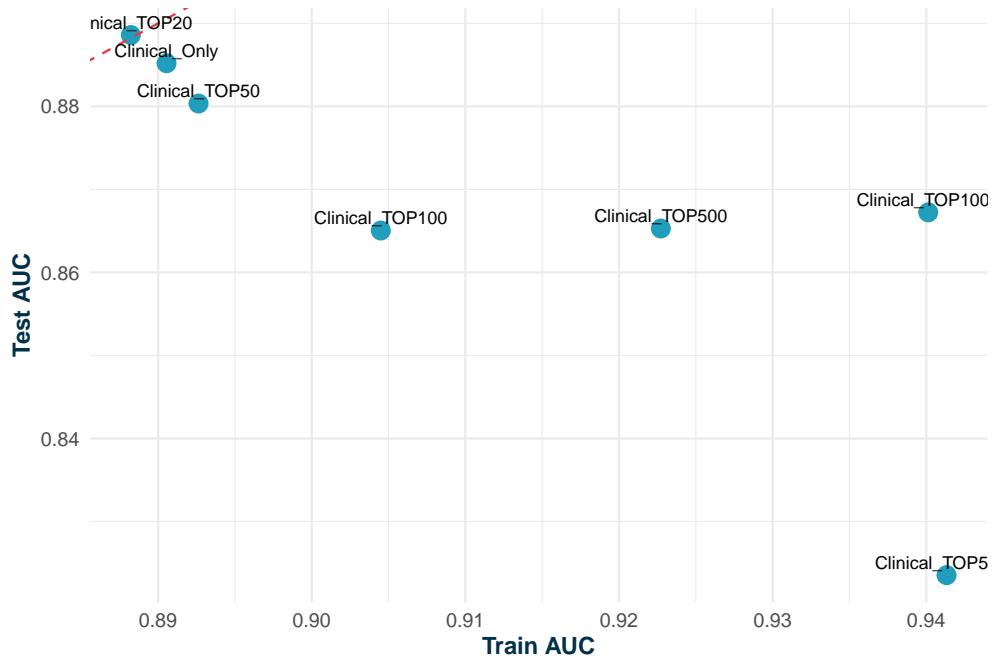
```

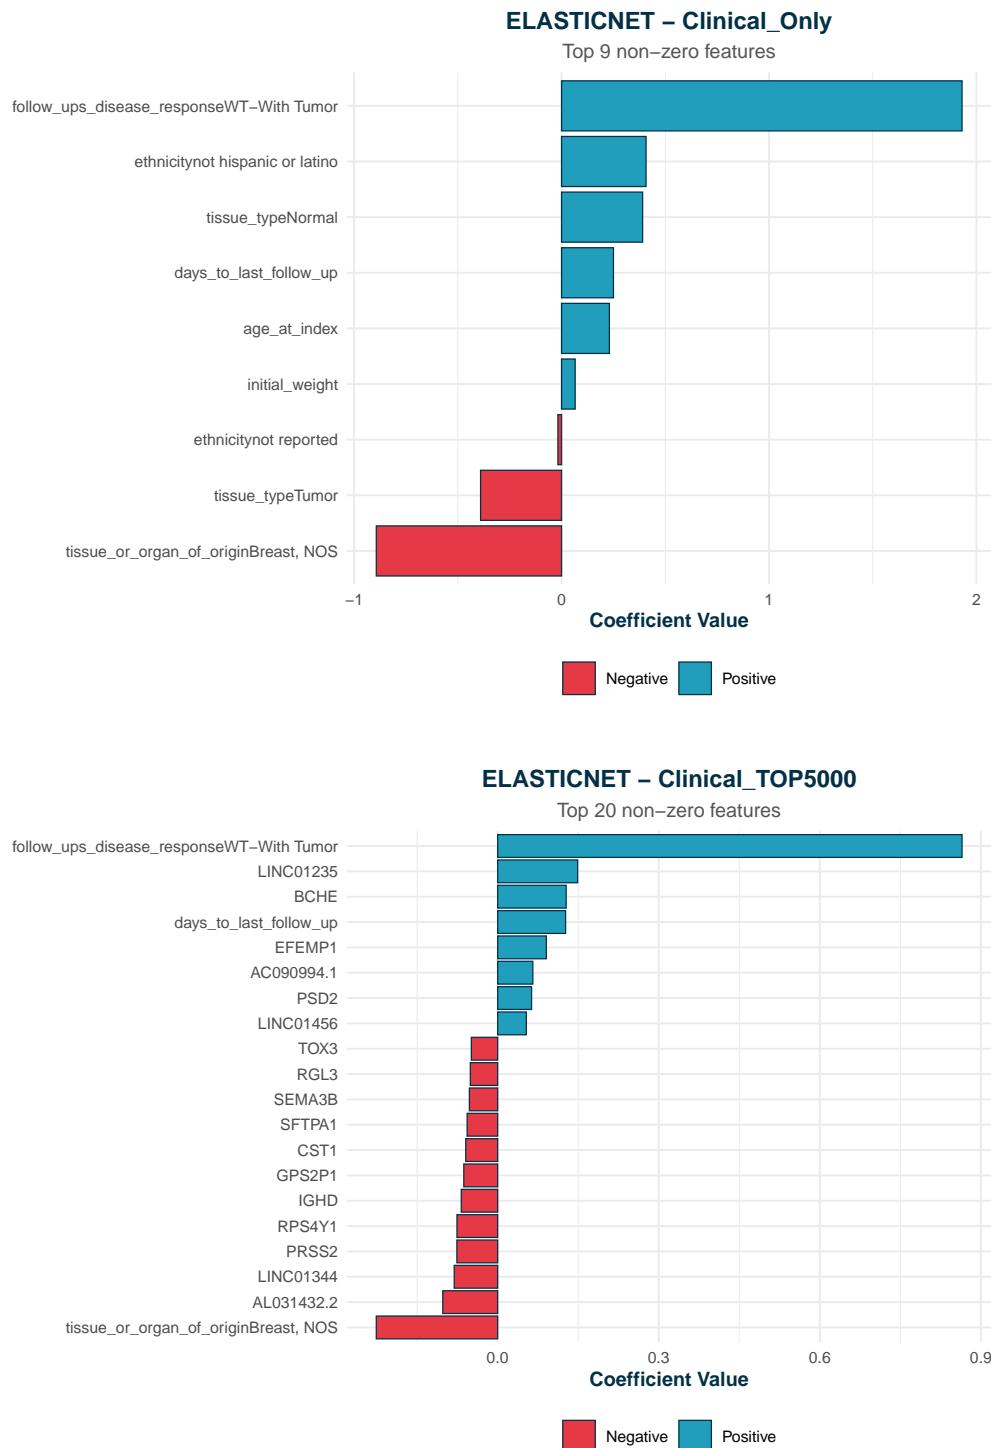


ELASTICNET – Selected Features



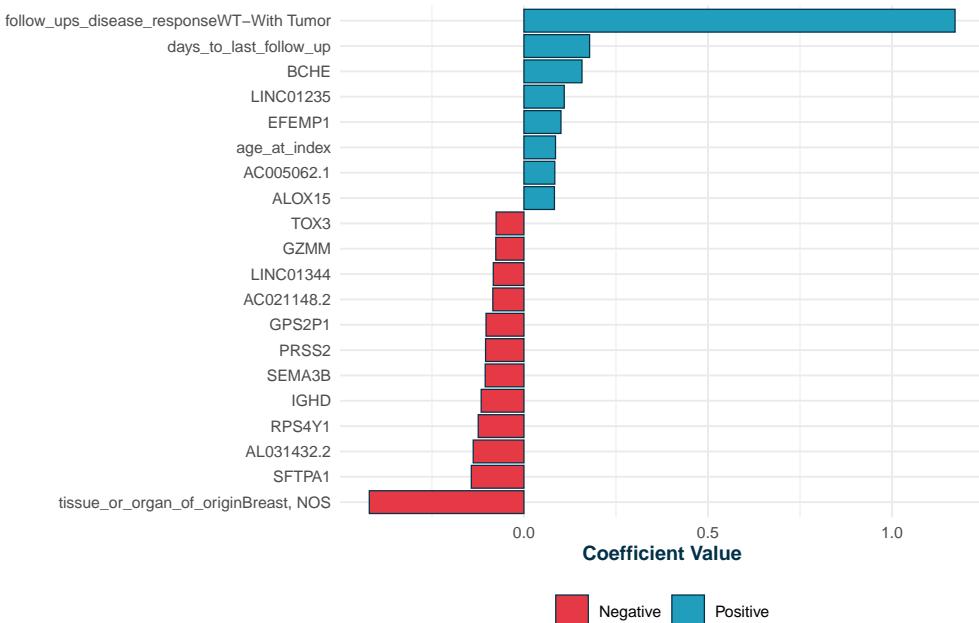
ELASTICNET – Train vs Test AUC





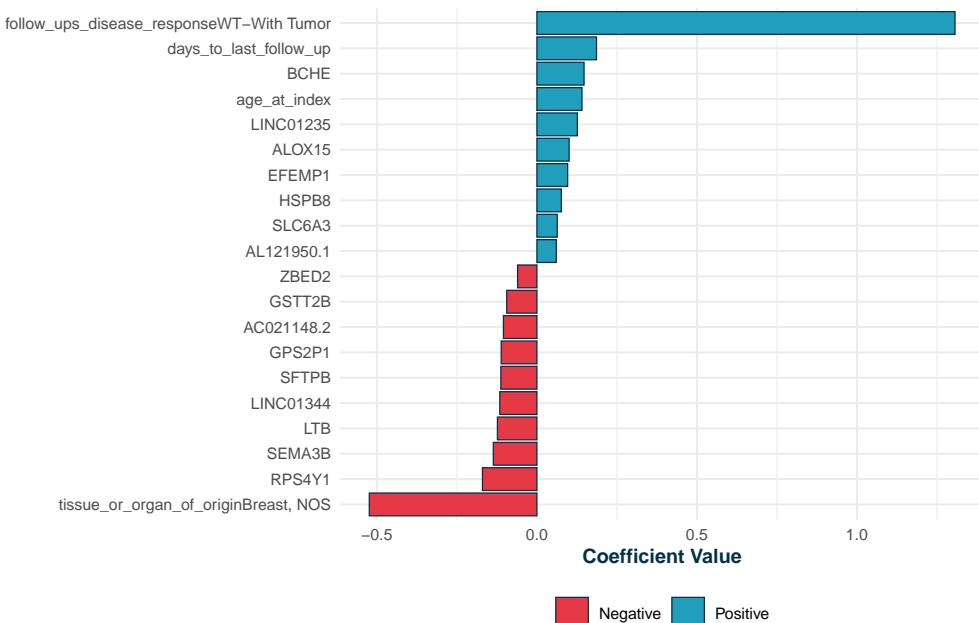
ELASTICNET – Clinical_TOP1000

Top 20 non-zero features



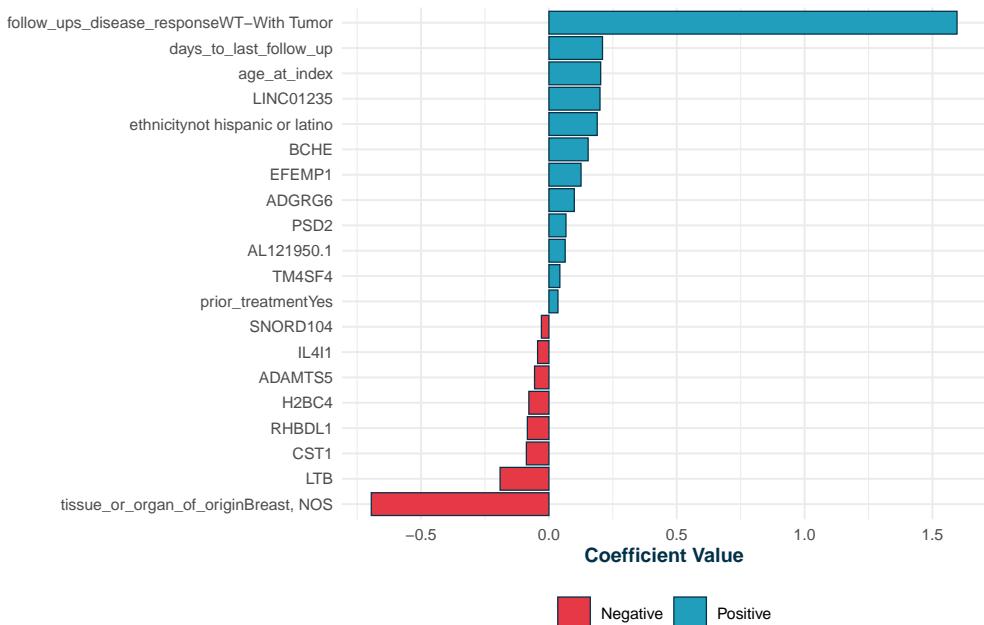
ELASTICNET – Clinical_TOP500

Top 20 non-zero features



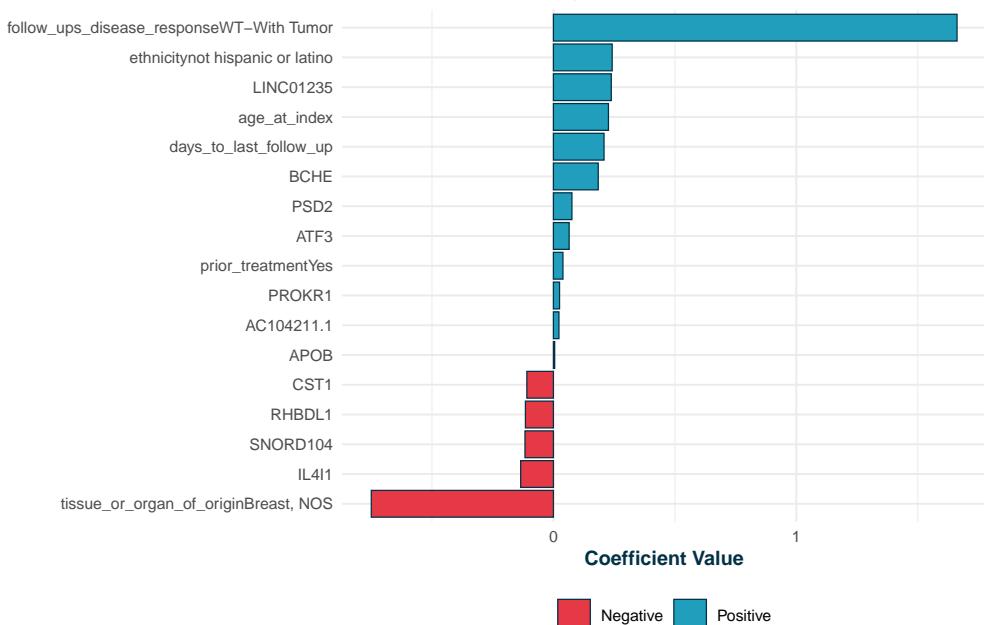
ELASTICNET – Clinical_TOP100

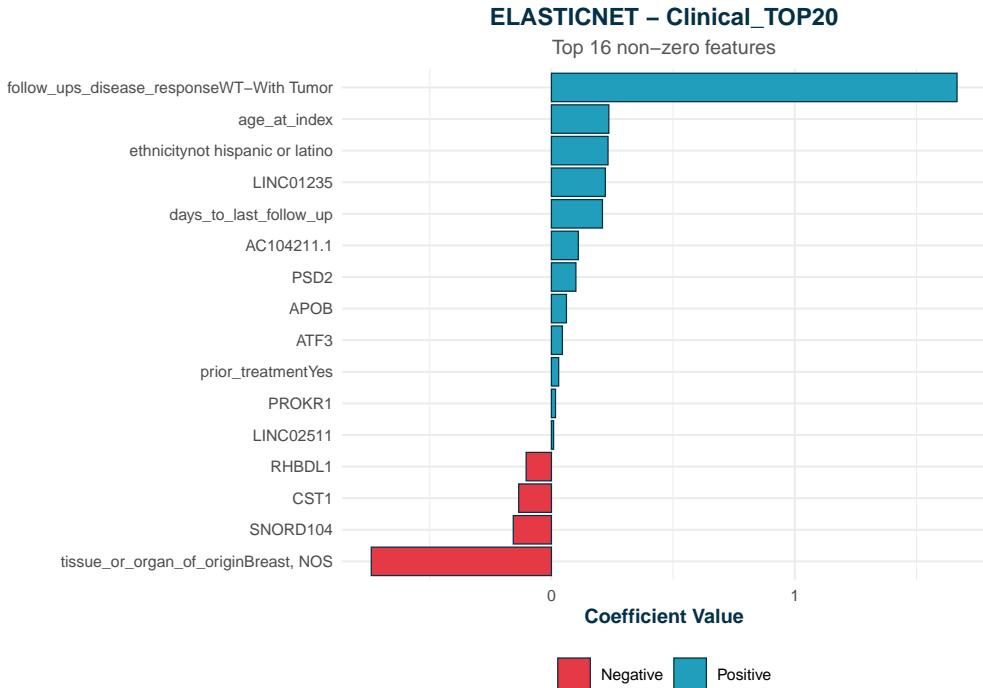
Top 20 non-zero features



ELASTICNET – Clinical_TOP50

Top 17 non-zero features





```

elasticnet_metrics <- plot_classification_metrics_single(elasticnet_results
, threshold = 0.5
, csv_filename = "elasticnet_classification_me"

## 
## === CLASSIFICATION METRICS ===

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_Only:
##   TP=10 TN=203 FP=3 FN=30
##   Accuracy=0.866 Precision=0.769 Recall=0.250 F1=0.377 AUC=0.885

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP5000:
##   TP=2 TN=205 FP=1 FN=38
##   Accuracy=0.841 Precision=0.667 Recall=0.050 F1=0.093 AUC=0.824

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP1000:
##   TP=5 TN=205 FP=1 FN=35
##   Accuracy=0.854 Precision=0.833 Recall=0.125 F1=0.217 AUC=0.867

```

```

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP500:
##   TP=9 TN=205 FP=1 FN=31
##   Accuracy=0.870 Precision=0.900 Recall=0.225 F1=0.360 AUC=0.865

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP100:
##   TP=12 TN=205 FP=1 FN=28
##   Accuracy=0.882 Precision=0.923 Recall=0.300 F1=0.453 AUC=0.865

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

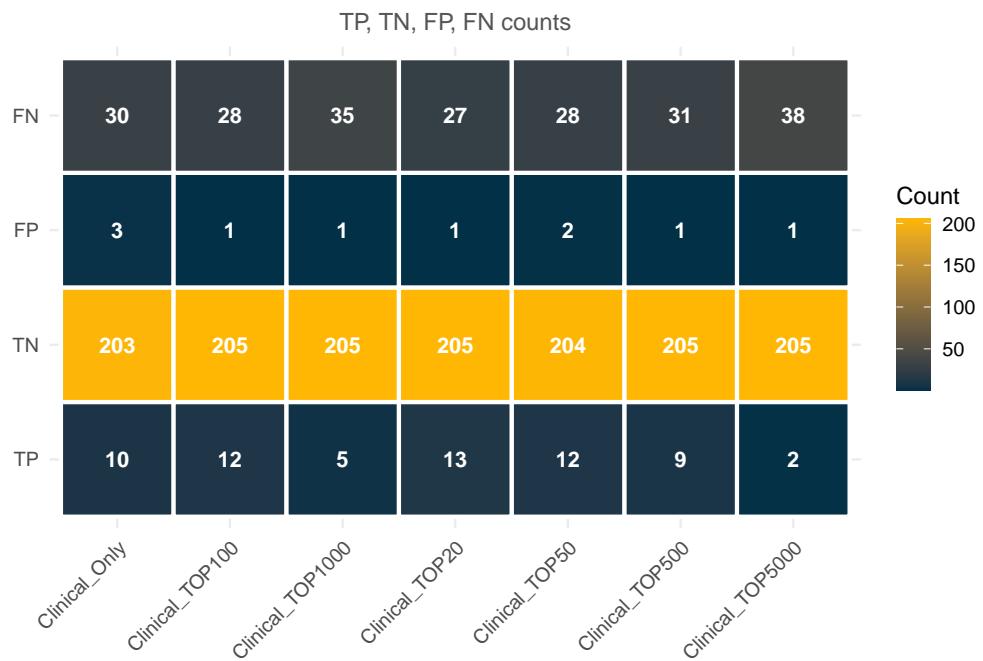
## Clinical_TOP50:
##   TP=12 TN=204 FP=2 FN=28
##   Accuracy=0.878 Precision=0.857 Recall=0.300 F1=0.444 AUC=0.880

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

## Clinical_TOP20:
##   TP=13 TN=205 FP=1 FN=27
##   Accuracy=0.886 Precision=0.929 Recall=0.325 F1=0.481 AUC=0.889

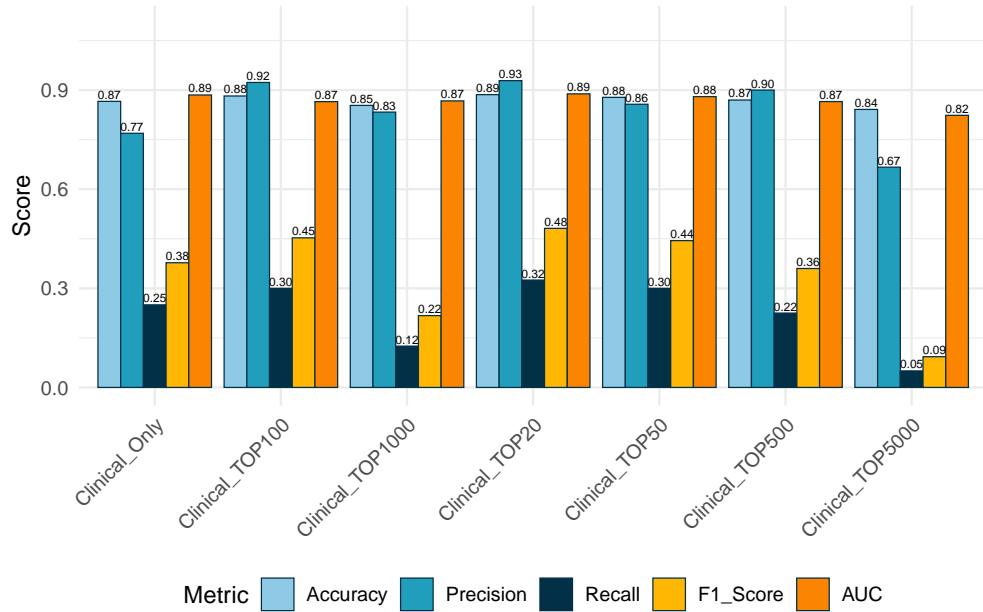
```

ELASTICNET – Confusion Matrix Across Feature Sets



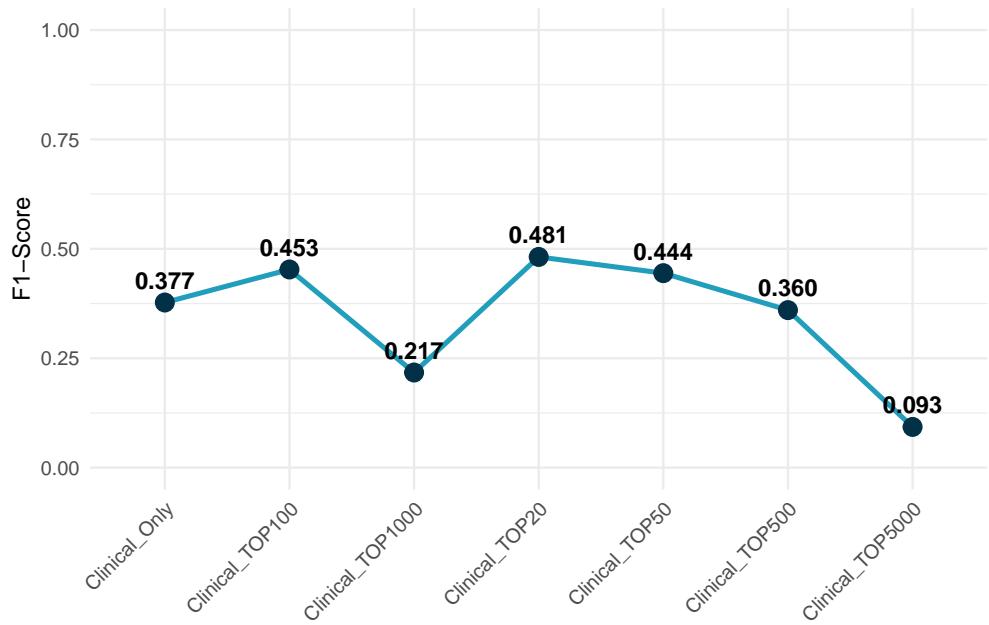
ELASTICNET – Classification Metrics

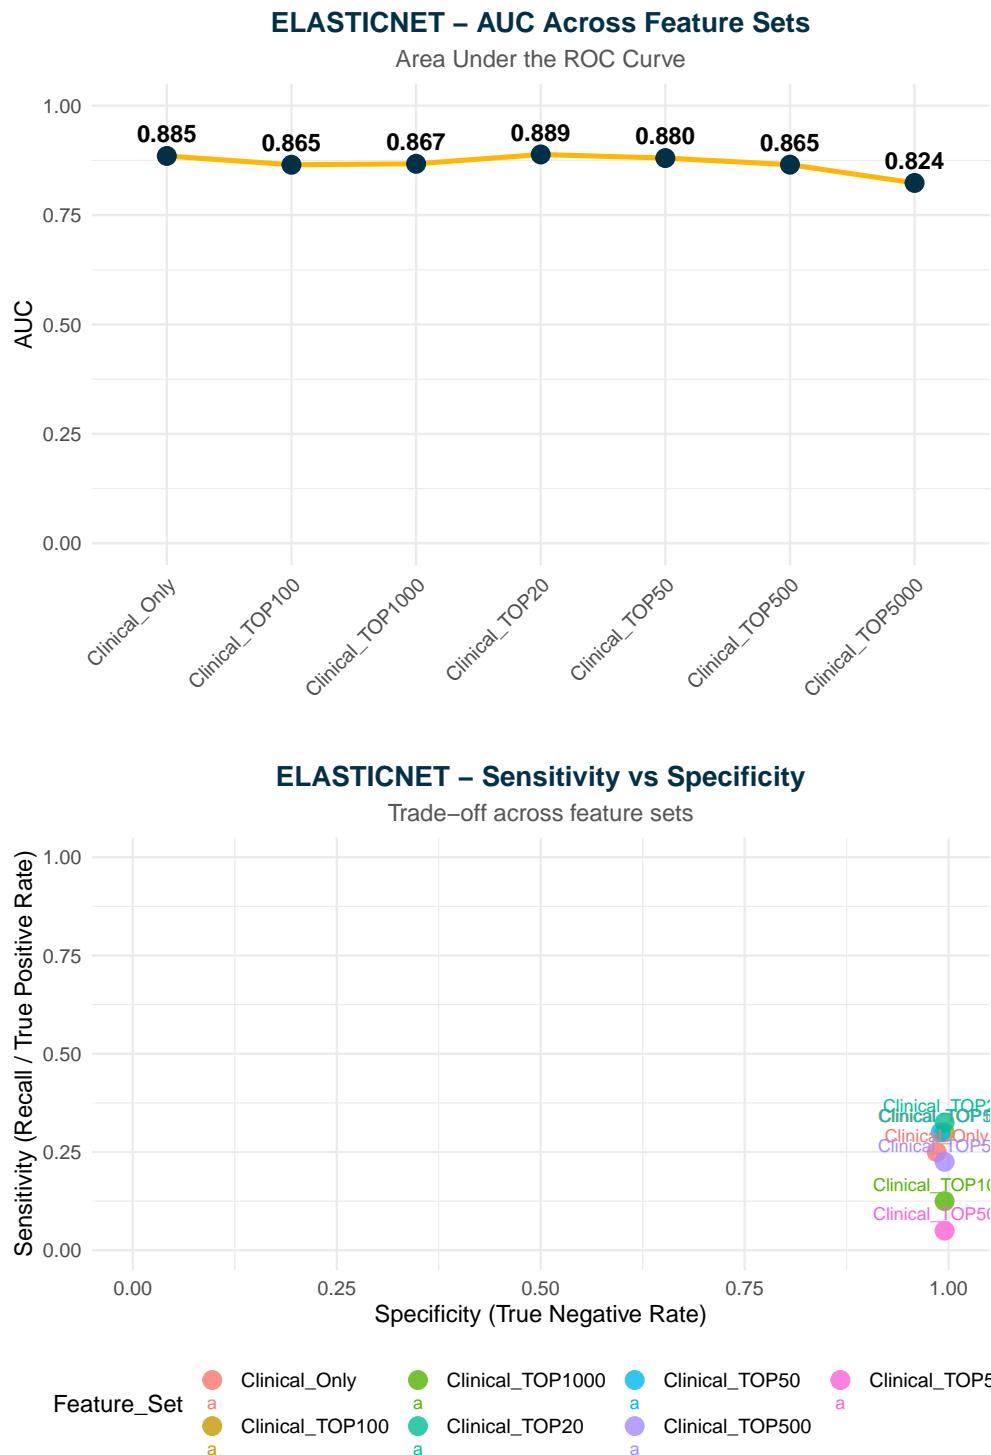
Accuracy, Precision, Recall, F1-Score, AUC



ELASTICNET – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1   Clinical_Only 10 203 3 30 0.8658537 0.7692308 0.250 0.9854369
## 2 Clinical_TOP5000  2 205 1 38 0.8414634 0.6666667 0.050 0.9951456
## 3 Clinical_TOP1000  5 205 1 35 0.8536585 0.8333333 0.125 0.9951456
## 4 Clinical_TOP500   9 205 1 31 0.8699187 0.9000000 0.225 0.9951456
```

```

## 5 Clinical_TOP100 12 205 1 28 0.8821138 0.9230769 0.300 0.9951456
## 6 Clinical_TOP50 12 204 2 28 0.8780488 0.8571429 0.300 0.9902913
## 7 Clinical_TOP20 13 205 1 27 0.8861789 0.9285714 0.325 0.9951456
## F1_Score AUC
## 1 0.37735849 0.8851942
## 2 0.09302326 0.8235437
## 3 0.21739130 0.8672330
## 4 0.36000000 0.8652913
## 5 0.45283019 0.8650485
## 6 0.44444444 0.8803398
## 7 0.48148148 0.8885922
##
## Exported classification metrics to: model_metrics/elasticnet_classification_metrics.csv

```

Elastic Net achieves consistently strong performance across all medium-sized gene sets, with test AUC values in the 0.86–0.89 range. It improves stability over Lasso when correlated genes are present and avoids the total collapse seen in Ridge when dimensionality is high. Precision is consistently high, while recall remains low but stable, reflecting conservative predictions in an imbalanced dataset. Elastic Net works best for gene subsets between 20 and 1000 genes, where it captures correlated genomic structure without being overwhelmed by noise.

Class Imbalance Handling with SMOTE

Current imbalance ratio: 5.12:1 (Alive:Dead)

Problem observed in baseline models: - Ridge: Recall = 0.00-0.27 (missing 73-100% of Dead patients) - Models biased toward majority class (Alive) - High Specificity (99%) but very low Recall (22%) - Clinical_TOP5000: Predicts 0 Dead patients (completely fails)

Why SMOTE: - Creates synthetic minority class samples (Dead patients) - Balances training data to ~1:1 ratio - Forces models to learn Dead patient patterns - No data loss (vs downsampling) - Prevents overfitting (vs simple upsampling)

Expected improvements: - Recall: 0.22 -> 0.50-0.70 (detect more Dead patients) - F1-Score: 0.36 -> 0.50+ (better balance) - Trade-off: Specificity may drop from 99% to 85-90% (acceptable)

Models selected for SMOTE testing: 1. Ridge - worst Recall performance, needs urgent fix 2. Lasso - feature selection sensitive to imbalance 3. ElasticNet - combination of L1/L2, middle priority

```

cat("== CLASS IMBALANCE ANALYSIS ==\n")

## == CLASS IMBALANCE ANALYSIS ==

cat("Training set imbalance:\n")

## Training set imbalance:

cat("  Alive:", sum(Y_train == 0), sprintf("(%.1f%%)\n", 100 * sum(Y_train == 0) / length(Y_train)))

##  Alive: 823 (83.6%)

```

```

cat("  Dead:", sum(Y_train == 1), sprintf("(%.1f%%)\n", 100 * sum(Y_train == 1) / length(Y_train)))

##  Dead: 161 (16.4%)

cat("  Ratio:", sprintf("%.2f:1\n\n", sum(Y_train == 0) / sum(Y_train == 1)))

##  Ratio: 5.11:1

smote_data <- apply_smote(X_train, Y_train, k = 5)

##
## === APPLYING SMOTE ===
## Before SMOTE:
##   Alive (0): 823
##   Dead (1): 161
##   Ratio: 5.11 :1
##
## After SMOTE:
##   Alive (0): 823
##   Dead (1): 805
##   Ratio: 1.02 :1
##   Total samples: 1628

```

Logistic with SMOTE

```

logistic_smote <- fit_single_model_across_features(
  model_type = "logistic"
  , X_train_all = smote_data$X_train
  , X_test_all = X_test
  , Y_train = smote_data$Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(100, 50, 20)
)

##
## === FITTING LOGISTIC ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...

## Fitting Clinical_TOP100...

## Fitting Clinical_TOP50...

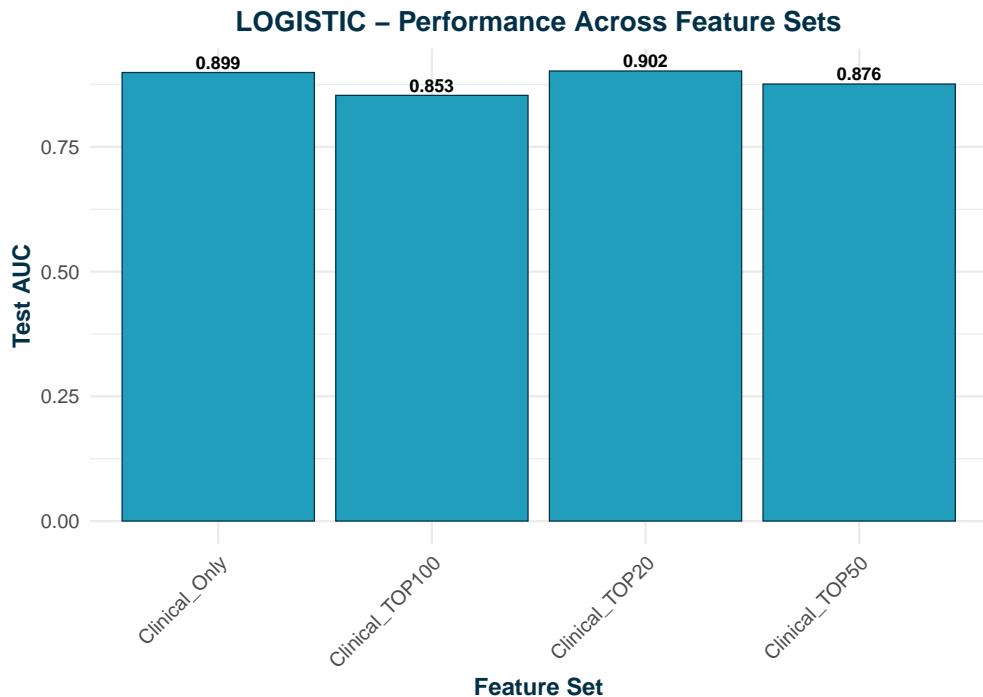
## Fitting Clinical_TOP20...

```

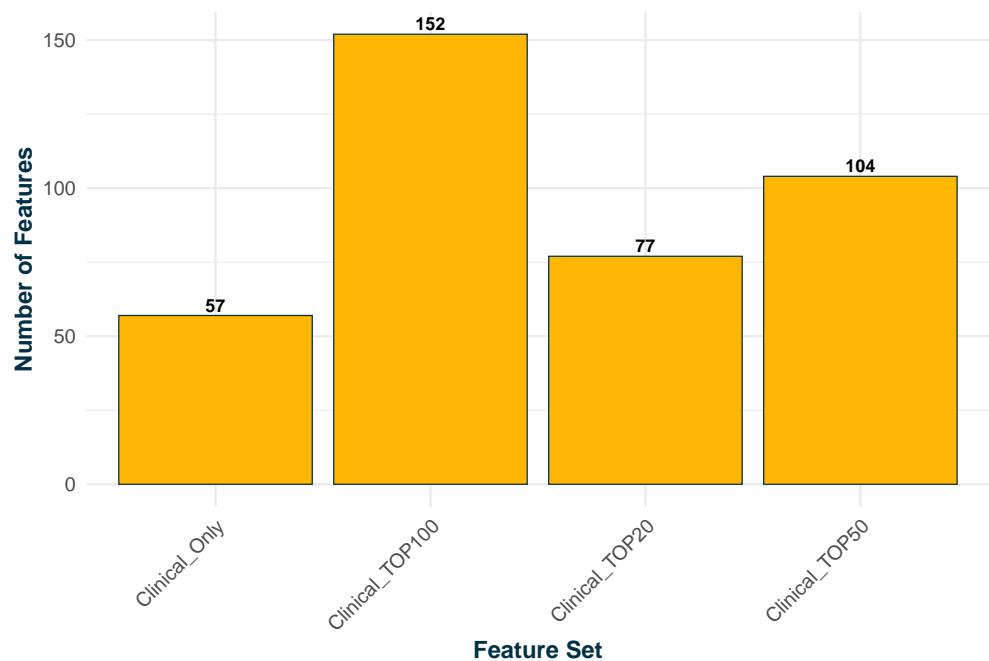
```

## 
## === SUMMARY TABLE ===
##   Feature_Set    Model Features Train_AUC  Test_AUC Test_Accuracy
## 1 Clinical_Only LOGISTIC      57 0.9573096 0.8990291 0.8821138
## 2 Clinical_TOP100 LOGISTIC    152 0.9845498 0.8531553 0.8373984
## 3 Clinical_TOP50 LOGISTIC    104 0.9716565 0.8759709 0.8617886
## 4 Clinical_TOP20 LOGISTIC    77 0.9639721 0.9020631 0.8943089
## Exported metrics to: model_metrics/logistic_across_features_metrics.csv

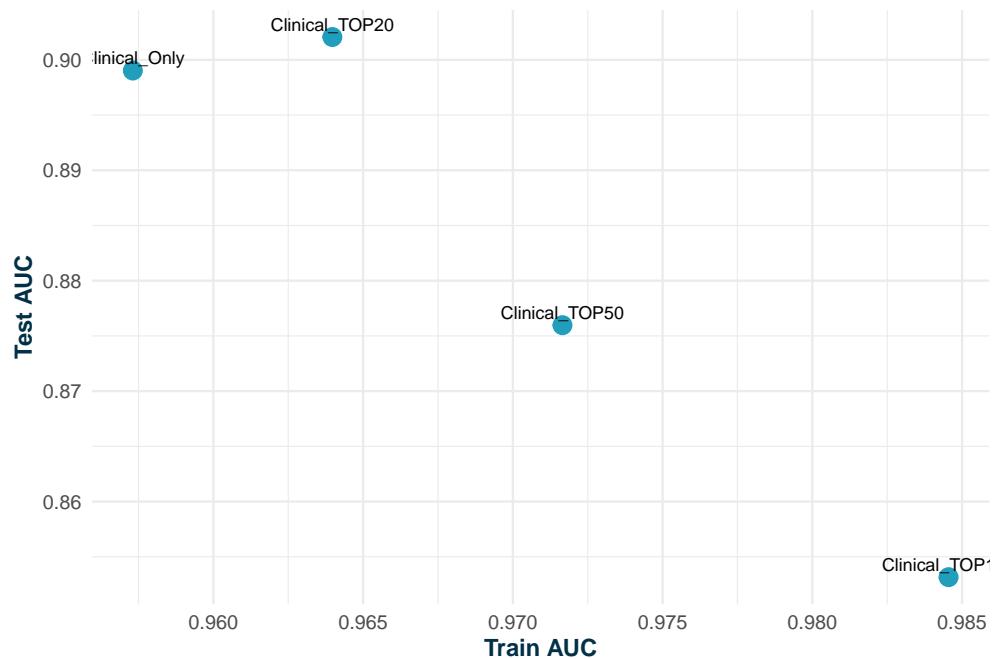
```



LOGISTIC – Selected Features

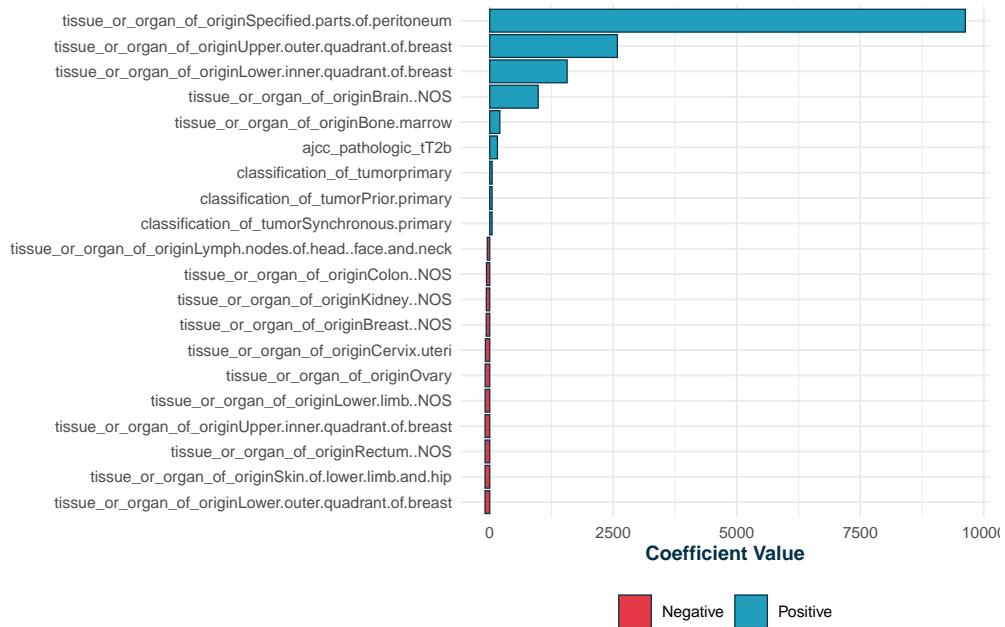


LOGISTIC – Train vs Test AUC



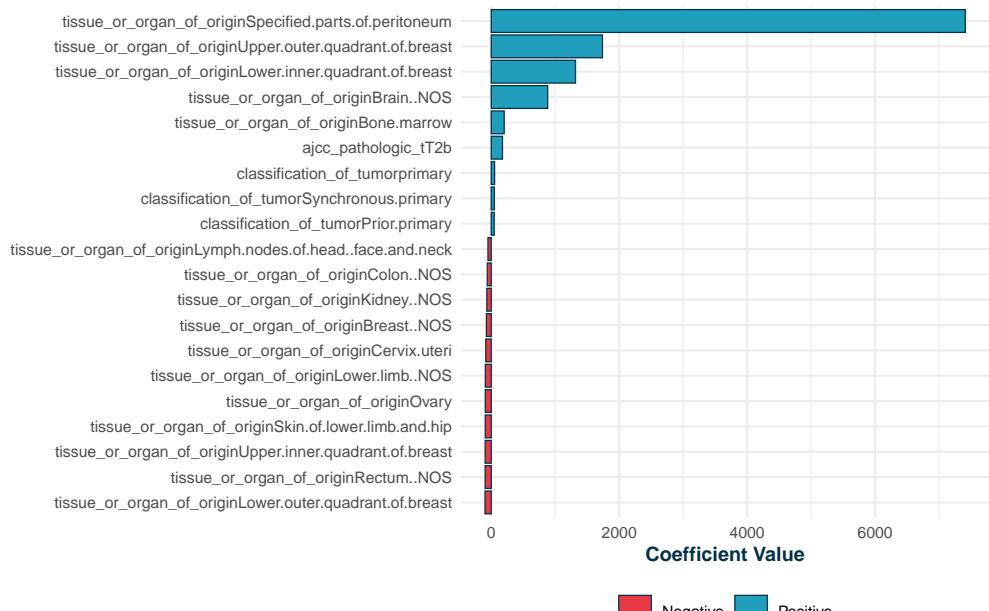
LOGISTIC – Clinical_Only

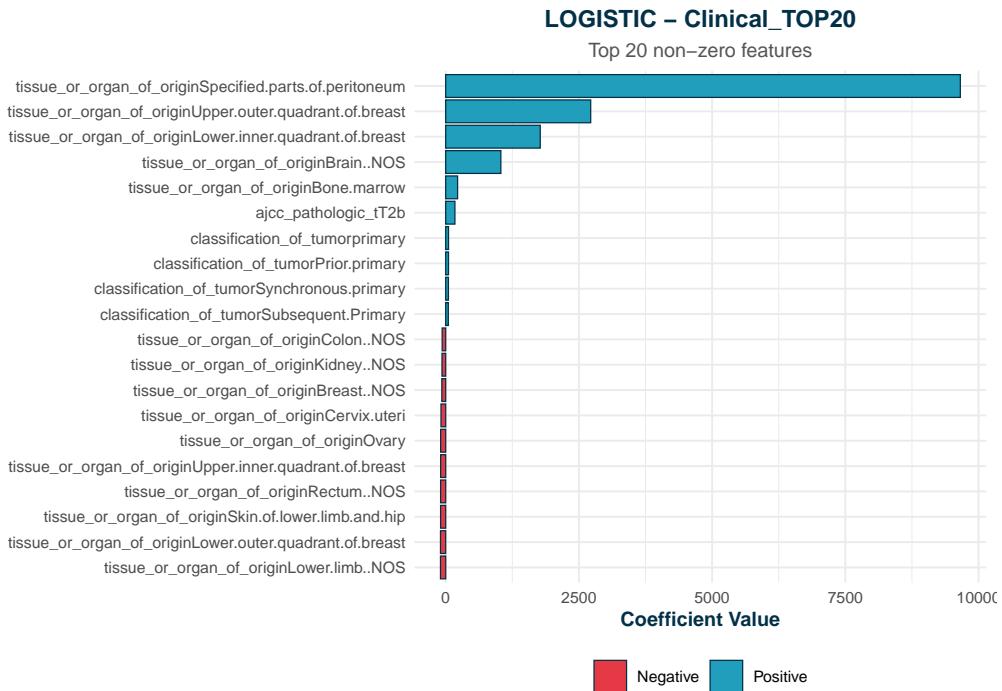
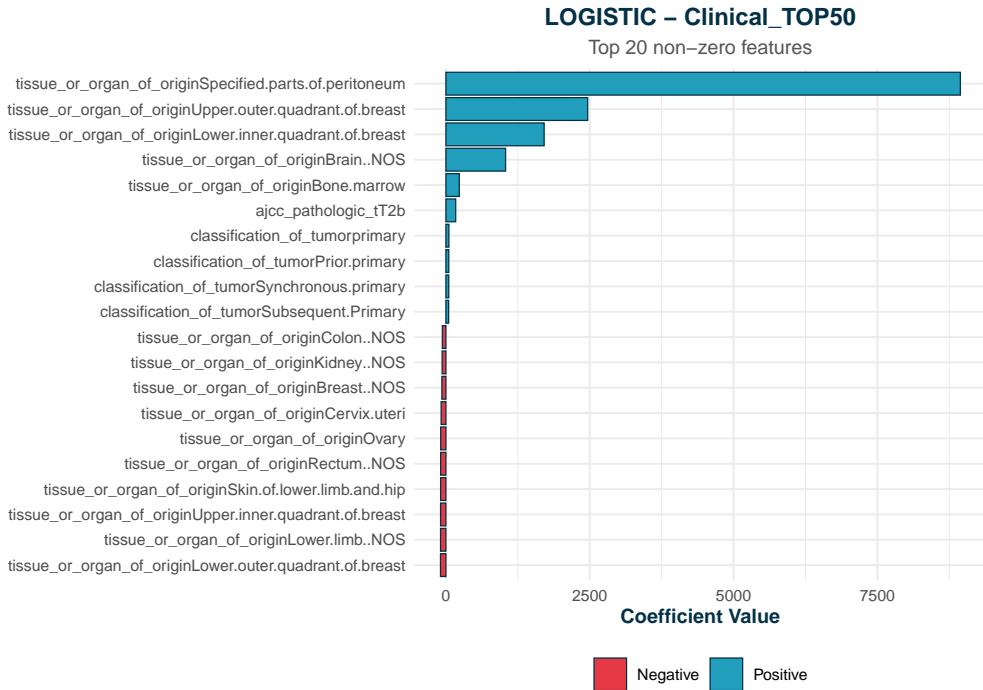
Top 20 non-zero features



LOGISTIC – Clinical_TOP100

Top 20 non-zero features





```
logistic_smote_metrics <- plot_classification_metrics_single(logistic_smote
, threshold = 0.5
, csv_filename = "logistic_smote_classification_metrics.csv")
```

```
## 
## === CLASSIFICATION METRICS ===
```

```

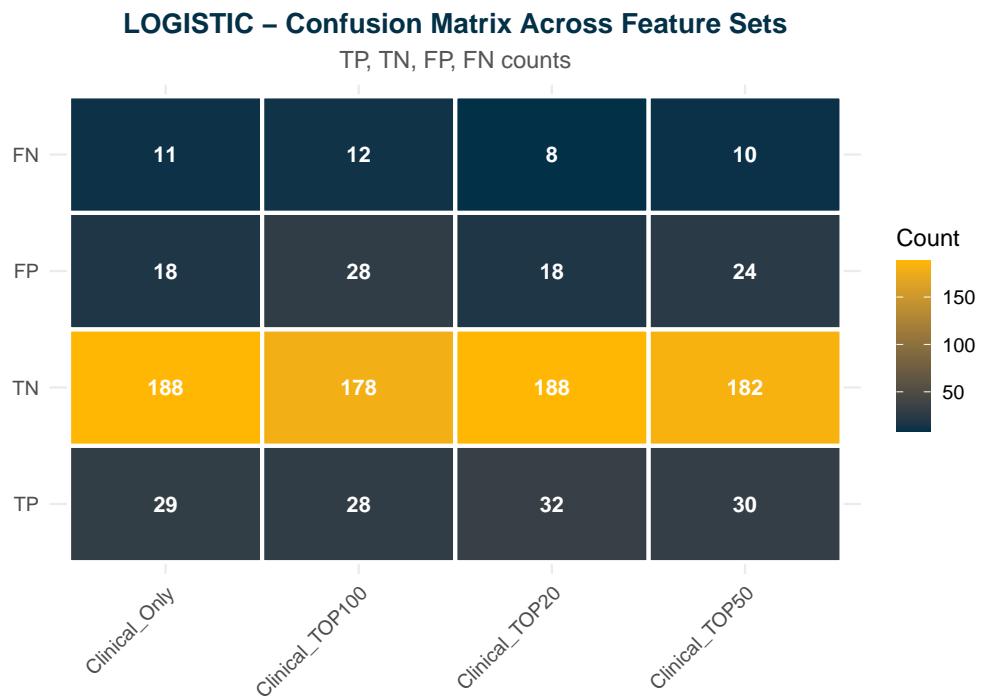
## Clinical_Only:
##   TP=29 TN=188 FP=18 FN=11
##   Accuracy=0.882 Precision=0.617 Recall=0.725 F1=0.667 AUC=0.899

## Clinical_TOP100:
##   TP=28 TN=178 FP=28 FN=12
##   Accuracy=0.837 Precision=0.500 Recall=0.700 F1=0.583 AUC=0.853

## Clinical_TOP50:
##   TP=30 TN=182 FP=24 FN=10
##   Accuracy=0.862 Precision=0.556 Recall=0.750 F1=0.638 AUC=0.876

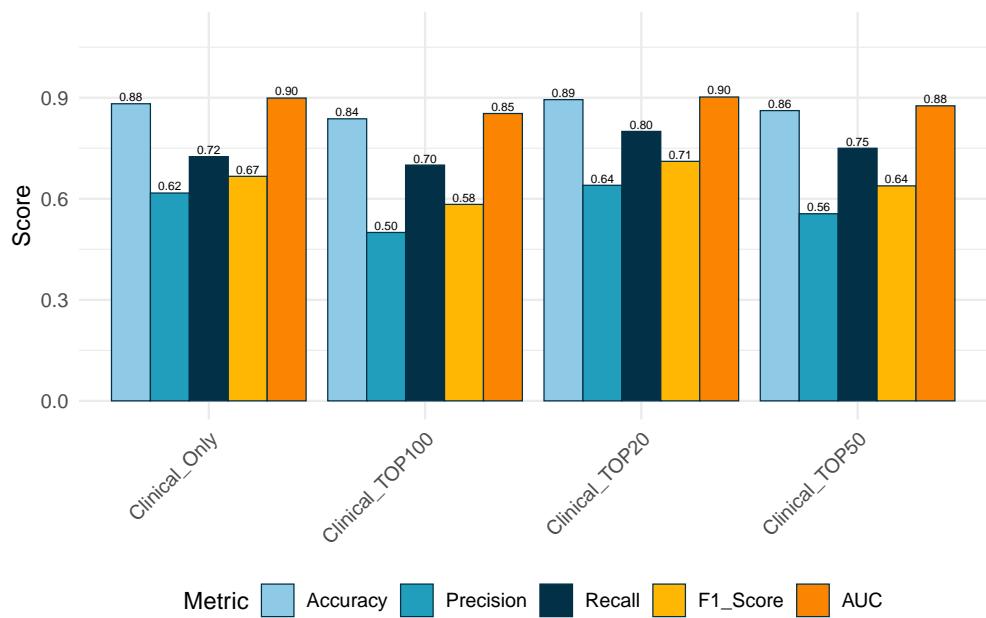
## Clinical_TOP20:
##   TP=32 TN=188 FP=18 FN=8
##   Accuracy=0.894 Precision=0.640 Recall=0.800 F1=0.711 AUC=0.902

```



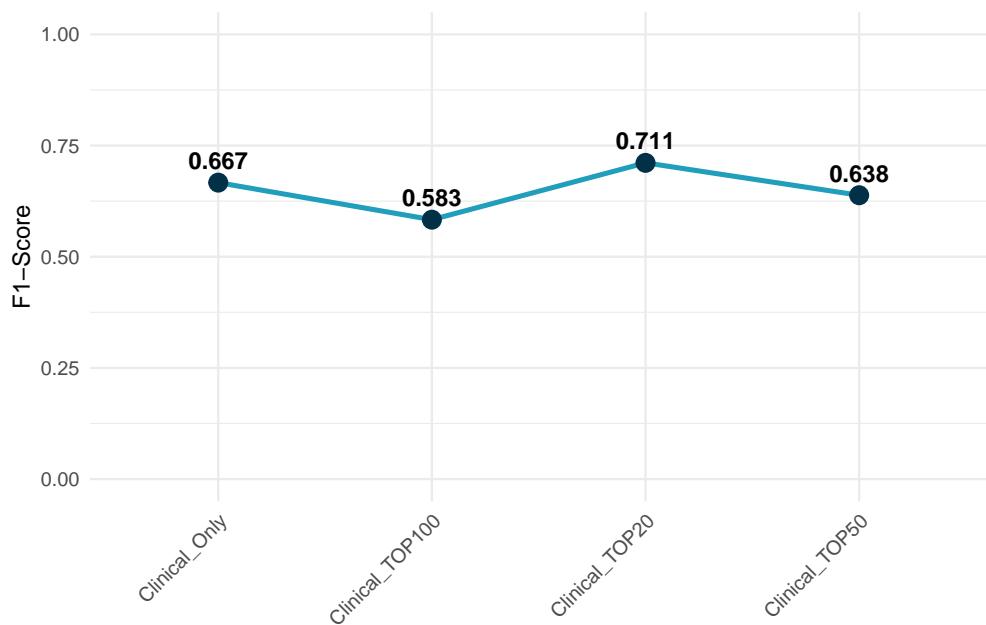
LOGISTIC – Classification Metrics

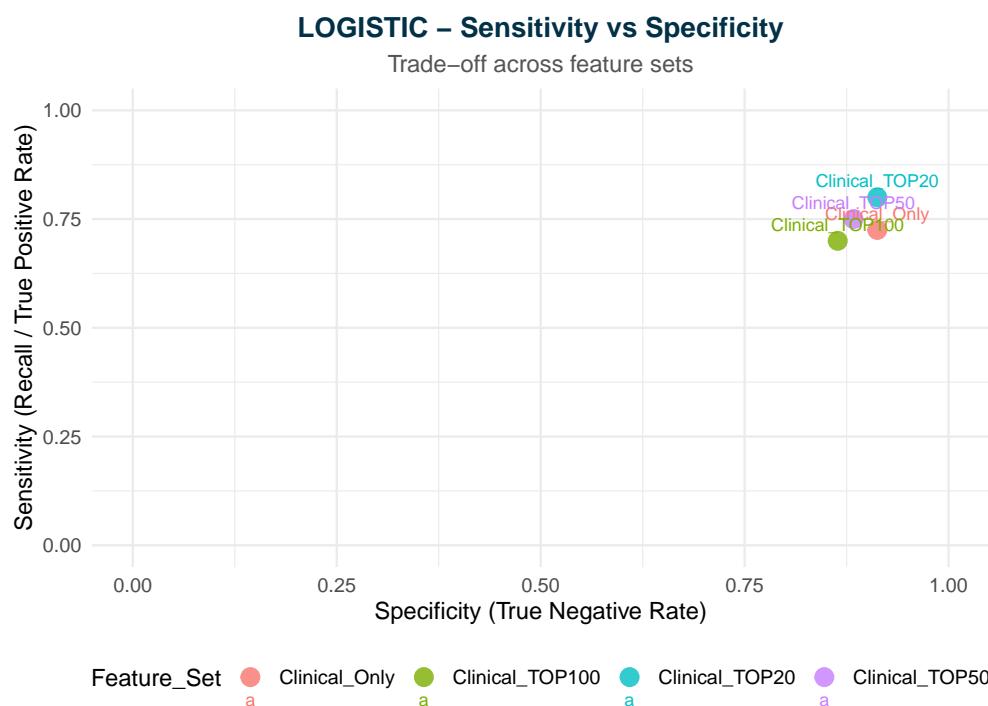
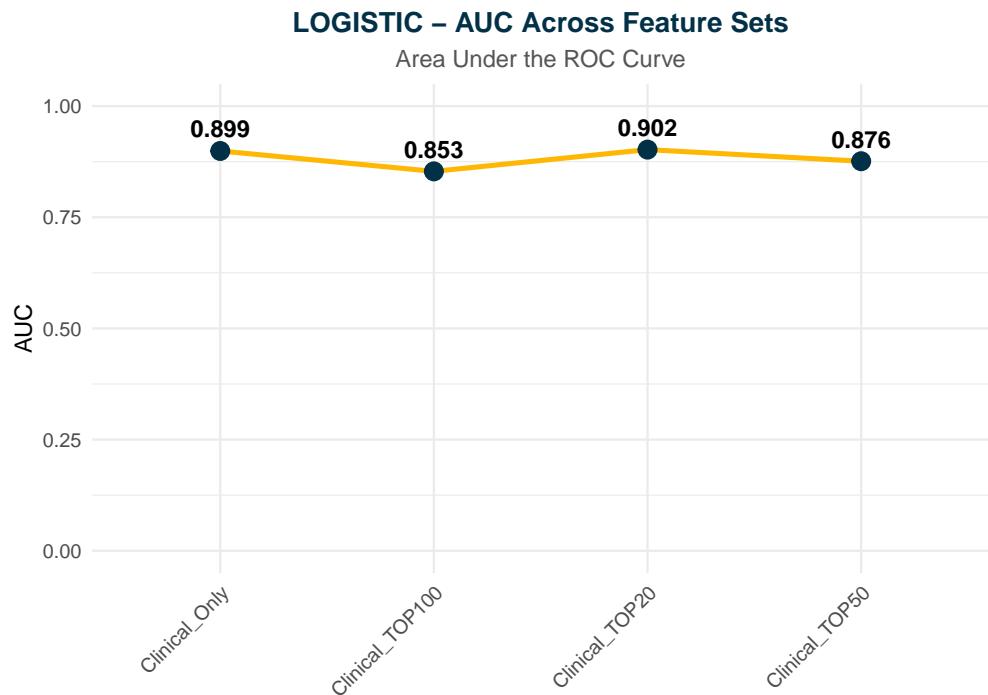
Accuracy, Precision, Recall, F1-Score, AUC



LOGISTIC – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP  TN  FP  FN  Accuracy Precision Recall Specificity F1_Score
## 1  Clinical_Only 29 188 18 11 0.8821138 0.6170213 0.725 0.9126214 0.6666667
## 2 Clinical_TOP100 28 178 28 12 0.8373984 0.5000000 0.700 0.8640777 0.5833333
## 3 Clinical_TOP50 30 182 24 10 0.8617886 0.5555556 0.750 0.8834951 0.6382979
## 4 Clinical_TOP20 32 188 18  8 0.8943089 0.6400000 0.800 0.9126214 0.7111111
```

```

##          AUC
## 1 0.8990291
## 2 0.8531553
## 3 0.8759709
## 4 0.9020631
##
## Exported classification metrics to: model_metrics/logistic_smote_classification_metrics.csv

```

Ridge with SMOTE

```

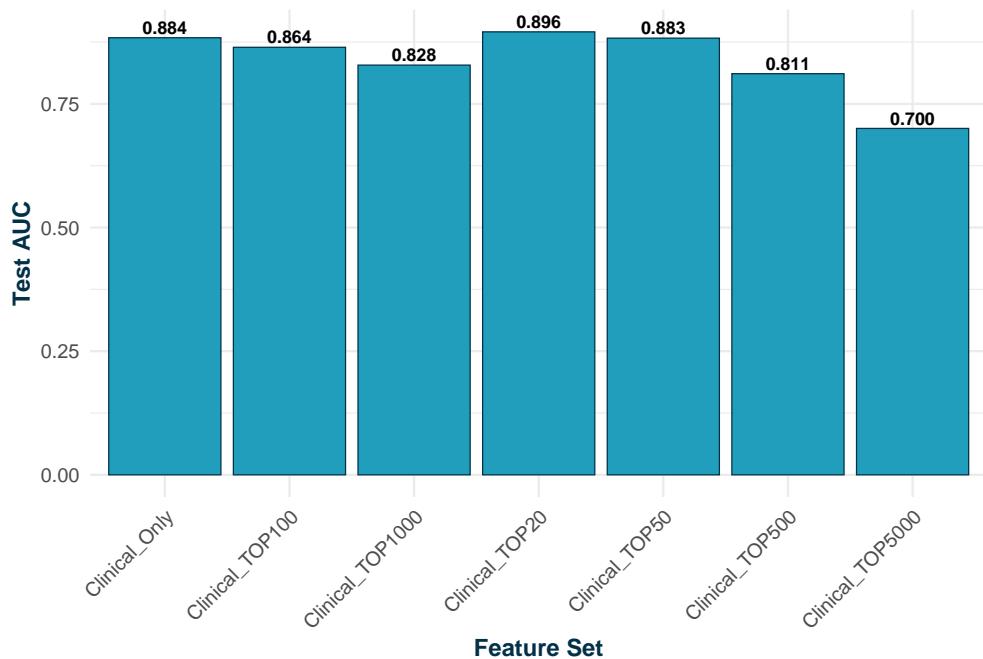
ridge_smote <- fit_single_model_across_features(
  model_type = "ridge"
  , X_train_all = smote_data$X_train
  , X_test_all = X_test
  , Y_train = smote_data$Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

##
## === FITTING RIDGE ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Fitting Clinical_TOP5000...
## Fitting Clinical_TOP1000...
## Fitting Clinical_TOP500...
## Fitting Clinical_TOP100...
## Fitting Clinical_TOP50...
## Fitting Clinical_TOP20...

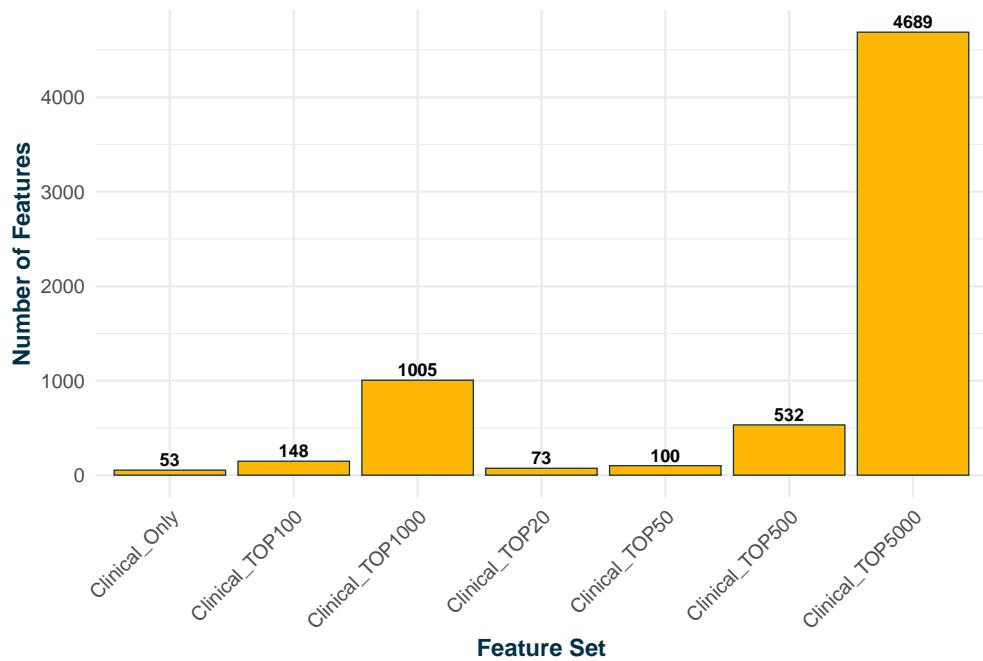
##
## === SUMMARY TABLE ===
##      Feature_Set Model Features Train_AUC Test_AUC Test_Accuracy
## 1    Clinical_Only RIDGE      53 0.9292484 0.8837379     0.8577236
## 2 Clinical_TOP5000 RIDGE    4689 0.9909693 0.7003641     0.7479675
## 3 Clinical_TOP1000 RIDGE   1005 0.9999125 0.8283981     0.8373984
## 4 Clinical_TOP500 RIDGE    532 0.9958703 0.8110437     0.8130081
## 5 Clinical_TOP100 RIDGE    148 0.9582802 0.8644417     0.8414634
## 6 Clinical_TOP50 RIDGE    100 0.9440571 0.8828883     0.8536585
## 7 Clinical_TOP20 RIDGE     73 0.9325766 0.8956311     0.8577236
## Exported metrics to: model_metrics/ridge_across_features_metrics.csv

```

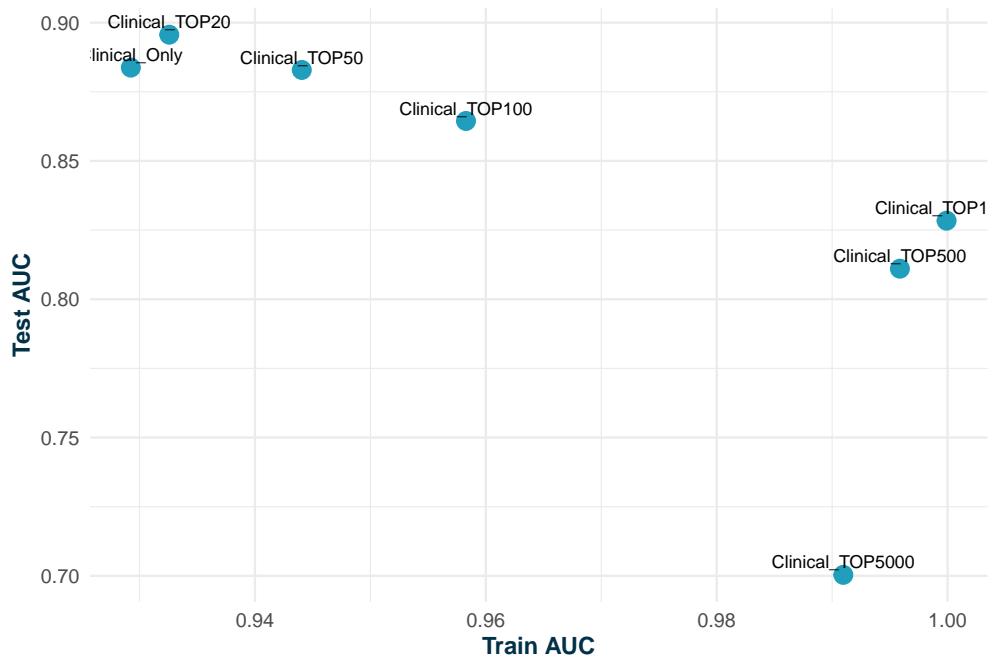
RIDGE – Performance Across Feature Sets



RIDGE – Selected Features

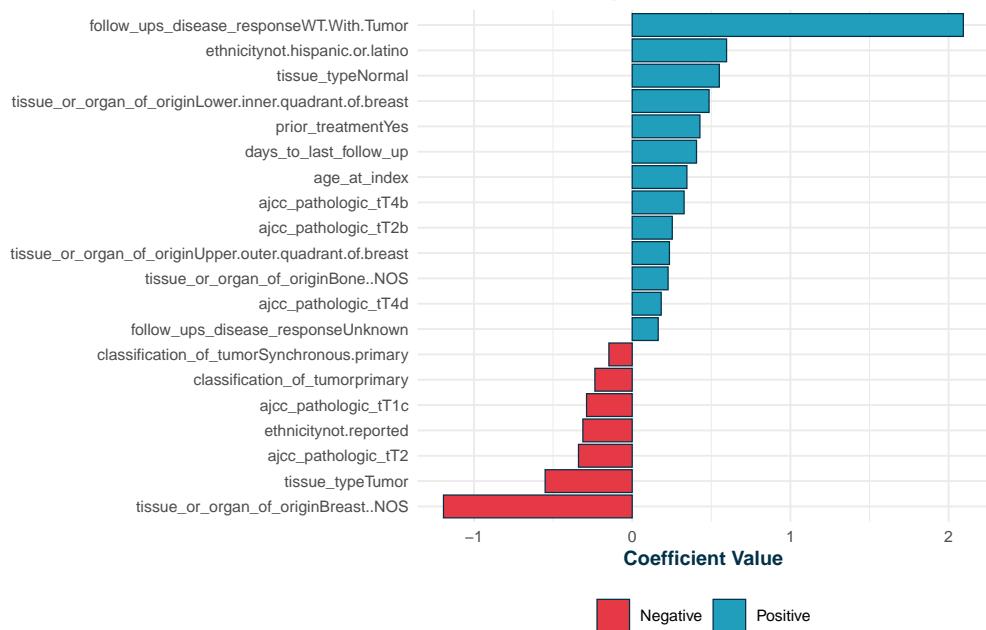


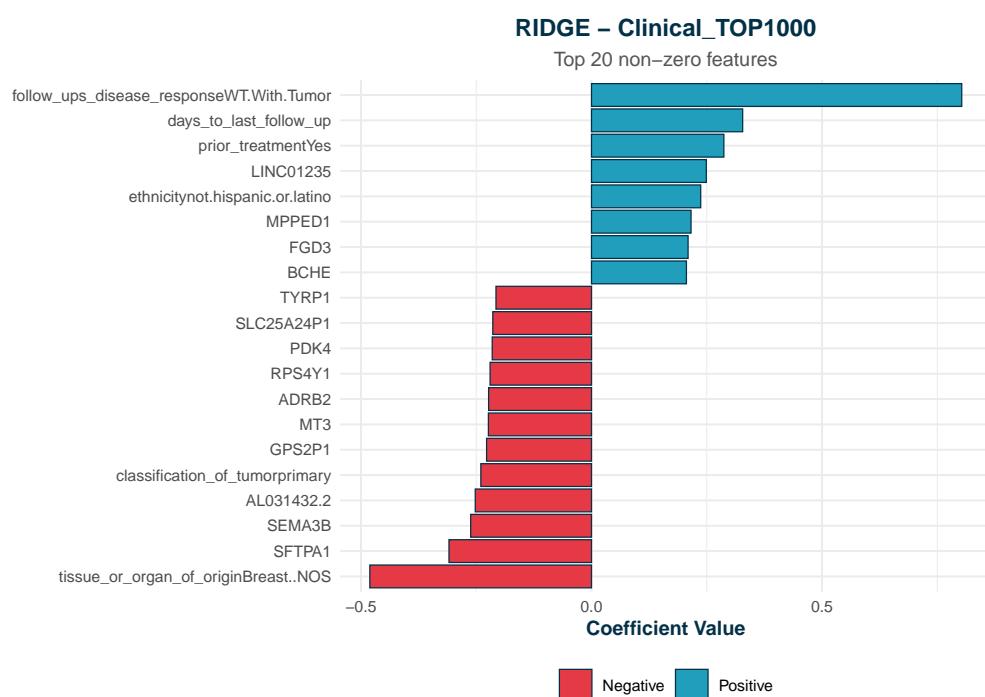
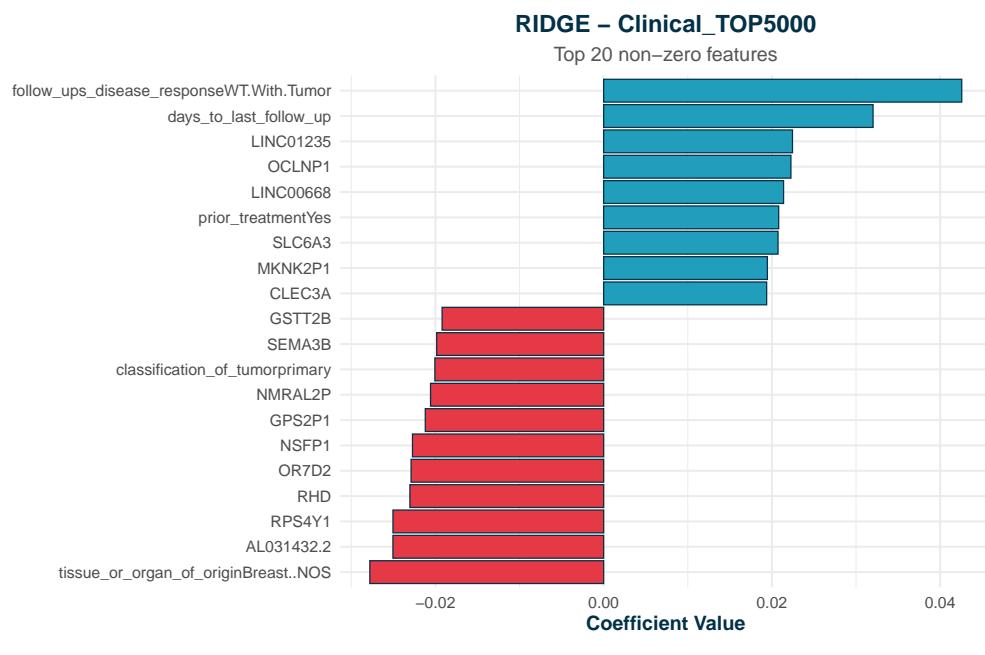
RIDGE – Train vs Test AUC



RIDGE – Clinical_Only

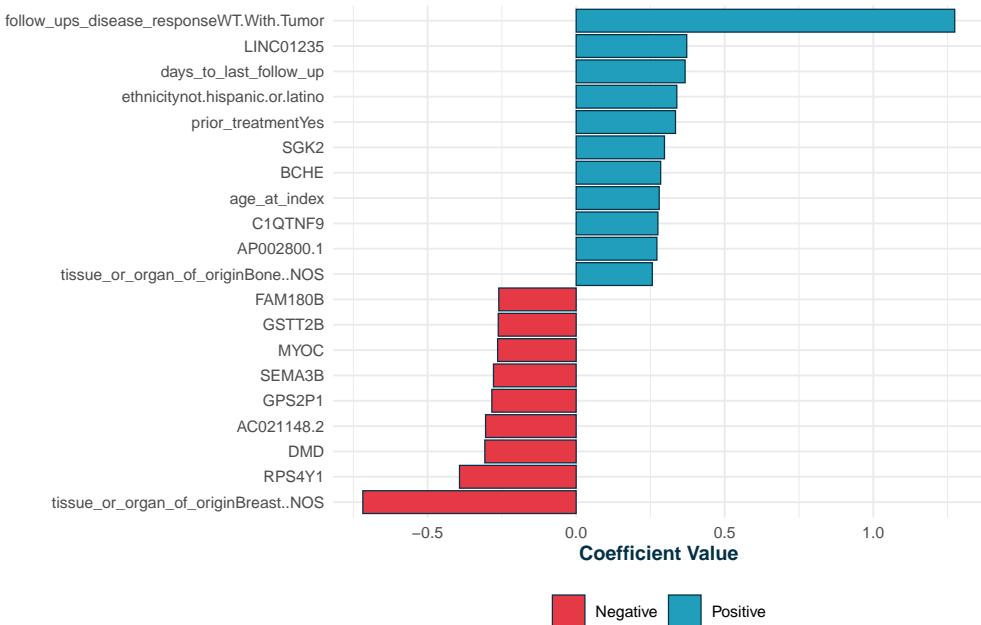
Top 20 non-zero features





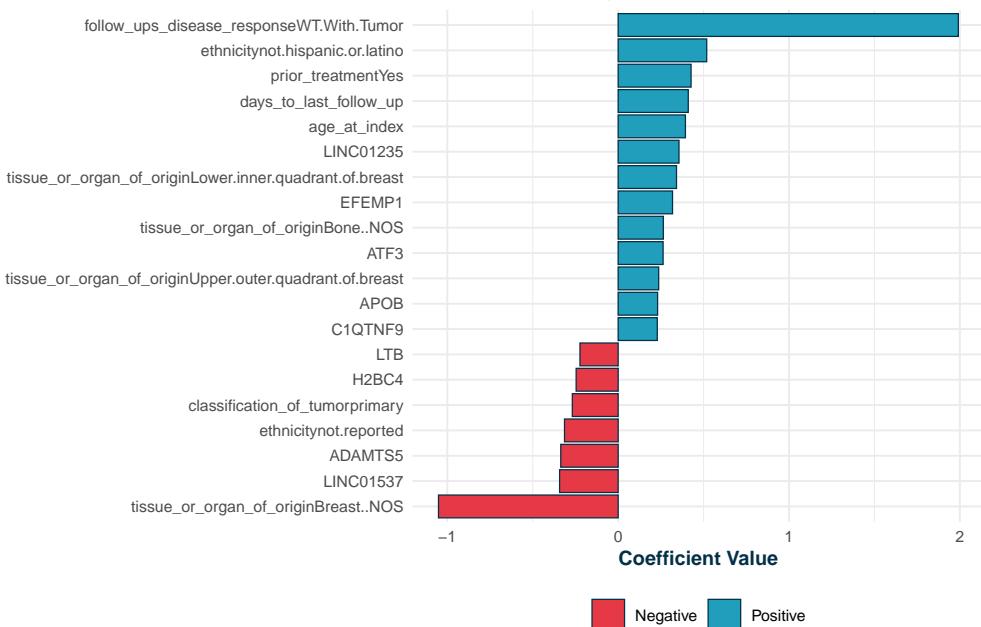
RIDGE – Clinical_TOP500

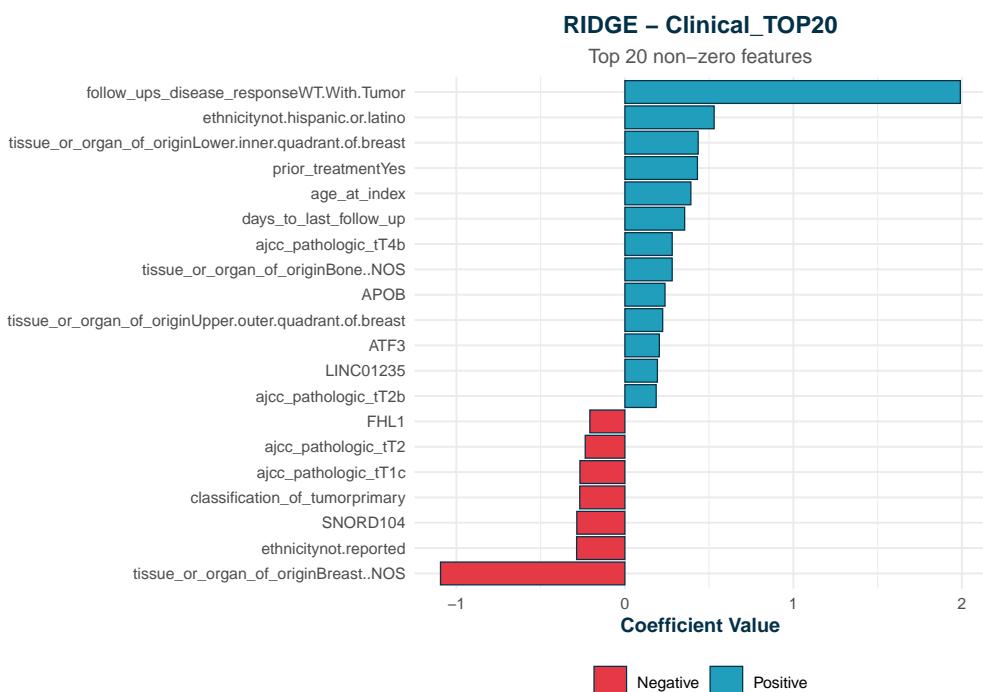
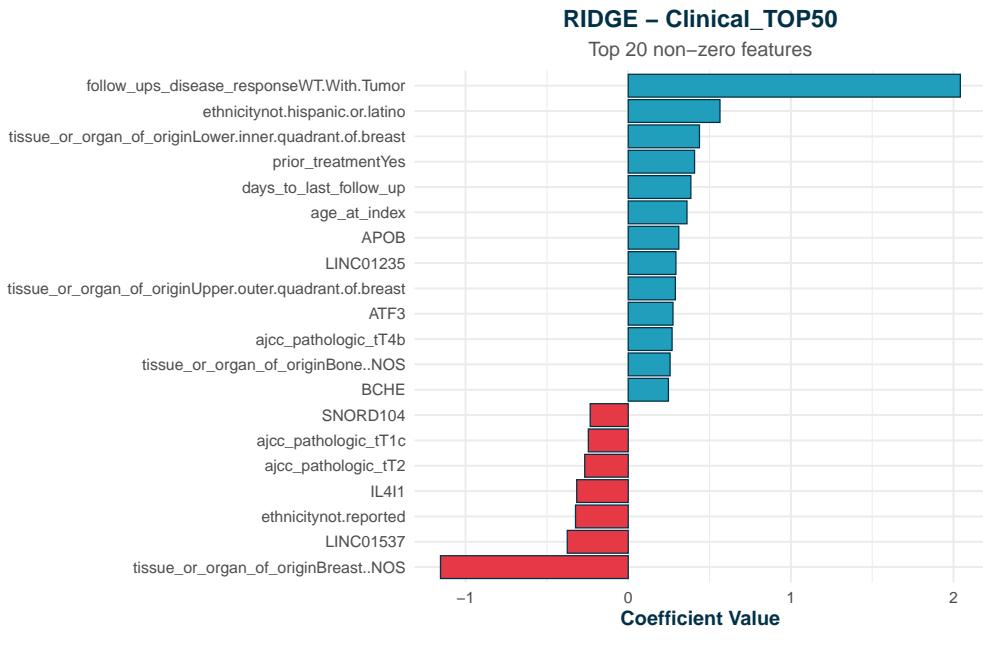
Top 20 non-zero features



RIDGE – Clinical_TOP100

Top 20 non-zero features





```
ridge_smote_metrics <- plot_classification_metrics_single(ridge_smote
  , threshold = 0.5
  , csv_filename = "ridge_smote_classification_metrics.csv")
```

```
##  
## === CLASSIFICATION METRICS ===
```

```

## Clinical_Only:
##   TP=29 TN=182 FP=24 FN=11
##   Accuracy=0.858 Precision=0.547 Recall=0.725 F1=0.624 AUC=0.884

## Clinical_TOP5000:
##   TP=17 TN=167 FP=39 FN=23
##   Accuracy=0.748 Precision=0.304 Recall=0.425 F1=0.354 AUC=0.700

## Clinical_TOP1000:
##   TP=24 TN=182 FP=24 FN=16
##   Accuracy=0.837 Precision=0.500 Recall=0.600 F1=0.545 AUC=0.828

## Clinical_TOP500:
##   TP=27 TN=173 FP=33 FN=13
##   Accuracy=0.813 Precision=0.450 Recall=0.675 F1=0.540 AUC=0.811

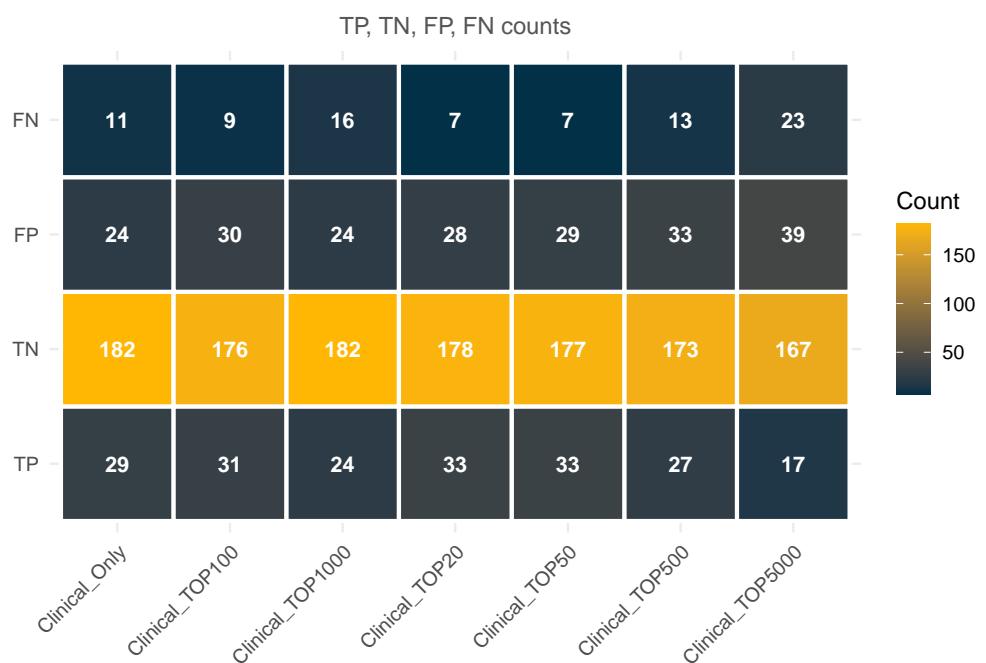
## Clinical_TOP100:
##   TP=31 TN=176 FP=30 FN=9
##   Accuracy=0.841 Precision=0.508 Recall=0.775 F1=0.614 AUC=0.864

## Clinical_TOP50:
##   TP=33 TN=177 FP=29 FN=7
##   Accuracy=0.854 Precision=0.532 Recall=0.825 F1=0.647 AUC=0.883

## Clinical_TOP20:
##   TP=33 TN=178 FP=28 FN=7
##   Accuracy=0.858 Precision=0.541 Recall=0.825 F1=0.653 AUC=0.896

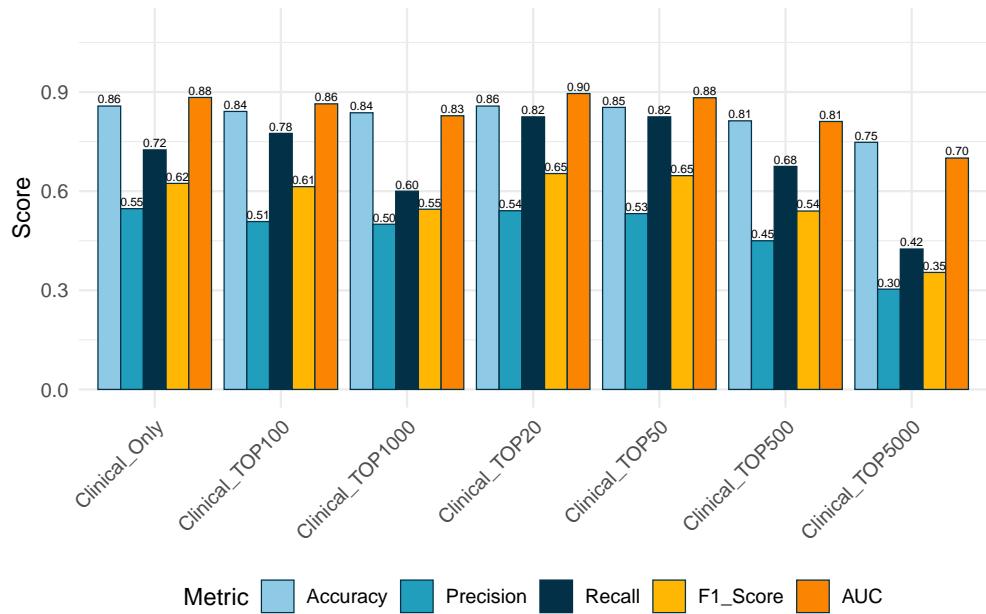
```

RIDGE – Confusion Matrix Across Feature Sets



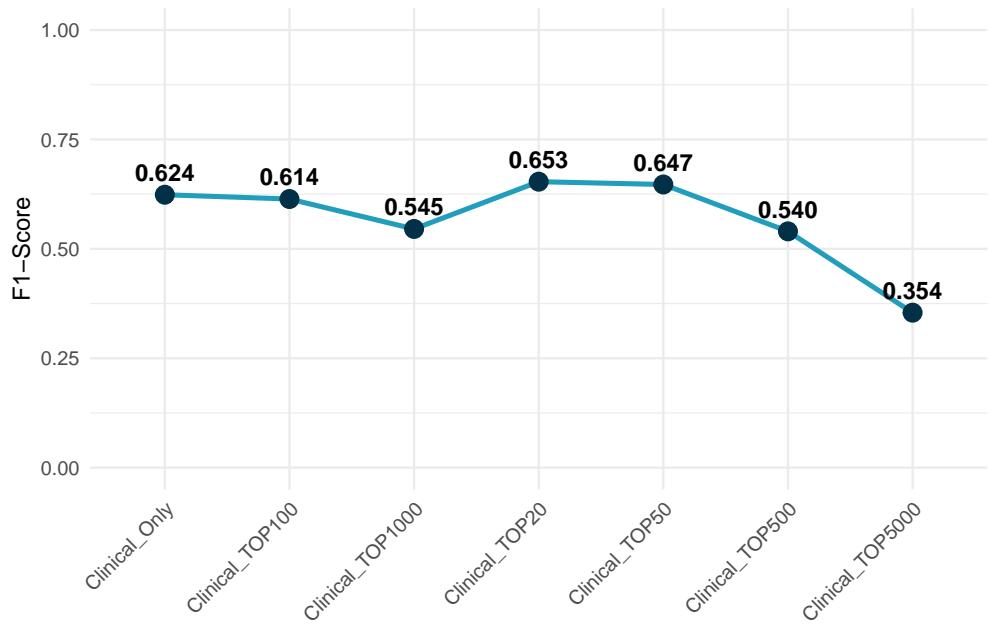
RIDGE – Classification Metrics

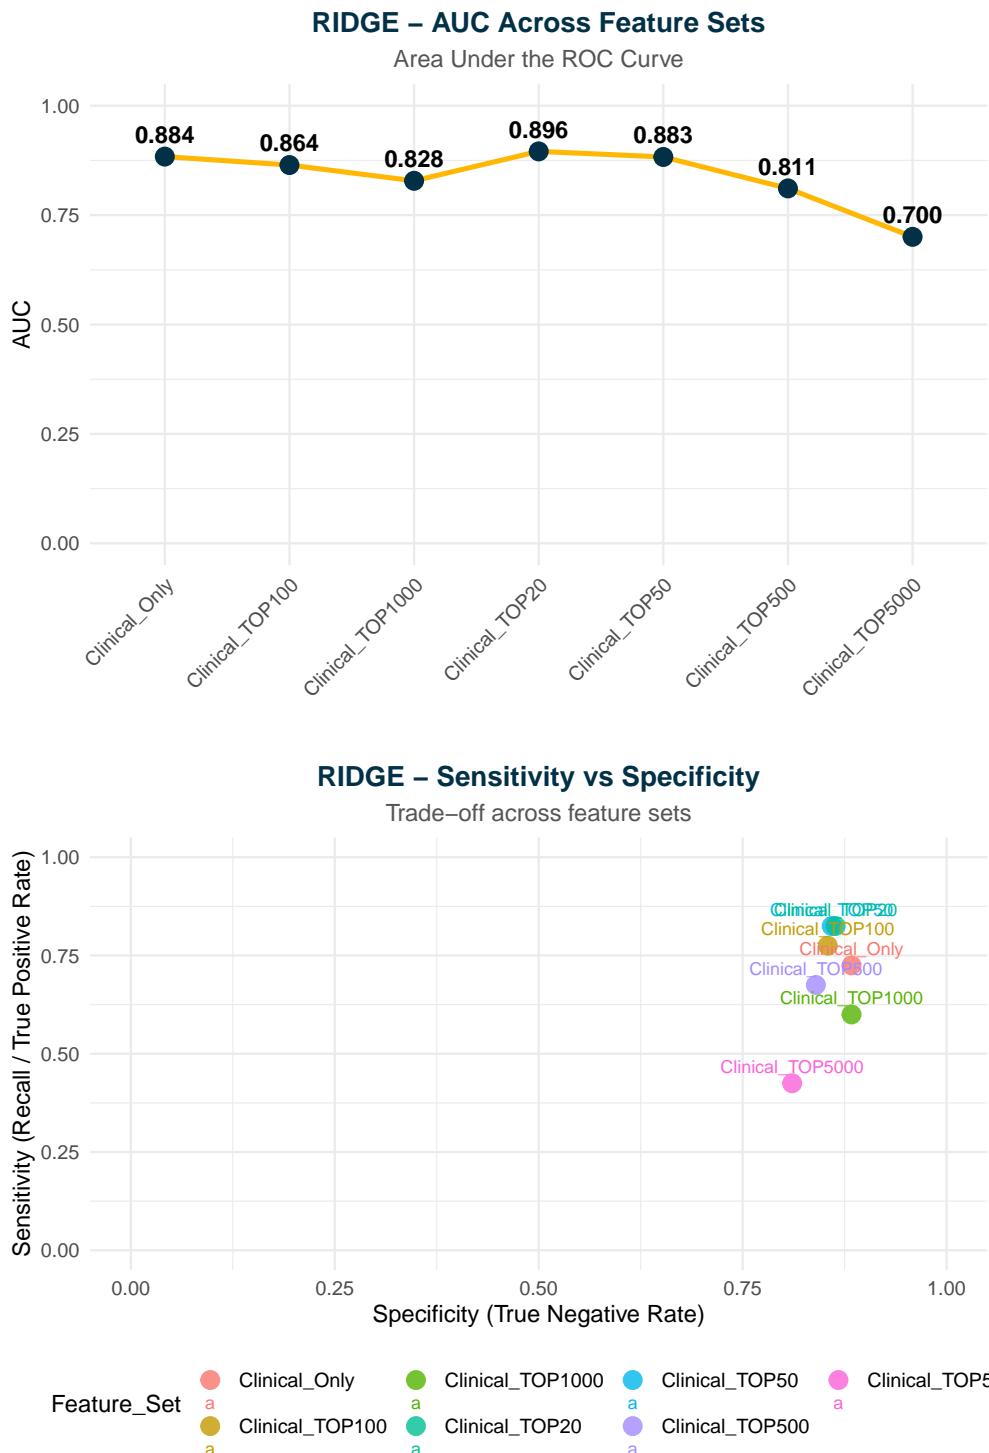
Accuracy, Precision, Recall, F1-Score, AUC



RIDGE – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1  Clinical_Only 29 182 24 11 0.8577236 0.5471698 0.725 0.8834951
## 2 Clinical_TOP5000 17 167 39 23 0.7479675 0.3035714 0.425 0.8106796
## 3 Clinical_TOP1000 24 182 24 16 0.8373984 0.5000000 0.600 0.8834951
## 4 Clinical_TOP500 27 173 33 13 0.8130081 0.4500000 0.675 0.8398058
```

```

## 5 Clinical_TOP100 31 176 30 9 0.8414634 0.5081967 0.775 0.8543689
## 6 Clinical_TOP50 33 177 29 7 0.8536585 0.5322581 0.825 0.8592233
## 7 Clinical_TOP20 33 178 28 7 0.8577236 0.5409836 0.825 0.8640777
## F1_Score AUC
## 1 0.6236559 0.8837379
## 2 0.3541667 0.7003641
## 3 0.5454545 0.8283981
## 4 0.5400000 0.8110437
## 5 0.6138614 0.8644417
## 6 0.6470588 0.8828883
## 7 0.6534653 0.8956311
##
## Exported classification metrics to: model_metrics/ridge_smote_classification_metrics.csv

```

Lasso with SMOTE

```

lasso_smote <- fit_single_model_across_features(
  model_type = "lasso"
  , X_train_all = smote_data$X_train
  , X_test_all = X_test
  , Y_train = smote_data$Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

##
## === FITTING LASSO ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Fitting Clinical_TOP5000...
## Fitting Clinical_TOP1000...
## Fitting Clinical_TOP500...
## Fitting Clinical_TOP100...
## Fitting Clinical_TOP50...
## Fitting Clinical_TOP20...

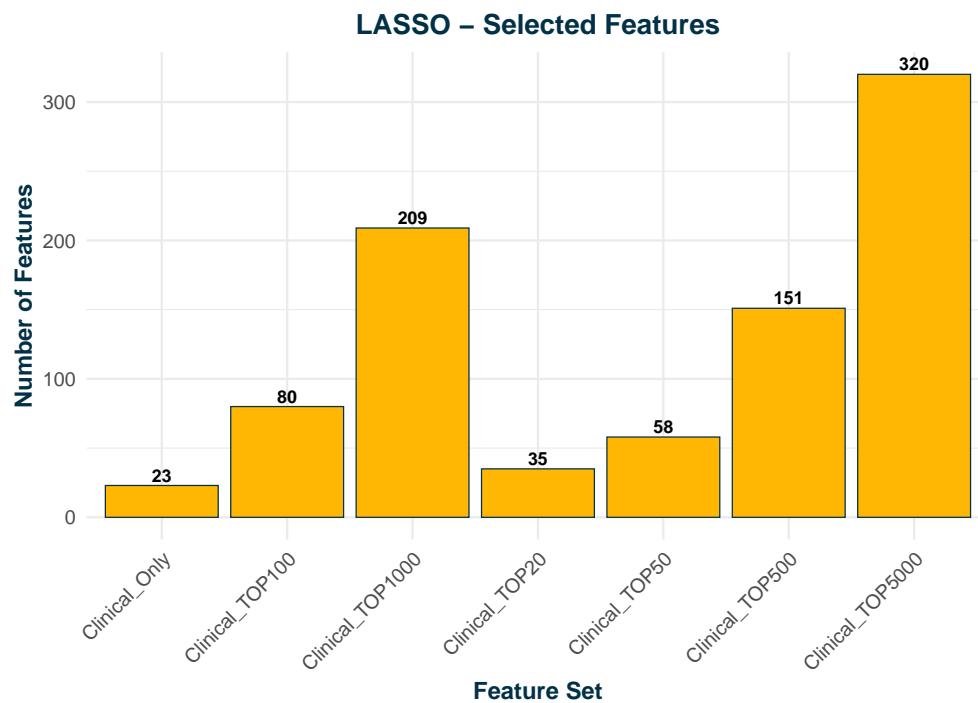
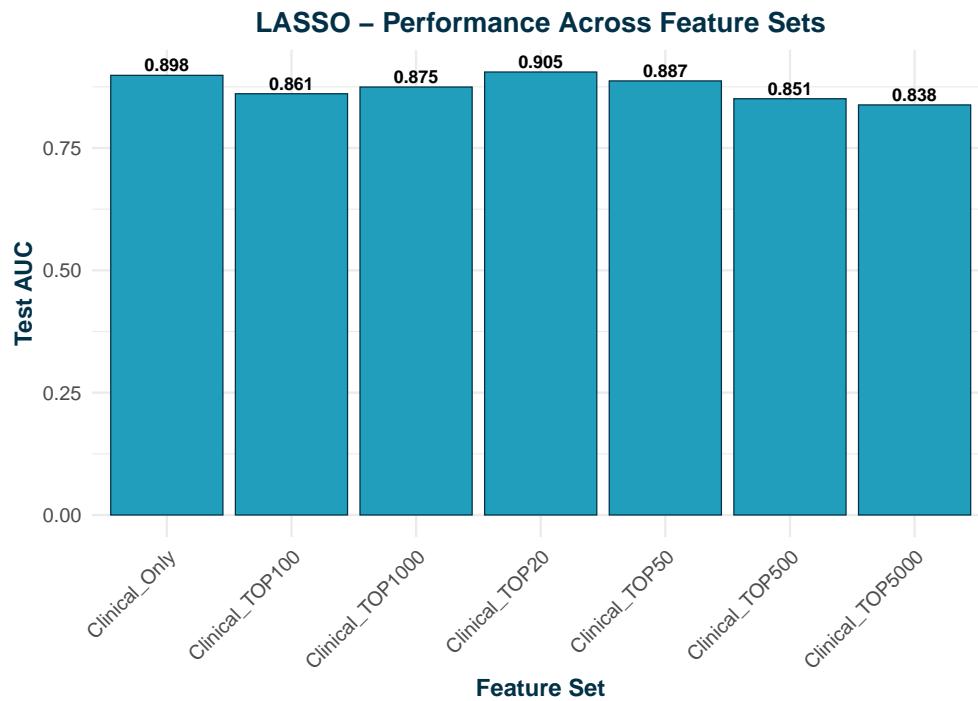
##
## === SUMMARY TABLE ===
##      Feature_Set Model Features Train_AUC Test_AUC Test_Accuracy
## 1    Clinical_Only LASSO        23 0.9422594 0.8984223 0.8739837
## 2 Clinical_TOP5000 LASSO       320 1.0000000 0.8379854 0.8739837
## 3 Clinical_TOP1000 LASSO      209 0.9975865 0.8745146 0.8699187

```

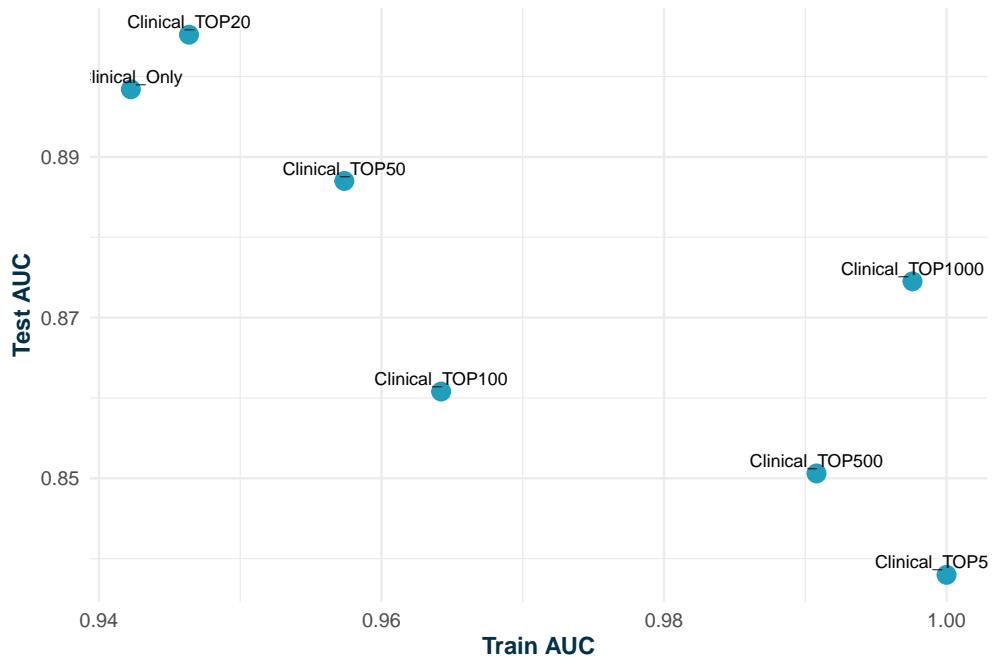
```

## 4 Clinical_TOP500 LASSO      151 0.9907942 0.8506068 0.8373984
## 5 Clinical_TOP100 LASSO      80 0.9642257 0.8608010 0.8739837
## 6 Clinical_TOP50 LASSO       58 0.9573700 0.8870146 0.8861789
## 7 Clinical_TOP20 LASSO       35 0.9463771 0.9052184 0.8861789
## Exported metrics to: model_metrics/lasso_across_features_metrics.csv

```

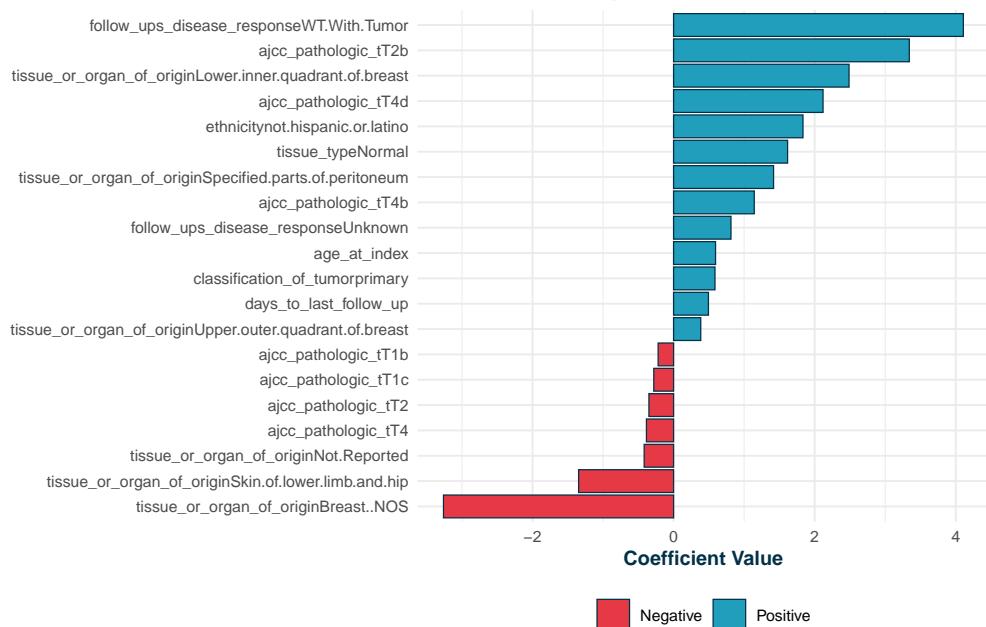


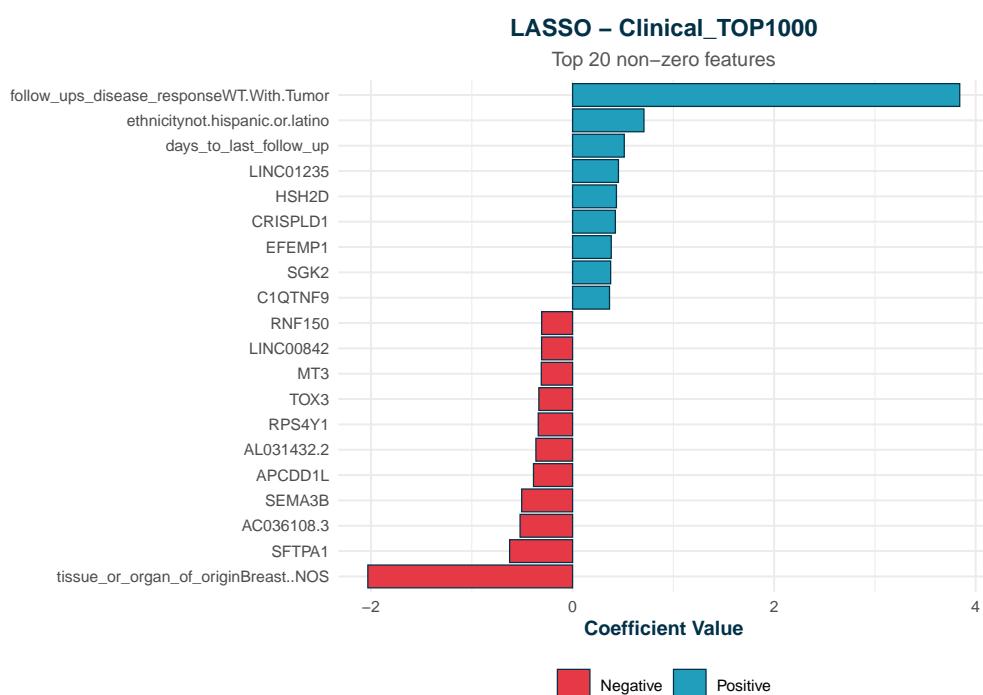
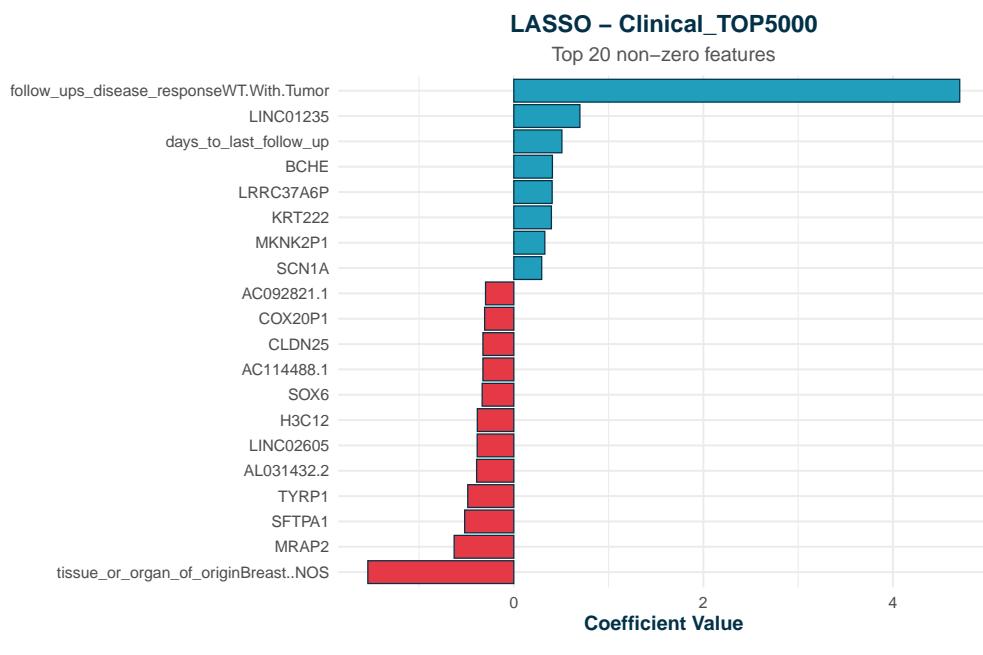
LASSO – Train vs Test AUC



LASSO – Clinical_Only

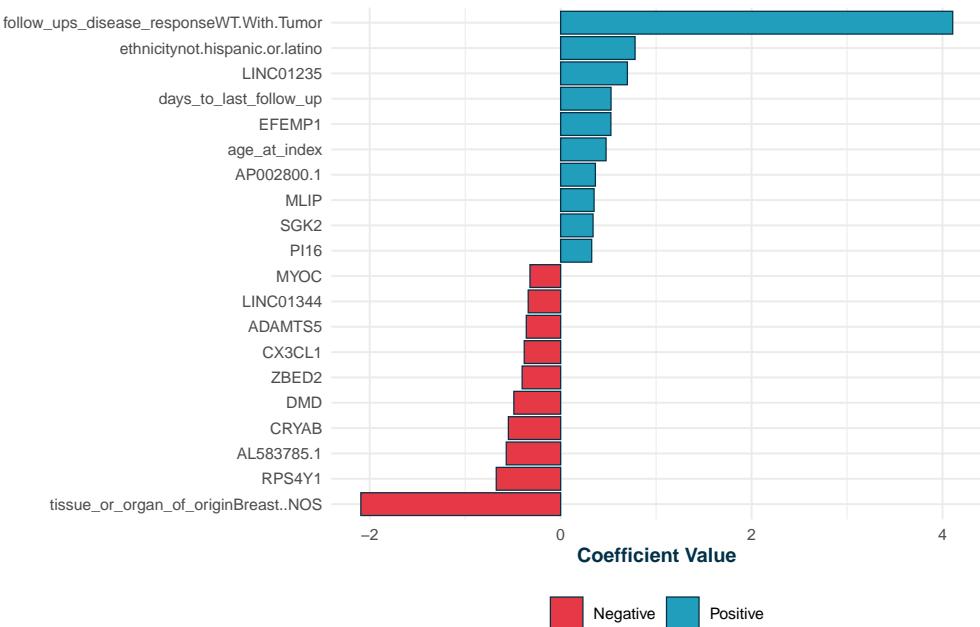
Top 20 non-zero features





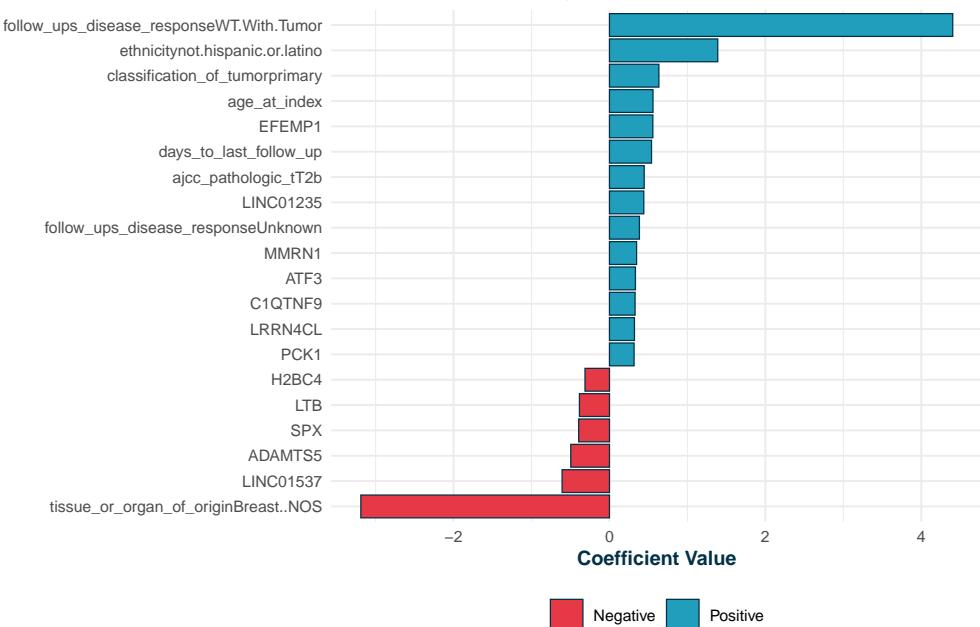
LASSO – Clinical_TOP500

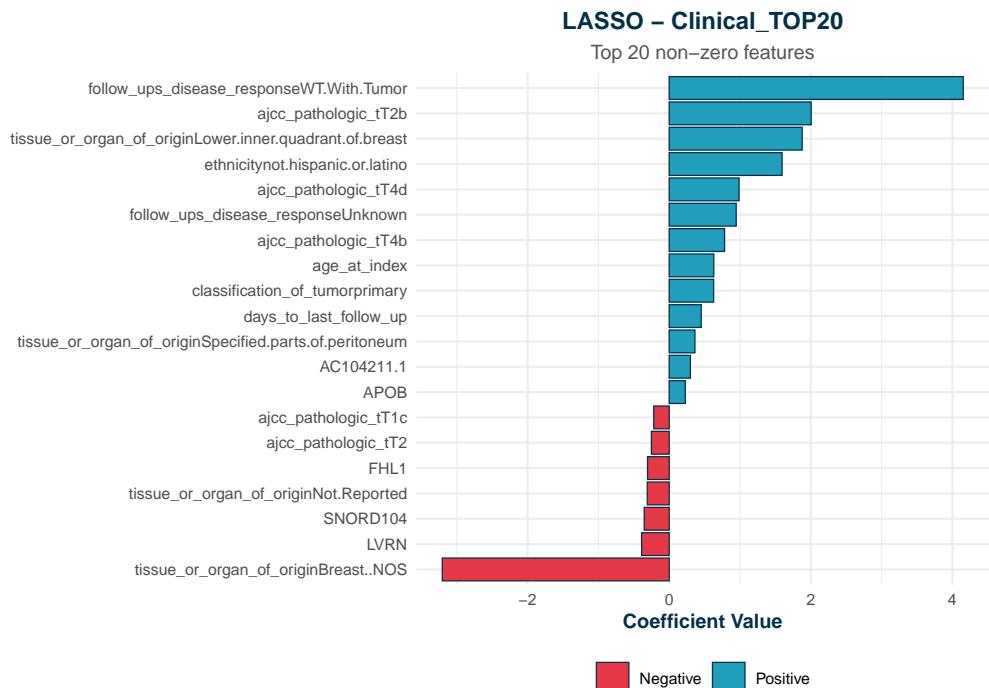
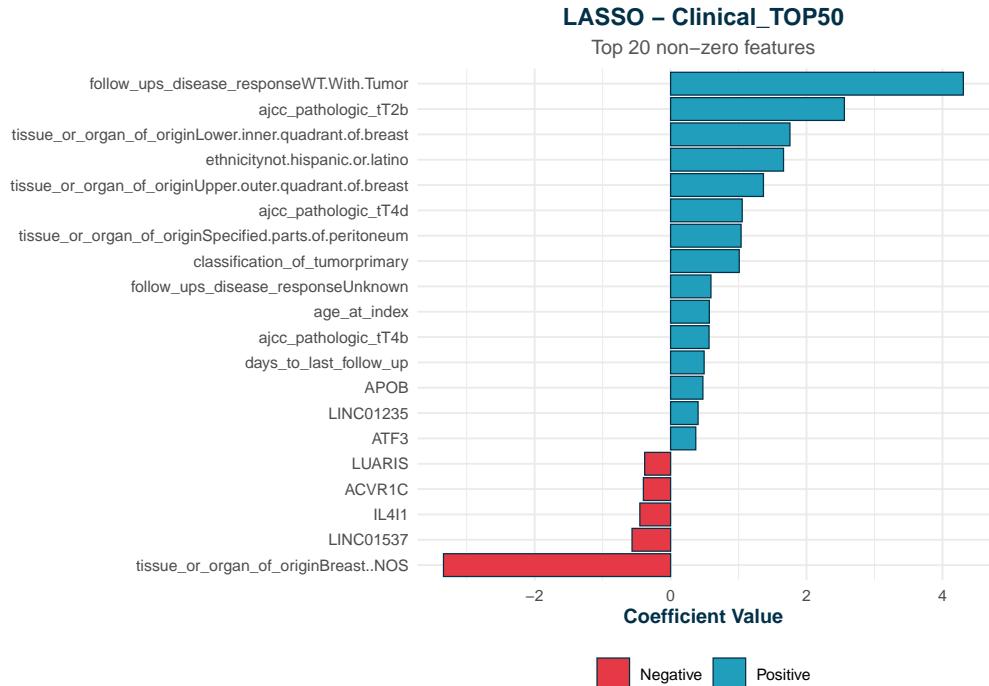
Top 20 non-zero features



LASSO – Clinical_TOP100

Top 20 non-zero features





```
lasso_smote_metrics <- plot_classification_metrics_single(lasso_smote
  , threshold = 0.5
  , csv_filename = "lasso_smote_classification_1.csv")
```

```
## 
## === CLASSIFICATION METRICS ===
```

```

## Clinical_Only:
##   TP=29 TN=186 FP=20 FN=11
##   Accuracy=0.874 Precision=0.592 Recall=0.725 F1=0.652 AUC=0.898

## Clinical_TOP5000:
##   TP=26 TN=189 FP=17 FN=14
##   Accuracy=0.874 Precision=0.605 Recall=0.650 F1=0.627 AUC=0.838

## Clinical_TOP1000:
##   TP=27 TN=187 FP=19 FN=13
##   Accuracy=0.870 Precision=0.587 Recall=0.675 F1=0.628 AUC=0.875

## Clinical_TOP500:
##   TP=25 TN=181 FP=25 FN=15
##   Accuracy=0.837 Precision=0.500 Recall=0.625 F1=0.556 AUC=0.851

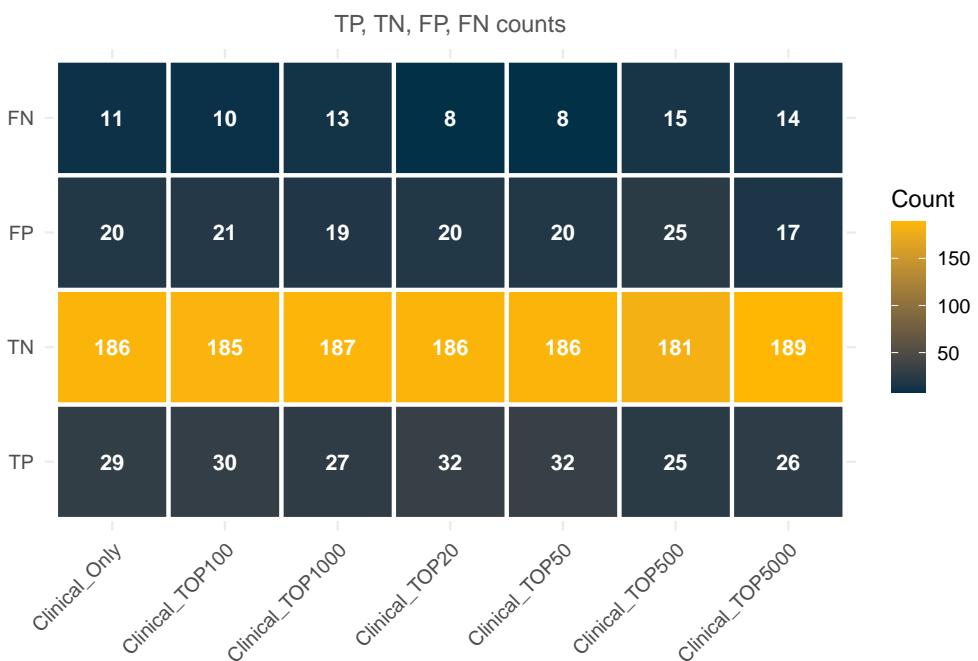
## Clinical_TOP100:
##   TP=30 TN=185 FP=21 FN=10
##   Accuracy=0.874 Precision=0.588 Recall=0.750 F1=0.659 AUC=0.861

## Clinical_TOP50:
##   TP=32 TN=186 FP=20 FN=8
##   Accuracy=0.886 Precision=0.615 Recall=0.800 F1=0.696 AUC=0.887

## Clinical_TOP20:
##   TP=32 TN=186 FP=20 FN=8
##   Accuracy=0.886 Precision=0.615 Recall=0.800 F1=0.696 AUC=0.905

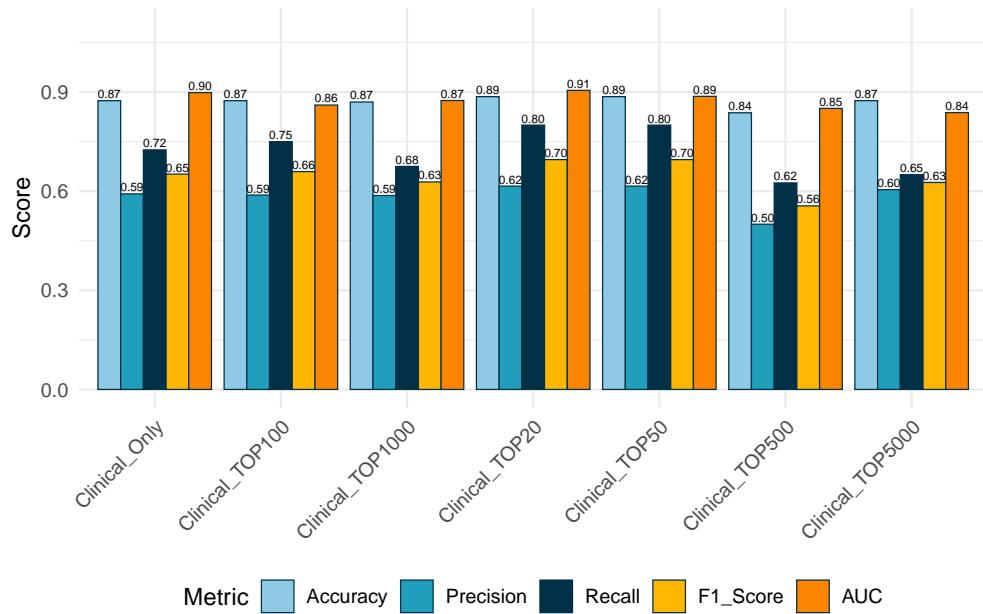
```

LASSO – Confusion Matrix Across Feature Sets



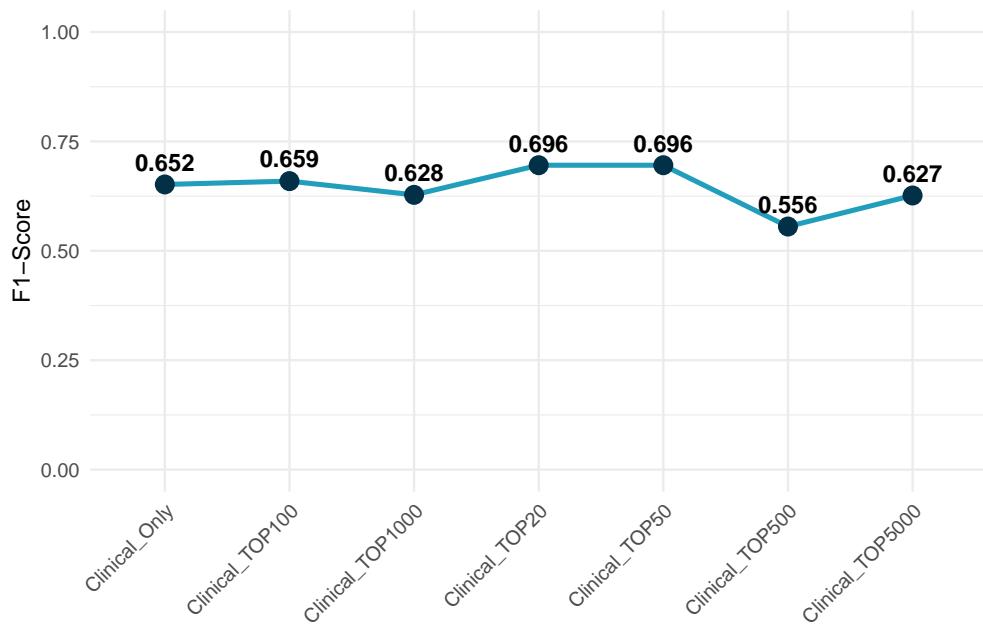
LASSO – Classification Metrics

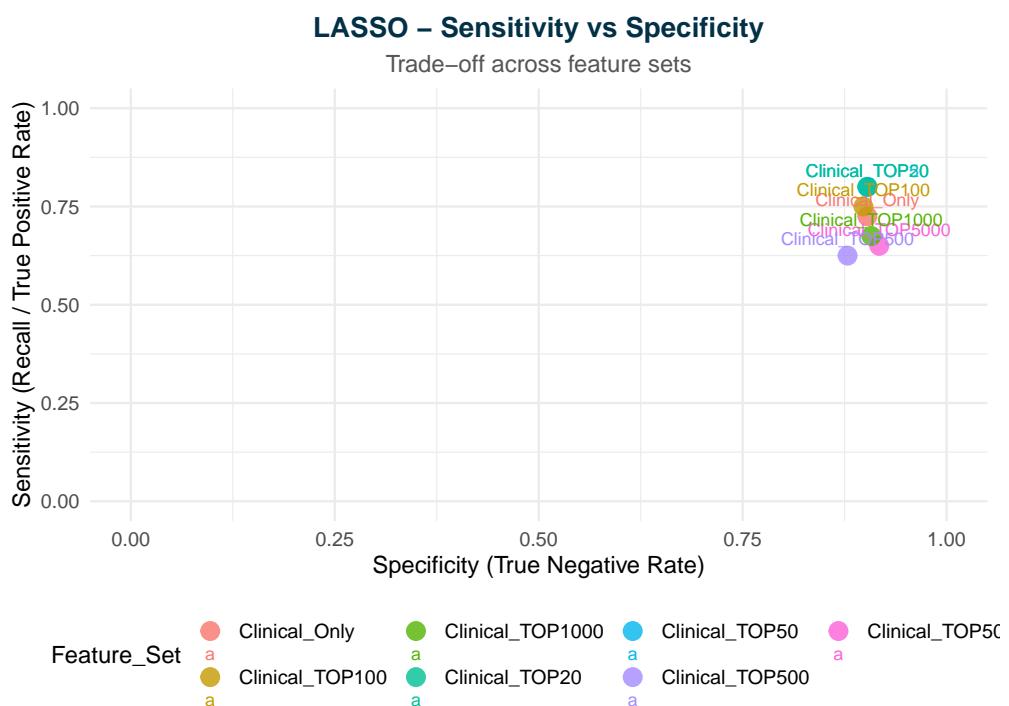
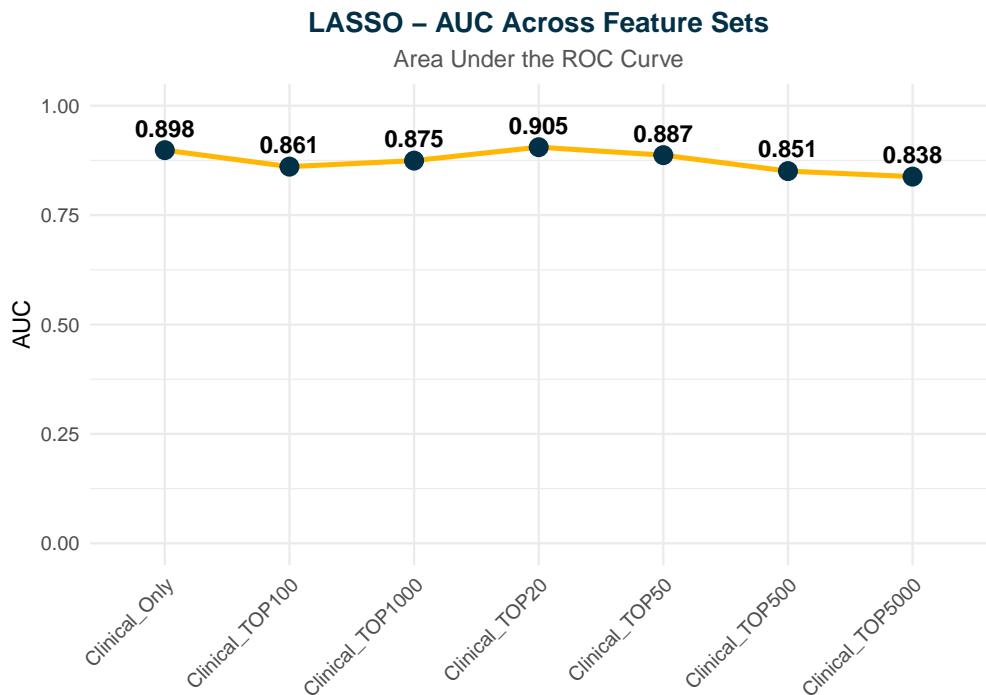
Accuracy, Precision, Recall, F1-Score, AUC



LASSO – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1   Clinical_Only 29 186 20 11 0.8739837 0.5918367 0.725 0.9029126
## 2 Clinical_TOP5000 26 189 17 14 0.8739837 0.6046512 0.650 0.9174757
## 3 Clinical_TOP1000 27 187 19 13 0.8699187 0.5869565 0.675 0.9077670
## 4 Clinical_TOP500 25 181 25 15 0.8373984 0.5000000 0.625 0.8786408
```

```

## 5 Clinical_TOP100 30 185 21 10 0.8739837 0.5882353 0.750 0.8980583
## 6 Clinical_TOP50 32 186 20 8 0.8861789 0.6153846 0.800 0.9029126
## 7 Clinical_TOP20 32 186 20 8 0.8861789 0.6153846 0.800 0.9029126
## F1_Score AUC
## 1 0.6516854 0.8984223
## 2 0.6265060 0.8379854
## 3 0.6279070 0.8745146
## 4 0.5555556 0.8506068
## 5 0.6593407 0.8608010
## 6 0.6956522 0.8870146
## 7 0.6956522 0.9052184
##
## Exported classification metrics to: model_metrics/lasso_smote_classification_metrics.csv

```

ElasticNet with SMOTE

```

elasticnet_smote <- fit_single_model_across_features(
  model_type = "elasticnet"
  , X_train_all = smote_data$X_train
  , X_test_all = X_test
  , Y_train = smote_data$Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

##
## === FITTING ELASTICNET ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Fitting Clinical_TOP5000...
## Fitting Clinical_TOP1000...
## Fitting Clinical_TOP500...
## Fitting Clinical_TOP100...
## Fitting Clinical_TOP50...
## Fitting Clinical_TOP20...

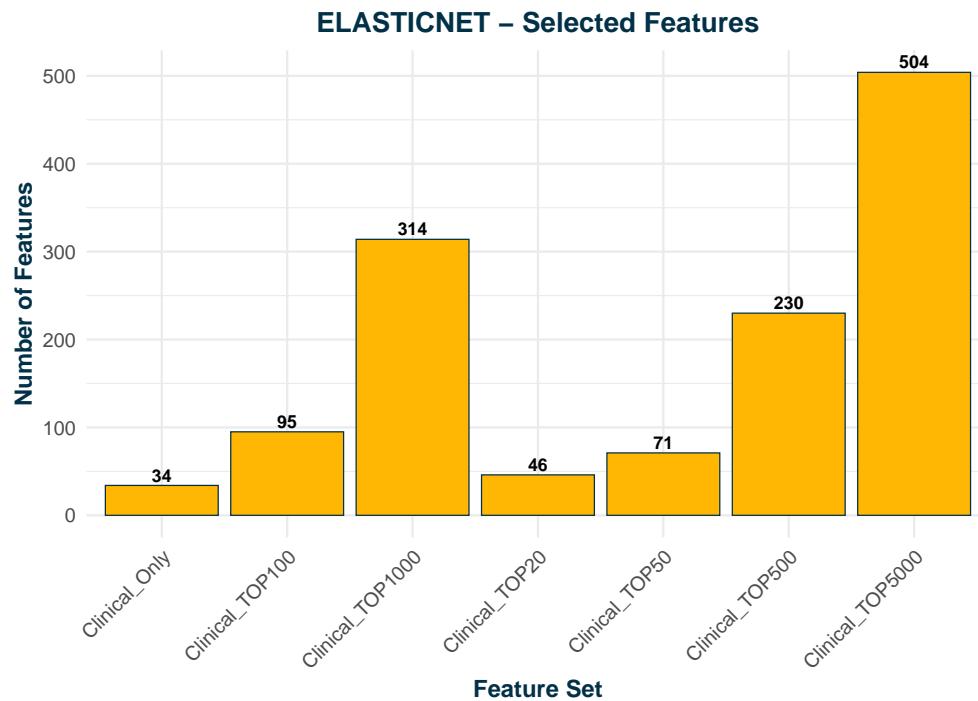
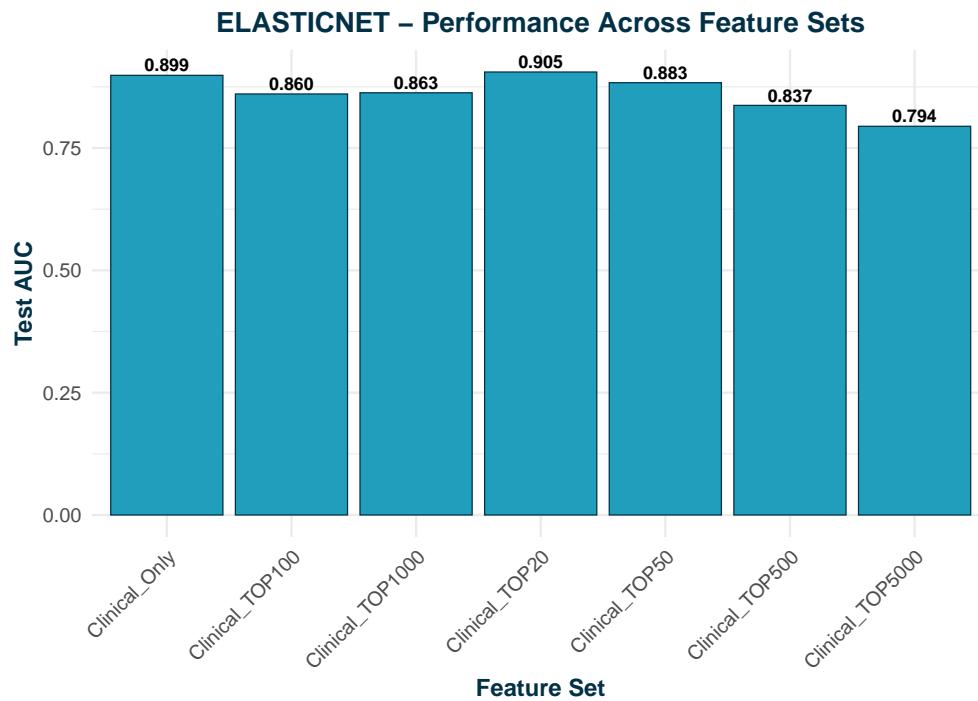
##
## === SUMMARY TABLE ===
##      Feature_Set      Model Features Train_AUC Test_AUC Test_Accuracy
## 1    Clinical_Only ELASTICNET      34 0.9428043 0.8985437 0.8699187
## 2 Clinical_TOP5000 ELASTICNET     504 1.0000000 0.7944175 0.8455285
## 3 Clinical_TOP1000 ELASTICNET     314 0.9995623 0.8628641 0.8699187

```

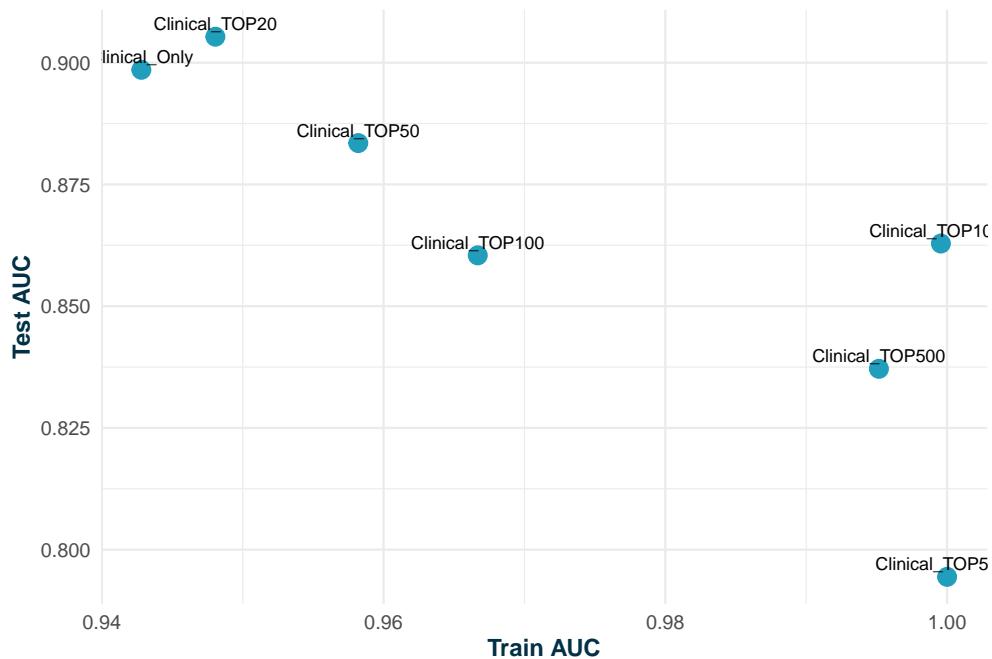
```

## 4 Clinical_TOP500 ELASTICNET      230 0.9951624 0.8371359 0.8373984
## 5 Clinical_TOP100 ELASTICNET      95 0.9666845 0.8604369 0.8739837
## 6 Clinical_TOP50 ELASTICNET      71 0.9582017 0.8834951 0.8739837
## 7 Clinical_TOP20 ELASTICNET      46 0.9480585 0.9053398 0.8943089
## Exported metrics to: model_metrics/elasticnet_across_features_metrics.csv

```

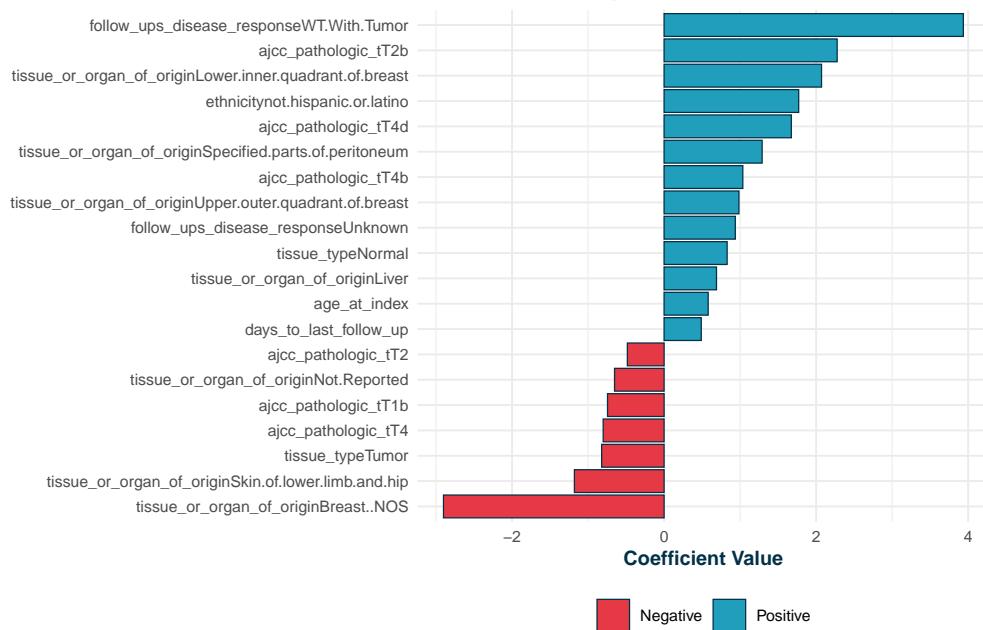


ELASTICNET – Train vs Test AUC



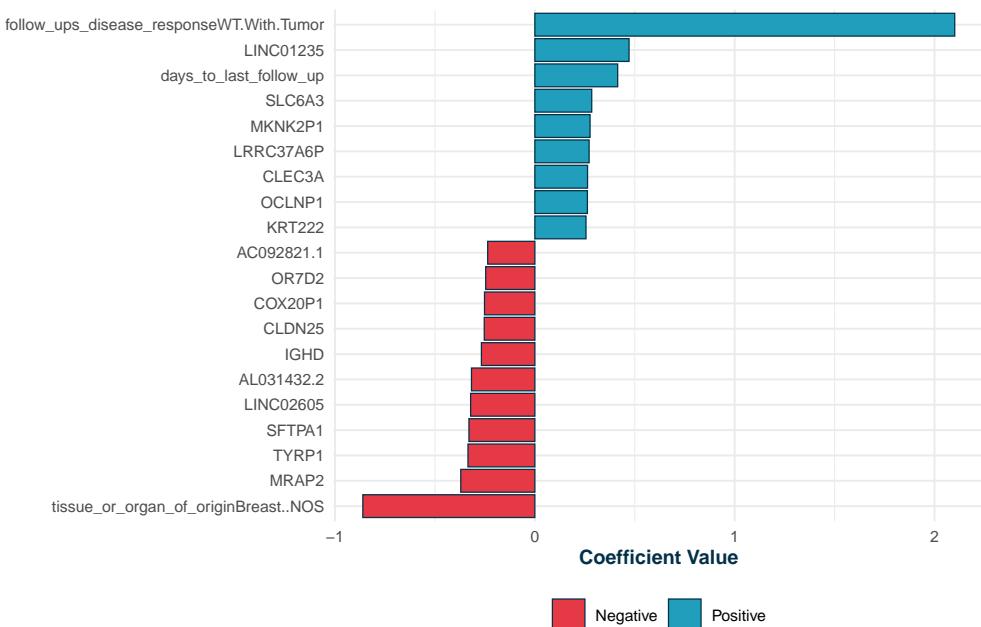
ELASTICNET – Clinical_Only

Top 20 non-zero features



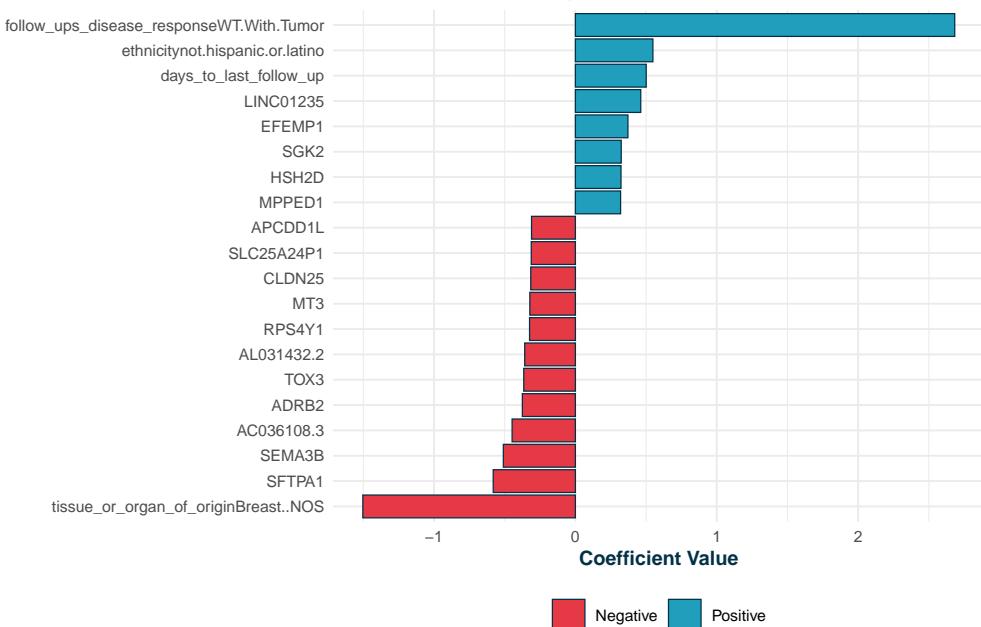
ELASTICNET – Clinical_TOP5000

Top 20 non-zero features



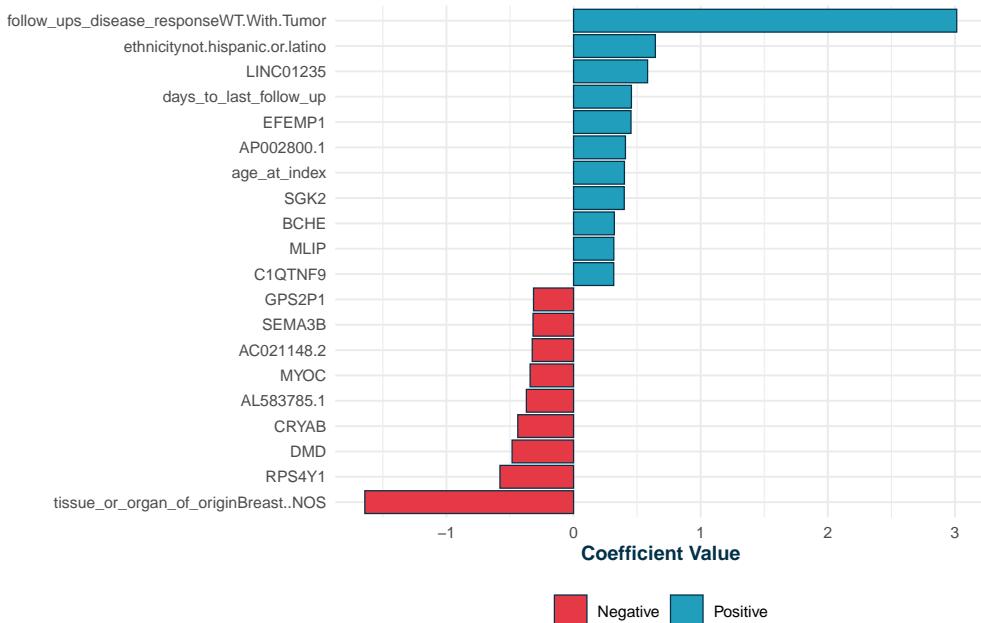
ELASTICNET – Clinical_TOP1000

Top 20 non-zero features



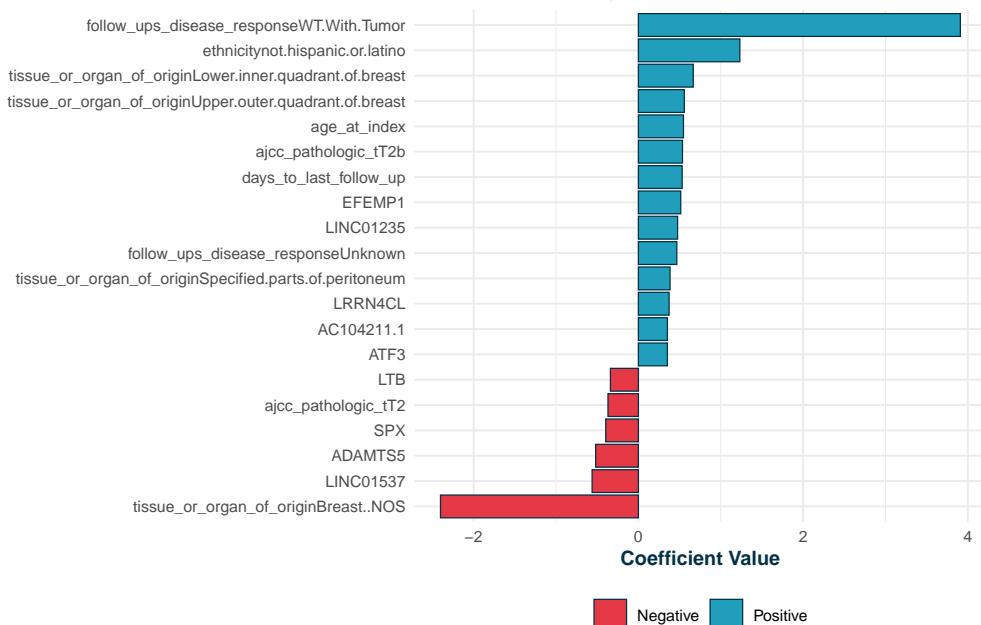
ELASTICNET – Clinical_TOP500

Top 20 non-zero features



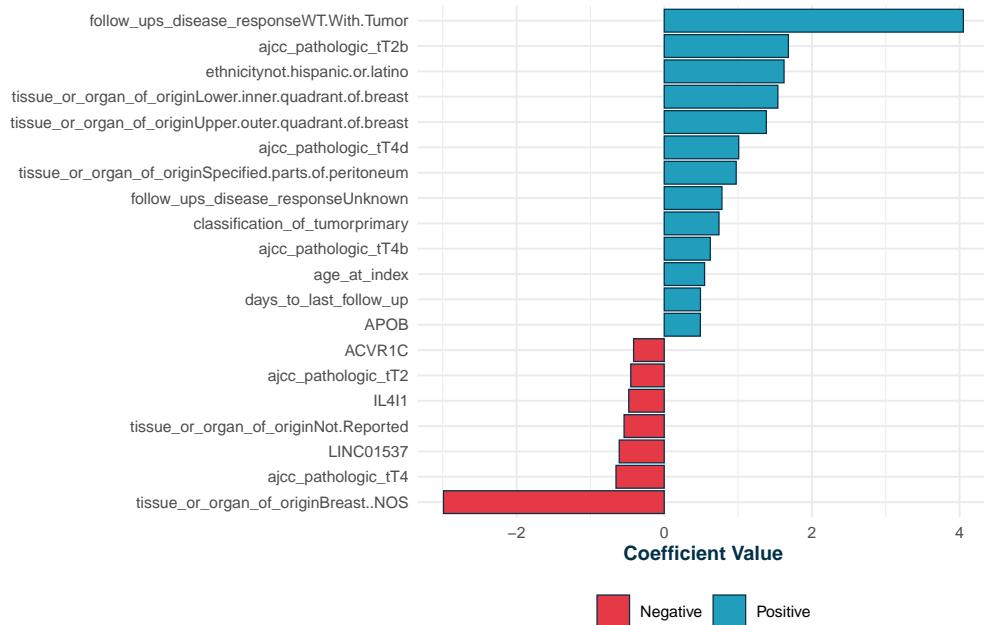
ELASTICNET – Clinical_TOP100

Top 20 non-zero features



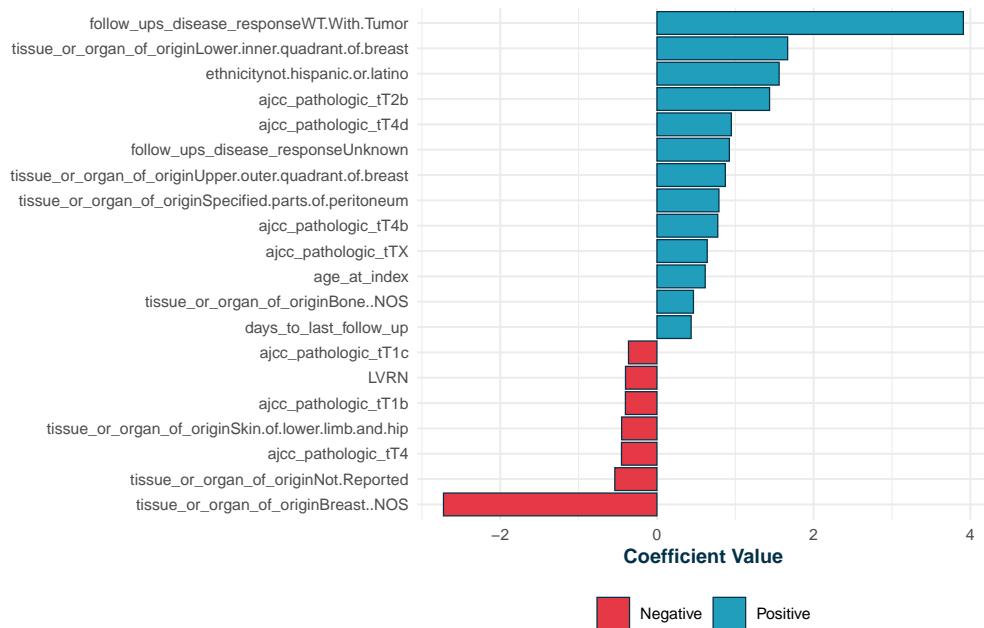
ELASTICNET – Clinical_TOP50

Top 20 non-zero features



ELASTICNET – Clinical_TOP20

Top 20 non-zero features



```
elasticnet_smote_metrics <- plot_classification_metrics_single(elasticnet_smote
, threshold = 0.5
, csv_filename = "elasticnet_smote_class"
```

```
##  
## === CLASSIFICATION METRICS ===
```

```

## Clinical_Only:
##   TP=29 TN=185 FP=21 FN=11
##   Accuracy=0.870 Precision=0.580 Recall=0.725 F1=0.644 AUC=0.899

## Clinical_TOP5000:
##   TP=20 TN=188 FP=18 FN=20
##   Accuracy=0.846 Precision=0.526 Recall=0.500 F1=0.513 AUC=0.794

## Clinical_TOP1000:
##   TP=27 TN=187 FP=19 FN=13
##   Accuracy=0.870 Precision=0.587 Recall=0.675 F1=0.628 AUC=0.863

## Clinical_TOP500:
##   TP=25 TN=181 FP=25 FN=15
##   Accuracy=0.837 Precision=0.500 Recall=0.625 F1=0.556 AUC=0.837

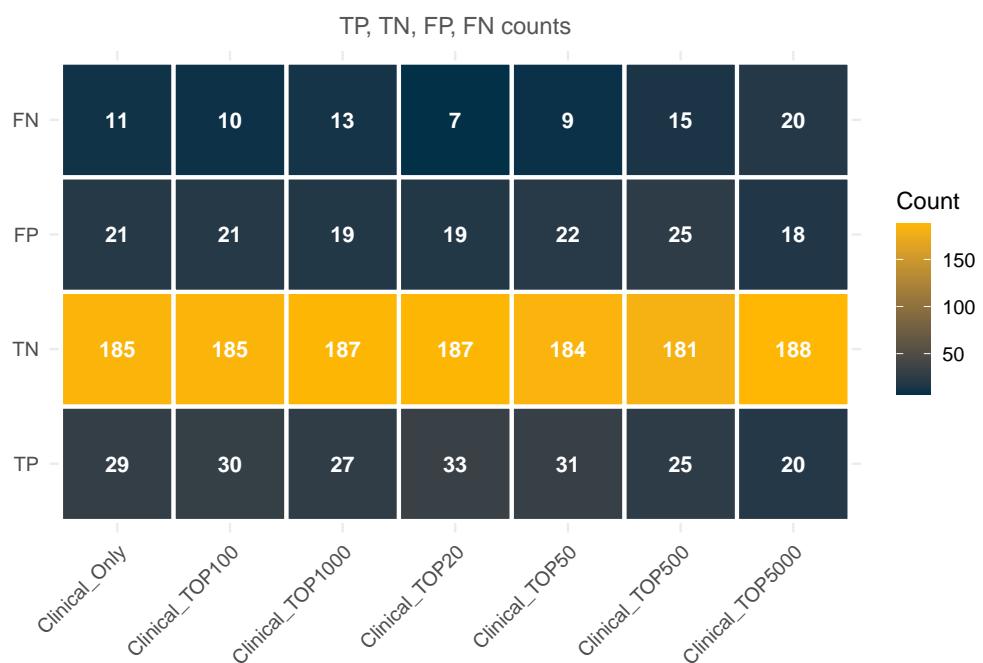
## Clinical_TOP100:
##   TP=30 TN=185 FP=21 FN=10
##   Accuracy=0.874 Precision=0.588 Recall=0.750 F1=0.659 AUC=0.860

## Clinical_TOP50:
##   TP=31 TN=184 FP=22 FN=9
##   Accuracy=0.874 Precision=0.585 Recall=0.775 F1=0.667 AUC=0.883

## Clinical_TOP20:
##   TP=33 TN=187 FP=19 FN=7
##   Accuracy=0.894 Precision=0.635 Recall=0.825 F1=0.717 AUC=0.905

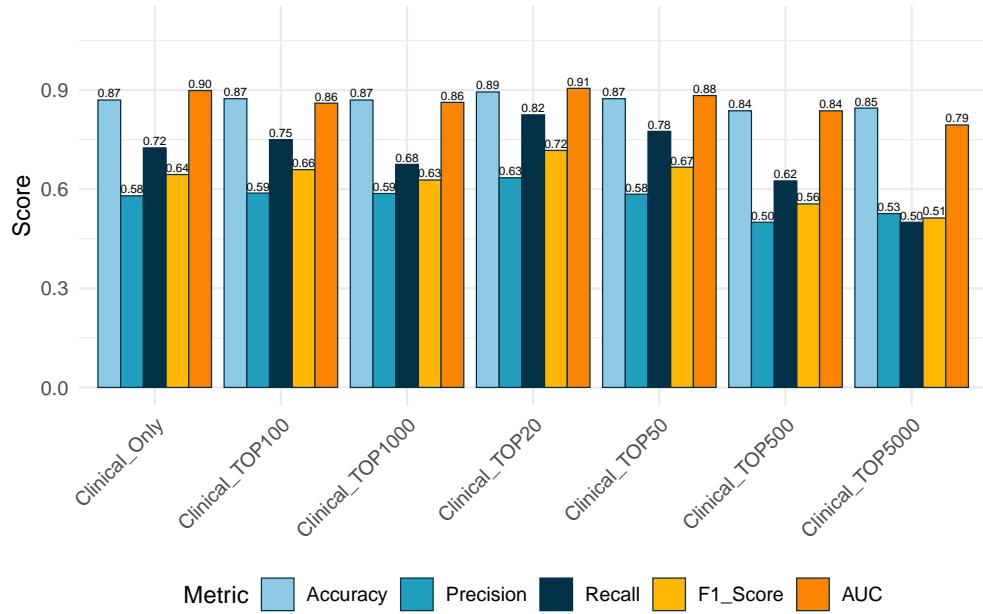
```

ELASTICNET – Confusion Matrix Across Feature Sets



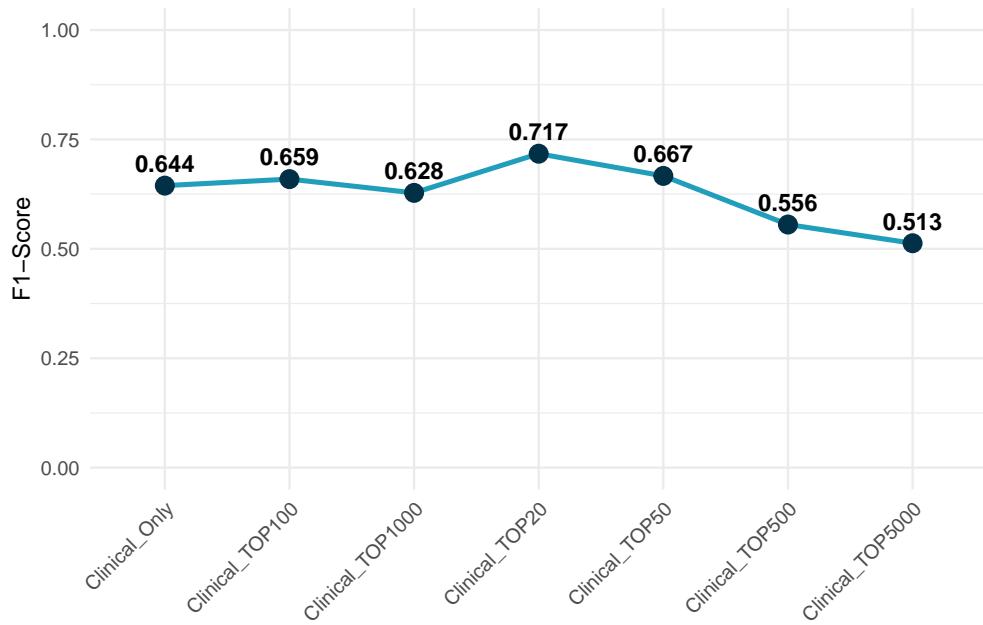
ELASTICNET – Classification Metrics

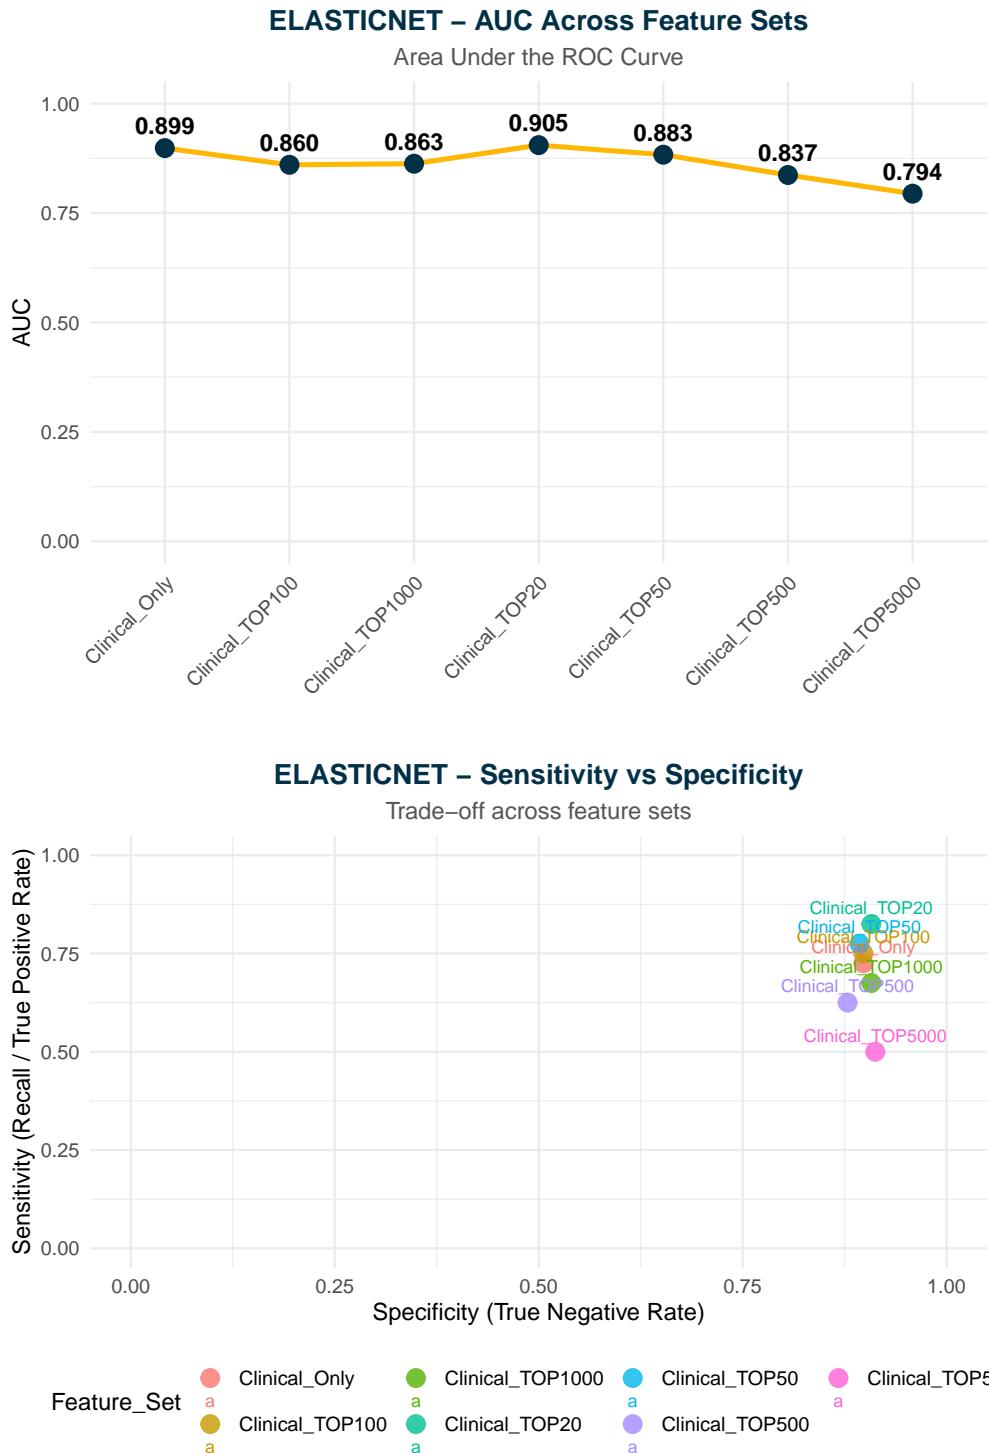
Accuracy, Precision, Recall, F1-Score, AUC



ELASTICNET – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1   Clinical_Only 29 185 21 11 0.8699187 0.5800000 0.725 0.8980583
## 2 Clinical_TOP5000 20 188 18 20 0.8455285 0.5263158 0.500 0.9126214
## 3 Clinical_TOP1000 27 187 19 13 0.8699187 0.5869565 0.675 0.9077670
## 4 Clinical_TOP500 25 181 25 15 0.8373984 0.5000000 0.625 0.8786408
```

```

## 5 Clinical_TOP100 30 185 21 10 0.8739837 0.5882353 0.750 0.8980583
## 6 Clinical_TOP50 31 184 22 9 0.8739837 0.5849057 0.775 0.8932039
## 7 Clinical_TOP20 33 187 19 7 0.8943089 0.6346154 0.825 0.9077670
## F1_Score AUC
## 1 0.6444444 0.8985437
## 2 0.5128205 0.7944175
## 3 0.6279070 0.8628641
## 4 0.5555556 0.8371359
## 5 0.6593407 0.8604369
## 6 0.6666667 0.8834951
## 7 0.7173913 0.9053398
##
## Exported classification metrics to: model_metrics/elasticnet_smote_classification_metrics.csv

```

Adaptive Lasso with SMOTE

```

adaptive_lasso_smote <- fit_single_model_across_features(
  model_type = "adaptive"
  , X_train_all = smote_data$X_train
  , X_test_all = X_test
  , Y_train = smote_data$Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

##
## === FITTING ADAPTIVE ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Fitting Clinical_TOP5000...
## Fitting Clinical_TOP1000...
## Fitting Clinical_TOP500...
## Fitting Clinical_TOP100...
## Fitting Clinical_TOP50...
## Fitting Clinical_TOP20...

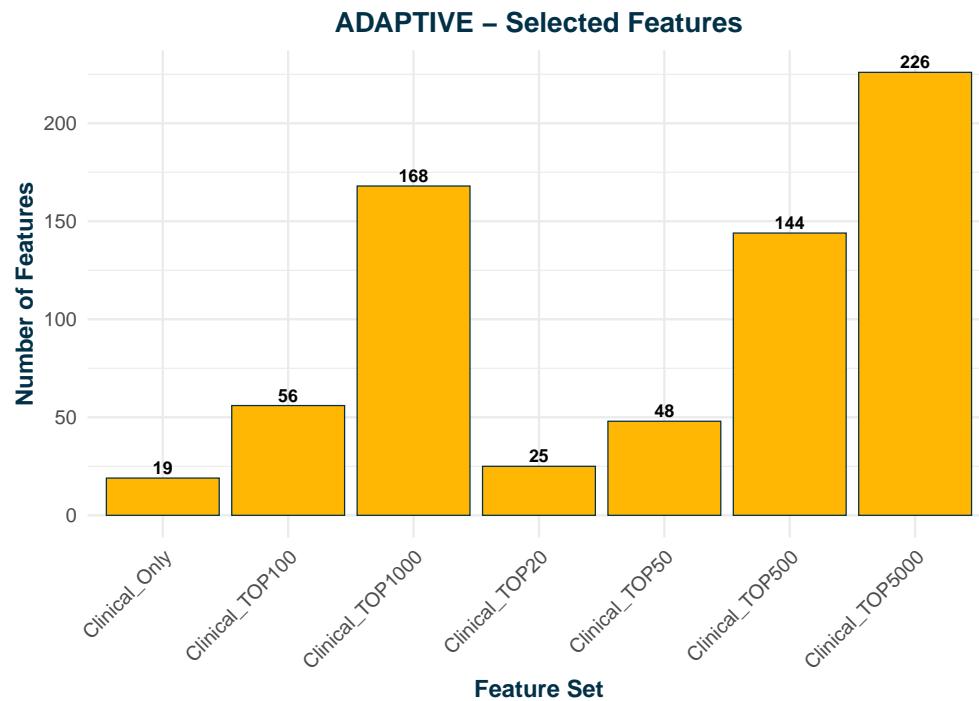
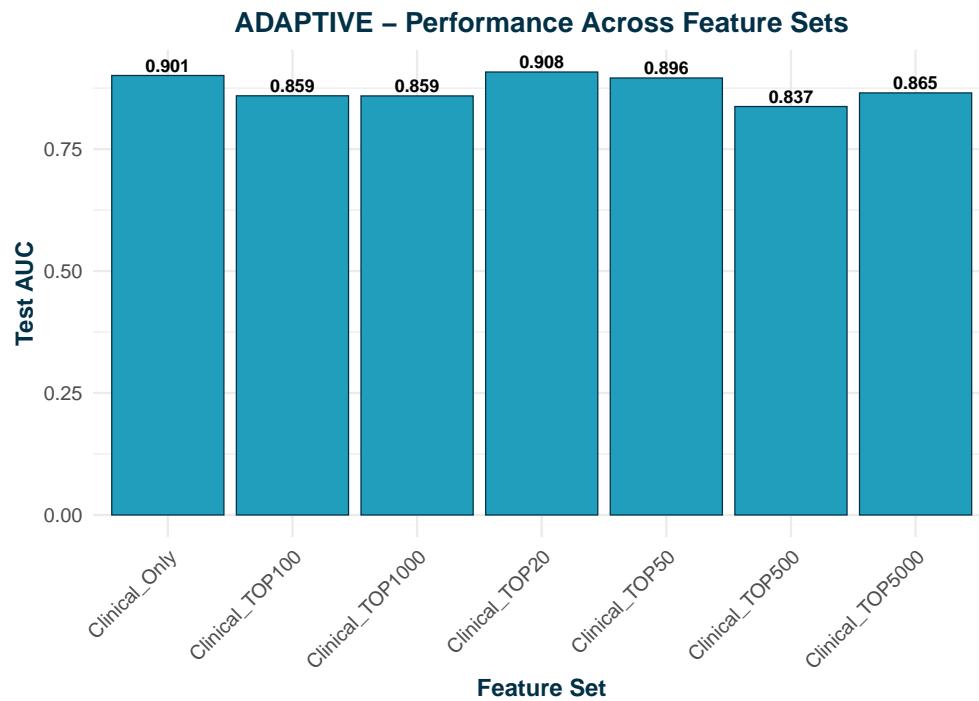
##
## === SUMMARY TABLE ===
##      Feature_Set Model Features Train_AUC Test_AUC Test_Accuracy
## 1 Clinical_Only ADAPTIVE      19 0.9425915 0.9009709 0.8821138
## 2 Clinical_TOP5000 ADAPTIVE     226 0.9999758 0.8652913 0.8861789
## 3 Clinical_TOP1000 ADAPTIVE     168 0.9988015 0.8591019 0.8739837

```

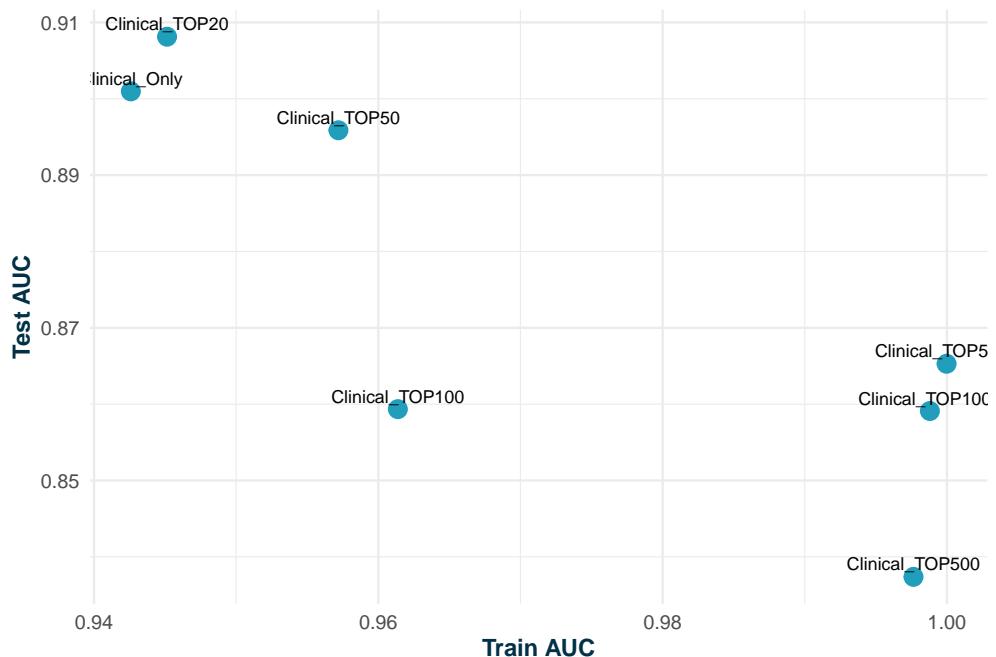
```

## 4 Clinical_TOP500 ADAPTIVE      144 0.9976499 0.8373786 0.8373984
## 5 Clinical_TOP100 ADAPTIVE      56 0.9613760 0.8593447 0.8699187
## 6 Clinical_TOP50 ADAPTIVE      48 0.9571949 0.8958738 0.8983740
## 7 Clinical_TOP20 ADAPTIVE      25 0.9451394 0.9081311 0.8943089
## Exported metrics to: model_metrics/adaptive_across_features_metrics.csv

```

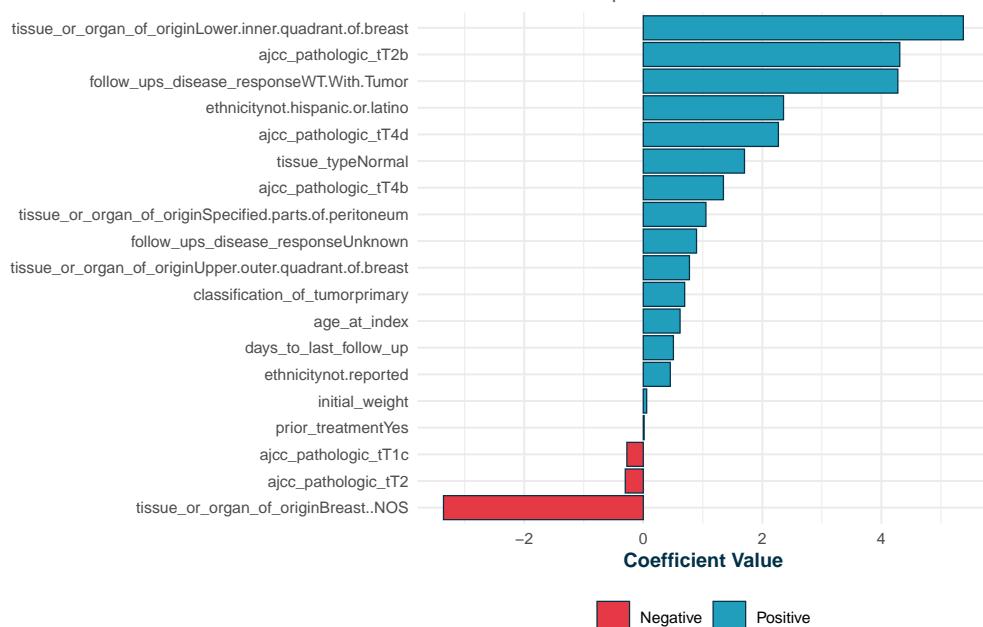


ADAPTIVE – Train vs Test AUC



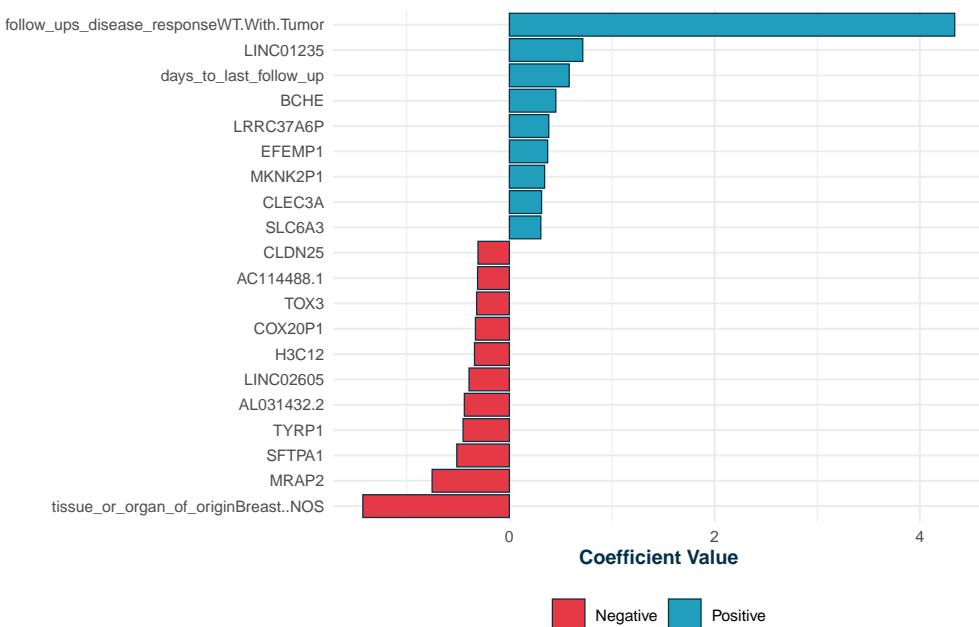
ADAPTIVE – Clinical_Only

Top 19 non-zero features



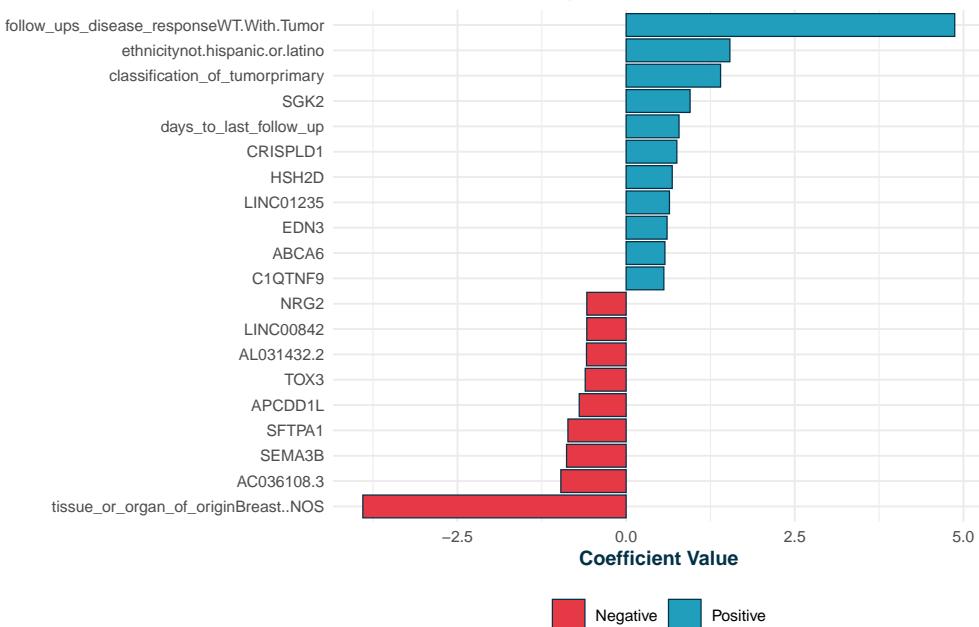
ADAPTIVE – Clinical_TOP5000

Top 20 non-zero features



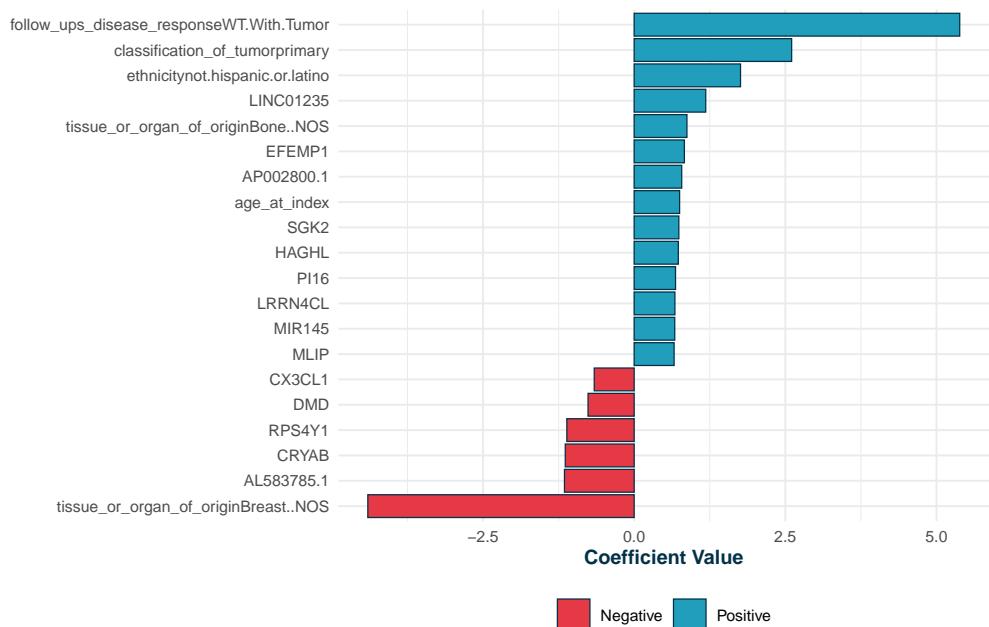
ADAPTIVE – Clinical_TOP1000

Top 20 non-zero features



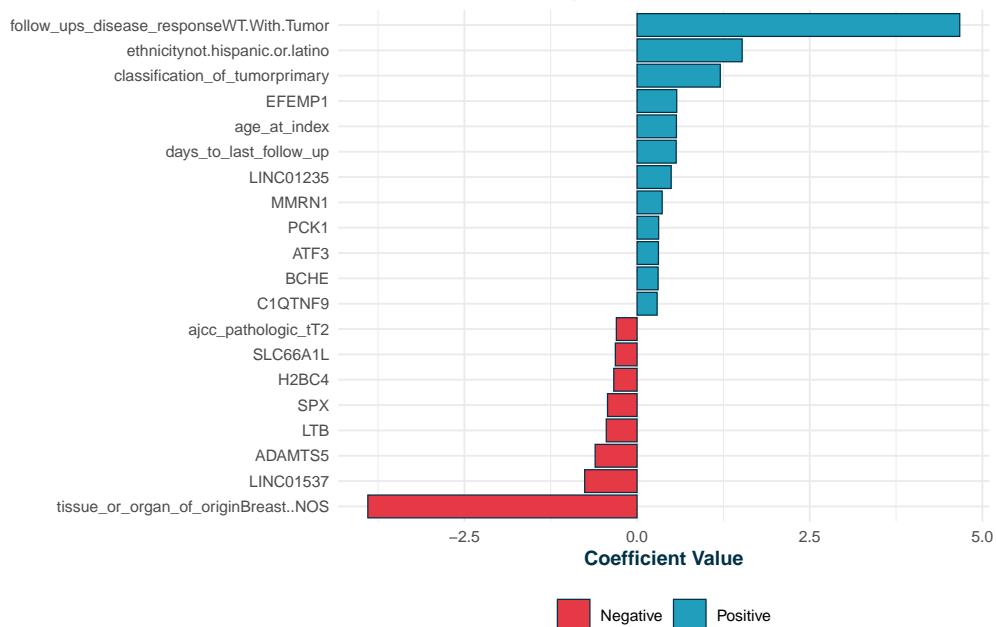
ADAPTIVE – Clinical_TOP500

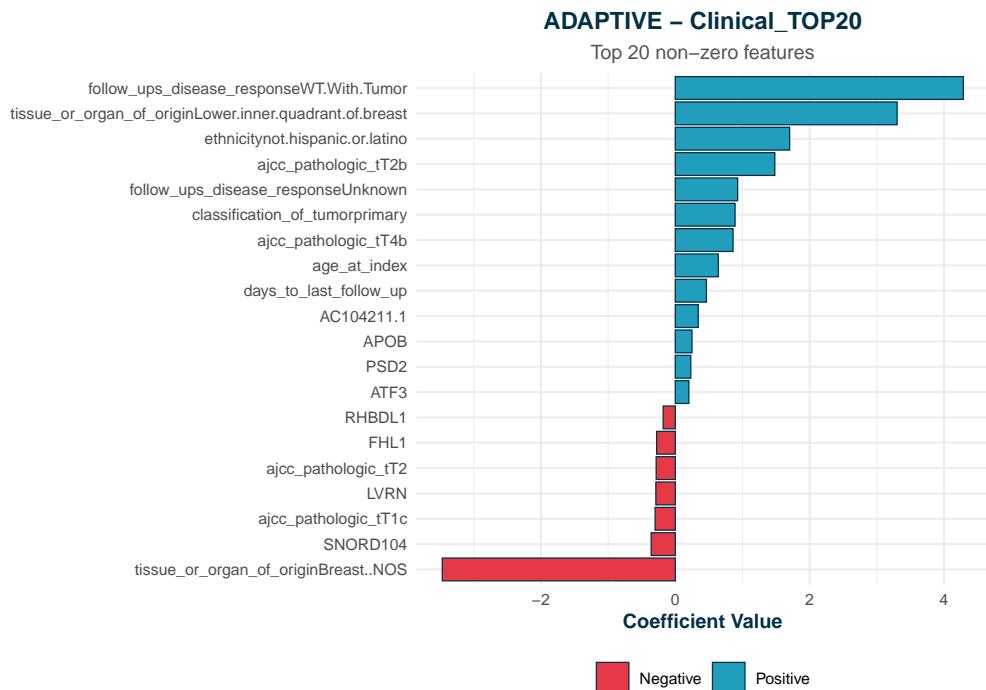
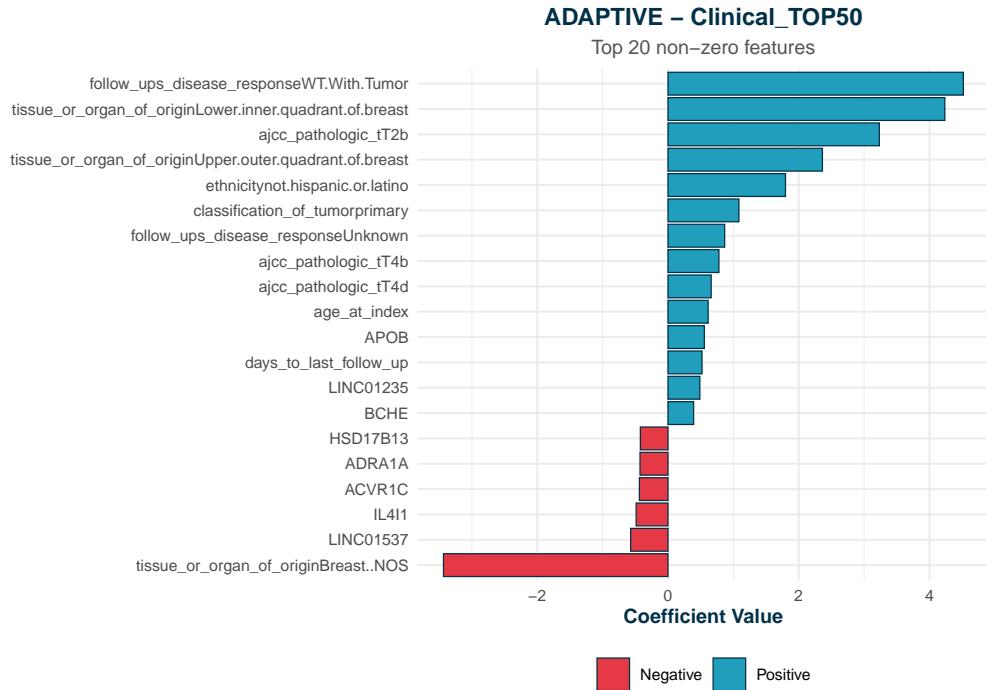
Top 20 non-zero features



ADAPTIVE – Clinical_TOP100

Top 20 non-zero features





```
adaptive_lasso_smote_metrics <- plot_classification_metrics_single(adaptive_lasso_smote
, threshold = 0.5
, csv_filename = "adaptive_lasso_smote.csv")
```

```
##  
## === CLASSIFICATION METRICS ===
```

```

## Clinical_Only:
##   TP=29 TN=188 FP=18 FN=11
##   Accuracy=0.882 Precision=0.617 Recall=0.725 F1=0.667 AUC=0.901

## Clinical_TOP5000:
##   TP=29 TN=189 FP=17 FN=11
##   Accuracy=0.886 Precision=0.630 Recall=0.725 F1=0.674 AUC=0.865

## Clinical_TOP1000:
##   TP=29 TN=186 FP=20 FN=11
##   Accuracy=0.874 Precision=0.592 Recall=0.725 F1=0.652 AUC=0.859

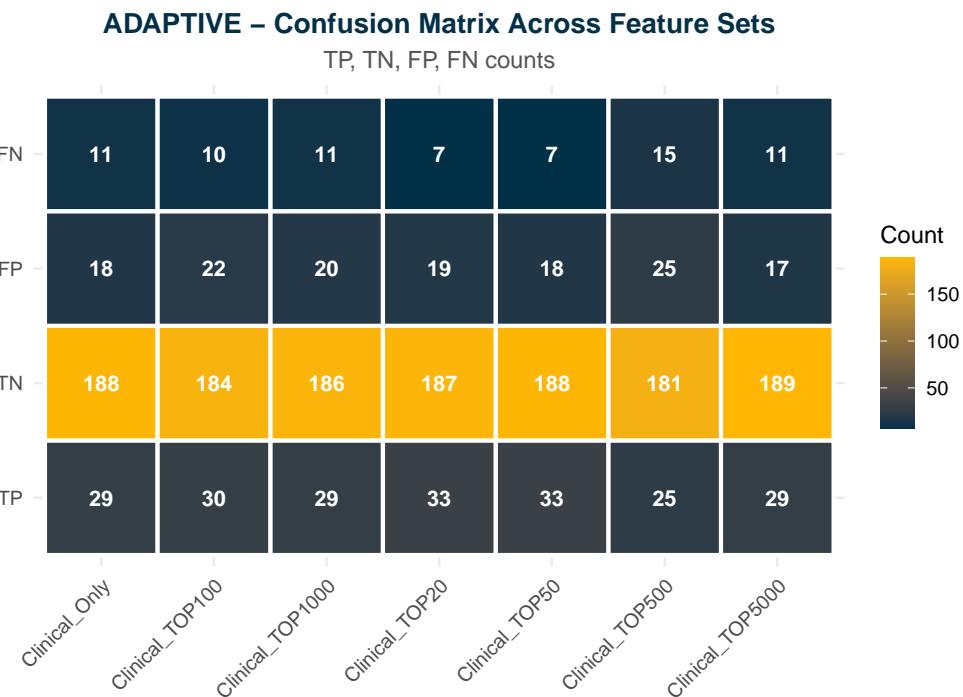
## Clinical_TOP500:
##   TP=25 TN=181 FP=25 FN=15
##   Accuracy=0.837 Precision=0.500 Recall=0.625 F1=0.556 AUC=0.837

## Clinical_TOP100:
##   TP=30 TN=184 FP=22 FN=10
##   Accuracy=0.870 Precision=0.577 Recall=0.750 F1=0.652 AUC=0.859

## Clinical_TOP50:
##   TP=33 TN=188 FP=18 FN=7
##   Accuracy=0.898 Precision=0.647 Recall=0.825 F1=0.725 AUC=0.896

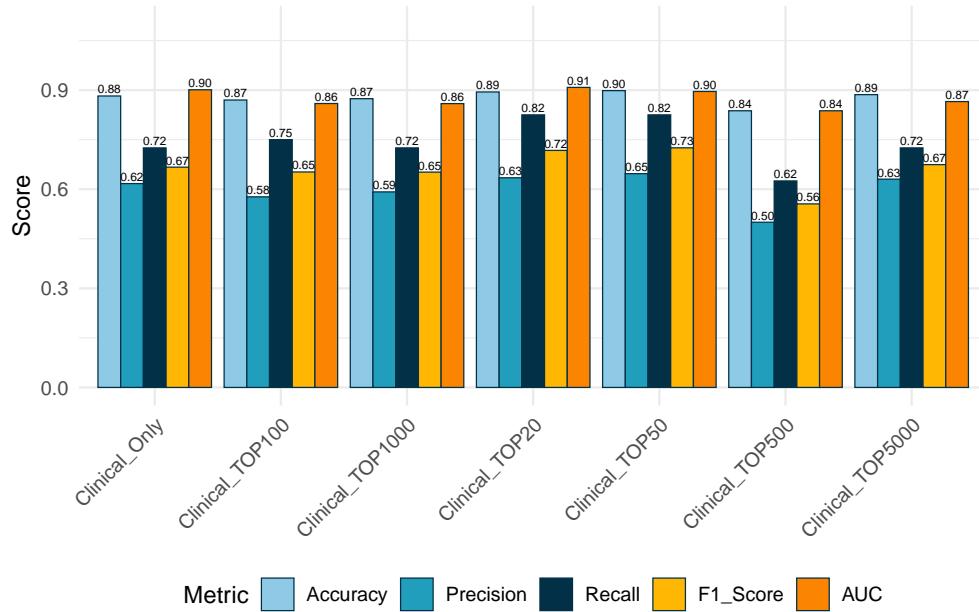
## Clinical_TOP20:
##   TP=33 TN=187 FP=19 FN=7
##   Accuracy=0.894 Precision=0.635 Recall=0.825 F1=0.717 AUC=0.908

```



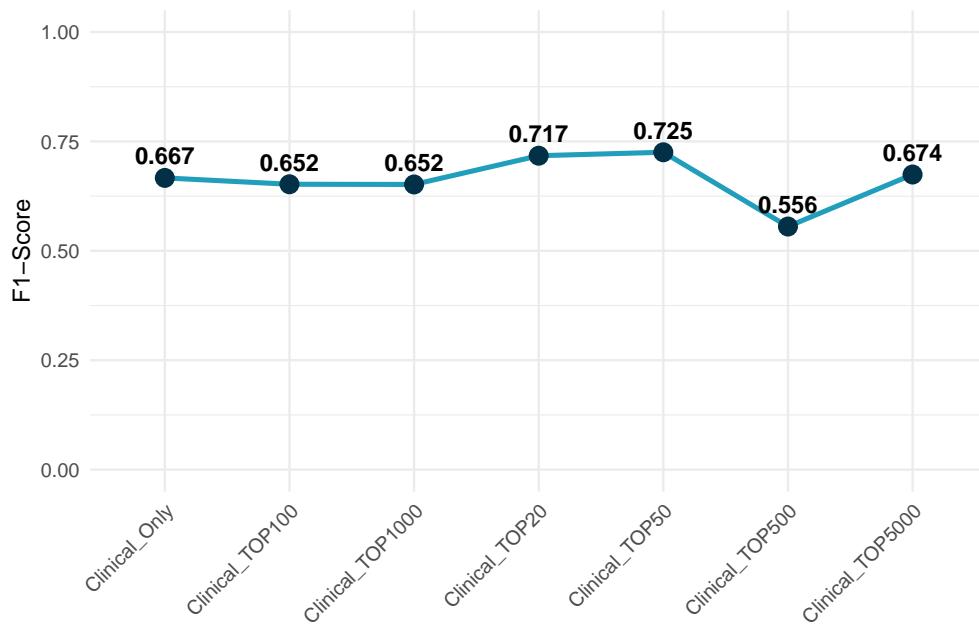
ADAPTIVE – Classification Metrics

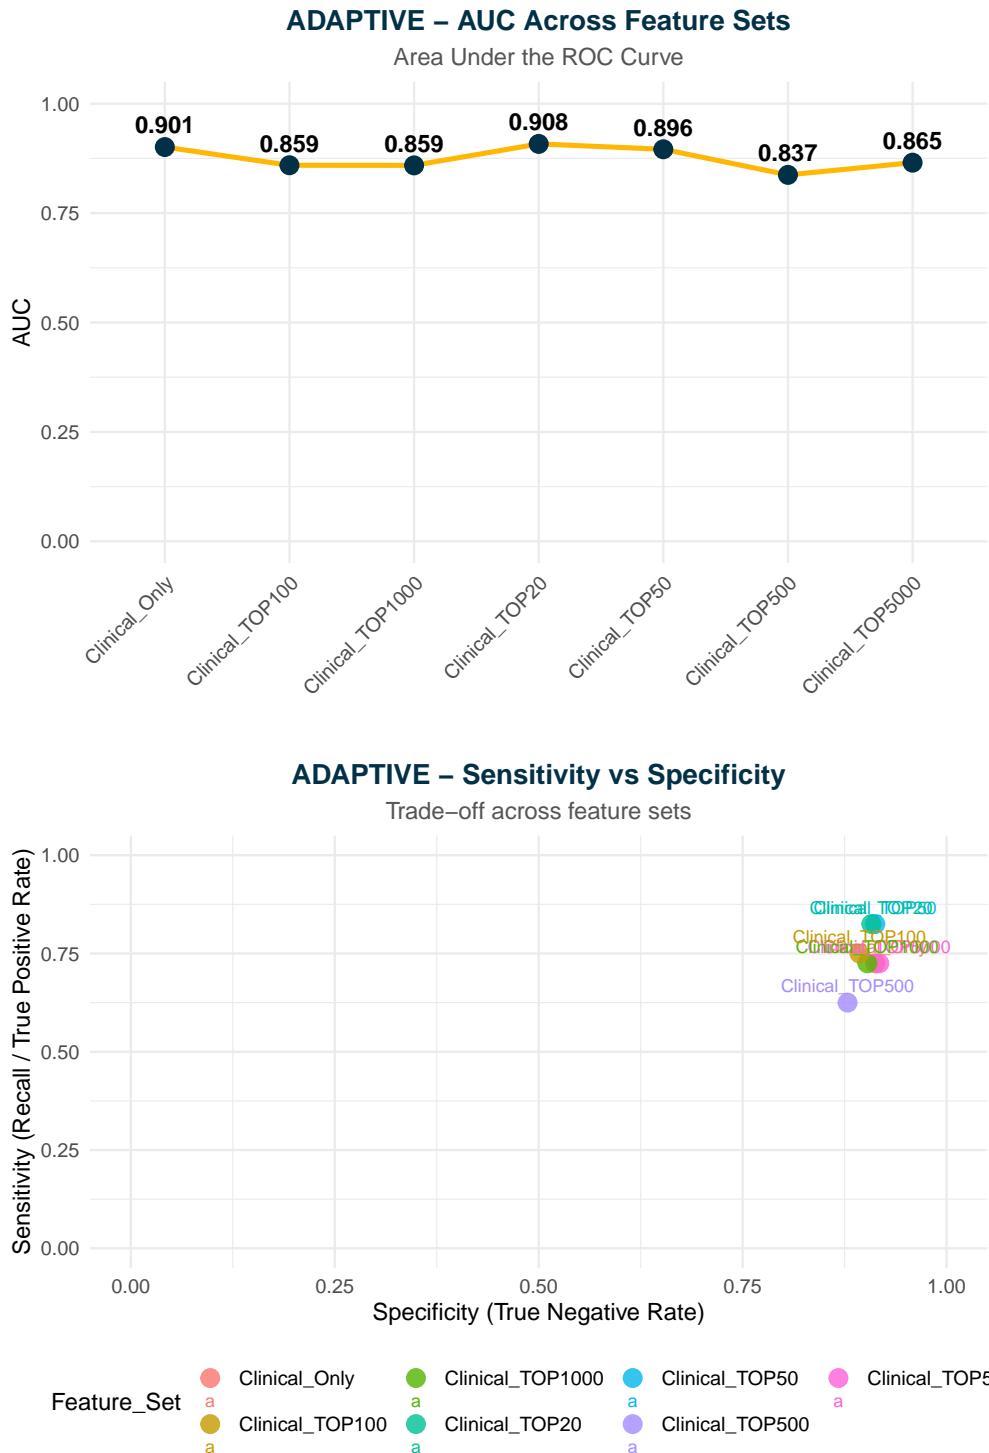
Accuracy, Precision, Recall, F1-Score, AUC



ADAPTIVE – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1   Clinical_Only 29 188 18 11 0.8821138 0.6170213 0.725 0.9126214
## 2 Clinical_TOP5000 29 189 17 11 0.8861789 0.6304348 0.725 0.9174757
## 3 Clinical_TOP1000 29 186 20 11 0.8739837 0.5918367 0.725 0.9029126
## 4   Clinical_TOP500 25 181 25 15 0.8373984 0.5000000 0.625 0.8786408
```

```

## 5 Clinical_TOP100 30 184 22 10 0.8699187 0.5769231 0.750 0.8932039
## 6 Clinical_TOP50 33 188 18 7 0.8983740 0.6470588 0.825 0.9126214
## 7 Clinical_TOP20 33 187 19 7 0.8943089 0.6346154 0.825 0.9077670
## F1_Score AUC
## 1 0.6666667 0.9009709
## 2 0.6744186 0.8652913
## 3 0.6516854 0.8591019
## 4 0.5555556 0.8373786
## 5 0.6521739 0.8593447
## 6 0.7252747 0.8958738
## 7 0.7173913 0.9081311
##
## Exported classification metrics to: model_metrics/adaptive_lasso_smote_classification_metrics.csv

```

UniLasso with SMOTE

```

unilasso_smote <- fit_single_model_across_features(
  model_type = "unilasso"
  , X_train_all = smote_data$X_train
  , X_test_all = X_test
  , Y_train = smote_data$Y_train
  , Y_test = Y_test
  , n_clinical = n_clinical
  , top_genes_ranked = top_genes
  , gene_sets = c(5000, 1000, 500, 100, 50, 20)
)

##
## === FITTING UNILASSO ACROSS FEATURE SETS ===
##
## Fitting Clinical_Only...
## Fitting Clinical_TOP5000...
## Fitting Clinical_TOP1000...
## Fitting Clinical_TOP500...
## Fitting Clinical_TOP100...
## Fitting Clinical_TOP50...
## Fitting Clinical_TOP20...

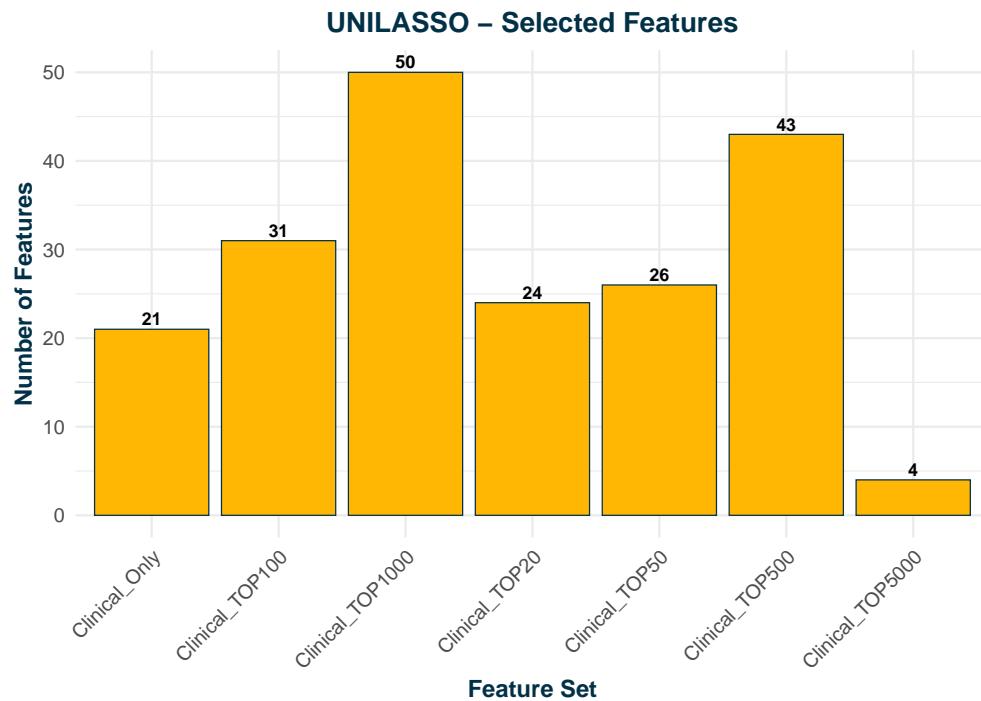
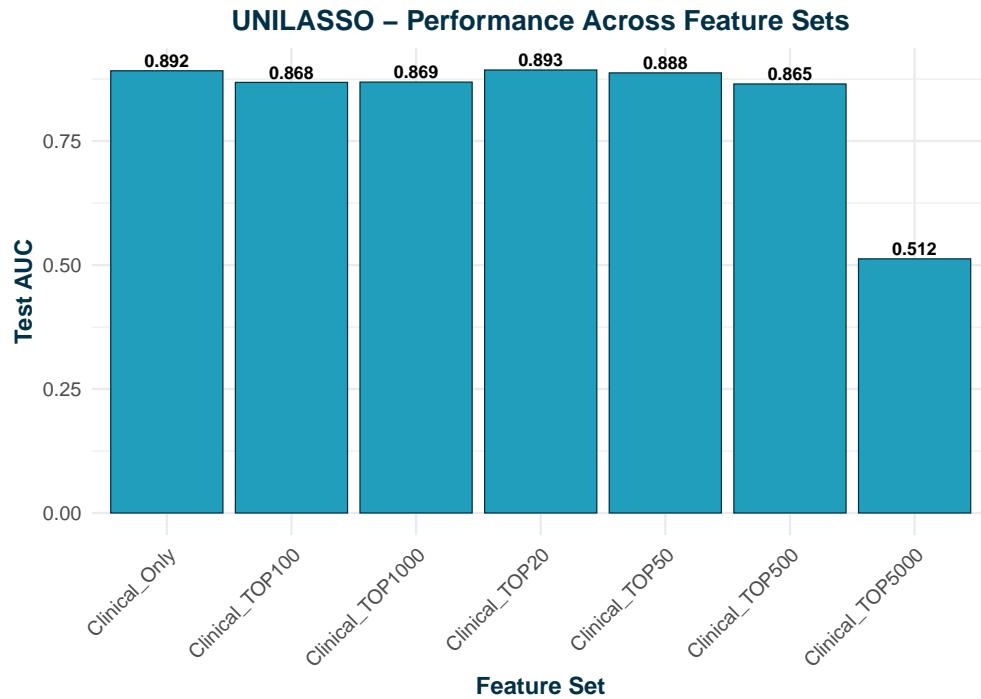
##
## === SUMMARY TABLE ===
##      Feature_Set Model Features Train_AUC Test_AUC Test_Accuracy
## 1 Clinical_Only UNILASSO        21 0.9490329 0.8917476    0.8821138
## 2 Clinical_TOP5000 UNILASSO        4 0.5950311 0.5125000    0.8414634
## 3 Clinical_TOP1000 UNILASSO       50 0.9685456 0.8689320    0.9065041

```

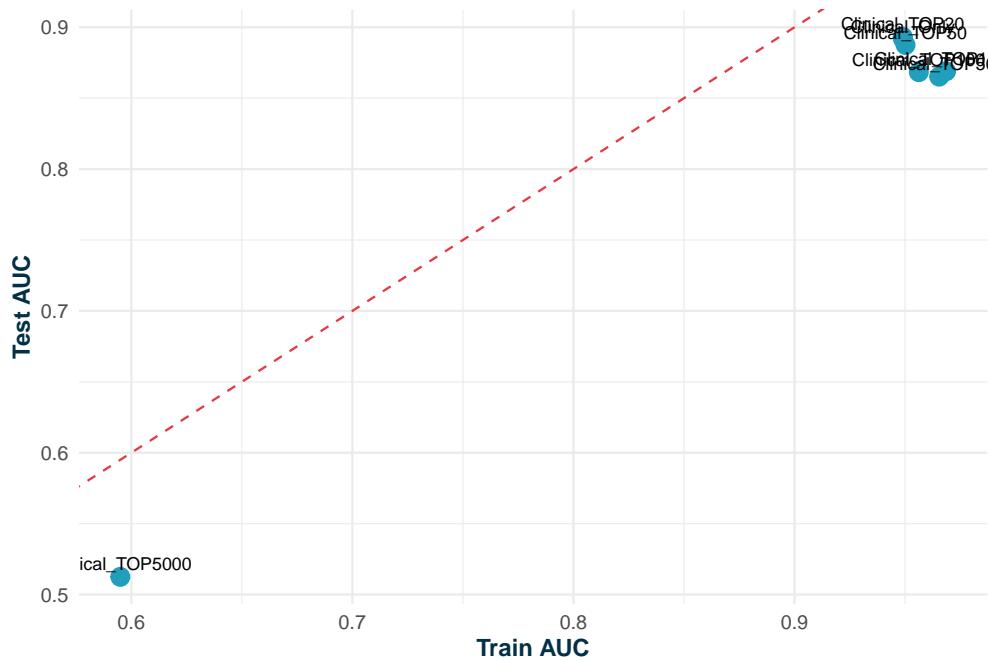
```

## 4 Clinical_TOP500 UNILASSO      43 0.9654393 0.8651699 0.9105691
## 5 Clinical_TOP100 UNILASSO     31 0.9562229 0.8683252 0.9065041
## 6 Clinical_TOP50 UNILASSO      26 0.9502049 0.8875000 0.9186992
## 7 Clinical_TOP20 UNILASSO      24 0.9489279 0.8933252 0.9186992
## Exported metrics to: model_metrics/unilasso_across_features_metrics.csv

```

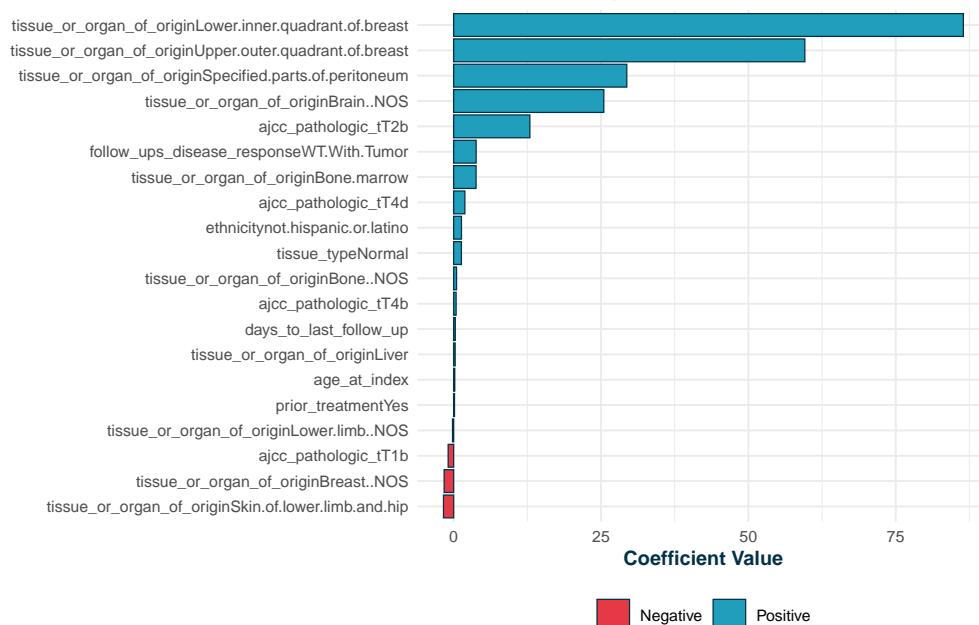


UNILASSO – Train vs Test AUC



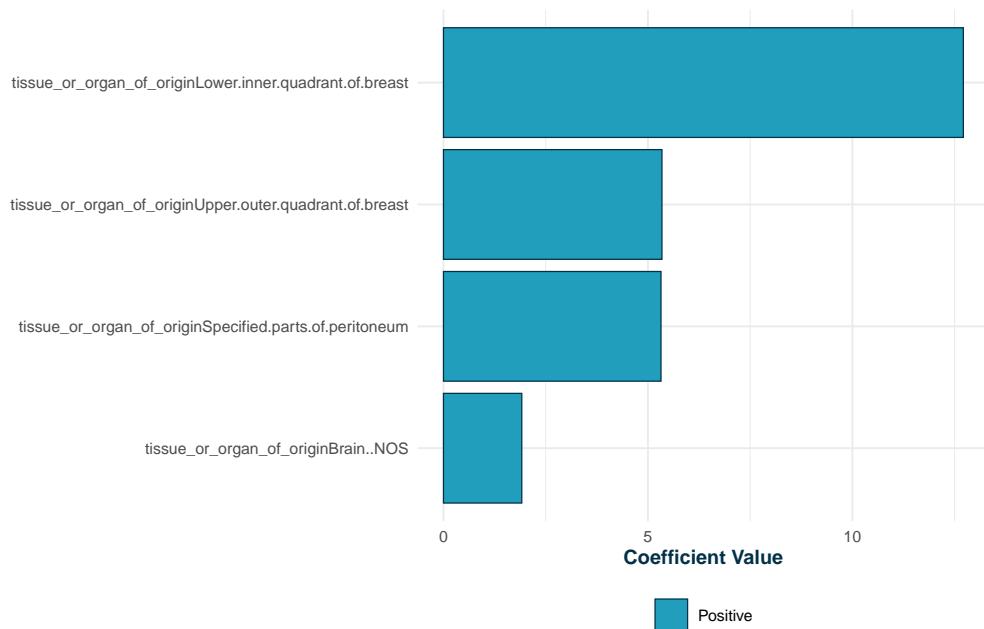
UNILASSO – Clinical_Only

Top 20 non-zero features



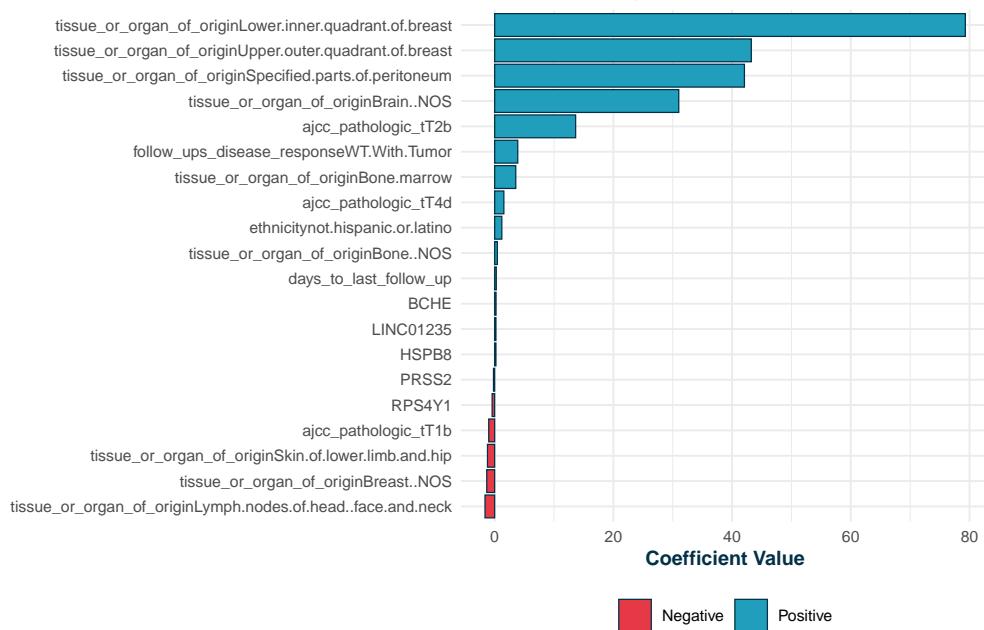
UNILASSO – Clinical_TOP5000

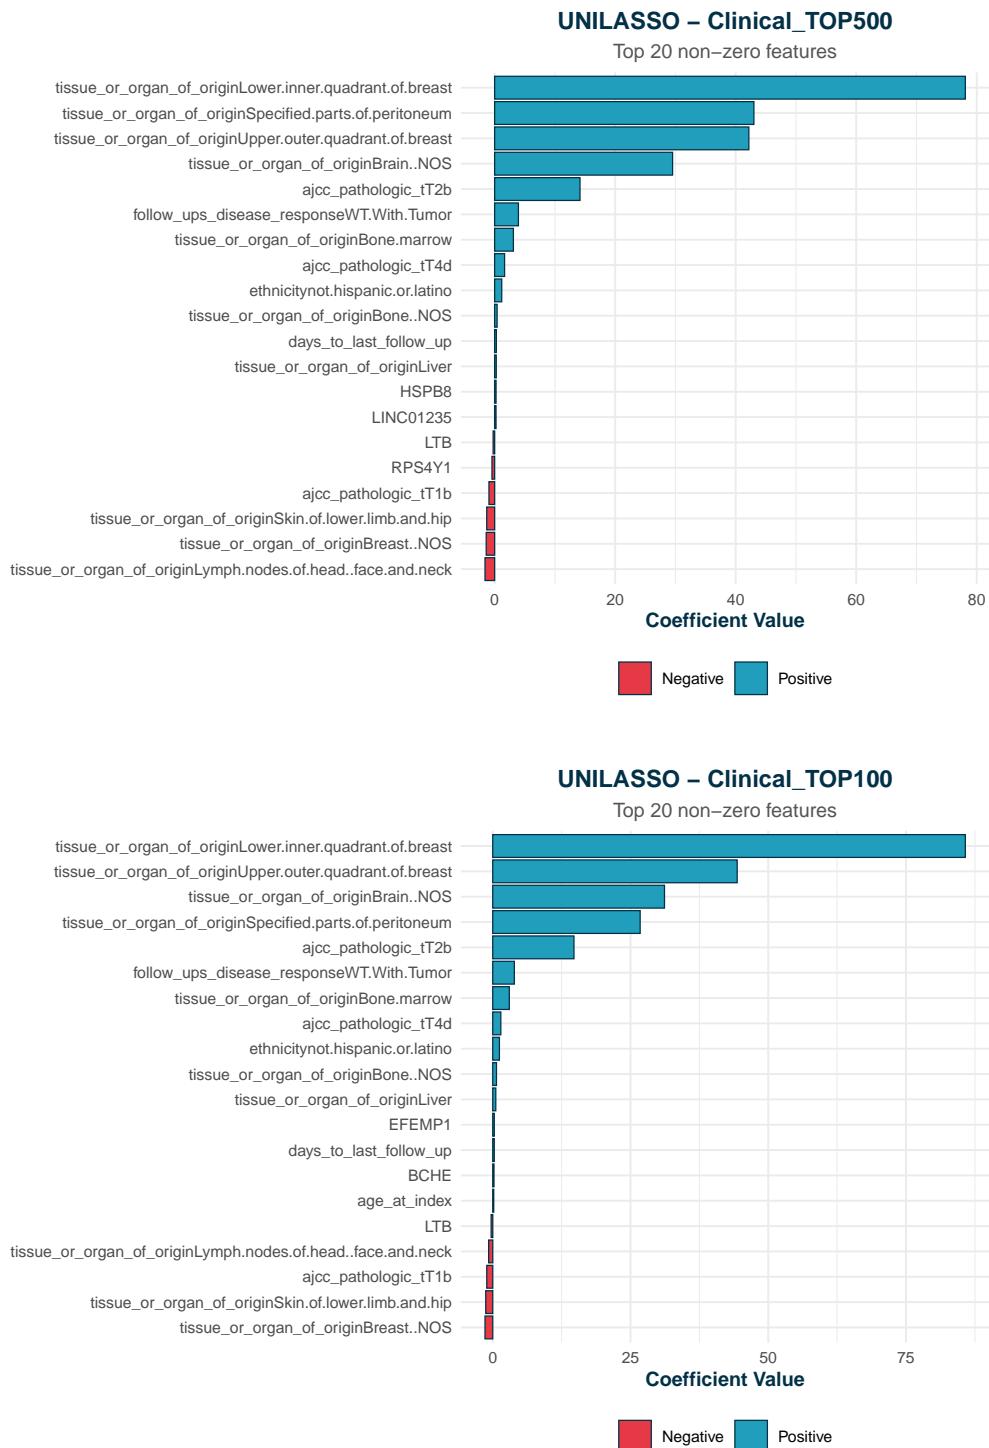
Top 4 non-zero features

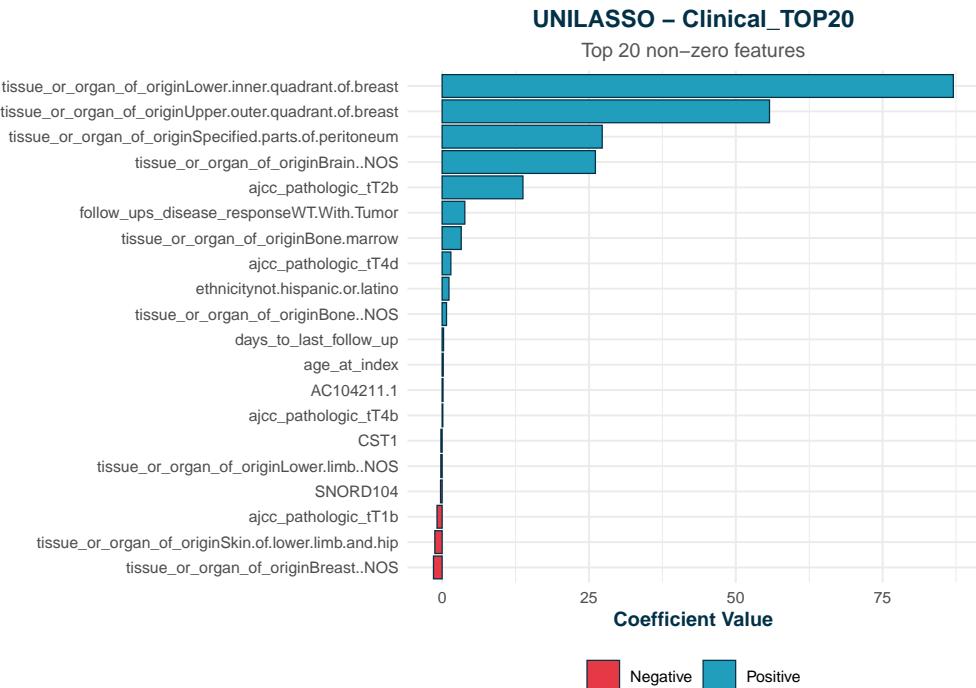
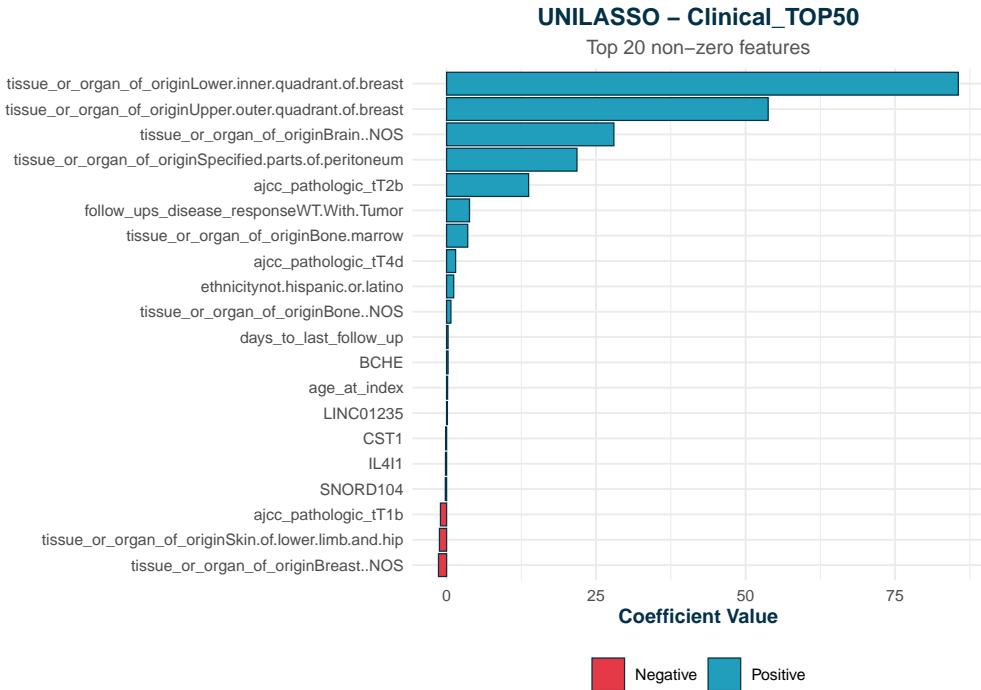


UNILASSO – Clinical_TOP1000

Top 20 non-zero features







```
unilasso_smote_metrics <- plot_classification_metrics_single(unilasso_smote
  , threshold = 0.5
  , csv_filename = "unilasso_smote_classification_metrics.csv")
```

```
##  
## === CLASSIFICATION METRICS ===
```

```

## Clinical_Only:
##   TP=26 TN=191 FP=15 FN=14
##   Accuracy=0.882 Precision=0.634 Recall=0.650 F1=0.642 AUC=0.892

## Clinical_TOP5000:
##   TP=1 TN=206 FP=0 FN=39
##   Accuracy=0.841 Precision=1.000 Recall=0.025 F1=0.049 AUC=0.512

## Clinical_TOP1000:
##   TP=28 TN=195 FP=11 FN=12
##   Accuracy=0.907 Precision=0.718 Recall=0.700 F1=0.709 AUC=0.869

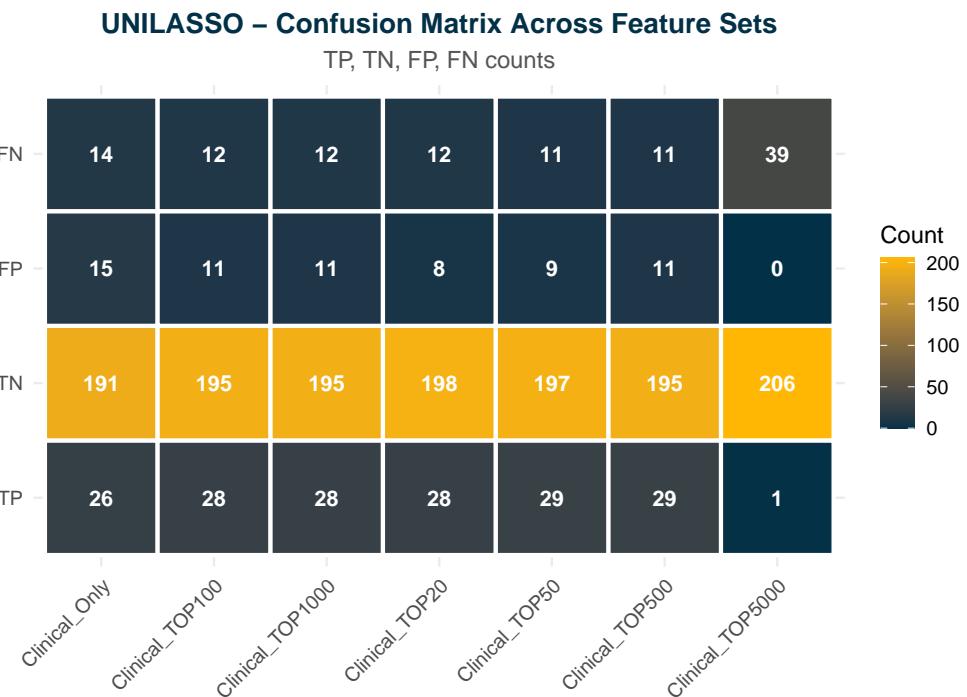
## Clinical_TOP500:
##   TP=29 TN=195 FP=11 FN=11
##   Accuracy=0.911 Precision=0.725 Recall=0.725 F1=0.725 AUC=0.865

## Clinical_TOP100:
##   TP=28 TN=195 FP=11 FN=12
##   Accuracy=0.907 Precision=0.718 Recall=0.700 F1=0.709 AUC=0.868

## Clinical_TOP50:
##   TP=29 TN=197 FP=9 FN=11
##   Accuracy=0.919 Precision=0.763 Recall=0.725 F1=0.744 AUC=0.888

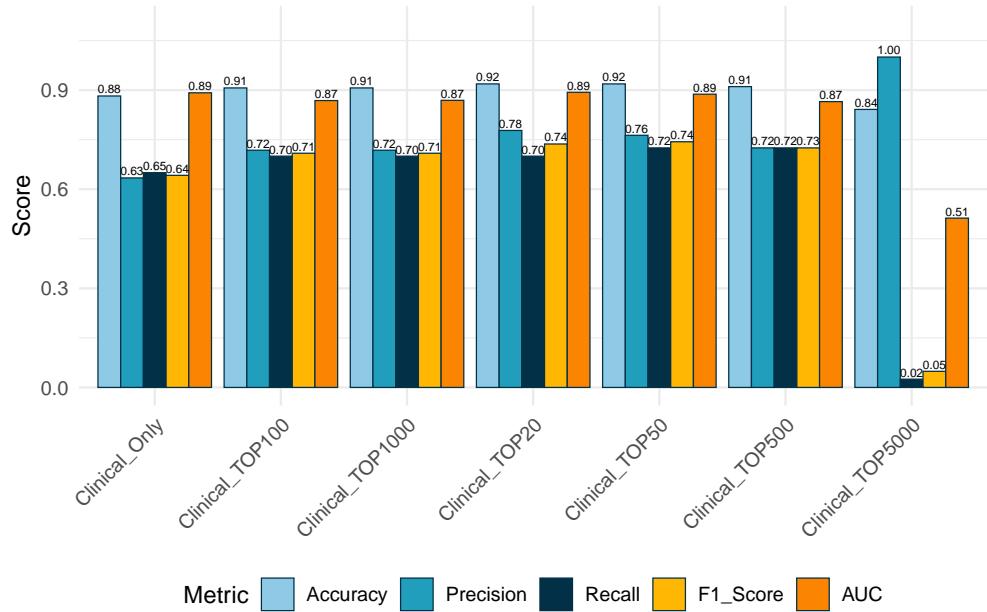
## Clinical_TOP20:
##   TP=28 TN=198 FP=8 FN=12
##   Accuracy=0.919 Precision=0.778 Recall=0.700 F1=0.737 AUC=0.893

```



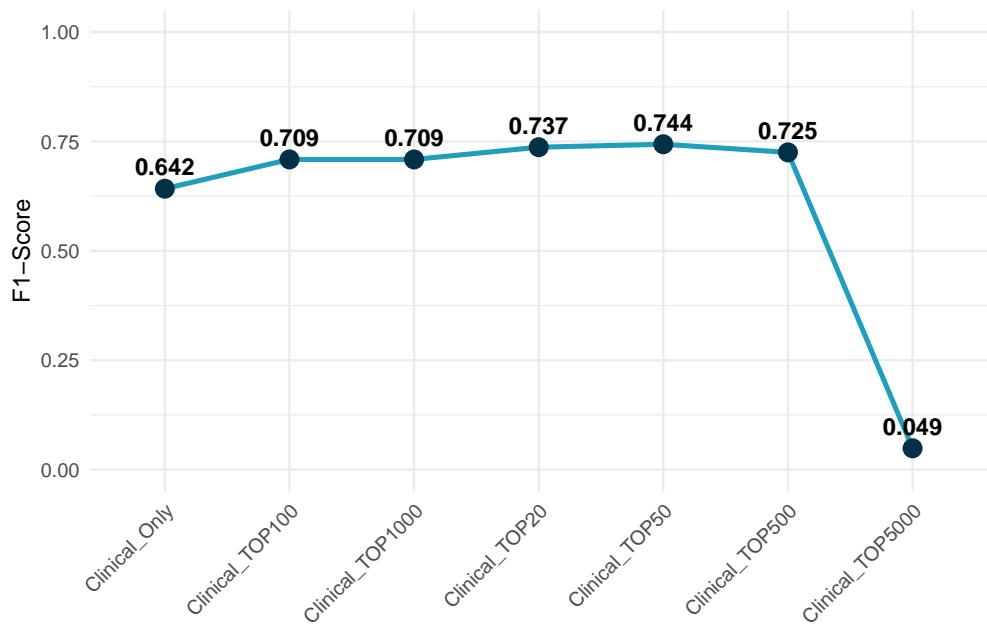
UNILASSO – Classification Metrics

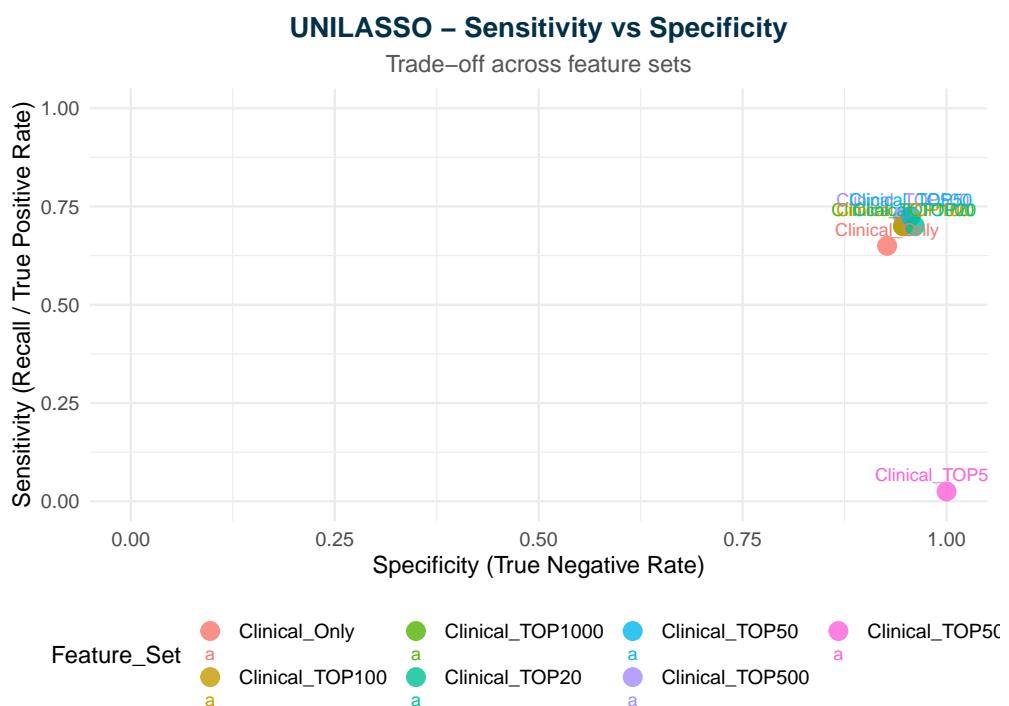
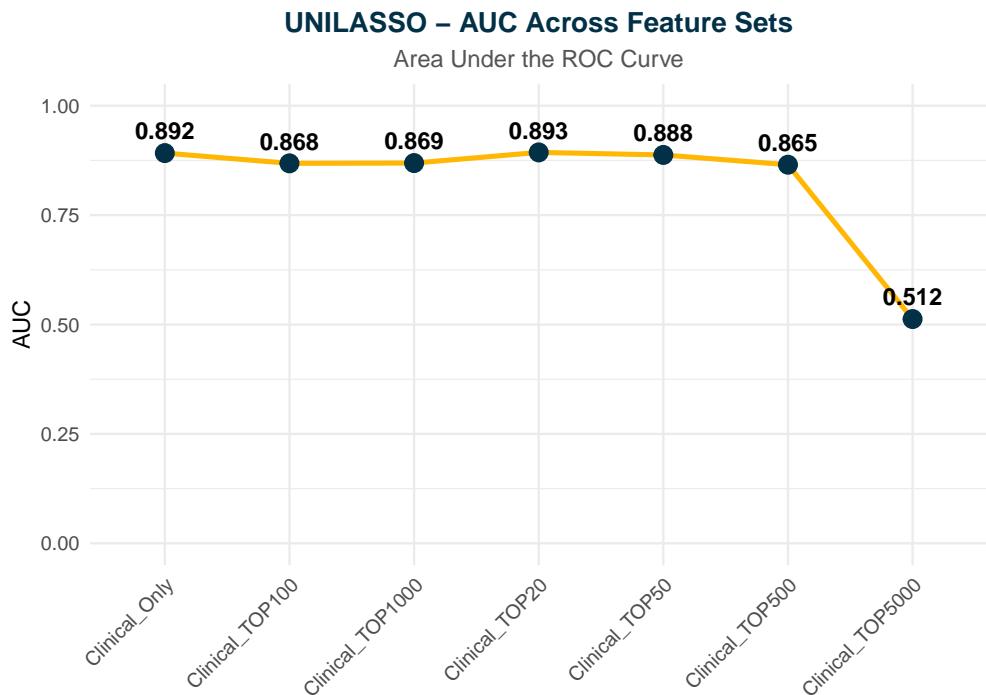
Accuracy, Precision, Recall, F1-Score, AUC



UNILASSO – F1-Score Across Feature Sets

Trend of model performance





```
##
## === SUMMARY TABLE ===
##      Feature_Set TP TN FP FN Accuracy Precision Recall Specificity
## 1  Clinical_Only 26 191 15 14 0.8821138 0.6341463 0.650 0.9271845
## 2 Clinical_TOP5000  1 206  0 39 0.8414634 1.0000000 0.025 1.0000000
## 3 Clinical_TOP1000 28 195 11 12 0.9065041 0.7179487 0.700 0.9466019
## 4  Clinical_TOP500 29 195 11 11 0.9105691 0.7250000 0.725 0.9466019
```

```

## 5 Clinical_TOP100 28 195 11 12 0.9065041 0.7179487 0.700 0.9466019
## 6 Clinical_TOP50 29 197 9 11 0.9186992 0.7631579 0.725 0.9563107
## 7 Clinical_TOP20 28 198 8 12 0.9186992 0.7777778 0.700 0.9611650
## F1_Score AUC
## 1 0.64197531 0.8917476
## 2 0.04878049 0.5125000
## 3 0.70886076 0.8689320
## 4 0.72500000 0.8651699
## 5 0.70886076 0.8683252
## 6 0.74358974 0.8875000
## 7 0.73684211 0.8933252
##
## Exported classification metrics to: model_metrics/unilasso_smote_classification_metrics.csv

```

SMOTE Impact Comparison

```
cat("\n==== RIDGE: SMOTE vs NO SMOTE COMPARISON ====\n\n")
```

Ridge Comparison

```

## 
## === RIDGE: SMOTE vs NO SMOTE COMPARISON ===

cat("Before SMOTE (Clinical_Only):\n")
```

```
## Before SMOTE (Clinical_Only):
```

```
cat("  Recall:", sprintf("%.3f", ridge_metrics$Recall[1]), "\n")
```

```
##   Recall: 0.175
```

```
cat("  Precision:", sprintf("%.3f", ridge_metrics$Precision[1]), "\n")
```

```
##   Precision: 0.778
```

```
cat("  F1-Score:", sprintf("%.3f", ridge_metrics$F1_Score[1]), "\n")
```

```
##   F1-Score: 0.286
```

```
cat("  AUC:", sprintf("%.3f", ridge_metrics$AUC[1]), "\n\n")
```

```
##   AUC: 0.883
```

```
cat("After SMOTE (Clinical_Only):\n")
```

```
## After SMOTE (Clinical_Only):
```

```

cat("  Recall:", sprintf("%.3f", ridge_smote_metrics$Recall[1]), "\n")

##   Recall: 0.725

cat("  Precision:", sprintf("%.3f", ridge_smote_metrics$Precision[1]), "\n")

##   Precision: 0.547

cat("  F1-Score:", sprintf("%.3f", ridge_smote_metrics$F1_Score[1]), "\n")

##   F1-Score: 0.624

cat("  AUC:", sprintf("%.3f", ridge_smote_metrics$AUC[1]), "\n\n")

##   AUC: 0.884

cat("Improvement:\n")

## Improvement:

cat("  Recall:", sprintf("%+.3f", ridge_smote_metrics$Recall[1] - ridge_metrics$Recall[1]), "\n")

##   Recall: +0.550

cat("  F1-Score:", sprintf("%+.3f", ridge_smote_metrics$F1_Score[1] - ridge_metrics$F1_Score[1]), "\n")

##   F1-Score: +0.338

comparison_ridge <- rbind(
  data.frame(Method = "No SMOTE", ridge_metrics[1, c("Feature_Set", "Recall", "Precision", "F1_Score")]),
  data.frame(Method = "SMOTE", ridge_smote_metrics[1, c("Feature_Set", "Recall", "Precision", "F1_Score")])
)

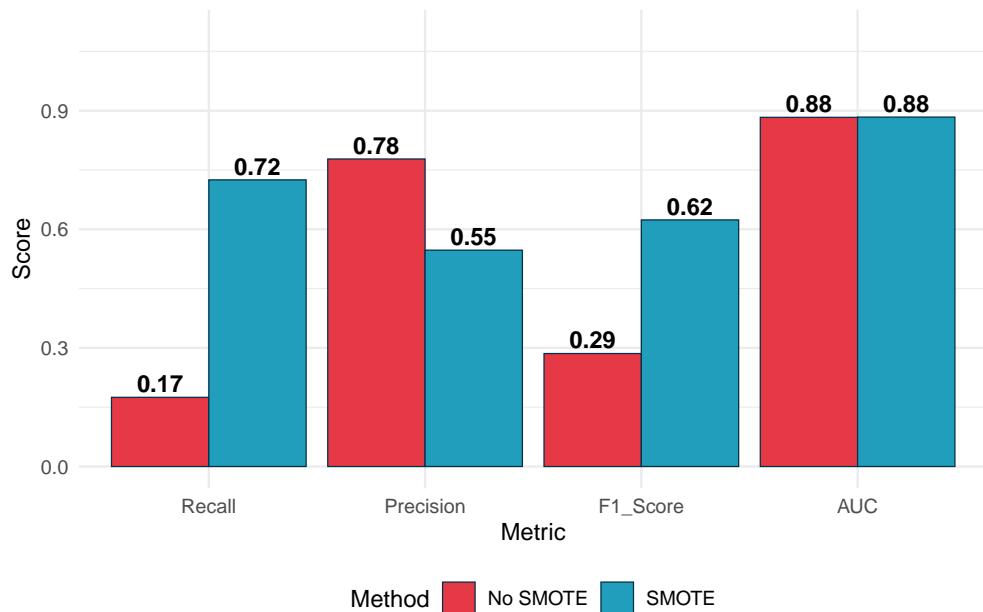
comp_ridge_long <- reshape2::melt(comparison_ridge, id.vars = c("Method", "Feature_Set"))

ggplot(comp_ridge_long, aes(x = variable, y = value, fill = Method)) +
  geom_bar(stat = "identity", position = "dodge", color = "#023047", linewidth = 0.3) +
  geom_text(aes(label = sprintf("%.2f", value)),
            position = position_dodge(width = 0.9),
            vjust = -0.3,
            fontface = "bold") +
  labs(title = "Ridge: Impact of SMOTE on Clinical_Only Model",
       subtitle = "Comparison of key metrics",
       x = "Metric",
       y = "Score") +
  theme_minimal(base_size = 12) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047"),
        plot.subtitle = element_text(hjust = 0.5, color = "#555555"),
        legend.position = "bottom") +
  scale_fill_manual(values = c("No SMOTE" = "#e63946", "SMOTE" = "#219ebc")) +
  ylim(0, 1.1)

```

Ridge: Impact of SMOTE on Clinical_Only Model

Comparison of key metrics



```
cat("\n==== LASSO: SMOTE vs NO SMOTE COMPARISON ===\n\n")
```

Lasso Comparison

```
##  
## === LASSO: SMOTE vs NO SMOTE COMPARISON ===
```

```
cat("Before SMOTE (Clinical_Only):\n")
```

```
## Before SMOTE (Clinical_Only):
```

```
cat("  Recall:", sprintf("%.3f", lasso_metrics$Recall[1]), "\n")
```

```
##  Recall: 0.400
```

```
cat("  Precision:", sprintf("%.3f", lasso_metrics$Precision[1]), "\n")
```

```
##  Precision: 0.762
```

```
cat("  F1-Score:", sprintf("%.3f", lasso_metrics$F1_Score[1]), "\n")
```

```
##  F1-Score: 0.525
```

```

cat("  AUC:", sprintf("%.3f", lasso_metrics$AUC[1]), "\n\n")

##   AUC: 0.886

cat("After SMOTE (Clinical_Only):\n")

## After SMOTE (Clinical_Only):

cat("  Recall:", sprintf("%.3f", lasso_smote_metrics$Recall[1]), "\n")

##   Recall: 0.725

cat("  Precision:", sprintf("%.3f", lasso_smote_metrics$Precision[1]), "\n")

##   Precision: 0.592

cat("  F1-Score:", sprintf("%.3f", lasso_smote_metrics$F1_Score[1]), "\n")

##   F1-Score: 0.652

cat("  AUC:", sprintf("%.3f", lasso_smote_metrics$AUC[1]), "\n\n")

##   AUC: 0.898

cat("Improvement:\n")

## Improvement:

cat("  Recall:", sprintf("%+.3f", lasso_smote_metrics$Recall[1] - lasso_metrics$Recall[1]), "\n")

##   Recall: +0.325

cat("  F1-Score:", sprintf("%+.3f", lasso_smote_metrics$F1_Score[1] - lasso_metrics$F1_Score[1]), "\n")

##   F1-Score: +0.127

comparison_lasso <- rbind(
  data.frame(Method = "No SMOTE", lasso_metrics[1, c("Feature_Set", "Recall", "Precision", "F1_Score")]),
  data.frame(Method = "SMOTE", lasso_smote_metrics[1, c("Feature_Set", "Recall", "Precision", "F1_Score")])
)

comp_lasso_long <- reshape2::melt(comparison_lasso, id.vars = c("Method", "Feature_Set"))

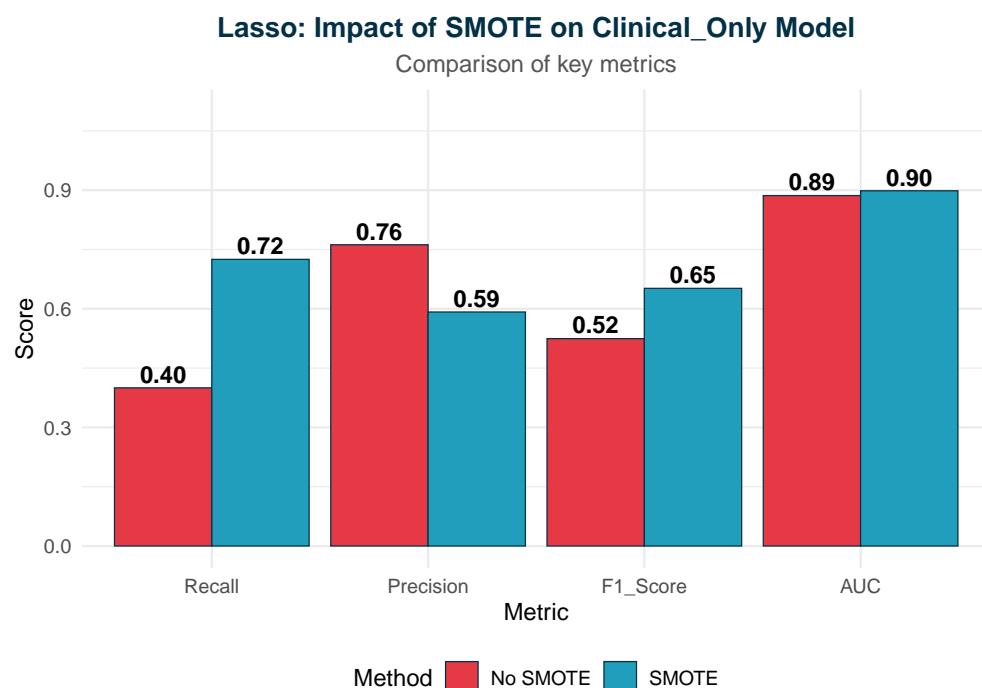
ggplot(comp_lasso_long, aes(x = variable, y = value, fill = Method)) +
  geom_bar(stat = "identity", position = "dodge", color = "#023047", linewidth = 0.3) +
  geom_text(aes(label = sprintf("%.2f", value)))

```

```

        , position = position_dodge(width = 0.9)
        , vjust = -0.3
        , fontface = "bold") +
  labs(title = "Lasso: Impact of SMOTE on Clinical_Only Model"
    , subtitle = "Comparison of key metrics"
    , x = "Metric"
    , y = "Score") +
  theme_minimal(base_size = 12) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047")
    , plot.subtitle = element_text(hjust = 0.5, color = "#555555")
    , legend.position = "bottom") +
  scale_fill_manual(values = c("No SMOTE" = "#e63946", "SMOTE" = "#219ebc")) +
  ylim(0, 1.1)

```



```
cat("\n==== ELASTICNET: SMOTE vs NO SMOTE COMPARISON ===\n\n")
```

ElasticNet Comparison

```

## 
## === ELASTICNET: SMOTE vs NO SMOTE COMPARISON ===

cat("Before SMOTE (Clinical_Only):\n")

## Before SMOTE (Clinical_Only):

```

```

cat("  Recall:", sprintf("%.3f", elasticnet_metrics$Recall[1]), "\n")

##   Recall: 0.250

cat("  Precision:", sprintf("%.3f", elasticnet_metrics$Precision[1]), "\n")

##   Precision: 0.769

cat("  F1-Score:", sprintf("%.3f", elasticnet_metrics$F1_Score[1]), "\n")

##   F1-Score: 0.377

cat("  AUC:", sprintf("%.3f", elasticnet_metrics$AUC[1]), "\n\n")

##   AUC: 0.885

cat("After SMOTE (Clinical_Only):\n")

## After SMOTE (Clinical_Only):

cat("  Recall:", sprintf("%.3f", elasticnet_smote_metrics$Recall[1]), "\n")

##   Recall: 0.725

cat("  Precision:", sprintf("%.3f", elasticnet_smote_metrics$Precision[1]), "\n")

##   Precision: 0.580

cat("  F1-Score:", sprintf("%.3f", elasticnet_smote_metrics$F1_Score[1]), "\n")

##   F1-Score: 0.644

cat("  AUC:", sprintf("%.3f", elasticnet_smote_metrics$AUC[1]), "\n\n")

##   AUC: 0.899

cat("Improvement:\n")

## Improvement:

cat("  Recall:", sprintf("%+.3f", elasticnet_smote_metrics$Recall[1] - elasticnet_metrics$Recall[1]), "\n")

##   Recall: +0.475

```

```

cat("  F1-Score:", sprintf("%+.3f", elasticnet_smote_metrics$F1_Score[1] - elasticnet_metrics$F1_Score[1]))

##   F1-Score: +0.267

comparison_elasticnet <- rbind(
  data.frame(Method = "No SMOTE", elasticnet_metrics[1, c("Feature_Set", "Recall", "Precision", "F1_Score")]),
  data.frame(Method = "SMOTE", elasticnet_smote_metrics[1, c("Feature_Set", "Recall", "Precision", "F1_Score")])
)

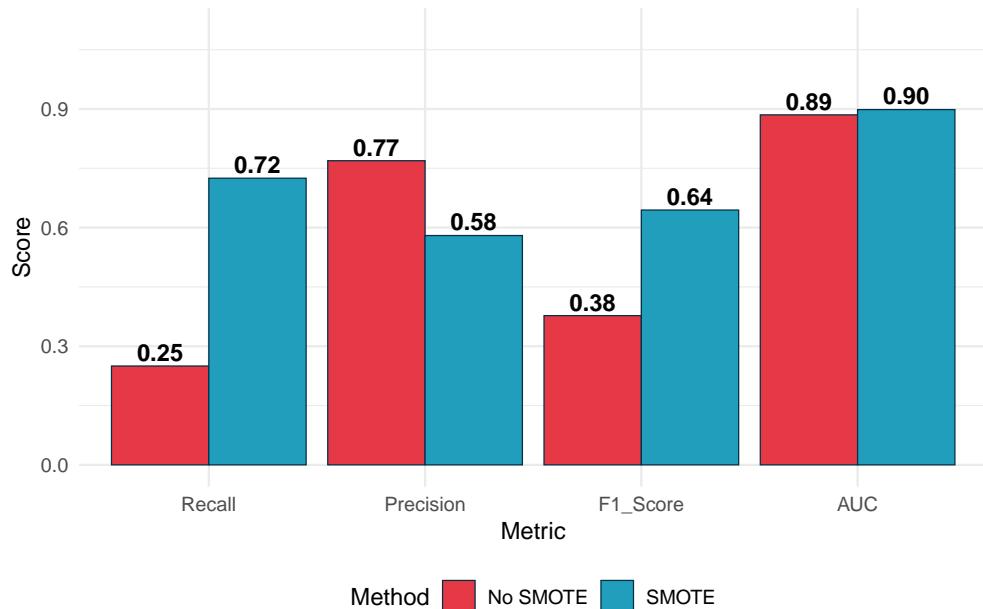
comp_elasticnet_long <- reshape2::melt(comparison_elasticnet, id.vars = c("Method", "Feature_Set"))

ggplot(comp_elasticnet_long, aes(x = variable, y = value, fill = Method)) +
  geom_bar(stat = "identity", position = "dodge", color = "#023047", linewidth = 0.3) +
  geom_text(aes(label = sprintf("%.2f", value)),
            position = position_dodge(width = 0.9),
            vjust = -0.3,
            fontface = "bold") +
  labs(title = "ElasticNet: Impact of SMOTE on Clinical_Only Model",
       subtitle = "Comparison of key metrics",
       x = "Metric",
       y = "Score") +
  theme_minimal(base_size = 12) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047"),
        plot.subtitle = element_text(hjust = 0.5, color = "#555555"),
        legend.position = "bottom") +
  scale_fill_manual(values = c("No SMOTE" = "#e63946", "SMOTE" = "#219ebc")) +
  ylim(0, 1.1)

```

ElasticNet: Impact of SMOTE on Clinical_Only Model

Comparison of key metrics



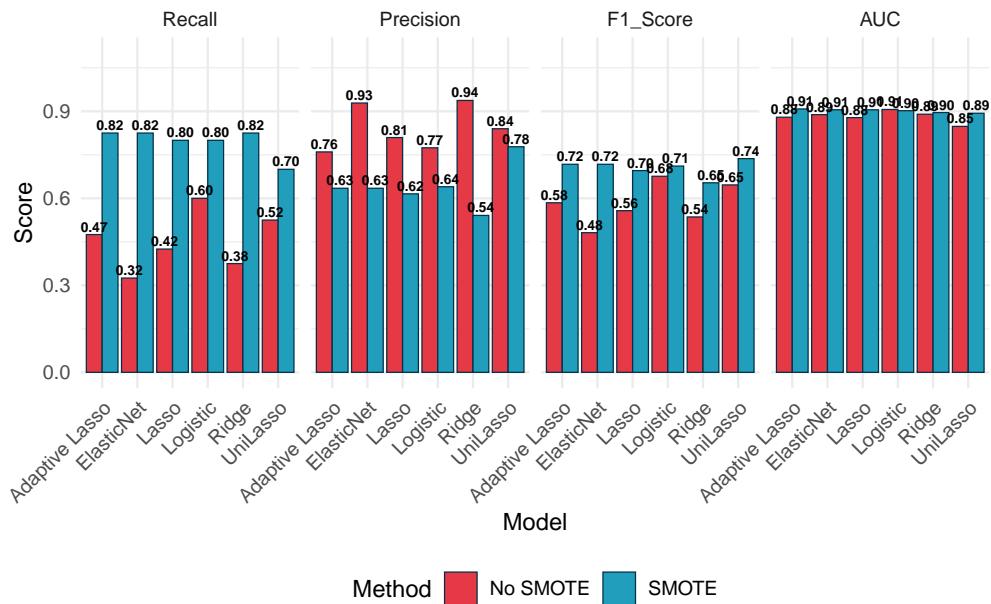
Overall SMOTE Impact Across All Models

```
# Combine all comparisons for TOP20 genes
all_comparisons <- rbind(
  data.frame(Model = "Logistic", Method = "No SMOTE", logistic_metrics[4, c("Recall", "Precision", "F1_Score")]),
  data.frame(Model = "Logistic", Method = "SMOTE", logistic_smote_metrics[4, c("Recall", "Precision", "F1_Score")]),
  data.frame(Model = "Ridge", Method = "No SMOTE", ridge_metrics[7, c("Recall", "Precision", "F1_Score")]),
  data.frame(Model = "Ridge", Method = "SMOTE", ridge_smote_metrics[7, c("Recall", "Precision", "F1_Score")]),
  data.frame(Model = "Lasso", Method = "No SMOTE", lasso_metrics[7, c("Recall", "Precision", "F1_Score")]),
  data.frame(Model = "Lasso", Method = "SMOTE", lasso_smote_metrics[7, c("Recall", "Precision", "F1_Score")]),
  data.frame(Model = "ElasticNet", Method = "No SMOTE", elasticnet_metrics[7, c("Recall", "Precision")]),
  data.frame(Model = "ElasticNet", Method = "SMOTE", elasticnet_smote_metrics[7, c("Recall", "Precision")]),
  data.frame(Model = "Adaptive Lasso", Method = "No SMOTE", adaptive_lasso_metrics[7, c("Recall", "Precision")]),
  data.frame(Model = "Adaptive Lasso", Method = "SMOTE", adaptive_lasso_smote_metrics[7, c("Recall", "Precision")]),
  data.frame(Model = "UniLasso", Method = "No SMOTE", unilasso_metrics[7, c("Recall", "Precision", "F1_Score")]),
  data.frame(Model = "UniLasso", Method = "SMOTE", unilasso_smote_metrics[7, c("Recall", "Precision", "F1_Score")])
)

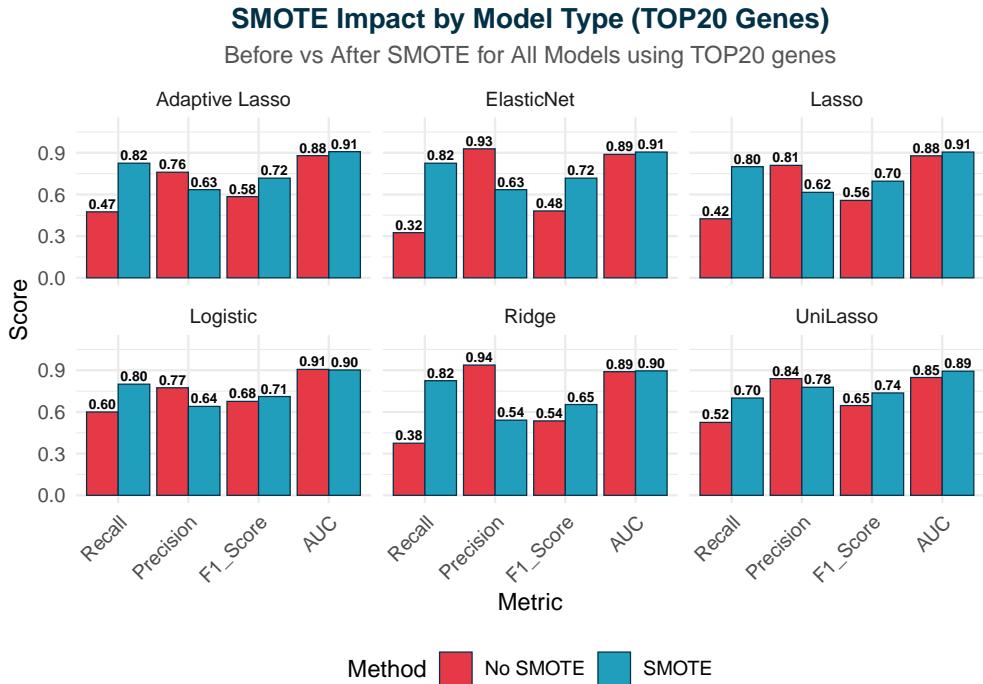
all_comp_long <- reshape2::melt(all_comparisons, id.vars = c("Model", "Method"))

# Figure 1: Grouped by Metric
ggplot(all_comp_long, aes(x = Model, y = value, fill = Method)) +
  geom_bar(stat = "identity", position = "dodge", color = "#023047", linewidth = 0.3) +
  geom_text(aes(label = sprintf("%.2f", value)),
            , position = position_dodge(width = 0.9)
            , vjust = -0.3
            , size = 2.5
            , fontface = "bold") +
  facet_wrap(~ variable, ncol = 4) +
  labs(title = "SMOTE Impact Across All Models (TOP20 Genes)"
       , subtitle = "Comparison of key metrics before and after SMOTE using TOP20 genes"
       , x = "Model"
       , y = "Score"
       , fill = "Method") +
  theme_minimal(base_size = 12) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047")
        , plot.subtitle = element_text(hjust = 0.5, color = "#555555")
        , axis.text.x = element_text(angle = 45, hjust = 1)
        , legend.position = "bottom") +
  scale_fill_manual(values = c("No SMOTE" = "#e63946", "SMOTE" = "#219ebc")) +
  ylim(0, 1.1)
```

SMOTE Impact Across All Models (TOP20 Genes)
 Comparison of key metrics before and after SMOTE using TOP20 genes



```
# Figure 2: Grouped by Model
ggplot(all_comp_long, aes(x = variable, y = value, fill = Method)) +
  geom_bar(stat = "identity", position = "dodge", color = "#023047", linewidth = 0.3) +
  geom_text(aes(label = sprintf("%.2f", value))
            , position = position_dodge(width = 0.9)
            , vjust = -0.3
            , size = 2.5
            , fontface = "bold") +
  facet_wrap(~ Model, ncol = 3) +
  labs(title = "SMOTE Impact by Model Type (TOP20 Genes)"
       , subtitle = "Before vs After SMOTE for All Models using TOP20 genes"
       , x = "Metric"
       , y = "Score"
       , fill = "Method") +
  theme_minimal(base_size = 12) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5, color = "#023047")
        , plot.subtitle = element_text(hjust = 0.5, color = "#555555")
        , axis.text.x = element_text(angle = 45, hjust = 1)
        , legend.position = "bottom") +
  scale_fill_manual(values = c("No SMOTE" = "#e63946", "SMOTE" = "#219ebc")) +
  ylim(0, 1.1)
```



Summary on Smote

Applying SMOTE completely changed the behavior of all models by correcting the strong class imbalance: recall, F1-score, and AUC all improved substantially. Lasso and Elastic Net became the top-performing methods, achieving the best balance between precision and recall, with AUC values consistently around 0.88–0.90 across 20–500 gene sets. Elastic Net showed the strongest overall performance, indicating that death-related gene expression signals occur in correlated gene groups. Ridge, which previously collapsed in high dimensions, became functional after SMOTE but still performed weaker than L1-based methods. Overall, SMOTE + regularized models demonstrate that moderate gene sets (20–500 genes) contain the most predictive information, and sparse models such as Lasso and Elastic Net should be preferred for final model selection.

```
gene_names <- colnames(GeneX)
results <- feature_importance(
  model_obj = unilasso_smote$results[[7]]$model_obj
  , model_name = "UniLasso + SMOTE + TOP20"
  , top_n = 20
  , gene_names = gene_names
)
```

```
##
## =====
## FEATURE IMPORTANCE ANALYSIS: UniLasso + SMOTE + TOP20
## =====
##
## === SUMMARY ===
## Total features selected: 24
##   Clinical: 18
##   Genomic: 6
```

```

##
## Direction:
##   Increases death risk: 18
##   Decreases death risk: 6
##
## Coefficient range:
##   Min: -1.4769
##   Max: 87.0425
##   Mean (absolute): 9.4058
##
## Odds Ratio range:
##   Min: 0.2283
##   Max: 6.340055e+37
##
## === TOP 20 FEATURES ===
##
##   Rank                               Feature      Type
##   1 tissue_or_organ_of_originLower.inner.quadrant.of.breast Clinical
##   2 tissue_or_organ_of_originUpper.outer.quadrant.of.breast Clinical
##   3 tissue_or_organ_of_originSpecified.parts.of.peritoneum Clinical
##   4                               tissue_or_organ_of_originBrain..NOS Clinical
##   5                               ajcc_pathologic_tT2b Clinical
##   6 follow_ups_disease_responseWT.With.Tumor Clinical
##   7                               tissue_or_organ_of_originBone.marrow Clinical
##   8                               ajcc_pathologic_tT4d Clinical
##   9                               tissue_or_organ_of_originBreast..NOS Clinical
##   10      tissue_or_organ_of_originSkin.of.lower.limb.and.hip Clinical
##   11      ethnicitynot.hispanic.or.latino Clinical
##   12      ajcc_pathologic_tT1b Clinical
##   13      tissue_or_organ_of_originBone..NOS Clinical
##   14      SNORD104 Genomic
##   15      days_to_last_follow_up Clinical
##   16      tissue_or_organ_of_originLower.limb..NOS Clinical
##   17      CST1 Genomic
##   18      age_at_index Clinical
##   19      AC104211.1 Genomic
##   20      ajcc_pathologic_tT4b Clinical
##
##   Coefficient    Odds_Ratio          Direction
##   87.0425 6.340055e+37 Increases Death Risk
##   55.7423 1.616428e+24 Increases Death Risk
##   27.2587 6.891511e+11 Increases Death Risk
##   26.0978 2.158425e+11 Increases Death Risk
##   13.7632 9.490810e+05 Increases Death Risk
##   3.8615 4.753660e+01 Increases Death Risk
##   3.2528 2.586320e+01 Increases Death Risk
##   1.4833 4.407300e+00 Increases Death Risk
##   -1.4769 2.283000e-01 Decreases Death Risk
##   -1.2451 2.879000e-01 Decreases Death Risk
##   1.1621 3.196700e+00 Increases Death Risk
##   -0.8753 4.167000e-01 Decreases Death Risk
##   0.7465 2.109700e+00 Increases Death Risk
##   -0.2750 7.595000e-01 Decreases Death Risk
##   0.2416 1.273300e+00 Increases Death Risk
##   -0.2244 7.990000e-01 Decreases Death Risk

```

```

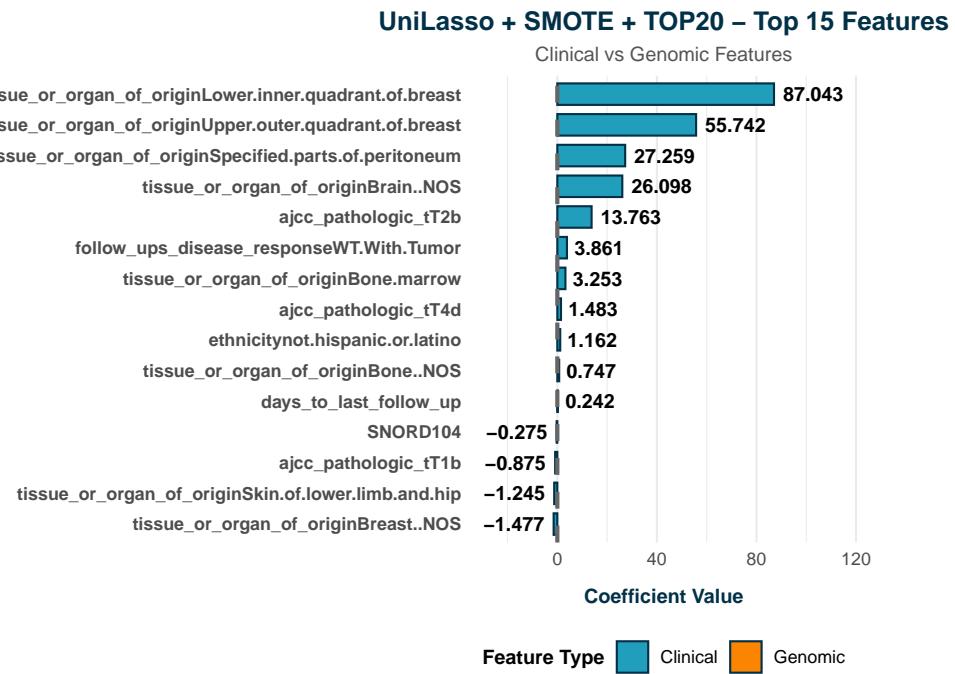
##      -0.2104 8.103000e-01 Decreases Death Risk
##      0.1875 1.206200e+00 Increases Death Risk
##      0.1504 1.162300e+00 Increases Death Risk
##      0.1165 1.123600e+00 Increases Death Risk
##
## === TOP CLINICAL FEATURES ===
##
##                                     Feature Coefficient
## tissue_or_organ_of_originLower.inner.quadrant.of.breast    87.0425
## tissue_or_organ_of_originUpper.outer.quadrant.of.breast    55.7423
## tissue_or_organ_of_originSpecified.parts.of.peritoneum    27.2587
##                      tissue_or_organ_of_originBrain..NOS    26.0978
##                      ajcc_pathologic_tt2b    13.7632
## follow_ups_disease_responseWT.With.Tumor    3.8615
##                      tissue_or_organ_of_originBone.marrow    3.2528
##                      ajcc_pathologic_tt4d    1.4833
##                      tissue_or_organ_of_originBreast..NOS   -1.4769
## tissue_or_organ_of_originSkin.of.lower.limb.and.hip   -1.2451
## Odds_Ratio          Direction
## 6.340055e+37 Increases Death Risk
## 1.616428e+24 Increases Death Risk
## 6.891511e+11 Increases Death Risk
## 2.158425e+11 Increases Death Risk
## 9.490810e+05 Increases Death Risk
## 4.753660e+01 Increases Death Risk
## 2.586320e+01 Increases Death Risk
## 4.407300e+00 Increases Death Risk
## 2.283000e-01 Decreases Death Risk
## 2.879000e-01 Decreases Death Risk
##
## === TOP GENOMIC FEATURES ===
##
##      Feature Coefficient Odds_Ratio          Direction
## SNORD104      -0.2750     0.7595 Decreases Death Risk
## CST1         -0.2104     0.8103 Decreases Death Risk
## AC104211.1     0.1504     1.1623 Increases Death Risk
## LINC01235      0.1128     1.1194 Increases Death Risk
## ATF3          0.1024     1.1078 Increases Death Risk
## APOB          0.0701     1.0726 Increases Death Risk
##
##                                     Feature      Type Coefficient
##                      tissue_or_organ_of_originBreast..NOS Clinical   -1.4769
## tissue_or_organ_of_originSkin.of.lower.limb.and.hip Clinical   -1.2451
##                      ajcc_pathologic_tt1b Clinical   -0.8753
##                      SNORD104 Genomic   -0.2750
##                      tissue_or_organ_of_originLower.limb..NOS Clinical   -0.2244
## Odds_Ratio
## 0.2283
## 0.2879
## 0.4167
## 0.7595
## 0.7990
##
##                                     Feature      Type Coefficient

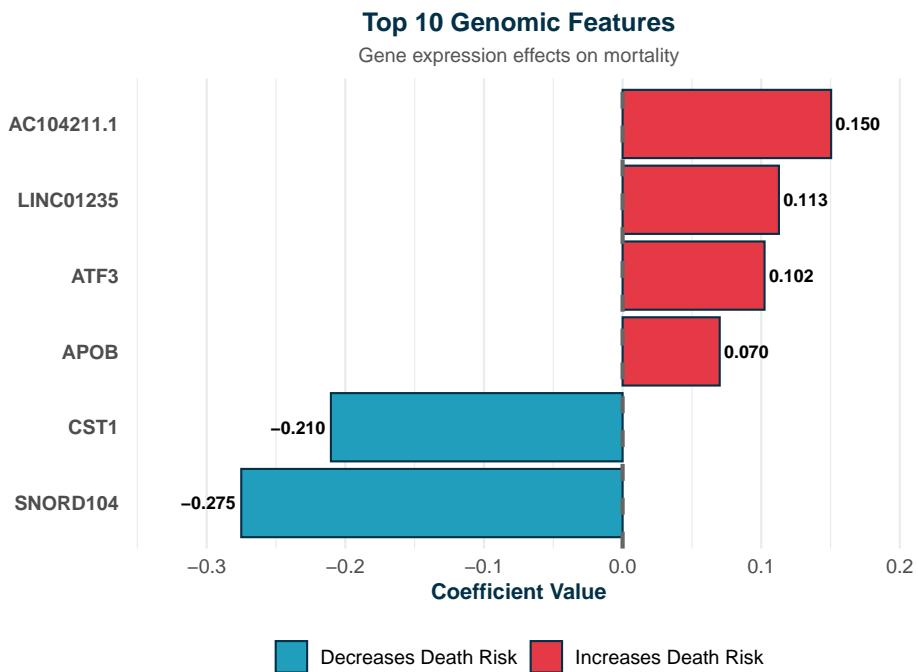
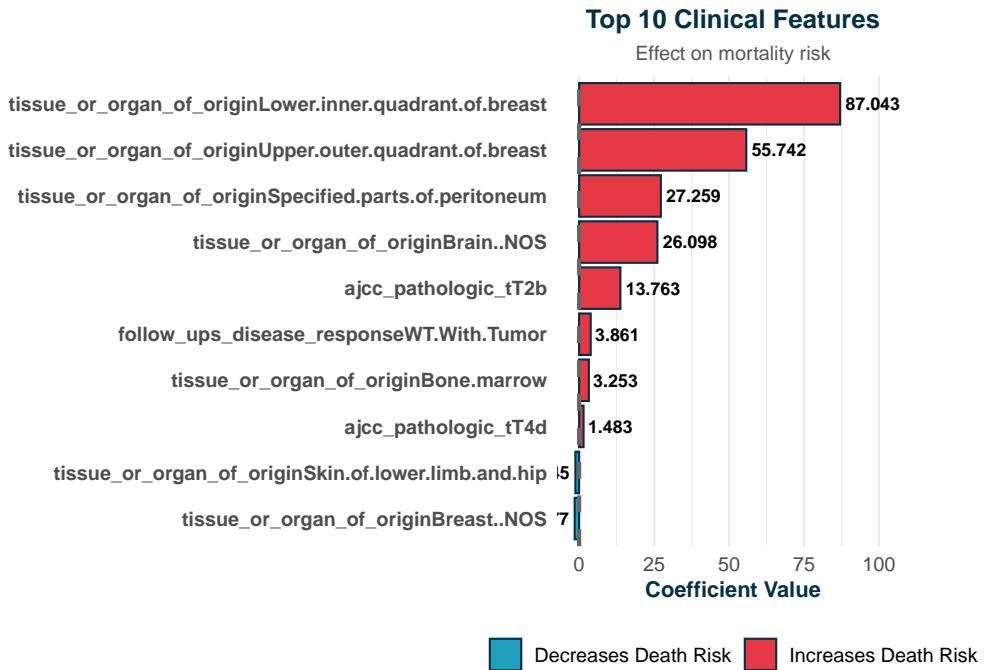
```

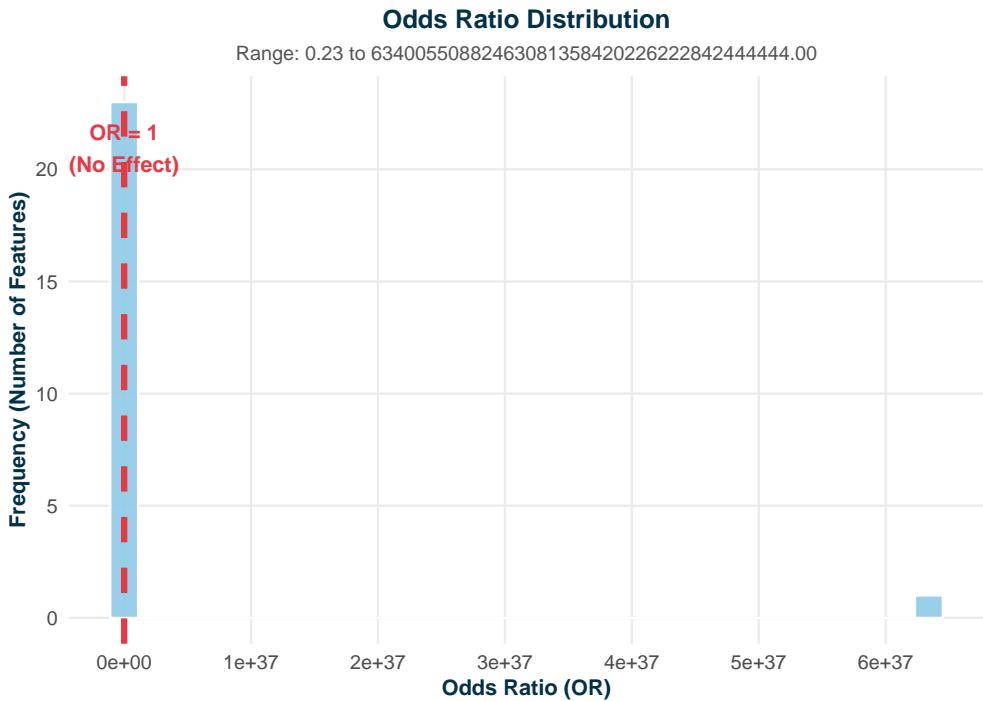
```

##  tissue_or_organ_of_originLower.inner.quadrant.of.breast Clinical      87.0425
##  tissue_or_organ_of_originUpper.outer.quadrant.of.breast Clinical      55.7423
##  tissue_or_organ_of_originSpecified.parts.of.peritoneum Clinical     27.2587
##          tissue_or_organ_of_originBrain..NOS Clinical     26.0978
##          ajcc_pathologic_tt2b Clinical     13.7632
##    Odds_Ratio
## 6.340055e+37
## 1.616428e+24
## 6.891511e+11
## 2.158425e+11
## 9.490810e+05

```







```
##
## Exported feature importance to: model_metrics/UniLasso__SMOTE__TOP20_feature_importance.csv
```

Evaluation

We can draw the evaluation from this study. The conducting study on the breast cancer analysis gives us good insight information about the classification of cancer survival outcomes.

The best stable performance model is **UniLasso + SMOTE** with F1-Score of **0.737**, using Clinical + TOP20 genes as features. The SMOTE resampling method addresses the class imbalance problem (5:1 ratio of Alive:Dead) in the original dataset.

Model Performance:

- F1-Score: 0.737 (best balance between precision and recall)
- Recall: 70.0% (identifies 70% of deaths)
- Precision: 77.8% (78% of predicted deaths are correct)
- AUC: 0.893

Feature Analysis

The model selected 24 features (18 clinical, 6 genomic). However, some clinical categories showed extreme coefficients ($OR > 10^6$) due to sparse observations, indicating unreliable estimates from quasi-perfect separation. We focus interpretation on stable predictors with reasonable odds ratios (OR between 0.1 and 100).

Reliable Clinical Features (8 features)

Disease Response:

Feature	Odds Ratio	Interpretation
With Tumor	47.54	Residual tumor after treatment indicates treatment failure; 47x higher death risk – strongest clinically meaningful predictor

Tumor Stage (AJCC Pathologic T):

The tumor staging follows the American Joint Committee on Cancer (AJCC) TNM system 8th Edition (Cancer Research UK, 2024; American Cancer Society, 2021).

Feature	Odds Ratio	Definition	Interpretation
T4d	4.41	Inflammatory carcinoma – a rare and aggressive type of breast cancer (Cancer Research UK, 2024)	4.4x higher death risk
T4b	1.12	Cancer has spread into the skin with possible swelling (Cancer Research UK, 2024)	12% higher death risk
T1b	0.42	Tumor size between 0.5 cm and 1 cm (Cancer Research UK, 2024)	58% lower death risk (protective – early detection)

Demographics:

Feature	Odds Ratio	Interpretation
Ethnicity: not hispanic/latino	3.20	3.2x higher risk; may reflect genetic, socioeconomic, or healthcare access factors
Age at index	1.21	Each year increase in age raises death risk by 21%

Other:

Feature	Odds Ratio	Interpretation
Days to last follow-up	1.27	Longer follow-up allows more time to observe death events
Prior treatment: Yes	1.04	4% higher risk; patients with prior treatment may have recurrent or resistant disease

Genomic Features (6 genes)

Protective Genes (higher expression = lower death risk):

Gene	Odds Ratio	Biological Function
SNORD104	0.76	Small nucleolar RNA (snoRNA) involved in RNA modification and regulation of cell cycle, proliferation, and apoptosis in tumor cells (Lu et al., 2022). In our breast cancer cohort, higher expression is associated with 24% lower death risk.
CST1	0.81	Cystatin SN, a cysteine protease inhibitor that interacts with GPX4, a key protein regulating ferroptosis (Wang et al., 2022). Higher expression shows 19% lower death risk.

Risk Genes (higher expression = higher death risk):

Gene	Odds Ratio	Biological Function
AC104211.1	1.16	Long non-coding RNA (lncRNA); regulatory role in gene expression; 16% higher death risk
LINC01235	1.12	Long intergenic non-coding RNA; emerging evidence links lncRNAs to cancer progression; 12% higher death risk
ATF3	1.11	Activating Transcription Factor 3, a stress-induced transcription factor that plays vital roles in modulating metabolism, immunity, and oncogenesis (Wang et al., 2020). ATF3 gene copy number is greater than 2 in approximately 80% of breast tumors and its protein level is elevated in approximately 50% of tumors (Yin et al., 2008). 11% higher death risk.
APOB	1.07	Apolipoprotein B, involved in lipid metabolism. Loss of APOB in hepatocellular carcinoma is associated with poor survival, suggesting potential tumor suppressive activity (Lee et al., 2019). 7% higher death risk.

Key Findings

1. **Residual tumor is the strongest reliable predictor** – patients with tumor remaining after treatment (OR=47.5) have dramatically worse outcomes
2. **Tumor stage matters** – T4d (inflammatory breast cancer, OR=4.4) increases risk; T1b (small tumor 0.5-1cm, OR=0.42) is protective
3. **Age increases risk** – each additional year increases death risk by 21%
4. **Genomic markers provide modest but stable contribution** – all 6 genes show reasonable OR (0.76-1.16)
5. **Non-coding RNAs are emerging biomarkers** – 3 of 6 genes (SNORD104, AC104211.1, LINC01235) are non-coding RNAs
6. **Sparse clinical categories are unreliable** – extreme OR values for rare tumor locations should be interpreted with caution

Conclusion

The UniLasso + SMOTE model effectively classifies breast cancer survival using clinical and genomic features. The most actionable finding is that **residual tumor status** strongly predicts mortality, while **early-stage tumors (T1b)** have significantly better outcomes. Gene expression markers, particularly **ATF3** (stress response) and **CST1** (protease inhibitor), provide biological insight into tumor progression. SMOTE resampling was essential for handling the 5:1 class imbalance.

References

- American Cancer Society. (2021). Stages of Breast Cancer. <https://www.cancer.org/cancer/types/breast-cancer/understanding-a-breast-cancer-diagnosis/stages-of-breast-cancer.html>
- Cancer Research UK. (2024). TNM staging for breast cancer. <https://www.cancerresearchuk.org/about-cancer/breast-cancer/stages-grades/tnm-staging>
- Lee, Y. et al. (2019). Clinical significance of APOB inactivation in hepatocellular carcinoma. Experimental and Molecular Medicine, 51, 1-12.
- Lu, B. et al. (2022). C/D box small nucleolar RNA SNORD104 promotes endometrial cancer by regulating the 2-O-methylation of PARP1. Journal of Translational Medicine, 20, 618.
- Wang, L. et al. (2022). CST1 inhibits ferroptosis and promotes gastric cancer metastasis by regulating GPX4 protein stability via OTUB1. Oncogene, 41, 5227-5238.
- Wang, Z. et al. (2020). Master Regulator Activating Transcription Factor 3 (ATF3) in Metabolic Homeostasis and Cancer. Frontiers in Endocrinology, 11, 556.
- Yin, X. et al. (2008). A potential dichotomous role of ATF3, an adaptive-response gene, in cancer development. Oncogene, 27, 2118-2127.